

Systems Thinking

From SE to STPA

C. Cabos

© 2026 C. Cabos

Table of Contents

1. Systems Thinking — From SE to STPA	10
1.1 Applying Systems Engineering Decomposition, Product Line Thinking, and STPA to Social and Technical Systems	10
1.2 Structure	10
1.3 The Core Insight	10
2. Introduction — What This Book Does and Why It Works	11
2.1 Start with a Puzzle	11
2.2 A Closer Example	11
2.3 The Pattern Appears Everywhere	12
2.4 What Engineering Knows	12
2.5 What This Book Does	13
2.6 What This Book Does Not Do	13
2.7 How to Read This Book	14
3. Part I — Foundations	15
4. What Is a System?	16
4.1 Before We Decompose, We Must Define	16
4.2 Theory I — Structure: Parts, Connections, and Goals	16
The Replaceability Hierarchy	16
Critical Assessment	16
4.3 Theory II — Behaviour: Emergence	17
Critical Assessment	17
4.4 Synthesis: The Working Definition	18
4.5 Why This Matters for Social Systems	18
5. The Five-Level SE Hierarchy	19
5.1 A Universal Decomposition Instrument	19
5.2 The Five Levels	19
5.3 Why Five Levels?	19
5.4 Traceability Rules	19
5.5 Example: The German Verein	20
6. Product Line Thinking	21
6.1 Platform, Variation Points, and Reuse	21
The Platform	21
Variation Points	21
Binding Decisions	21
The Reuse Gradient	22
The Coherence Rule	22

The Workaround Principle	22
7. STPA and STAMP	23
7.1 System-Theoretic Process Analysis	23
The STAMP Accident Model	23
The Four Steps of STPA	23
The Four UCA Types	23
From SE Decomposition to STPA	23
Rung tagging — the social-systems extension	24
References	24
8. Justificatory Rungs	25
8.1 What this chapter adds to the book	25
8.2 The seven rungs at a glance	25
8.3 The seven rungs in detail	26
Rung 0 — Power / coercion	26
Rung 1 — Authority / rhetoric / identity	26
Rung 2 — Formal consistency	26
Rung 3 — Empirical testability	27
Rung 4 — Cumulative evidence and consilience	27
Rung 5 — Meta-rational integration	27
Rung 6 — Normative legitimacy / deliberative ethics	27
What every controller carries on every loop	27
8.4 Why the ladder matters for STPA	28
Same rung is necessary for transmission	28
Same rung is not sufficient — and at low rungs it is dangerous	28
Different rungs on the two arrows of one loop	28
8.5 Three dangerous mismatches	28
Pattern A — Asymmetric loop	28
Pattern B — Claimed-rung inflation	28
Pattern C — Cross-loop rung imposition	29
8.6 Claimed rung vs operating rung	29
8.7 A normative caveat	29
8.8 Forward references	29
9. Part II — Social Systems	31
9.1 The Systems	31
Governance Family	31
Organization Family	31
Intimate Systems Family	31
9.2 Chapters	31

9.3 Additional worked examples	31
10. Religion — SE Decomposition	32
10.1 The Founding Example	32
10.2 The Decomposition	34
10.3 Key Structural Observations	34
10.4 See Also	35
11. Kingdom & Republic — A Reuse Analysis	36
11.1 The Constitutional Monarchy as a Product-Line Variant	36
11.2 Reuse Classification	37
Fully Reusable — Adopt Directly	37
Partially Reusable — Adapt with Modifications	38
Incompatible — Cannot Reuse Without Abolishing the Monarchy	38
12. Product Line Thinking for Social Systems	39
12.1 Applying Systems Engineering Decomposition and Platform Reuse to Human Institutions	39
12.2 Abstract	39
12.3 1. Introduction: Why Product Lines?	39
1.1 The Problem of Institutional Comparison	39
1.2 The Product Line Concept in Systems Engineering Terms	39
1.3 The Five-Level SE Hierarchy	40
12.4 2. The Decompositions: Ten Social Systems	40
12.5 3. The Social System Platform	40
12.6 4. The Variation Points	41
12.7 5. Cross-Family Reuse: Case Studies	41
12.8 6. Implications and Principles	41
12.9 7. Conclusion	41
12.10 Interactive Explorers	42
13. Governance as a Single Platform	43
13.1 Kingdom, Republic, Theocracy, and One-Party State as Four Bindings of One System	43
13.2 1. The Polity Platform	43
13.3 2. The Four VP1 Bindings	44
13.4 3. Reading the Pathologies off the Bindings	44
13.5 4. The Reuse Gradient Across Polities	45
13.6 5. What This Synthesis Buys	45
13.7 See Also	45
14. Example: Der Deutsche Verein (e.V.)	46
14.1 SE Decomposition	47
14.2 Variation Point Bindings	47
14.3 The Structural Pathology of Voluntary Democratic Participation	48

14.4	STPA Highlights	48
14.5	Fixes That Worked	48
14.6	Platform Mapping	49
15.	Example: Democracy (Modern Republic)	50
15.1	SE Decomposition	51
15.2	Variation Point Bindings	51
15.3	Why This System Is Structurally Unusual	52
15.4	STPA Highlights	52
15.5	Fixes That Worked: Germany's Basic Law (1949)	53
15.6	Platform Mapping	53
16.	Example: Theocracy	54
16.1	SE Decomposition	55
16.2	Variation Point Bindings	55
16.3	The Structural Pathology of VP1 = Divine Mandate	56
16.4	STPA Highlights	56
16.5	Fixes That Worked	56
16.6	Platform Mapping	57
17.	Example: Corporation (For-Profit)	58
17.1	SE Decomposition	59
17.2	Variation Point Bindings	59
17.3	The Structural Pathology of VP1 = Capital Ownership	60
17.4	STPA Highlights	60
17.5	Fixes That Worked	60
17.6	Platform Mapping	61
18.	Example: University	62
18.1	SE Decomposition	63
18.2	Variation Point Bindings	63
18.3	The Structural Pathology of VP1 = Scholarly Merit	63
18.4	STPA Highlights	64
18.5	Fixes That Worked	64
18.6	Platform Mapping	65
19.	Example: Military	66
19.1	SE Decomposition	67
19.2	Variation Point Bindings	67
19.3	The Structural Pathology of VP3 = Command Order	68
19.4	STPA Highlights	68
19.5	Fixes That Worked	68
19.6	Platform Mapping	69

20. Example: The Family	70
20.1 SE Decomposition	71
20.2 Variation Point Bindings	71
20.3 The Structural Pathology of Parental Authority over Non-Consenting Members	71
20.4 STPA Highlights	72
20.5 Fixes That Worked	72
20.6 Platform Mapping	73
21. Example: One-Party State	74
21.1 SE Decomposition	75
21.2 Variation Point Bindings	75
21.3 The Structural Pathology of VP3 = Party Committee Decision	75
21.4 STPA Highlights	76
21.5 Fixes That Worked	76
21.6 Platform Mapping	77
22. Example: The School (Secondary Education)	78
22.1 SE Decomposition	80
22.2 Variation Point Bindings	80
22.3 The Structural Pathology of Proxy-Metric Substitution	80
22.4 STPA Highlights	80
22.5 Fixes That Worked	81
22.6 Platform Mapping	82
23. Example: The Hospital (Acute Healthcare)	83
23.1 SE Decomposition	85
23.2 Variation Point Bindings	85
23.3 The Structural Pathology of Triadic Membership	85
23.4 STPA Highlights	85
23.5 Fixes That Worked	86
23.6 Platform Mapping	87
24. Example: Wikipedia / Open-Source Project	88
24.1 SE Decomposition	89
24.2 Variation Point Bindings	90
24.3 The Structural Pathology of Rough Consensus	90
24.4 STPA Highlights	90
24.5 Fixes That Worked	91
24.6 Platform Mapping	92
25. Part III — Development Frameworks	93
25.1 The Frameworks	93
25.2 Chapters	93

26. Product Line Thinking for Development Frameworks	94
26.1 How Waterfall, Scrum, Kanban, V-Model, PRINCE2, Design Thinking, DevOps, and SAFe Share a Common Platform — and How to Merge Them Better	94
26.2 Abstract	94
26.3 1. Introduction	94
1.1 The False Dichotomy	94
1.2 Scope: Eight Frameworks	95
26.4 2. The Decompositions	95
26.5 3. The Development Framework Platform	95
26.6 4. The Variation Points	96
26.7 5. Reuse Analysis Across Frameworks	96
26.8 6. Three Proposed Hybrid Architectures	96
26.9 7. Principles for Framework Merging	96
26.10 8. Conclusion	97
27. The Framework Platform	98
27.1 Ten Universal Functional Slots in Development	98
27.2 Four Variation Points	98
28. Hybrid Architectures	99
28.1 Three Principled Mergers	99
Hybrid A: "Governed Agile"	99
Hybrid B: "Discovery-Driven Flow"	99
Hybrid C: "Scaled Governed Flow"	99
28.2 Five Design Principles	99
29. Part IV — STPA Applied	100
29.1 Chapters	100
30. STPA Analysis of Religion as a Social System	101
30.1 System-Theoretic Process Analysis Applied to the Systems Engineering Decomposition of Religion	101
30.2 Abstract	101
30.3 1. Background: The SE Decomposition of Religion	102
1.1 The Five-Level Hierarchy	102
1.2 Why STPA?	102
30.4 2. STPA Step 1: Define Purpose of the Analysis	103
2.1 System-Level Losses	103
2.2 System-Level Hazards	103
2.3 System-Level Constraints	104
30.5 3. STPA Step 2: Model the Control Structure	104
3.1 The STAMP Control Structure	104
3.2 Why these control loops — and why only these?	105

3.3 Key Observations about the Control Structure	107
30.6 4. STPA Step 3: Identify Unsafe Control Actions	108
4.1 UCAs for Central Doctrinal Authority → Clergy	108
4.2 UCAs for Clergy → Congregation	108
4.3 UCAs for Congregation → Individual	109
4.4 UCAs for Individual (Self-Control / Internalized Beliefs)	109
30.7 5. STPA Step 4: Identify Loss Scenarios	110
5.1 Loss Scenario LS-1: Doctrine Fails to Self-Correct (UCA-1, UCA-6 → H1 → L1, L6)	110
5.2 Loss Scenario LS-2: Abuse Cover-Up (UCA-2, UCA-15 → H2, H5 → L2, L4, L5)	110
5.3 Loss Scenario LS-3: Radicalization through In-Group/Out-Group Escalation (UCA-3, UCA-10 → H3 → L3)	111
5.4 Loss Scenario LS-4: Spiritual Coercion and Psychological Damage (UCA-11, UCA-16, UCA-17, UCA-20 → H4, H8 → L1, L4)	111
5.5 Loss Scenario LS-5: Financial Exploitation (UCA-8, UCA-13 → H7 → L2, L5)	112
5.6 Loss Scenario LS-6: Suppression of Individual Conscience (UCA-4, UCA-7, UCA-21 → H1, H4 → L1, L6)	112
30.8 6. Cross-Cutting Structural Findings	113
6.1 The Feedback Problem Is Universal	113
6.2 The Self-Sealing Model Problem	113
6.3 The Conflicting-Goals Problem	113
6.4 The rung mismatch behind the four findings	114
30.9 7. Recommendations: Architectural Improvements	114
7.1 Close the Feedback Loop	114
7.2 Make Process Models Updatable	114
7.3 Separate Control Functions	114
7.4 Install Circuit Breakers	115
7.5 Protect Vulnerable Populations	115
30.10 8. Conclusion	115
References	115
31. STPA Analysis of Frontier AI as a Sociotechnical System	116
31.1 What Will AI Do to Us? A Systems-Engineering Reading	116
31.2 Abstract	116
31.3 1. The questions, and why the usual answers do not satisfy	116
31.4 2. Two ways of thinking about the system	117
31.5 3. The system	118
3.1 The first mismatch: lab safety frameworks claim more legitimacy than they operate	119
3.2 The second mismatch: capital cannot evaluate the laboratories it funds	119
3.3 The third mismatch: the labour-policy cycle cannot keep up with the capability-deployment cycle	120
3.4 The fourth mismatch: incidents inside the labs do not reach outside observers	120
3.5 The fifth mismatch: antitrust looks at mergers; concentration happens in contracts	120

3.6 What having five mismatches in hand changes	121
31.6 4. The nine ways the system can fail	121
4.1 Concentration entrenches at the infrastructure layer	121
4.2 Lab safety frameworks become ceremonial	122
4.3 Workers cannot keep up with capability deployment	122
4.4 The productivity surplus accrues to capital and stays there	123
4.5 The open-weights price-setting mechanism collapses	123
4.6 Compute competition escalates to kinetic conflict	124
4.7 The Chinese mirror	124
4.8 The information environment degrades faster than authentication scales	124
4.9 Rules calibrated to capability tiers already obsolete	125
4.10 What having nine scenarios gives us	125
31.7 5. What can be done	125
5.1 The four-phase sequence	126
5.2 The non-state actor's leverage	127
5.3 Reversibility	127
31.8 6. What the analysis does not do	128
31.9 7. Coda: the question reformulated	128
32. Control Structures in Social Systems	130
32.1 The Universal Control Pattern	130
32.2 The Ten Systems: Control Structure Profiles	130
32.3 The Four Most Dangerous Structural Patterns	131
33. Architectural Remedies	132
33.1 The Analysis Generates the Fix	132
33.2 The Five Worked Cases	132
1. Democracy (Republic) — the Weimar → Basic Law transition	132
2. Corporation — Sarbanes-Oxley, Mitbestimmung, and Cadbury	132
3. Religion — mandatory reporting and Vatican II	132
4. University — pre-registration, Registered Reports, and DORA	133
5. Military — parliamentary authorisation and independent prosecution	133
33.3 Cross-System Pattern	133
33.4 Reading the Remedies as Rung-Elevation Moves	133
34. About This Book	135
34.1 Origin	135
34.2 Approach	135
34.3 Technology	135
34.4 License	135

1. Systems Thinking — From SE to STPA

1.1 Applying Systems Engineering Decomposition, Product Line Thinking, and STPA to Social and Technical Systems

A place to start

In 1975, Kodak built the world's first digital camera — and buried it. Twenty years later, digital photography destroyed the company that invented it. The engineers were excellent. The technology was real. The failure was structural: the organisation was wired to protect film revenues, and it reliably smothered anything that threatened them. No personnel change could have fixed it. The design had to change. Most institutional failures look like this. **The Introduction** develops this argument through several cases before any framework or equation appears.

This book explores a single, powerful idea: **the tools we use to engineer safe aircraft and reliable software apply with surprising precision to human institutions** — political regimes, religions, corporations, families, and development methodologies.

We use three frameworks, layered on top of each other:

1. **The Five-Level SE Hierarchy** — decompose any system into Goals → Requirements → Functions → Logical Architecture → Physical Implementation, with full traceability.
2. **Product Line Thinking** — compare systems by identifying their shared platform (what's common) and their variation points (where they diverge), then study which elements can be reused across system types.
3. **STPA (System-Theoretic Process Analysis)** — model the system as a hierarchical control structure and systematically identify unsafe control actions that produce outcomes opposite to the system's stated goals.

1.2 Structure

The book is organized in four parts:

	Focus
Introduction	Why structural analysis beats moral exhortation — argued through concrete cases
Part I — Foundations	What is a system? The three frameworks: SE hierarchy, product lines, STPA
Part II — Social Systems	Ten social systems decomposed and compared
Part III — Development Frameworks	Waterfall, Scrum, Kanban, DevOps, and others — analysed as a product family
Part IV — STPA Applied	STPA applied to religion and other social systems

1.3 The Core Insight

Social systems look more different from each other than they actually are. A kingdom and a republic, a church and a corporation, a family and a Verein — these seem like categorically different things. But when decomposed through a rigorous systems engineering hierarchy, they reveal a **shared platform** of functional slots, and they differ primarily at a small number of well-defined **variation points**.

Product line thinking does not reduce human institutions to interchangeable parts. It does the opposite: it makes the genuine differences visible by stripping away the illusion that everything is different.

And STPA reveals that the most harmful outcomes of social systems are not random malfunctions but **predictable consequences of their control architecture** — consequences that can be addressed through architectural redesign, not just moral exhortation.

2. Introduction — What This Book Does and Why It Works

2.1 Start with a Puzzle

In the 1970s, Kodak employed more than 60,000 people in Rochester, New York. The company was profitable, technically excellent, and entirely dominant in its market. It also, in 1975, built the world's first digital camera. The engineer who built it, Steve Sasson, presented the prototype to management. The response was cautious appreciation and a quiet burial. The product was not released. Twenty years later, digital photography destroyed Kodak's business, and the company filed for bankruptcy in 2012.

What happened? The simple story is that Kodak's managers were short-sighted or afraid. But that explanation predicts the wrong fix. If you replace short-sighted managers with visionary ones, and nothing else changes, you get the same outcome — because the problem was not the people. Kodak had a different problem, one that no personnel change could solve.

Kodak's film business generated enormous profit margins. Every internal reporting line, every budget process, every promotion criterion was calibrated to protect and extend that revenue stream. The digital camera project reported into the film division. Its budget competed against film products. Its success was measured against metrics designed for film. The structure of the organisation — the connections between its parts and the goals those connections served — guaranteed that any technology threatening film would be smothered, regardless of who was in charge. Kodak did not fail because of bad people. It failed because its organisational structure contained an architectural flaw that turned a threat into an impossibility.

This is a systems problem. And systems problems have systems solutions.

2.2 A Closer Example

Consider something more familiar: a secondary school.

The school's declared goal is to educate students — to develop their capacity to think, reason, create, and contribute. Nobody disputes this. Teachers enter the profession because they believe in it. Administrators genuinely want students to succeed. Parents care about their children's development. The people are good and their intentions are aligned.

And yet, in many such schools, the following happens: students who understand nothing about history can pass a history exam. Teachers who are brilliant at inspiring curiosity get poor performance reviews. Students optimise for grades rather than understanding, teachers optimise for exam results rather than education, and school administrators optimise for league table rankings rather than learning. Everybody does their job correctly. The system produces the wrong outcomes.

Why? Because the connections in the system — the feedback loops that link behaviour to reward — do not connect to the declared goal. They connect to a measurable proxy for the goal. Exam scores are measurable; understanding is not. The moment you attach resources, careers, and prestige to exam scores, the whole system reorients toward producing exam scores. This is not cynicism — it is the mechanical consequence of the control structure.

The fix is not to motivate teachers more, threaten students more, or exhort administrators to care about real learning. None of those interventions touch the structure. The fix requires changing what is measured, how it feeds back, and what it controls — which means changing the connections. That is an architectural intervention, not a moral one.

But what does that architectural fix look like, concretely? This is where the engineering approach earns its keep. Once you have traced the causal chain — teacher evaluation derives from class exam scores, school ranking derives from published test data, university admission derives from GPA — each link in that chain becomes a candidate for redesign:

Misconfigured connection	Architectural remedy
Teacher appraisal tied only to exam averages	Add peer review of lesson quality and longitudinal tracking of how former students actually perform
School ranking uses only test results	Require portfolio assessment, inspection reports, and employer or university feedback as parallel signals
Single high-stakes exam controls university entry	Reduce the stakes of any one instrument; distribute the selection signal across multiple independent measures

The remedies are not exotic. Several school systems have implemented versions of them. Finland removed published school rankings and standardised national testing before age 16; teacher evaluation shifted to professional peer review; the feedback loops changed. The outcomes changed with them — not because Finnish children or teachers are different, but because the control structure is different.

This is the pattern the methods in this book are designed to produce: not merely a diagnosis of what went wrong, but a precise identification of which connection to change — and confidence that changing it will change the outcome. The analysis generates the remedy. The two are the same work.

2.3 The Pattern Appears Everywhere

Once you see this pattern, you find it in every kind of institution.

Religions often declare a goal of compassion, care for the vulnerable, and guidance toward transcendence. The institutional structure — hierarchy, doctrine, territory, resources — has its own logic. When institutional survival conflicts with the declared goal, the structure tends to win. The result is not hypocrisy by individuals but an architectural conflict between declared goals and operative feedback loops. The Catholic Church's response to the sexual abuse crisis is inexplicable as individual moral failure at scale; it is entirely explicable as the output of a control structure in which the preservation of institutional reputation was connected, through many feedback loops, to the decision-making of bishops — while the protection of victims was not.

Democracies are designed on the premise that distributing power across competing institutions prevents any single actor from accumulating enough control to subvert the system. When democracies fail, the failure is rarely a sudden coup. It is typically a sequence of individually legal changes — each one defensible in isolation — that progressively weakens the connections between institutions, concentrates decision-making, and removes the feedback mechanisms (free press, independent courts, competitive elections) that generate corrective signals. The structure changes before anyone notices the system has changed.

Corporations fail not because they hire bad people but because they build control structures that reward the wrong behaviour. The 2008 financial crisis did not require criminal intent at every level. It required a mortgage originator whose bonus was tied to volume, not quality; a securitiser whose fee was tied to deal flow, not risk; a rating agency whose revenue depended on the issuers it rated; and a regulator whose authority did not extend to the instruments being created. Each actor responded rationally to the incentives their position created. The system produced catastrophe.

Families can be analysed the same way. A family that communicates only through the children, where no direct channel exists between parents, has a structural problem. Replacing the children with different children — or telling the existing children to communicate more clearly — does not address it. The connection pattern needs to change.

In each case, the same diagnosis applies: the declared goals of the system diverge from the operative goals encoded in its feedback structure, and the people inside it respond to the operative goals. The fix requires architectural change — changing connections, changing what the system measures, changing what triggers what — not just changing the people or the rhetoric.

2.4 What Engineering Knows

Engineers have been building complex systems — aircraft, nuclear plants, air traffic control, spacecraft — for decades, and those systems must work. The cost of failure is immediate and physical. This pressure has produced a set of analytical disciplines that are ruthlessly practical: they do not ask whether the people involved are good; they assume people will respond to the structure they are in, and they ask whether the structure is safe.

Three of those disciplines are central to this book.

Systems Engineering decomposition takes any system — a jet engine, a hospital, a political party — and breaks it into five layers: the goals it pursues, the requirements those goals generate, the functions it must perform, the logical architecture that organises those functions, and the physical components that realise that architecture. This decomposition makes the structure explicit and legible. You cannot analyse what you cannot see, and most social systems are analysed at the level of individuals and events rather than at the level of structure. The SE hierarchy changes that.

Product Line Thinking observes that many systems that look very different are actually variations on a common underlying architecture. Boeing builds multiple aircraft types from a shared engineering platform. IKEA produces thousands of products from a shared manufacturing and distribution system. This book argues that kingdoms and republics, corporations and universities, religions and militaries, are variations on a shared social system platform — and that understanding which elements are shared and which are genuinely different is more useful than treating each system as entirely unique.

STPA — System-Theoretic Process Analysis — was developed by Nancy Leveson at MIT after studying catastrophic failures in complex socio-technical systems. Its central insight is that most catastrophic failures in modern systems do not result from component breakdown; they result from control failures — situations where the commands, feedback, and authority relationships between parts of the system produce unsafe actions even when each part is functioning normally. STPA does not ask "what broke?" It asks "what could the control structure do that it should not, and what is it not doing that it should?" Applied to social systems, it produces a systematic map of the ways an institution can harm the people it was built to serve — not through individual malice, but through structural design.

2.5 What This Book Does

This book takes those three engineering tools and applies them to the social systems most of us live inside.

We decompose ten social systems — the kingdom, the republic, the theocracy, the one-party state, the corporation, the university, the military, the family, the church, and the Verein — through the five-level SE hierarchy. Each decomposition produces a table: goals at the top, physical implementation at the bottom, with traceable connections between every level. The decomposition makes visible what is usually invisible: what the system is actually for, how it is actually organised, and where the critical connections are.

We then apply Product Line Thinking to compare these decompositions. The result is striking: beneath the obvious differences in flag, creed, and custom, these ten systems share a common architectural platform of ten functional slots that every human institution must fill. They differ at a small number of variation points — source of authority, membership boundary, decision-making mechanism, succession mechanism, legitimation narrative, and norm enforcement. Understanding the variation points tells you what can be changed and what cannot: institutional reform that crosses variation-point constraints will fail, not because of bad intentions, but because the structural constraints are real.

Finally, we apply STPA to trace how these systems can turn against their own stated goals. We show how the control structure of a religion can produce exactly the outcomes the religion forbids. We show how the control structure of a democracy can produce authoritarian outcomes if the feedback mechanisms are progressively dismantled. And we identify, for each type of system, the architectural features — the design choices — that make harmful outcomes more or less likely.

2.6 What This Book Does Not Do

It does not argue that all institutions are secretly the same, or that cultural and historical particularity does not matter. The variation points are real and consequential. A theocracy and a republic are genuinely different systems, and those differences explain a great deal about how they behave.

It does not argue that structural analysis replaces moral and political judgement. It is a precondition for it. You cannot judge a system's design wisely if you cannot see the design.

And it does not promise that structural reform is easy. The Kodak story illustrates the problem: the people inside a system are rewarded by the system's operative goals. Those who propose changing the structure threaten the operative goals of the people the structure rewards. Structural reform is resisted not because people are perverse but because the structure makes resistance rational.

What this book does promise is that structural analysis makes the conversation more honest. When we understand that harmful outcomes are predictable consequences of identifiable architectural features, we can argue about design instead of character. That is a more productive argument — and a fairer one.

2.7 How to Read This Book

If you want the examples first — continue reading here and move directly to [Part II](#), which applies all three frameworks to ten social systems. The theoretical foundations are in Part I, but the examples in Part II are self-contained.

If you want the theory first — [Part I](#) develops the three frameworks systematically, starting with a precise definition of what a system is and ending with the full STPA methodology.

If you want the most unsettling chapter — go directly to the [STPA analysis of religion](#) in Part IV. It is the worked example that brings all three frameworks together and shows, in formal detail, how a system designed to serve the transcendent can be structurally configured to serve itself instead.

The introduction you just read is the argument in compressed form. The rest of the book is the evidence.

3. Part I — Foundations

This part introduces the foundational definition of a system and the three analytical frameworks that form the backbone of the book. Each framework is well-established in engineering practice; our contribution is applying them systematically to social and organizational systems.

- **What Is a System?** — the foundational definition
- **The Five-Level SE Hierarchy** — the decomposition instrument
- **Product Line Thinking** — the comparative methodology
- **STPA and STAMP** — the safety and control analysis
- **Justificatory Rungs** — the epistemic-standard extension to STPA for social systems

4. What Is a System?

Canonical reference

The normative statement of the foundational concepts in this chapter — the working definition, the replaceability hierarchy, emergence, boundaries, and the distinction between technical and non-technical systems — is maintained as the project's knowledge base under [knowledge/foundations/](#). This chapter is the narrative discussion; the knowledge base is the reference.

4.1 Before We Decompose, We Must Define

Every analytical framework rests on an ontology — a theory of what kinds of things exist. Before applying the SE hierarchy, product line thinking, or STPA, we need a precise answer to the question that precedes all of them: what makes something a system in the first place?

Two families of answers dominate the literature. They are not mutually exclusive; this chapter argues they are complementary and that understanding the tension between them clarifies what our analytical tools are actually doing.

4.2 Theory I — Structure: Parts, Connections, and Goals

The first and most widely held view holds that a system is constituted by three elements:

1. **Parts** — the distinguishable components that can, in principle, be identified and listed
2. **Connections** — the relationships, flows, and dependencies between parts
3. **Goals** — the purpose or purposes the system pursues, whether declared or operative

The Replaceability Hierarchy

A significant observation follows from this ontology: the three elements are not equally essential to system identity, and they differ in how easily they can be replaced without the system ceasing to be the same system.

Element	Replaceability	Consequence of Change
Parts	High — most easily swapped	System continues; this is routine maintenance
Connections	Moderate — harder to change	System is restructured; behaviour shifts significantly
Goals	Low — very difficult to change	System is transformed; what was there before is effectively a different system

A human body replaces most of its cells over years and remains the same organism. A corporation can change its entire workforce and remain recognisably the same corporation. But a corporation whose goal shifts from profit to public service is, in any meaningful sense, a different institution wearing the same legal clothes. The Roman Catholic Church that abandoned the salvation of souls as its operative goal would not be a reformed Catholic Church — it would be something else entirely occupying the same physical infrastructure.

This hierarchy maps directly onto the SE decomposition used throughout this book. Physical implementations (P) correspond to parts. Logical architecture and functions (L, F) encode the connection patterns. Goals and requirements (G, R) encode purpose. Institutional reform that touches only P-level elements is maintenance. Reform that reaches G-level is transformation — and the historical record shows it is both rare and destabilising.

Critical Assessment

The structural theory is productive and largely correct, but it has four significant gaps.

1. Goals are often contested, latent, or split. The framework assumes a system has goals. Real social systems typically have declared goals and operative goals that diverge substantially. A regulatory agency declared to protect the public may operate primarily to protect the regulated industry. A religion declared to save souls may operate primarily to perpetuate its own institutional authority. The structural theory needs to distinguish these two types of goal, or it will mistake the label for the contents. This book addresses the divergence directly in the STPA analysis (Part IV), where unsafe control actions typically arise from exactly this split.

2. Connections may outrank goals as the determinant of system identity. Donella Meadows argued that the feedback structure of a system — its loops, delays, and reinforcing dynamics — determines its long-term behaviour more than any stated goal. A system whose goals are replaced but whose feedback loops remain intact will tend to revert to its former behaviour over time. The Soviet Union changed its declared goals repeatedly; the feedback structure of centralised authority, information suppression, and reward for loyalty produced remarkably consistent outcomes across six decades and several ideological pivots. The replaceability ordering may need revision: connections as structure may be more fundamental to identity than goals as declaration.

3. Every part is itself a system. The framework is valid but incomplete: it defers the definition rather than grounding it. A part is a system with its own parts, connections, and goals. This recursion is not a defect — it reflects the genuine nested structure of complex systems — but it means the theory requires a boundary condition: a decision about where to stop decomposing. That decision is not given by the theory itself.

4. The boundary question is missing. The most important definitional act for any system is drawing its boundary: what is inside and what is outside. For social systems this is particularly fraught. Who is "in" a state? Citizens only? Residents? Who counts as a member of a religion — the registered faithful, the active practitioner, the nominal cultural adherent? The boundary choice determines what counts as a part, what counts as an internal connection, and what counts as an external relationship. A theory of systems that omits the boundary is like a theory of cells that omits the membrane.

4.3 Theory II — Behaviour: Emergence

The second view defines systems not by what they are made of but by what they do that their parts cannot. A system is characterised by **emergent properties** — properties that are present in the organised whole but absent from any of the individual parts considered in isolation.

No individual neuron is conscious. No individual water molecule is wet. No individual trader sets market prices. Consciousness, wetness, and market prices are real; they exist; but they exist only at the level of the organised system, not at the level of its parts. If emergence is present, something is a system. If it is absent — if the whole is simply the sum of its parts — what we have is an aggregate, not a system.

This provides an observational test: can the behaviour of the whole be predicted by summing the behaviours of the parts? If yes: aggregate. If no: system.

Critical Assessment

The emergence criterion captures something real and important. But it has three significant weaknesses that limit its usefulness as a definition of system.

1. Emergence is descriptive, not constitutive. The emergence criterion tells you what to observe once you have already decided you are looking at a system; it does not tell you how to identify or engineer one. You cannot look for emergence before you have decided what to treat as the parts — and that decision already presupposes a structural theory. Emergence is the output of a system's structure, not the criterion that defines it. For an SE-based approach, this is a practical limitation: you design structure and observe emergence as a consequence; you cannot design emergence directly.

2. The criterion is trivially satisfied by aggregates. A pile of sand has a height and an angle of repose that no individual grain possesses. A crowd has a density that no individual person has. These are properties present in the aggregate but absent from the parts — but no one seriously claims a pile of sand is a system in the relevant sense. The emergence criterion must be qualified: we need to specify which emergent properties matter. The relevant ones are goal-directed or self-maintaining behaviours — and as soon as we say this, we have smuggled goal back in through the back door, returning to Theory I.

3. The strong/weak emergence distinction matters. Philosophers distinguish weak emergence (in principle reducible to lower-level description, but practically too complex to compute) from strong emergence (genuinely irreducible, a new causal power not

explainable even in principle from the parts). Almost all emergence in social systems is weak: complex, non-linear, and surprising, but not metaphysically mysterious. Treating social emergence as strong emergence — as if "the system" were a force beyond analysis — is an error that forecloses the very analysis this book attempts.

4.4 Synthesis: The Working Definition

The two theories are not rivals. They are the **inside view** and the **outside view** of the same reality.

- **From the inside** (Theory I): a system is a structure — organised parts, patterned connections, directing goals, within a defined boundary.
- **From the outside** (Theory II): a system is an actor — something whose boundary-level behaviour cannot be read off from its parts in isolation.

Neither is sufficient alone. A structure with no emergent behaviour is a machine, not a system in the interesting sense. An emergent phenomenon with no describable structure is magic, not analysis.

Adding the boundary gives the working definition this book operates with throughout:

Working Definition

A system is a **bounded structure of parts and connections organised by goals**, whose behaviour at the boundary **cannot be predicted from its parts in isolation**.

Four corollaries follow from this definition and recur throughout the analysis:

1. **Changing parts is maintenance.** It preserves the system.
2. **Changing connections restructures the system.** It changes what the system does, often dramatically.
3. **Changing goals transforms the system.** The result is a different system with inherited infrastructure.
4. **Changing the boundary redefines the system.** What counts as internal, external, and relational changes entirely.

The SE hierarchy in the next chapter is a formal instrument for making this structure explicit and tractable. Goals become the G-level. Connections between functional components become the L-level. Parts become the P-level. The boundary is the decision to start the decomposition at all.

4.5 Why This Matters for Social Systems

Technical systems — aircraft, software, power grids — are designed with explicit goals, documented connections, and identifiable parts. Their structure is legible because engineers made it so.

Social systems — states, religions, corporations, families — were not designed in the same deliberate sense. Their goals are often implicit, their connections informal, and their boundaries contested. The structural theory predicts that this makes them harder to reform: if the connection patterns are not documented, you cannot know what you are changing. If operative goals are not distinguished from declared goals, reform efforts target the label and leave the engine untouched.

The emergent-behaviour theory adds a further warning: social systems produce outcomes that no participant intended and no individual can fully understand. This is not mysticism — it is the practical consequence of non-linear feedback in complex structures. It means that moral responsibility is always distributed across the structure, never fully located in any single part.

Both warnings are taken seriously in what follows.

5. The Five-Level SE Hierarchy

Canonical reference

The normative definition of the five-level hierarchy, its traceability rules, ID conventions, decomposition order, and failure pathologies is maintained as the project's knowledge base under [knowledge/se-techniques/goals-requirements-hierarchy/](#). This chapter is the narrative introduction; the knowledge base is the reference.

5.1 A Universal Decomposition Instrument

Systems engineering decomposes any system — technical or social — into five levels of abstraction. Each level answers a different question, and each item at one level traces to items at the next level down, forming a directed acyclic graph with full traceability.

5.2 The Five Levels

Level	Prefix	Question Answered
Goals	G	Why does this system exist? What overarching objectives does it pursue?
Requirements	R	What must be true? What constraints, criteria, and conditions are derived from the goals?
Functions	F	What does the system do? What operations does it perform to meet requirements?
Logical Architecture	L	How is it organized? How are functions grouped into interacting subsystems?
Physical Implementation	P	What is it made of? What concrete people, buildings, documents, tools, and processes realize the architecture?

5.3 Why Five Levels?

The five levels are not arbitrary. They correspond to the fundamental decisions that any system designer — whether an engineer or a lawmaker — must make, in order:

1. **Purpose before structure.** You cannot design a system without knowing what it is for (Goals).
2. **Constraints before mechanisms.** You cannot choose mechanisms without knowing the constraints they must satisfy (Requirements).
3. **Functions before grouping.** You cannot group things into subsystems without knowing what functions need to exist (Functions).
4. **Logical before physical.** You should decide how subsystems interact before choosing which specific technologies, people, or buildings implement them (Logical → Physical).

Skipping a level — for example, jumping from goals directly to physical implementation — produces systems that "work" but are fragile, incoherent, and resistant to change. This is as true of a parliament as it is of a software architecture.

5.4 Traceability Rules

Every item in the hierarchy must be connected:

- Every Goal links to at least one Requirement
- Every Requirement links to at least one Function
- Every Function links to at least one Logical Architecture element
- Every Logical Architecture element links to at least one Physical Implementation
- A child can have multiple parents (many-to-many is valid)
- No orphan nodes — every item must be traceable upward to a goal and downward to an implementation

These rules ensure **completeness** (every goal is implemented) and **necessity** (every implementation serves a goal). When applied to social systems, they reveal a common pathology: physical implementations that have lost their connection to any goal ("we do this because we've always done it") and goals that have no implementation ("we believe in this but have no mechanism to achieve it").

5.5 Example: The German Verein

A brief example to make the abstraction concrete:

Level	ID	Description
Goal	G1	Enable citizens to pursue a shared non-commercial purpose collectively
Req.	R1	Written Satzung defining the purpose
Req.	R3	Democratic decision-making with equal voting rights
Func.	F3	Democratic assembly — Mitgliederversammlungen, elect Vorstand
Logical	L1	Governance Subsystem — Satzung, Wahlordnung, Geschäftsordnung
Physical	P2	Mitgliederversammlung (annual meeting), Protokolle, Kassenprüfung

The traceability chain: G1 → R3 → F3 → L1 → P2. Every physical artifact (the annual meeting, the minutes, the financial audit) traces back to a goal. If you removed the annual meeting, you would break the trace to G1 — the system would lose its mechanism for collective purpose.

This is the instrument. In the following chapters, we apply it to systems that are not usually analyzed this way.

6. Product Line Thinking



Canonical reference

The normative definitions of platform, variation points, binding decisions, the reuse gradient, the coherence rule, the workaround principle, and the product-line architect's toolkit are maintained as the project's knowledge base under [knowledge/se-techniques/product-line-engineering/](#). This chapter is the narrative introduction; the knowledge base is the reference.

6.1 Platform, Variation Points, and Reuse

A **product line** (or product family) is a set of systems that share a **common platform** but differ at defined **variation points**. The engineering discipline has three core ideas:

The Platform

The platform is the set of architectural elements — goals, requirements, functions, subsystems, and physical components — that are shared across all (or most) members of the family. In automotive terms, this is the chassis. In social system terms, this is the governance function, the membership mechanism, the resource allocation subsystem.

Variation Points

Variation points are the locations in the architecture where the platform deliberately leaves a decision open, because different family members will make different choices. A variation point has a defined interface (what connects to it) and a set of variants (what can be plugged in).



Automotive Analogy

Interface: "Here, a powertrain connects to the drivetrain."

Variants: 2.0L diesel · electric motor · plug-in hybrid

Everything below the interface (suspension, brakes, steering) works with any variant.



Political Analogy

Interface: "Here, executive authority is sourced and legitimated."

Variants: hereditary succession · popular election · party appointment · clerical selection

Everything below the interface (courts, civil service, tax authority) works with any variant.

Binding Decisions

Binding decisions are the specific choices made at each variation point for a given product. Once a binding is made, it constrains which other bindings are compatible. Not all combinations are coherent.

The Reuse Gradient

Our analysis reveals a general principle: **reuse feasibility decreases as you move upward in the SE hierarchy.**

SE Level	Reuse Potential
Physical Implementation	Highest — a tax authority works under any regime
Logical Architecture	High — subsystems transfer with adaptation
Functions	Medium — mechanisms may need interface adjustment
Requirements	Low — must be checked for goal compatibility
Goals	Lowest — goals define the system type; changing them means changing the system

This gradient has a practical corollary: **start reform at the bottom.** Modernize physical implementations first (professional civil service, independent courts, transparent budgets), then work upward. Top-down reform — imposing new goals before building the implementing infrastructure — is the recipe for failed states and institutional collapse.

The Coherence Rule

A system is coherent when its binding decisions at each variation point are mutually compatible, and when every element at every SE level traces consistently to the goals above it and the implementations below it. Incoherent systems — those with contradictory bindings — are unstable.

The Workaround Principle

When a desirable module from another system type conflicts with a binding decision, the most successful reforms create a **workaround** — a modified version that preserves the function while respecting the constraint. The constitutional monarchy is the paradigm case.

These concepts are applied extensively in [Part II](#) (social systems) and [Part III](#) (development frameworks).

7. STPA and STAMP

Canonical reference

The normative statements of STAMP (the accident model) and STPA (the four-step hazard analysis procedure, the four UCA types, and the rules for translating an SE decomposition into STPA inputs) are maintained as the project's knowledge base under [knowledge/se-techniques/stamp/](#) and [knowledge/se-techniques/stpa/](#). This chapter is the narrative introduction; the knowledge base is the reference.

7.1 System-Theoretic Process Analysis

The STAMP Accident Model

STAMP (Systems-Theoretic Accident Model and Processes), developed by Nancy Leveson at MIT, treats accidents not as the result of component failures but as **emergent properties of complex systems** — caused by inadequate control, including failures between components that may each be functioning exactly as designed.

This is precisely the right model for social systems, where the most devastating outcomes (inquisitions, abuse cover-ups, financial crises, failed reforms) typically occur not because the system has "broken" but because its control mechanisms are operating exactly as designed — just in a way that produces harm.

The Four Steps of STPA

Step	What It Produces
Step 1: Define Purpose	Losses (unacceptable outcomes), Hazards (system states leading to losses), System-level Constraints
Step 2: Model Control Structure	Hierarchical diagram of controllers, control actions, feedback channels, and process models
Step 3: Identify Unsafe Control Actions	Systematic enumeration of UCAs across four types
Step 4: Identify Loss Scenarios	Causal analysis: why each UCA occurs, what contextual factors enable it

The Four UCA Types

For every control action in the system, STPA asks four questions:

1. **Not providing** — What happens if this control action is not issued when it should be?
2. **Providing** — What happens if this control action is issued when it shouldn't be, or is the wrong action?
3. **Too early / too late / wrong order** — What happens if the timing is wrong?
4. **Applied too long / stopped too soon** — What happens if the duration is wrong?

From SE Decomposition to STPA

The SE hierarchy feeds directly into STPA:

- **Goals** define what the system must achieve → STPA **Losses** are the negation of goals
- **Requirements** define constraints → STPA **System-level Constraints** are safety requirements
- **Logical Architecture** defines subsystems → STPA **Control Structure** maps controllers and controlled processes
- **Physical Implementation** identifies actors → STPA identifies who issues which control actions

This integration is applied in detail in [Part IV](#).

Rung tagging — the social-systems extension

For social systems, Step 2 acquires one more move: tag every control action and every feedback channel with the **justificatory rung** at which it operates (rung 0 coercion through rung 6 deliberative legitimacy), and tag each controller with both its claimed and operating rung.

Rung information is not derivable from the SE decomposition — it is its own elicitation activity, performed during Step 2 — and it is required by Step 3 to identify rung-mismatch UCAs (over-rung command, under-rung command, asymmetric loop). The previous chapter [Justificatory Rungs](#) introduces the ladder and the three dangerous mismatch patterns; the canonical integration with STPA Steps 1-4 is at [knowledge/se-techniques/justification-rungs/application-to-stpa.md](#).

For engineering systems where every loop operates at rung 3-4 throughout, rung tagging adds nothing and is omitted. Use this extension where the system is social, institutional, or policy-shaped — where the question being asked is whether the loop's rung fits its function.

References

- Leveson, N.G. (2012). Engineering a Safer World: Systems Thinking Applied to Safety. MIT Press.
- Leveson, N.G. & Thomas, J.P. (2018). STPA Handbook. MIT Partnership for Systems Approaches to Safety and Security.
- Leveson, N.G. (2004). "A New Accident Model for Engineering Safer Systems." Safety Science, 42(4), 237-270.

8. Justificatory Rungs

At a glance

What this chapter adds: a seven-rung ladder of standards under which people accept a claim as a reason to act (rung 0 coercion → rung 6 deliberative legitimacy), and three structural patterns — Asymmetric Loop, Claimed-Rung Inflation, Cross-Loop Rung Imposition — where the rungs do not match and the system fails predictably.

What you need to read it: Part I chapters 1–4 (system definition, SE hierarchy, product lines, STPA).

Skip to §3 if you already know what a justificatory standard is and only want the three mismatch patterns. **Skip to §5** if you came from Part IV and just need the rung-tag conventions used in the STPA chapters.

Canonical reference

The normative statement of the seven-rung ladder, its three rung-mismatch patterns, and the seven hooks into STPA Steps 1–4 are maintained as the project's knowledge base under knowledge/se-techniques/justification-rungs/. This chapter is the narrative introduction; the knowledge base is the reference.

8.1 What this chapter adds to the book

Parts I and IV introduce STPA — the safety-and-control analysis at the heart of this book. STPA tells us which control actions are unsafe. It does not tell us why a control action that should land does not. That second question is the everyday experience of social systems: the bishop's sermon does not sway the empirically-minded congregant, the scientist's data does not move the believer, the parliament's resolution does not constrain the autocrat. In each case the control action is issued, transmitted, and received — but the receiver does not accept it as a reason to change behaviour.

The reason is that human beings act on different **standards** for what counts as a decisive reason. There are seven of them in this book's analysis. They form a ladder, and most of the catastrophic failures of social systems are traceable to mismatches on that ladder: between the two ends of a control loop, or between what a controller claims and what it operates on.

8.2 The seven rungs at a glance

Rung	Standard	Where you meet it
0	Power / coercion ("do it or else")	Coups, monopolies, threats
1	Authority / rhetoric / identity	Slogans, scripture, "as a parent..."
2	Formal consistency	Math proofs, court syllogisms
3	Empirical testability	Randomised trials, falsification
4	Cumulative evidence and consilience	Meta-analyses, IPCC reports
5	Meta-rational integration	Risk engineering, AI-safety, complex-systems policy
6	Normative legitimacy / deliberative ethics	Constitutional courts, citizens' assemblies

Each later rung adds a check that earlier rungs lack:

- **Rung 0 → Rung 1.** The receiver moves voluntarily, without a threat, because they recognise the source as legitimate. Rung 1 has no check on truth or consistency — but it has stopped using raw force.
- **Rung 1 → Rung 2.** A consistency check appears: even an authority cannot make an inconsistent argument win. Contradictions lose. This is the first internal check.
- **Rung 2 → Rung 3.** An external check appears: data can sink a consistent theory if its premises do not match the world.
- **Rung 3 → Rung 4.** The check is broadened against single-study flukes, publication bias, and method artefacts. Findings have to hold up across replications and methods.
- **Rung 4 → Rung 5.** The receiver becomes able to step outside any single method, recognise its limits, and combine methods under uncertainty. Rung 5 is rare, fragile, and cognitively costly.
- **Rung 5 → Rung 6.** A second kind of question appears that no amount of rung-5 sharpening can answer: what should we do? Rung 6 is the check that surviving open inclusive discourse imposes on claims about oughts.

The full rung definitions and examples are in knowledge/se-techniques/justification-rungs/rungs.md.

8.3 The seven rungs in detail

Each rung admits four questions: what is the standard, where do you meet it, what does the rung add over the previous one, and what does it still lack — and one extra question that matters for STPA: how does the rung show up in a control structure when something goes wrong.

Rung 0 — Power / coercion

The standard is "do it or else." Compliance is enforced by threat of sanction, violence, or exclusion; there is no shared criterion for what is true, valid, or fair. Two parties at rung 0 cannot disagree productively — they can only enforce or submit. Examples range from coups and occupation regimes to organised crime and parental "because I said so."

What rung 0 lacks: any common standard at all. In an STPA control structure this shows up as control actions that land without persuasion and feedback that either flatters the controller or is suppressed. Loss scenarios concentrate around the absence of corrective feedback — the controller cannot be wrong because the controller cannot be questioned.

Rung 1 — Authority / rhetoric / identity

A claim is decisive because of who makes it, what tradition it comes from, or what group the listener identifies with. The listener moves voluntarily, but without an obligation to check truth or consistency. Examples: campaign slogans, scripture, "as a parent...", brand loyalty, expert testimony where the listener cannot evaluate the substance, charismatic leadership.

What rung 1 adds over rung 0: voluntary acceptance — the listener moves without a sanction being threatened. What rung 1 lacks: any internal check. Two contradictory rung-1 authorities cannot adjudicate the contradiction without escalating to rung 0 (force) or rung 2+ (some standard outside personal authority). In STPA terms, rung-1 controllers treat who said it as the deciding factor; feedback that comes from the wrong source is dismissed regardless of content. The self-sealing process models catalogued in the next chapter are most often rung-1 process models.

Rung 2 — Formal consistency

A conclusion follows validly from premises both parties already accept. The argument can be checked step by step; contradictions lose. Examples: mathematical proofs, court syllogisms, formal logic, rule-based bureaucratic decisions, legal reasoning within an established code.

What rung 2 adds over rung 1: an internal check. Even an authority cannot make an inconsistent argument win — a rung-2 challenge can defeat a rung-1 claim by exhibiting contradiction within the claim's own premises. What rung 2 lacks: an external check. Premises themselves are not tested against the world. Loops that operate at rung 2 produce internally coherent decisions that may diverge from external reality — the characteristic failure is the system that "follows the rules" straight into harm.

Rung 3 — Empirical testability

A hypothesis is decisive only if it survives tests against rivals — controlled experiments, falsification, prediction against unseen data. Examples: randomised trials, falsification in physical sciences, A/B tests, engineering acceptance tests, audited financial statements.

What rung 3 adds: an external check. Data can sink a neat theory; premises themselves become accountable to observation. What rung 3 lacks: protection against one-off flukes, publication bias, p-hacking, and the small-N problem. A single positive trial does not yet make a robust claim. Rung-3 feedback channels are powerful but fragile — a single biased sample can defeat them — so loops that operate at rung 3 require active maintenance of the channel's integrity (independence of the auditor, pre-registration of trials, replication budgets).

Rung 4 — Cumulative evidence and consilience

A claim is decisive when results survive replication and line up across multiple methods, populations, and disciplines. Meta-analyses, systematic reviews, and cross-method triangulation beat single studies. Examples: evidence-based medicine, IPCC assessment reports, well-replicated psychology effects, cross-discipline syntheses, regulatory science.

What rung 4 adds: robustness against flukes, biases, and method artefacts. A finding that holds up across independent methods is harder to explain away. What rung 4 lacks: a way to handle situations in which the methods themselves are in dispute, or where no single method applies, or where the question is partly normative. Rung-4 feedback is slow but durable; systems that depend on it (climate policy, drug regulation) have characteristic failure modes around delay — the rung-4 signal arrives only after irreversible action.

Rung 5 — Meta-rational integration

A claim is decisive when the speaker can show why a method works, when it breaks, and how to combine multiple methods under uncertainty. The speaker steps outside any single formalism to choose the right one for the question. Examples: risk engineering, AI-safety research, complex-systems policy labs, integrated assessment in climate policy, intelligence analysis at its best.

What rung 5 adds: an acknowledgement that no single method covers all messy domains, and a discipline for choosing among methods rather than defaulting to one. Rung 5 is where one becomes able to talk sensibly about model failure, unknown unknowns, and the limits of one's own evidence. What rung 5 lacks: any answer to what should be done once the empirics are settled. Rung 5 sharpens facts; it does not deliver oughts. Rung-5 controllers can hold multiple feedback channels in superposition — a rung-3 RCT, a rung-1 stakeholder intuition, a rung-2 legal constraint — and weight them appropriately. This is rare, fragile, and cognitively costly.

Rung 6 — Normative legitimacy / deliberative ethics

A claim is decisive when it survives open, inclusive discourse and meets normative criteria — fairness, reversibility, universalisability, respect for the autonomy of the affected. Even perfect facts and models do not tell us oughts; society still has to justify rules to all affected by them. Examples: constitutional courts, citizens' assemblies, Habermasian deliberative democracy, well-functioning parliamentary deliberation, bioethics committees.

What rung 6 adds: a way to settle questions about what should be done that cannot be reduced to fact — distribution, risk acceptance, value tradeoffs, intergenerational duties. Rung 6 is the only rung that handles value pluralism without smuggling one party's values in as facts. What rung 6 lacks: speed and reach. Rung-6 deliberation is slow, expensive, and does not scale to many decisions per unit time, so most operational control stays at lower rungs. The characteristic failure mode is displacement: when the rung-6 channel is suppressed, the unresolved normative question accumulates pressure that eventually escapes through rung 0 (revolution).

What every controller carries on every loop

Drawing these rungs onto an STPA control structure means every arrow gets a tag, and every controller acquires two extra attributes: the rung at which it accepts commands from above and the rung at which it accepts feedback from below. The two are often different on the same controller, and that asymmetry is where most rung-related failures live. The next section catalogues the three patterns of mismatch.

8.4 Why the ladder matters for STPA

STPA Step 2 produces a control structure: controllers issuing control actions to controlled processes, with feedback returning. For social systems, every arrow in that structure operates at some rung. The bishop-to-priest arrow is rung 1 (sacred authority). The scientist-to-policymaker arrow is rung 3 (empirical testability). The court-to-parliament arrow is rung 2 (formal consistency) backed by rung 6 (constitutional principle).

Tagging the rungs makes three failure modes visible:

Same rung is necessary for transmission

If the controller speaks at a rung the receiver does not operate on, the control action does not land. A scientist briefing a charismatic authority at rung 3 ("the data shows...") to a rung-1 audience does not produce behavioural change — the audience cannot process data as a reason. This failure mode is silent: the controller may believe the action was issued, transmitted, and obeyed.

Same rung is not sufficient — and at low rungs it is dangerous

Two actors both at rung 1 transmit perfectly but have no internal check. The system efficiently propagates whatever the upper controller wants, and the same low-rung filter applies to feedback. The corrective signal that would expose error is never registered as evidence — it is registered as disloyalty, hostility, or sin.

This is the dominant failure mode of Religion (the Part IV worked example): rung-1 control downward is honest and stable, but rung-3 reality (abuse evidence, scientific findings, demographic harm) tries to enter the upward channel and is reclassified as rung-1 noise.

Different rungs on the two arrows of one loop

The most subtle pattern: the same loop operates at different rungs in each direction. The controller issues commands at rung 1 but expects feedback at rung 3. Or it issues commands at rung 3 but receives feedback only at rung 1. Either way the loop is structurally broken: information flows but does not correct.

This is what the cross-system catalogue identifies in religion, the unreformed military justice system, board-captured corporations, and family systems where the privacy norm absorbs all upward signal into a rung-1 filter.

8.5 Three dangerous mismatches

The seven rungs admit many possible per-arrow combinations. Three recur across the catalogued social systems and account for most of the rung-related failures.

Pattern A — Asymmetric loop

Rung-1 control downward, rung-3 (or higher) feedback expected upward. The controller's process model rejects the higher-rung feedback unless it arrives wrapped in rung-1 packaging — which strips the empirical content.

Generic remedy: insert an independent rung-3 channel that parallels (does not replace) the rung-1 channel and reaches the apex of control directly. Mandatory external audit, independent ombudspersons, pre-registered studies, and constitutional courts hearing rung-3 evidence against rung-1 sovereign claims are all implementations.

Pattern B — Claimed-rung inflation

The system publicly claims a high rung (rung 6 deliberation, rung 4 cumulative evidence) but operates at rung 0/1. The catastrophic systems of the twentieth century — one-party states claiming "people's democracy," theocracies claiming sacralised ultimate legitimacy — are inflations of this kind.

Generic remedy: close the gap. Either lower the claim to match operation, or raise the operation to match the claim. The first is rare because high claims confer near-term legitimacy; the second is what most catalogued architectural remedies attempt.

Pattern C — Cross-loop rung imposition

A controller legitimate at one rung in its own loop tries to extend that rung's authority into adjacent loops where a different rung is appropriate. Rung-1 religious authority claiming jurisdiction over rung-3 scientific questions; rung-3 scientific authority ruling on rung-6 normative questions; markets pricing questions whose decisive considerations sit at rung 6.

Generic remedy: domain-appropriate rung containment via subsidiarity rules and standing forums for the rung the question requires.

The full catalogue is in [knowledge/se-techniques/justification-rungs/dangerous-mismatches.md](#).

8.6 Claimed rung vs operating rung

Every controller has two rungs. The **claimed rung** is the standard the controller's public legitimacy appeals to — what the system tells participants and outsiders it operates on. The **operating rung** is the standard at which control actions and feedback actually run.

The gap between the two is itself a diagnostic. The widest gaps produce the most catastrophic failures, because participants experience the gap empirically and the system's legitimacy collapses faster than the system can correct.

Across the ten social systems in this book's catalogue, the gaps distribute as follows:

System	Claimed rung	Operating rung	Gap
Kingdom	1	0/1	Small
Republic	6	3–6 mixed	Variable; opens with capture
Theocracy	6 (sacralised)	1	Large
One-Party State	6	0/1	Extreme
Corporation	3	1–3 mixed	Moderate
University	3–4	1–3 mixed	Moderate
Military	1+2+6 (modern)	0/1	Small in tradition; widens under stress
Religion	1+6 (sacralised)	1	Large
Family	Variable	Variable	The gap is the diagnostic
Verein	6	6+2	Smallest

The full per-system analysis is in [knowledge/system-catalogues/social-systems/cross-system/justification-rungs-by-system.md](#).

8.7 A normative caveat

STPA in its original form is deliberately neutral about which controllers are "better" — it describes control loops without ranking them. The rung ladder is openly normative: rung 6 sits above rung 0 in a way that rung 0 cannot reciprocate. We flag this explicitly rather than smuggle it in.

The ladder is a useful add-on for social systems where epistemic standards matter. It is not a generic STPA tool. Engineering systems mostly operate at rung 3–4 throughout, the rungs do not mismatch, and the ladder adds nothing. Use this technique where the question being asked is whether the loop's rung fits its function — not where the question is purely informational.

8.8 Forward references

- The technique is applied to the religion case study in [STPA on Religion](#). Religion is the cleanest example of Pattern A asymmetric-loop failure.
- The dangerous patterns and their remedies generalise across all ten social systems — see [Control Structures in Social Systems](#) and [Architectural Remedies](#).

- The full canonical reference is at knowledge/se-techniques/justification-rungs/.

9. Part II — Social Systems

This part applies the SE hierarchy and product line framework to ten social systems across three families.

9.1 The Systems

Governance Family

Systems that rule a territory and its population:

- Kingdom (Traditional Monarchy)
- Republic (Modern Democratic)
- Theocracy
- One-Party State

Organization Family

Systems that coordinate collective effort toward a declared purpose:

- Corporation (For-Profit)
- University
- Military
- Religion (Church / Religious Organisation)

Intimate Systems Family

Systems at the scale of direct human relationships:

- The Family
- The Verein (Voluntary Association)

9.2 Chapters

- **Religion** — The founding decomposition that started this analysis
- **Kingdom & Republic** — Side-by-side comparison with reuse analysis
- **Ten Social Systems Compared** — Full decompositions, platform analysis, variation points, cross-family reuse, and interactive explorers
- **Governance as a Single Platform** — A four-system zoom: kingdom, republic, theocracy, one-party state as four bindings on one polity platform

9.3 Additional worked examples

Beyond the ten governance and organisation systems, the Worked Examples section also analyses:

- **The School** — Completes the worked case opened in the introduction (proxy-metric pathology)
- **The Hospital** — Triadic membership (patient, clinician, payer) and its structural failure modes
- **Wikipedia / Open Source** — Authority by demonstrated contribution against a rung-3 standard

10. Religion — SE Decomposition

10.1 The Founding Example

Religion was the first social system we decomposed through the five-level SE hierarchy. It serves as the worked example throughout this book and as the input to the STPA analysis in [Part IV](#).



religion — Graph

LEVEL	ID	DESCRIPTION	↑	↓
🎯 Goals	G1	Provide existential meaning and purpose to human life	—	R1 R2
	G2	Establish moral and ethical orientation for individuals and communities	—	R3 R4
	G3	Create social cohesion and collective identity	—	R5 R6
	G4	Offer psychological comfort in the face of suffering and mortality	—	R2 R7
💎 Requirements	R1	Coherent cosmological narrative (creation, afterlife, transcendence)	G1	F1
	R2	Framework for interpreting personal experience as meaningful	G1 G4	F1 F2
	R3	Codified behavioral norms (commandments, precepts, dharma)	G2	F3
	R4	Authority structure to interpret and enforce norms	G2	F3 F4
	R5	Shared symbols, narratives, and practices for group bonding	G3	F5
	R6	Mechanism for distinguishing in-group from out-group	G3	F4 F6
	R7	Rituals addressing grief, loss, and life transitions	G4	F2 F5
🔹 Functions	F1	Mythological narration — transmit origin and eschatological stories	R1 R2	L1
	F2	Pastoral care — counsel individuals through crisis and life stages	R2 R7	L2
	F3	Moral legislation — define permissible and forbidden behavior	R3 R4	L3
	F4	Doctrinal governance — interpret scripture, settle disputes	R4 R6	L3
	F5	Ritual performance — enact shared ceremonies (prayer, sacraments, festivals)	R5 R7	L1 L2
	F6	Missionary expansion — propagate belief system to new populations	R6	L4
🔸 Logical Architecture	L1	Liturgical Subsystem — scripture reading, preaching, ceremonial calendar	F1 F5	P1 P2 P3
	L2	Pastoral Subsystem — counseling, rites of passage, community support	F2 F5	P3 P4
	L3	Doctrinal/Juridical Subsystem — theology, canon law, ethics boards	F3 F4	P2 P5 P6
	L4	Outreach Subsystem — evangelization, education, charity, media	F6	P7 P8
🏢 Physical Implementation	P1	Churches, mosques, temples, synagogues, gurdwaras	L1	—
	P2	Sacred texts — Bible, Quran, Torah, Vedas, Tripitaka	L1 L3	—
	P3	Clergy — priests, imams, rabbis, pastors, monks, gurus	L1 L2	—
	P4	Parish networks, sanghas, congregations, umma structures	L2	—
	P5	Central authorities — Vatican, Al-Azhar, ecumenical councils	L3	—
	P6	Legal systems — canon law, fatwa apparatus, halakhic courts, Sharia boards	L3	—
	P7	Missionary organizations, madrasas, Sunday schools, yeshivas	L4	—
	P8	Religious media — TV networks, podcasts, publishing houses	L4	—

religion — Table

10.2 The Decomposition

Level	ID	Description
Goal	G1	Provide existential meaning and purpose to human life
Goal	G2	Establish moral and ethical orientation for individuals and communities
Goal	G3	Create social cohesion and collective identity
Goal	G4	Offer psychological comfort in the face of suffering and mortality
Req.	R1	Coherent cosmological narrative (creation, afterlife, transcendence)
Req.	R2	Framework for interpreting personal experience as meaningful
Req.	R3	Codified behavioral norms (commandments, precepts, dharma)
Req.	R4	Authority structure to interpret and enforce norms
Req.	R5	Shared symbols, narratives, and practices for group bonding
Req.	R6	Mechanism for distinguishing in-group from out-group
Req.	R7	Rituals addressing grief, loss, and life transitions
Func.	F1	Mythological narration — transmit origin and eschatological stories
Func.	F2	Pastoral care — counsel individuals through crisis and life stages
Func.	F3	Moral legislation — define permissible and forbidden behavior
Func.	F4	Doctrinal governance — interpret scripture, settle disputes
Func.	F5	Ritual performance — enact shared ceremonies (prayer, sacraments, festivals)
Func.	F6	Missionary expansion — propagate belief system to new populations
Log.	L1	Liturgical Subsystem — scripture reading, preaching, ceremonial calendar
Log.	L2	Pastoral Subsystem — counseling, rites of passage, community support
Log.	L3	Doctrinal/Juridical Subsystem — theology, canon law, ethics boards
Log.	L4	Outreach Subsystem — evangelization, education, charity, media
Phys.	P1	Churches, mosques, temples, synagogues, gurdwaras
Phys.	P2	Sacred texts — Bible, Quran, Torah, Vedas, Tripitaka
Phys.	P3	Clergy — priests, imams, rabbis, pastors, monks, gurus
Phys.	P4	Parish networks, sanghas, congregations, umma structures
Phys.	P5	Central authorities — Vatican, Al-Azhar, ecumenical councils
Phys.	P6	Legal systems — canon law, fatwa apparatus, halakhic courts, Sharia boards
Phys.	P7	Missionary organizations, madrasas, Sunday schools, yeshivas
Phys.	P8	Religious media — TV networks, podcasts, publishing houses

10.3 Key Structural Observations

R4 (authority) is the most consequential requirement. It enables F3 (moral legislation) and F4 (doctrinal governance), which together determine how the system interprets and enforces its own norms. When R4 becomes an end in itself — authority for authority's sake — the system begins to produce outcomes opposite to its goals. This dynamic is analyzed in detail in the [STPA analysis](#).

R6 (in-group/out-group) is a double-edged mechanism. It serves G3 (social cohesion) but contains no built-in limit. The same mechanism that creates belonging can escalate to dehumanization. The absence of a structural "circuit breaker" is a key hazard identified in the STPA.

The Liturgical and Pastoral subsystems (L1, L2) serve different goals that can conflict. L1 (liturgy) serves G1 (meaning) and G3 (cohesion) through shared ritual. L2 (pastoral care) serves G4 (comfort) through individual attention. When institutional resources are scarce, L1 tends to win — communal performance is more visible than individual care.

10.4 See Also

The political system that fuses religious authority with state governance shares many structural elements with this decomposition. Explore it in the [Theocracy Interactive Explorer](#).

The accompanying STPA analysis built on this decomposition lives in [STPA on Religion](#).

11. Kingdom & Republic — A Reuse Analysis

11.1 The Constitutional Monarchy as a Product-Line Variant



democracy — Graph

LEVEL	ID	DESCRIPTION	↑	↓
🎯 Goals	G1	Derive governing authority from consent of the governed	—	R1 R2
	G2	Protect individual rights and prevent tyranny	—	R2 R3 R4
	G3	Ensure order, security, and territorial integrity	—	R3 R5
	G4	Enable peaceful, periodic transfer of executive power	—	R1 R6
🔧 Requirements	R1	Free, periodic elections with universal suffrage	G1 G4	F1 F4
	R2	Codified constitution with enforceable rights	G1 G2	F1 F3
	R3	Separation of powers — legislative, executive, judicial	G2 G3	F1 F2 F3
	R4	Independent judiciary with constitutional review	G2	F3
	R5	Professional civil service	G3	F2
	R6	Free press and public accountability	G4	F4
⬠ Functions	F1	Legislative process — enact laws through parliament	R1 R2 R3	L1
	F2	Executive governance — implement policy, manage ministries	R3 R5	L2
	F3	Constitutional adjudication — protect rights	R2 R3 R4	L3
	F4	Electoral & accountability administration	R1 R6	L4
🔗 Logical Architecture	L1	Legislative Subsystem — parliament, committees	F1	P1 P2
	L2	Executive Subsystem — head of government, ministries	F2	P3
	L3	Judicial Subsystem — constitutional court, courts	F3	P4
	L4	Electoral & Transparency Subsystem	F4	P5
🏢 Physical Implementation	P1	Parliament (Bundestag, Congress)	L1	—
	P2	Constitution, bill of rights	L1	—
	P3	President/Chancellor, cabinet, civil servants	L2	—
	P4	Constitutional court (BVerfG, Supreme Court)	L3	—
	P5	Election commission, polling stations, free press	L4	—

democracy — Table

Key Finding

A constitutional monarchy reuses approximately **80% of the republican physical implementation** and **60% of its logical architecture**, while preserving the one goal-level requirement (dynastic continuity) that defines it as a distinct system variant.

11.2 Reuse Classification

Fully Reusable — Adopt Directly

Republican Element	Kingdom Equivalent	Historical Example
Independent judiciary	Royal courts → constitutional courts	UK Constitutional Reform Act 2005
Professional civil service	Patronage officials → merit-based Beamte	Stein-Hardenberg reforms, Prussia 1807
Modern fiscal system	Royal treasury → transparent tax authority + audit	Prussian fiscal reforms 1807-1815
Free press	No equivalent → introduce press freedom	Sweden's Freedom of the Press Act 1766

Partially Reusable — Adapt with Modifications

Republican Element	Adaptation Required	Historical Example
Parliament	Co-legislative body alongside crown, not replacing it	UK Parliament Act 1911
Written constitution	Limits royal power to enumerated prerogatives	Belgian Constitution 1831
Elections	For legislature only — head of state remains hereditary	Every constitutional monarchy
Separation of powers	Partial — monarch bridges branches as constitutional organ	Norwegian Grunnloven 1814

Incompatible — Cannot Reuse Without Abolishing the Monarchy

Republican Element	Why Incompatible
Popular sovereignty as sole source of authority	Eliminates hereditary authority by definition
Periodic election of head of state	Replaces hereditary succession
Electoral administration for choosing the executive	Structurally replaces dynastic succession

The full decomposition of both systems and the detailed cross-system comparison are in [Ten Social Systems Compared](#).

12. Product Line Thinking for Social Systems

12.1 Applying Systems Engineering Decomposition and Platform Reuse to Human Institutions

12.2 Abstract

Product line engineering — the discipline of designing a family of systems that share a common platform while differing at defined variation points — is a mature practice in automotive, aerospace, and software engineering. This chapter argues that the same framework applies with surprising precision to social systems: political regimes, organizations, religious institutions, and even families. By decomposing ten social systems through a five-level systems engineering hierarchy (Goals → Requirements → Functions → Logical Architecture → Physical Implementation), we reveal that human institutions share far more structural DNA than their surface differences suggest. We identify a "social system platform" — a common architecture present across nearly all institutions — and map the variation points that produce distinct system types. The result is a conceptual toolkit for understanding institutional reform, cross-cultural comparison, and the structural reasons why some institutions modernize successfully while others collapse under attempted change.

Canonical reference

The normative versions of the ten decompositions, the social-system platform, and the cross-system analyses are maintained as the project's knowledge base under [knowledge/system-catalogues/social-systems/](#). This chapter is the narrative overview; the knowledge base is the reference.

12.3 1. Introduction: Why Product Lines?

1.1 The Problem of Institutional Comparison

When we compare a kingdom to a republic, a corporation to a university, or a church to a military, we tend to focus on what makes them different. A king inherits power; a president is elected. A corporation maximizes shareholder value; a university pursues knowledge. These differences are real and important. But they obscure a deeper structural truth: these systems are far more alike than they are different.

Consider: every one of these systems needs to make collective decisions, allocate resources, resolve disputes, recruit and manage members, maintain legitimacy, and adapt to external change. The mechanisms differ, but the functional slots are nearly identical. This is precisely the situation that product line engineering was invented to exploit.

1.2 The Product Line Concept in Systems Engineering Terms

A product line is a set of systems that share a **common platform** but differ at defined **variation points**. The platform is the set of architectural elements — goals, requirements, functions, subsystems, and physical components — that are shared across all (or most) members of the family. In automotive terms, this is the chassis. In social-system terms, this is the governance function, the membership mechanism, the resource allocation subsystem.

Variation points are the locations in the architecture where the platform deliberately leaves a decision open. A variation point has a defined interface and a set of variants. In political terms: "here, executive authority is sourced and legitimated" is the interface; "hereditary succession" or "popular election" are the variants.

Binding decisions are the specific choices made at each variation point for a given product. Once a binding is made, it constrains which other bindings are compatible.

The full technique — platform, variation points, binding decisions, reuse gradient, coherence rule, workaround principle — is defined in [knowledge/se-techniques/product-line-engineering/](#).

1.3 The Five-Level SE Hierarchy

Throughout this chapter, we decompose every system using a five-level hierarchy: Goals → Requirements → Functions → Logical Architecture → Physical Implementation. Each item at one level traces to items at the next level down. This traceability is what makes the product line analysis possible: we can identify which elements are shared and which are variant, at every level of abstraction.

The level definitions and traceability rules are in <knowledge/se-techniques/goals-requirements-hierarchy/>.

12.4 2. The Decompositions: Ten Social Systems

We decompose ten social systems, grouped into three families:

- **Governance family** — Kingdom (Traditional Monarchy), Republic (Modern Democracy), Theocracy, One-Party State.
- **Organization family** — Corporation (For-Profit), University, Military, Church / Religious Organization.
- **Intimate systems family** — Family, Verein (Voluntary Association).

Each decomposition names the system's goals, requirements, functions, logical architecture, and physical implementation, with traceability across all five levels. The full tables are maintained in the knowledge base:

System	Decomposition file
Kingdom	knowledge/system-catalogues/social-systems/kingdom/se-decomposition.md
Democracy / Republic	knowledge/system-catalogues/social-systems/democracy/se-decomposition.md
Theocracy	knowledge/system-catalogues/social-systems/theocracy/se-decomposition.md
One-Party State	knowledge/system-catalogues/social-systems/one-party-state/se-decomposition.md
Corporation	knowledge/system-catalogues/social-systems/corporation/se-decomposition.md
University	knowledge/system-catalogues/social-systems/university/se-decomposition.md
Military	knowledge/system-catalogues/social-systems/military/se-decomposition.md
Religion	knowledge/system-catalogues/social-systems/religion/
Family	knowledge/system-catalogues/social-systems/family/se-decomposition.md
Verein	knowledge/system-catalogues/social-systems/verein/se-decomposition.md

Each file presents the full five-level hierarchy with IDs, descriptions, and cross-level traceability. The interactive explorers linked at the bottom of this chapter render the same decompositions as browsable graphs.

12.5 3. The Social System Platform

Once you align the ten decompositions level by level, a common core emerges: a set of functional slots that every system fills, however differently.

The slots are: **Authority, Membership, Resource Allocation, Norms, Dispute Resolution, Legitimation, Succession, External Representation, Socialisation, and Activity Delivery**. Every social system — whether a monarchy, a corporation, or a family — contains all ten slots, because they correspond to problems every collective must solve.

The full platform description (goals, requirements, functions, logical, and physical elements shared across all ten systems) is in <knowledge/system-catalogues/social-systems/cross-system/platform.md>.

The platform is not an abstract ideal — it is the specific structural residue that appears when you decompose enough systems the same way and look at what is left in common.

12.6 4. The Variation Points

What varies across the family is concentrated at a small number of well-defined variation points. Six have been identified:

- **VP1 — Source of Authority** (hereditary · popular · party · divine · capital · voluntary · statutory · ...)
- **VP2 — Membership Boundary** (birth · contract · profession · faith · election · ...)
- **VP3 — Decision-Making Mechanism** (command · vote · consensus · ritual · market · ...)
- **VP4 — Succession Mechanism** (inheritance · election · appointment · nomination · ...)
- **VP5 — Legitimation Narrative** (tradition · law · results · doctrine · contract · ...)
- **VP6 — Norm Enforcement Mechanism** (coercion · social sanction · ritual · employment · ...)

Variants are not freely combinable: some combinations are coherent, others are structurally contradictory. The full set of variation points, observed variants, and interaction constraints is in [knowledge/system-catalogues/social-systems/cross-system/variation-points.md](#).

12.7 5. Cross-Family Reuse: Case Studies

The binding decisions at each variation point constrain which elements can be transplanted between systems. Five worked cases illustrate this:

1. **Kingdom adopting republican structures** — the constitutional monarchy as a workaround that preserves hereditary legitimation while grafting on parliamentary legislation.
2. **Church adopting corporate structures** — institutional professionalization as partial reuse of corporate management under a divine-mandate top.
3. **University adopting corporate structures** — New Public Management as an amber-zone reuse with documented side effects.
4. **One-party state adopting republican fiscal structures** — China's reform model as selective transplantation below the authority variation point.
5. **Family adopting Verein structures** — democratic parenting as controlled opening of parental authority.

All five case studies are in [knowledge/system-catalogues/social-systems/cross-system/reuse-case-studies.md](#).

12.8 6. Implications and Principles

Four principles recur across the case studies:

- **The Reuse Gradient.** Reuse feasibility decreases as you move upward in the SE hierarchy. Physical-level reuse is easy; goal-level reuse is almost impossible without transforming the system.
- **The Coherence Rule.** A system's binding decisions must be mutually compatible. Incoherent systems either resolve the contradiction (reform or revolution) or persist in institutional dysfunction.
- **The Workaround Principle.** When a desirable element conflicts with a binding decision, the most successful reforms create a modified version that preserves the function while respecting the constraint.
- **The Product Line Architect's Toolkit.** An eight-step procedure for designing or reforming a system: decompose both source and target, identify the platform, map the variation points, check binding constraints, classify reuse potential, start at the bottom, design workarounds, accept red items as system-defining.

Full statements of all four principles are in [knowledge/system-catalogues/social-systems/cross-system/principles.md](#).

12.9 7. Conclusion

Social systems look more different from each other than they actually are. A kingdom and a republic, a church and a corporation, a family and a Verein — these seem like categorically different things. But when decomposed through a rigorous systems engineering

hierarchy, they reveal a shared platform of functional slots, and they differ primarily at a small number of well-defined variation points.

Product line thinking does not reduce human institutions to interchangeable parts. It does the opposite: it makes the genuine differences visible by stripping away the illusion that everything is different. Once you see that a kingdom's judicial system and a republic's judicial system fill the same functional slot with similar physical implementations, you can focus your analytical energy on the actual variation point — the source of executive authority — and reason clearly about which reforms are compatible with that binding and which are not.

This framework is not merely academic. Every successful institutional modernization in history — the Stein-Hardenberg reforms, the Meiji Restoration, the post-WWII German constitution, Vatican II, China's economic opening — can be understood as a product-line reuse exercise: identify what can be transplanted, adapt what must be adapted, and respect what cannot be changed without changing the system's identity.

The five-level SE decomposition is the analytical instrument. The variation point map is the strategic guide. And the coherence rule is the warning system. Together, they form a toolkit for anyone who wants to reform an institution — or simply to understand why institutions are the way they are.

12.10 Interactive Explorers

Explore each social system's SE decomposition interactively:

System	Explorer
Corporation	→ Open
Democracy (Republic)	→ Open
Family	→ Open
Hospital	→ Open
Military	→ Open
One-Party State	→ Open
Religion	→ Open
School	→ Open
Theocracy	→ Open
University	→ Open
Verein	→ Open
Wikipedia / Open Source	→ Open

13. Governance as a Single Platform

13.1 Kingdom, Republic, Theocracy, and One-Party State as Four Bindings of One System

The four governance systems in this book — kingdom (traditional monarchy), republic (modern democracy), theocracy, and one-party state — are usually treated as fundamentally different political systems. The systems-engineering reading is that they are four binding decisions on the same platform. Once you align their SE decompositions level by level, the platform is identical at L (logical architecture) and largely identical at F (functions). What differs is concentrated at a small number of variation points, all of which trace upward to a single binding decision at VP1: the source of authority.

This chapter is the synthesis that the four worked examples in this part imply but do not state. It is short by design; the underlying decompositions live in [Kingdom & Republic](#), [Theocracy](#), and [One-Party State](#), and the cross-system platform reference is at [knowledge/system-catalogues/social-systems/cross-system/platform.md](#).

At a glance

What this chapter does: treats kingdom/republic/theocracy/one-party-state as one polity-platform with four VP1 bindings. Reads each governance pathology from the binding it follows from. Gives the architect's checklist for what reforms can be transplanted between bindings and what cannot.

What you need: the four individual chapters (Kingdom & Republic, Theocracy, One-Party State) or at least their variation-point bindings. The cross-system platform definition is helpful but not required.

13.2 1. The Polity Platform

Every governance system fills the same ten platform slots already identified in [Ten Social Systems Compared](#). The four governance variants implement them with strikingly similar logical-architecture-level subsystems:

Functional Slot	Logical subsystem (all four polities)
Authority & Decision-Making	Supreme authority + advisory body
Membership & Belonging	Citizenship / subjecthood + boundary enforcement
Resource Allocation	Treasury / fiscal apparatus
Norm Setting & Enforcement	Legal system + enforcement apparatus
Dispute Resolution	Court system
Legitimation	Legitimation narrative + ritual
Succession & Continuity	Succession mechanism
External Representation	Diplomatic apparatus + military
Socialisation	Education + state media
Activity Delivery	Civil service + public services

The subsystems are not identical in their physical implementation, but they are identical in their role. A republic's parliament and a kingdom's privy council are different physical implementations of the same logical subsystem ("supreme deliberative body"). A republic's election commission and a one-party state's organisation department are different physical implementations of the same logical subsystem ("succession mechanism"). The structural identity of the slots is what makes the comparison possible.

13.3 2. The Four VP1 Bindings

What distinguishes the four polities is, fundamentally, a single binding decision at VP1 (source of authority) plus the cascade of constraints that binding imposes downstream:

Variant	VP1 binding	Cascades to
Kingdom	Hereditary / dynastic	VP3 = royal will / advised; VP4 = dynastic succession; VP5 = tradition + sacral kingship
Republic	Popular sovereignty	VP3 = competitive election + parliamentary deliberation; VP4 = periodic election; VP5 = constitution + consent
Theocracy	Divine mandate	VP3 = clerical interpretation; VP4 = clerical selection; VP5 = sacred text
One-Party State	Party ideology	VP3 = party committee; VP4 = cadre promotion; VP5 = ideological doctrine

The cascade direction is downward in the SE hierarchy. VP1 constrains what bindings are coherent at VP3, VP4, and VP5. A binding incoherent with VP1 is not just suboptimal; it is structurally unstable, and the system either resolves the incoherence (reform), maintains it as institutional dysfunction (the "façade" pattern), or collapses.

13.4 3. Reading the Pathologies off the Bindings

The structural pathologies documented in the individual chapters are not coincidences. They follow predictably from the VP1 binding.

Kingdom — succession crisis. Hereditary succession (VP4) produces a leadership pool of one (the heir). The pool may turn out to be incompetent, infirm, infant, or absent. The structural fix — adoption (Roman emperors), elected monarchy (Holy Roman Empire), regency mechanism — is a workaround that preserves VP1 = dynastic while substituting a less rigid VP4.

Republic — vulnerability to anti-system parties. Competitive election (VP3) admits any party that meets the formal entry requirements, including parties whose programme is to abolish competitive election. The structural fix — *streitbare Demokratie* (Germany, Basic Law Article 21), party-banning constitutional provisions, judicial review — installs a rung-2/6 backstop at exactly the level VP3 leaves open.

Theocracy — doctrine cannot revise. Divine mandate (VP1) cannot, by its own logic, be revised by evidence; doing so would dissolve the legitimating claim. The structural fix is either a workaround (admit rung-3 evidence in bounded domains while preserving VP1 in the religious-political core — China's socialist market economy binding pattern works the same way for a different VP1) or a workaround at VP3 (a "constitutional theocracy" with elected institutions ratified by clerical authority — Iran's structural model).

One-Party State — no error-correction channel. Party committee decision (VP3) eliminates by design every external channel that would carry error-correcting evidence to the apex. The structural fix is either narrow rung-3 channel openings (China post-1978 economic, Vietnam *Đổi Mới*) or partial dissolution of VP1 (Gorbachev's *glasnost*, which proved structurally unstable and ended the system).

In each case the diagnostic insight is the same: the pathology is what it is because the binding is what it is. Read the binding, predict the pathology.

13.5 4. The Reuse Gradient Across Polities

The Reuse Gradient from [Part I](#) applies with unusual clarity to the polity platform. Reuse between governance variants is asymmetric depending on the SE level:

SE level	Reuse across polity variants
Physical	High — civil service, tax authority, courts, post, police, military: all four polities use structurally similar implementations
Logical	High — the ten platform subsystems are identical in role across all four
Functions	Medium — same functions, but the assignment to controllers differs (a republic's legislature ≠ a kingdom's privy council in who exercises legislative function)
Requirements	Low — derived requirements differ because they derive from different goals
Goals	Lowest — goals encode VP1, and VP1 is what differentiates the variants

This is why the Stein-Hardenberg reforms (Prussia 1807) succeeded: they reformed Prussia at the physical and logical levels (civil service, fiscal system, education, military command) without touching VP1 (Prussian monarchy preserved). It is also why Gorbachev's reforms did not: they attempted goal-level reform (glasnost, multi-candidate elections) on a one-party-state platform, which destabilised VP1 itself.

The **Workaround Principle** also reads naturally here. The constitutional monarchy is the paradigm: take the desirable elements of the republic (parliament, constitution, independent judiciary, civil service) and graft them onto a kingdom platform whose VP1 (hereditary) is preserved by making the monarch a constitutional organ rather than a sovereign. Norway 1814, Belgium 1831, the UK 1911-1949 are all examples of the same architectural pattern executed with different parameter settings.

13.6 5. What This Synthesis Buys

Treating the four polities as one platform with four VP1 bindings does three things the four separate analyses do not.

First, it makes the **reform question precise**. "Can Country X democratise?" is a vague question that admits no engineering answer. "Can Country X transplant Republican L-level subsystems while keeping its current VP1 binding?" is a precise question that admits specific answers (Stein-Hardenberg-pattern: yes; full goal-level rebuild: rarely).

Second, it makes the **comparison legitimate**. Kingdom and republic can be compared on platform terms without falling into the false dichotomy that they are fundamentally different things. They are different bindings on the same platform; the comparison is between the bindings, not the platforms.

Third, it surfaces the **shared failure modes**. The Accountability Void, the Self-Sealing Process Model, and the Proxy Metric ([Part IV, Control Structures](#)) appear in all four variants; the form differs but the structural pattern is the same. A reader who has understood one polity's accountability void has understood the others' too.

13.7 See Also

- [Kingdom & Republic — A Reuse Analysis](#) — the worked case of transplanting between the first two bindings.
- [Theocracy, One-Party State](#) — the two cases not covered in detail by kingdom-republic.
- [Ten Social Systems Compared](#) — the full ten-system platform analysis of which this chapter is a four-system zoom.
- [knowledge/system-catalogues/social-systems/cross-system/platform.md](#) — canonical definition of the social-system platform.
- [knowledge/system-catalogues/social-systems/cross-system/variation-points.md](#) — canonical catalogue of the six variation points.

14. Example: Der Deutsche Verein (e.V.)

The Verein is Germany's universal civic institution — over 600,000 registered associations covering everything from football to beekeeping. It is the simplest social system that instantiates all ten platform functional slots, making it an ideal worked example.



verein — Graph

LEVEL	ID	DESCRIPTION	↑	↓
🎯 Goals	G1	Enable citizens to pursue a shared non-commercial purpose collectively	—	R1 R2
	G2	Provide a democratic, self-governing legal structure	—	R3 R4
	G3	Create social cohesion and community identity	—	R5 R6
	G4	Ensure institutional continuity independent of individuals	—	R2 R4 R7
💎 Requirements	R1	Written Satzung defining the non-commercial purpose	G1	F1
	R2	Legal personality through Vereinsregister (e.V.)	G1 G4	F1 F2
	R3	Democratic decision-making with equal voting rights	G2	F3
	R4	Elected Vorstand accountable to membership	G2 G4	F3 F4
	R5	Regular activities bringing members together	G3	F5
	R6	Open, voluntary membership	G3	F2
	R7	Financial sustainability through dues, donations, grants	G4	F6
◇ Functions	F1	Constitutional governance — draft and enforce Satzung	R1 R2	L1
	F2	Membership administration — admit, track, manage	R2 R6	L2
	F3	Democratic assembly — Mitgliederversammlungen, elections	R3 R4	L1
	F4	Executive management — represent Verein externally	R4	L3
	F5	Activity delivery — events, training, competitions	R5	L4
	F6	Financial management — dues, bookkeeping, Gemeinnützigkeit	R7	L2 L3
○ Logical Architecture	L1	Governance — Satzung, Versammlung, Wahlordnung	F1 F3	P1 P2
	L2	Administration — member database, dues, Vereinsregister	F2 F6	P3 P4
	L3	Executive — Vorstand, contracts, bank account	F4 F6	P5 P6
	L4	Operations — Abteilungen, coaches, facilities	F5	P7 P8
🏢 Physical Implementation	P1	Satzung, Geschäftsordnung, Beitragsordnung	L1	—
	P2	Mitgliederversammlung, Protokolle, Kassenprüfung	L1	—
	P3	Membership software (SEWOBE, Vereinsflieger)	L2	—
	P4	Vereinsregister entry, Finanzamt registration	L2	—
	P5	Vorstand (Vorsitzender, Schatzmeister, Schriftführer)	L3	—
	P6	Vereinskonto, DATEV, Zuwendungsbestätigungen	L3	—
	P7	Vereinsheim, sports fields, rehearsal rooms	L4	—
	P8	Abteilungen (Fußball, Turnen, Chor), Übungsleiter	L4	—

verein — Table

14.1 SE Decomposition

The full five-level decomposition is shown in the interactive visualization linked above. Click any node to see its description, parent links, and child links. Use the Table view for the complete traceability matrix.

14.2 Variation Point Bindings

The Verein binds VP1 (authority) to voluntary democratic participation, VP2 (membership) to open and voluntary, VP3 (decisions) to one-person-one-vote, and VP4 (succession) to periodic election. This makes it structurally closest to a republic, but at micro-scale and with a non-commercial purpose constraint.

14.3 The Structural Pathology of Voluntary Democratic Participation

The Verein's binding produces a structurally lovely system: every member can vote, every euro is approved by the assembly, every officer is elected. The matching weakness is that the binding depends on the membership actually showing up. When attendance at the Mitgliederversammlung falls below an effective participation threshold — and in most large Vereine it does — the assembly mechanism becomes a fig leaf for an effective oligarchy of the Vorstand and a small core of long-standing members.

The result is the "Vorstand-only Verein": legally democratic, structurally a board-run organisation, with the membership relating to the Verein as customers of an activity rather than members of a self-governing body. The Satzung promises rung-6 deliberation; in practice the Verein operates at rung 1 (Vorstand authority unchallenged) with rung-6 ratification at most once per year by twenty people.

A second pathology is **entryism**: because membership is open and voting power is per-head rather than per-stake, an organised faction can join a Verein in numbers sufficient to capture its assembly. The pattern is well documented in political parties (the early German Pirate Party, multiple shooting clubs taken over by extreme-right networks) but applies to any Verein whose decision-mechanism cannot distinguish a long-standing supporter from a recent recruit.

These are not failures of voluntarism. They are structural consequences of the VP1 binding when participation is low and membership is open.

14.4 STPA Highlights

UCA-V1: The Vorstand acts on substantial decisions without meaningful Mitgliederversammlung input because attendance is below quorum or engagement is below effective oversight.

Causal factor: the same volunteer board prepares the agenda, runs the meeting, presents the accounts, recommends decisions, and counts the votes. Without an actively engaged member base, the assembly is structurally captured by its preparers — not because anyone is acting in bad faith, but because the workload of running a Verein concentrates in the people willing to do it.

UCA-V2: The Kassenprüfer is socially close to the Vorstand and cannot perform an independent audit.

Causal factor: Kassenprüfer are usually fellow members elected from the same general assembly that has elected the Vorstand. The role is statutory but the social independence is not enforceable; in small Vereine the Kassenprüfer is often the only person who knows the bookkeeping system because the treasurer trained them.

UCA-V3: An organised faction joins the Verein in numbers sufficient to capture its assembly.

Causal factor: the open-membership binding has no mechanism to distinguish committed supporters of the Verein's purpose from organised entrants whose loyalty is to an outside agenda. The defence — admission control by the Vorstand under §39 BGB — collides with the open-membership norm and is rarely used at scale until the takeover is already underway.

14.5 Fixes That Worked

Vereinsregister + Finanzamt Gemeinnützigkeit audit — partial structural defence against UCA-V1 and UCA-V2.

Registered Vereine must file evidence of properly conducted elections with the Vereinsregister, and gemeinnützige Vereine must satisfy the Finanzamt's recurring audit that their activities and finances actually serve the registered purpose. Evidence: the audit is an external rung-3 check on the rung-1 internal claim, and is the most common point at which structural problems in a Verein become visible to the wider system — annual reports of loss of Gemeinnützigkeit reach the low thousands across the German Verein population.

§43 BGB statutory Vorstand liability — addressing UCA-V1.

The Vorstand's personal liability for breaches of duty (incompetent management, conflict of interest, embezzlement) operates without depending on the assembly to notice. Evidence: the threat of personal liability is the strongest non-assembly correction in the legal architecture; the existence of standard Vereins-Haftpflicht insurance products is the market's pricing of the residual risk.

Documented two-step admission for politically exposed Vereine — addressing UCA-V3.

Vereine in domains where takeover risk is recognised (political parties, Schützenvereine after 2010s extremism cases, sports federations) increasingly adopt a probationary membership period during which the new member has activity rights but no voting rights, followed by Vorstand confirmation of full membership. This is a workaround in product-line terms: it preserves the open-membership ideal while inserting a finite delay that an organised takeover must survive.

**The Design Principle**

The Verein's legal architecture is well designed for the steady state of an engaged democratic association. Its failure modes appear when the steady-state assumption breaks — low participation, low engagement, or hostile entry. The fixes that work are not changes to the democratic core; they are external channels (registry, tax audit, statutory liability) and procedural workarounds (probationary admission) that catch the failure paths the democratic core cannot catch by itself.

14.6 Platform Mapping

This system fills all ten universal functional slots identified in the [Ten Social Systems Compared](#):

Functional Slot	How This System Fills It
Authority & Decision-Making	Mitgliederversammlung (general assembly) as sovereign body; elected Vorstand handles day-to-day execution
Membership & Belonging	Voluntary application; admission approved by Vorstand; conditions and fees defined in the Satzung
Resource Allocation	Membership fees (Beiträge); donations; assembly-approved annual budget; independent Kassenprüfer audit
Norm Setting & Enforcement	Satzung (articles of association); Geschäftsordnung (procedural rules); group social norms
Dispute Resolution	Ehrengericht (honour court) or Vorstand mediation; expulsion requires assembly vote by absolute majority
Legitimation	Shared non-commercial purpose; democratic participation; voluntary membership as constitutive feature
Succession & Continuity	Annual or biennial election of Vorstand at Mitgliederversammlung; written Satzung ensures institutional memory
External Representation	1. Vorsitzender (chair) as sole legal representative (Vertretungsberechtigung); umbrella organisation membership
Socialisation	Club culture and shared practices; activity rituals; mentoring of new members into group norms
Activity Delivery	Regular activities (training sessions, meetings, events, competitions) serving the stated purpose

Navigate to the [interactive visualization](#) for the full graph and table.

15. Example: Democracy (Modern Republic)

The democratic republic is the governance family's most complex variant, with the richest set of checks, balances, and feedback mechanisms of all ten systems analysed in this book. Its SE decomposition reveals why it is also the most structurally resilient — it has more redundant control paths, more independent feedback channels, and more explicit separation of authority than any other social system architecture.



democracy — Graph

LEVEL	ID	DESCRIPTION	↑	↓
🎯 Goals	G1	Derive governing authority from consent of the governed	—	R1 R2
	G2	Protect individual rights and prevent tyranny	—	R2 R3 R4
	G3	Ensure order, security, and territorial integrity	—	R3 R5
	G4	Enable peaceful, periodic transfer of executive power	—	R1 R6
🔧 Requirements	R1	Free, periodic elections with universal suffrage	G1 G4	F1 F4
	R2	Codified constitution with enforceable rights	G1 G2	F1 F3
	R3	Separation of powers — legislative, executive, judicial	G2 G3	F1 F2 F3
	R4	Independent judiciary with constitutional review	G2	F3
	R5	Professional civil service	G3	F2
	R6	Free press and public accountability	G4	F4
🔗 Functions	F1	Legislative process — enact laws through parliament	R1 R2 R3	L1
	F2	Executive governance — implement policy, manage ministries	R3 R5	L2
	F3	Constitutional adjudication — protect rights	R2 R3 R4	L3
	F4	Electoral & accountability administration	R1 R6	L4
🗑 Logical Architecture	L1	Legislative Subsystem — parliament, committees	F1	P1 P2
	L2	Executive Subsystem — head of government, ministries	F2	P3
	L3	Judicial Subsystem — constitutional court, courts	F3	P4
	L4	Electoral & Transparency Subsystem	F4	P5
🏢 Physical Implementation	P1	Parliament (Bundestag, Congress)	L1	—
	P2	Constitution, bill of rights	L1	—
	P3	President/Chancellor, cabinet, civil servants	L2	—
	P4	Constitutional court (BVerfG, Supreme Court)	L3	—
	P5	Election commission, polling stations, free press	L4	—

democracy — Table

15.1 SE Decomposition

The full five-level decomposition is shown in the interactive visualization above. Click any node to see its description and connections. Use the Table view for the complete traceability matrix.

15.2 Variation Point Bindings

Variation Point	Binding
VP1: Source of Authority	Popular sovereignty — the only legitimate source of governing power is the consent of the governed
VP2: Membership Boundary	Citizenship by territory, birth, or naturalisation
VP3: Decision-Making	Representative democratic vote — competitive elections with universal suffrage
VP4: Succession	Periodic election — no hereditary or self-perpetuating leadership

The key constraint: VP1 (popular sovereignty) structurally forces VP3 and VP4 to include genuine competitive elections. A system that calls itself democratic but controls elections, eliminates competition, or makes succession non-periodic has changed VP1 — it is no longer a democracy in the product-line sense, whatever its name.

15.3 Why This System Is Structurally Unusual

Every other system in this book concentrates authority at the top and manages the risk of that concentration through informal norms, cultural expectations, or individual virtue. The republic does something structurally different: it distributes authority across institutions that are designed to check each other, and it builds the checking mechanism into the constitution rather than leaving it to goodwill.

The result is the highest feedback richness of any system: elections connect government behaviour to electoral consequences; independent courts connect government action to constitutional standards; a free press connects official narrative to observable reality; civil society connects individual grievance to collective voice. These channels operate in parallel and are deliberately independent of one another — so that capturing one does not disable the others.

This redundancy is not accidental. It is the product of historical learning from democratic failures — most clearly documented in the engineering of Germany's Basic Law after the Weimar Republic's collapse.

15.4 STPA Highlights

STPA identifies the conditions under which democracy's redundant architecture degrades. The most important unsafe control actions are:

UCA-D1: Legislature enacts legislation that dismantles constitutional democracy, with no body able to invalidate it.

The Weimar Republic had no mechanism for judicial review of constitutionality. The Enabling Act of 1933 was passed by a formally valid parliamentary vote. There was no circuit breaker. The causal factor is structural, not moral: even a legislature of good people cannot prevent constitutional erosion if no external authority can constrain it.

UCA-D2: Government collapses without a functioning successor, creating conditions for emergency rule.

Weimar's Article 48 combined easily dismissed coalition governments with presidential emergency powers. The path from parliamentary deadlock to authoritarian rule was structurally open. Fourteen chancellors in fourteen years exhausted the system's legitimacy before it was overthrown.

UCA-D3: An anti-constitutional party uses democratic procedures to eliminate democracy.

When a system's rules contain no defence against participants who intend to abolish the rules, the system can be closed by the actors it most needs to exclude. The structural openness that is democracy's greatest strength becomes a vulnerability without a corresponding structural defence.

UCA-D4: The information environment is captured, severing corrective feedback between government behaviour and electoral accountability.

When citizens cannot see what the government is doing, they cannot vote against it. Press freedom and media independence are not decorative features of democracy — they are the feedback channel that makes elections meaningful.

15.5 Fixes That Worked: Germany's Basic Law (1949)

The Federal Republic of Germany is the most thoroughly documented case of constitutional engineering explicitly designed to prevent identified structural failure modes. Its designers worked from what was in effect a structural post-mortem of Weimar — identifying each failure pattern and designing a structural countermeasure.

UCA	Weimar structural failure	Basic Law remedy
UCA-D1	No judicial review of constitutionality	Bundesverfassungsgericht — independent constitutional court with power to strike down unconstitutional legislation. Has struck down hundreds of laws since 1951.
UCA-D2	Easy government collapse + emergency powers	Constructive vote of no confidence — the Bundestag can only remove a Chancellor by simultaneously electing a successor. No more leadership vacuum. Invoked twice in 75 years, both times appropriately.
UCA-D3	No defence against anti-constitutional parties	Article 21 GG — parties seeking to undermine the free democratic basic order can be banned by the Constitutional Court. SRP banned 1952; KPD 1956. The mechanism functions as a structural deterrent.
UCA-D4	Centralised or capturable press	Public broadcasting (ARD, ZDF) governed by Rundfunkräte — councils of civil society representatives, not government officials. Financially independent.

The Design Principle

Germany's post-Weimar architects did not trust good intentions. They trusted structure. The Basic Law has supported a functioning democracy for 75 years through reunification, financial crises, and significant polarisation — not because Germans are inherently different, but because the control structure was deliberately engineered to be resilient to the failure modes that destroyed Weimar.

15.6 Platform Mapping

Functional Slot	How This System Fills It
Authority & Decision-Making	Popular sovereignty expressed through elected representative institutions
Membership & Belonging	Citizenship by birth, territory, or naturalisation; universal suffrage
Resource Allocation	Parliamentary budget process; taxation by elected legislature
Norm Setting & Enforcement	Legislative process with constitutional review; independent judiciary
Dispute Resolution	Independent courts; constitutional court for fundamental rights
Legitimation	Consent of the governed; periodic competitive elections
Succession & Continuity	Periodic elections; constitutional transfer of power
External Representation	Elected executive; treaty ratification by legislature
Socialisation	Public education; civic education; free press
Activity Delivery	Civil service; public services; rule of law

For the full SE decomposition: [interactive visualization](#). For the comparative analysis across all ten systems: [Ten Social Systems Compared](#). For the cross-system control structure analysis: [Control Structures](#).

16. Example: Theocracy

The theocracy derives all authority from a sacred source, making it the governance type with the strongest legitimization narrative — and the weakest feedback mechanisms. Its SE decomposition reveals a structurally closed system where doctrine, law, education, and enforcement form a self-reinforcing loop.



theocracy — Graph

LEVEL	ID	DESCRIPTION	↑	↓
⊙ Goals	G1	Implement divine law as foundation of governance	—	R1 R2
	G2	Ensure spiritual salvation and moral purity	—	R3 R6
	G3	Maintain order, security, and territory	—	R4 R5
	G4	Preserve clerical authority as sole divine interpreter	—	R2 R3
◆ Requirements	R1	Sacred text as supreme legal source	G1	F1
	R2	Clerical qualification for governance	G1 G4	F1 F4
	R3	Religious conformity enforcement	G2 G4	F3
	R4	Military loyal to the theocratic order	G3	F4
	R5	Revenue system (religious taxes, endowments)	G3	F5
	R6	Education transmitting doctrine as civic education	G2	F6
◇ Functions	F1	Divine legislation — derive laws from scripture	R1 R2	L1
	F2	Religious adjudication — disputes per sacred law	—	L2
	F3	Moral enforcement — police virtue, suppress dissent	R3	L3
	F4	Military command and territorial defense	R2 R4	L4
	F5	Revenue collection and treasury	R5	L5
	F6	Doctrinal education	R6	L6
○ Logical Architecture	L1	Clerical-Legislative — supreme authority, councils	F1	P1 P2
	L2	Religious-Judicial — Sharia courts, tribunals	F2	P3
	L3	Enforcement — morality police, censorship	F3	P4
	L4	Military — army, revolutionary guard	F4	P5
	L5	Fiscal — treasury, religious taxes	F5	P6
	L6	Education — seminaries, madrasas	F6	P7
■ Physical Implementation	P1	Supreme Leader / Council of Guardians	L1	—
	P2	Sacred text as constitutional document	L1	—
	P3	Religious courts, Sharia judges	L2	—
	P4	Morality police, censorship boards	L3	—
	P5	Revolutionary Guard / military forces	L4	—
	P6	Religious treasury, zakat, endowments (waqf)	L5	—
	P7	Seminaries (Hawza), madrasas	L6	—

theocracy — Table

16.1 SE Decomposition

The full five-level decomposition is shown in the interactive visualization linked above. Click any node to see its description, parent links, and child links. Use the Table view for the complete traceability matrix.

16.2 Variation Point Bindings

VP1 = divine mandate, VP2 = territorial + religious conformity, VP3 = doctrinal interpretation by clergy, VP4 = clerical selection (conclave, council, or dynastic). The binding at VP1 constrains VP3 so severely that popular input to governance is structurally impossible.

16.3 The Structural Pathology of VP1 = Divine Mandate

The theocracy binds VP1 (source of authority) to a claim the system cannot, by its own logic, allow to be revised. This produces the strongest legitimization narrative in the entire social-systems catalogue — and the structurally weakest error-correction. Doctrine cannot be revised by empirical evidence because revising it would dissolve the legitimating claim. Authority cannot be reviewed externally because no external rung is recognised as competent to review divine mandate.

The result is **Pattern B — Claimed-Rung Inflation** in its purest form: the system claims rung 6 (deliberative + sacred legitimacy) while operating at rung 1 (clerical authority). The gap is enforced doctrinally rather than closed. Where a republic's separation of powers admits rung-3 evidence about the executive's conduct via independent courts and a free press, a theocracy treats the same evidence as rung-1 hostility — an attack on doctrine rather than information about practice.

Doctrine, law, education, and enforcement form a self-reinforcing loop: the clerical hierarchy interprets the sacred text, codifies it into law, teaches it as civic education, and enforces it through the morality apparatus. Each subsystem reinforces the others; none provides corrective feedback to the central doctrinal authority. There is, in STPA terms, no controller external to the loop authorised to register that the loop is producing harms.

These are not failures of clerical character. They are structural consequences of the VP1 binding.

16.4 STPA Highlights

UCA-T1: Religious courts apply divine law uniformly without contextual judgement.

Causal factor: the controller (religious court) and the source of authority (sacred text) are connected by a single chain of clerical interpretation. Contextual mitigation — circumstance, intent, proportionality — competes against doctrinal fidelity, and the rung-1 framing makes fidelity unambiguously the higher value. The result is the documented pattern of harsh, decontextualised sentences.

UCA-T2: Morality police enforces conformity instead of voluntary virtue.

Causal factor: G2 (moral orientation) becomes operationalised as compliance with visible behavioural norms (dress, prayer, gender separation, public conduct). The compliance is observable and so the enforcement mechanism can be built around it; voluntary internalised virtue, which is what the goal actually requires, is not observable and so cannot be enforced. The system substitutes the observable proxy for the unobservable goal — the classic proxy-metric pathology, in religious form.

UCA-T3: Clerical succession concentrates power in a self-perpetuating body.

Causal factor: VP4 = clerical selection (conclave, council, dynastic) means the body that selects the next supreme authority is itself the body that derives its authority from the supreme authority. There is no external rung at which the selection process can be challenged. The system tends over time toward concentration in whichever faction can secure the selection mechanism.

16.5 Fixes That Worked

Sustained reforms in active theocracies are rare for the same architectural reason that one-party state reforms are rare: the binding that produces the failure is also what produces the system's identity. The reforms that have worked have done so by **constraining the theocratic binding without abandoning it**.

Iranian constitutional architecture (1979-2009) — partial structural constraint on UCA-T3.

The Islamic Republic combines the rung-1 supreme religious authority (the Velayat-e Faqih) with elected institutions (parliament, presidency, assemblies). The structural intent was a workaround: preserve clerical legitimacy at the top while admitting rung-6 deliberation below. Evidence (positive): reformist movements 1997 and 2009 used the elected channels to surface significant rung-3 evidence (election irregularities, economic mismanagement, women's rights) that reached the public sphere. Evidence (negative): the rung-1 Guardian Council retains veto power over candidates and laws, and has used it to neutralise the channel openings. The reform is partial because the underlying binding has not changed.

Saudi Vision 2030 (2016 onward) — domain-narrowed channel opening.

Same pattern as China post-1978: open rung-3 economic and technical channels (foreign direct investment, tourism, social-policy modernisation) while preserving VP1. Evidence: significant rapid social reform on questions previously closed (women's driving,

public entertainment) accompanied by sustained closure on political-religious dissent. The structural argument: economic rung-3 evidence is admissible to the apex; political-religious rung-3 evidence is not.

Vatican II (1962-1965) — the most complete documented case.

The Catholic Church is not a theocracy, but the structural pattern of clerical doctrinal authority is the same. The Council was a structural opening: liturgical reform admitted vernacular languages (admitting laity to doctrinal participation), ecumenical engagement admitted other Christian rung-1 traditions as conversation partners, and modern biblical scholarship was admitted as a rung-3 input to doctrinal interpretation. Evidence: fifty years of demonstrable ecclesial change. Evidence of partial reversal: successive pontificates have restricted some of the openings (Latin Mass reauthorisation, doctrinal congregation interventions). The structural lesson is precisely what the religion-STPA chapter spells out: rung-elevated channels work where they are maintained, and are reversible by the same rung-1 hierarchy that authorised them unless the elevation is made structurally independent.

**The Design Principle**

The structurally honest reform of a theocracy is the one Vatican II attempted: admit rung-3 evidence and rung-6 deliberation in carefully bounded domains, do not contest VP1 directly, and accept that the reform will be partially reversible until the channel opening is given an independence the original hierarchy cannot rescind. Most attempted theocratic reforms fail at exactly this last step.

16.6 Platform Mapping

This system fills all ten universal functional slots identified in the [Ten Social Systems Compared](#):

Functional Slot	How This System Fills It
Authority & Decision-Making	Supreme religious authority (Pope, Supreme Leader, Caliph); clerical council; no independent secular check
Membership & Belonging	Territorial citizenship combined with required religious conformity; apostasy carries legal or social penalty
Resource Allocation	Religious endowments (waqf, church estates); state budget under clerical control; no separation of sacred and fiscal authority
Norm Setting & Enforcement	Religious law (canon law, sharia, halakha); clerical courts; religious police where applicable
Dispute Resolution	Religious courts; clerical arbitration; all dispute resolution channels pass through the religious hierarchy
Legitimation	Divine mandate; sacred texts as constitutional authority; clerical monopoly on authoritative interpretation
Succession & Continuity	Clerical selection: conclave, council deliberation, or dynastic clerical succession
External Representation	Supreme religious leader conducts state diplomacy; religious and foreign policy are indistinguishable
Socialisation	Mandatory religious education (madrasas, seminaries, religious schools); ritual participation as civic obligation
Activity Delivery	Religious services; welfare distribution (zakat, tithing); education; governance fused with religious practice

Navigate to the [interactive visualization](#) for the full graph and table.

17. Example: Corporation (For-Profit)

The for-profit corporation is the organizational family's most studied variant. Its SE decomposition reveals a system optimised for a single metric — profit — with governance designed to serve capital holders. This binding produces both extraordinary efficiency in capital allocation and a set of well-documented structural pathologies that follow directly from the architecture, not from individual moral failure.



corporation — Graph

LEVEL	ID	DESCRIPTION	↑	↓
⊙ Goals	G1	Generate sustainable profit and shareholder value	—	R1 R3
	G2	Deliver products/services satisfying market demand	—	R3 R5
	G3	Provide livelihood and purpose to employees	—	R4 R5
	G4	Ensure organizational continuity	—	R1 R2
◆ Requirements	R1	Legal incorporation with limited liability	G1 G4	F1 F5
	R2	Board accountable to shareholders	G4	F1
	R3	Revenue exceeding costs	G1 G2	F2 F3
	R4	Hierarchical management with clear authority	G3	F1 F4
	R5	Employment contracts	G2 G3	F4
◇ Functions	F1	Strategic direction — vision, budgets, appointments	R1 R2 R4	L1
	F2	Product/service delivery	R3	L2
	F3	Financial management — accounting, treasury	R3	L3
	F4	HR management — recruit, develop, compensate	R4 R5	L4
	F5	Legal and compliance	R1	L5
○ Logical Architecture	L1	Governance — board, shareholders meeting, audit	F1	P1 P2
	L2	Operations — product dev, manufacturing, supply chain	F2	P3
	L3	Financial — accounting, treasury, controlling	F3	P4
	L4	HR — recruiting, payroll, training	F4	P5
	L5	Legal & Compliance — counsel, regulatory	F5	P1
■ Physical Implementation	P1	Articles of incorporation, board resolutions	L1 L5	—
	P2	CEO, C-suite, management hierarchy	L1	—
	P3	Factories, offices, IT, supply chain	L2	—
	P4	ERP (SAP/Oracle), bank accounts, annual report	L3	—
	P5	Employment contracts, HR software, training	L4	—

corporation — Table

17.1 SE Decomposition

The full five-level decomposition is shown in the interactive visualization above. Click any node to see its description and connections. Use the Table view for the complete traceability matrix.

17.2 Variation Point Bindings

Variation Point	Binding
VP1: Source of Authority	Capital ownership — authority is proportional to shareholding
VP2: Membership Boundary	Employment contract (voluntary, conditional, terminable)
VP3: Decision-Making	Board vote weighted by share ownership
VP4: Succession	Board appointment — not election, not merit review by peers

The key structural feature: employees are the system's primary operators — they hold the most accurate process models of what the system actually does — but they have no formal governance authority. The people who know most about the corporation's real functioning are structurally excluded from its decision-making.

17.3 The Structural Pathology of VP1 = Capital Ownership

Every variation-point binding creates trade-offs. VP1 = capital ownership has two well-documented consequences.

Time horizon misalignment. Shareholders can exit in milliseconds (liquid markets). Employees, suppliers, and communities cannot exit as quickly. The governance structure therefore weights the interests of the fastest-exit stakeholder most heavily — which creates systematic pressure toward decisions that serve 90-day earnings cycles over multi-year value creation.

Operational knowledge exclusion. The people with the most accurate information about product quality, customer relationships, safety risks, and long-term capabilities — the workers — have no formal channel through which that information reaches the decision-making authority. Executives receive information that has been filtered upward through layers of management with their own incentives. The board receives information filtered through the executives. The process model at the governance level is systematically less accurate than the process model held by the operators.

These are not failures of bad management. They are structural consequences of the VP1 binding.

17.4 STPA Highlights

UCA-C1: The external auditor certifies materially false financial statements.

Causal factor: Arthur Andersen earned \$25 million in audit fees and \$27 million in consulting fees from Enron in 2000. The auditor's business revenue depended on client satisfaction. Accurate qualification of the accounts would end the relationship; passing the accounts continued it. The control structure made honest auditing financially irrational for the auditor — not because Arthur Andersen's people were unusually corrupt, but because the incentive structure made the corrupted outcome individually optimal.

UCA-C2: The board approves compensation that rewards short-term stock price over long-term value creation.

Causal factor: Options vesting over 3–4 years, combined with quarterly earnings reports and average executive tenure of 5 years, encodes a time horizon into the compensation structure that is shorter than the time horizon of the decisions being made. Leveraged acquisitions, capacity-depleting cost cuts, and aggressive accounting are individually rational under this incentive structure and collectively destructive at scale.

UCA-C3: Workers have no voice in decisions that directly affect them.

Causal factor: VP1 = capital ownership structurally excludes the operational knowledge held by workers from the governance process. Decisions about working conditions, investment priorities, and cost structures are made by people whose information about operational reality is systematically less complete than the people executing those decisions.

17.5 Fixes That Worked

Sarbanes-Oxley Act 2002 — addressing UCA-C1

After Enron and WorldCom, the US Congress made these structural changes:

- Prohibited audit firms from providing most non-audit services (consulting, internal audit outsourcing) to their audit clients — severing the revenue dependency that created the conflict.
- Required CEO and CFO to personally certify financial statement accuracy under criminal penalty, removing deniability.
- Created the Public Company Accounting Oversight Board (PCAOB) as an independent inspectorate, replacing self-regulation by the audit profession.

Evidence: No Enron-scale accounting fraud has recurred in US-listed companies at comparable magnitude in the 22 years since SOX. The structural conflict of interest was directly addressed.

German co-determination — Mitbestimmung (1976) — addressing UCA-C2 and UCA-C3

Germany's Mitbestimmungsgesetz requires companies with more than 2,000 employees to give workers equal representation on the supervisory board (Aufsichtsrat). In Germany's two-tier structure, the Aufsichtsrat oversees management strategy without running day-to-day operations. Worker representatives bring:

- A different time horizon (employment stability, long-term skill investment, not quarterly share price)
- Operational knowledge that capital holders structurally lack
- Incentive to flag risks visible to the workforce before they surface in financial results

Evidence: German DAX companies maintained significantly more stable employment during the 2008–2009 financial crisis and the 2020 COVID recession than comparable US firms. German industrial companies (Volkswagen, Siemens, BASF, Bosch) have sustained R&D investment cycles measured in decades and maintained global technological leadership through multiple economic downturns. The Mitbestimmung structure is a contributing — not the sole — causal factor; it works in combination with Germany's long-term banking relationships and patient capital model.

UK Corporate Governance Code (post-Cadbury 1992) — addressing UCA-C2

Following UK corporate scandals (Maxwell, BCCI, Polly Peck), the Cadbury Report mandated:

- Separation of CEO and Chairman roles — preventing a single person from both running the company and overseeing themselves.
- Majority independent non-executive directors on audit, remuneration, and nomination committees.
- Annual shareholder vote on executive compensation (say-on-pay).

Evidence: UK listed companies report >90% compliance. The structural concentration of unchecked personal power that characterised the Maxwell case has been architecturally eliminated as a governance norm for listed companies.

The Design Principle

None of these remedies asks executives or auditors to be more virtuous. They change the structure so that the incentive points in a different direction: severing the auditor's financial dependence on the auditee; changing the time horizon encoded in the compensation structure; adding voices with different interests and different knowledge to the decision room. Structural change, not moral exhortation.

17.6 Platform Mapping

Functional Slot	How This System Fills It
Authority & Decision-Making	Board of Directors; weighted shareholder vote
Membership & Belonging	Employment contract; shareholder register
Resource Allocation	Capital budgeting; executive discretion within board-approved limits
Norm Setting & Enforcement	HR policies; compliance function; employment law
Dispute Resolution	Employment tribunals; shareholder litigation; regulatory enforcement
Legitimation	Profit generation; shareholder returns; employment provision
Succession & Continuity	Board appointment; executive succession planning
External Representation	CEO; investor relations; legal and regulatory affairs
Socialisation	Onboarding; training; corporate culture
Activity Delivery	Operations; product/service development and delivery

For the full SE decomposition: [interactive visualization](#). For the comparative analysis across all ten systems: [Ten Social Systems Compared](#). For the cross-system control structure analysis: [Control Structures](#).

18. Example: University

The university is the organization family's most internally contested variant — it must simultaneously serve truth (research), education (teaching), and self-governance (academic freedom), goals that regularly conflict. Its SE decomposition exposes why corporatization threatens the system's core logic.



university — Graph

LEVEL	ID	DESCRIPTION	↑	↓
⊙ Goals	G1	Advance knowledge through research	—	R1 R5
	G2	Educate and credential the next generation	—	R2 R6
	G3	Serve society through knowledge transfer	—	R1 R2
	G4	Preserve academic freedom and autonomy	—	R3 R5
◆ Requirements	R1	Peer review as knowledge validation standard	G1 G3	F1
	R2	Structured curricula leading to degrees	G2 G3	F2
	R3	Academic self-governance — senates, elected rectors	G4	F3
	R4	Sustainable funding — tuition, grants, state	—	F5
	R5	Tenure protecting freedom of inquiry	G1 G4	F1 F3
	R6	Accreditation by external bodies	G2	F4
◇ Functions	F1	Research — conduct and publish scholarship	R1 R5	L2
	F2	Teaching — lectures, seminars, examinations	R2	L3
	F3	Academic governance — senate, appointments	R3 R5	L1
	F4	Student administration — admissions, graduation	R6	L4
	F5	Financial management — budgets, grants	R4	L4
○ Logical Architecture	L1	Governance — senate, rector, faculty councils	F3	P1 P2
	L2	Research — departments, institutes, labs, libraries	F1	P3
	L3	Teaching — curricula, lecture halls, exams	F2	P3 P4
	L4	Administration — student office, finance, HR	F4 F5	P4 P5
■ Physical Implementation	P1	University charter, Grundordnung	L1	—
	P2	Rector, deans, senate, faculty councils	L1	—
	P3	Labs, libraries, lecture halls, institutes	L2 L3	—
	P4	Student registry (HISinOne), exam office	L3 L4	—
	P5	Endowment, state funding, DFG/ERC grants	L4	—

university — Table

18.1 SE Decomposition

The full five-level decomposition is shown in the interactive visualization linked above. Click any node to see its description, parent links, and child links. Use the Table view for the complete traceability matrix.

18.2 Variation Point Bindings

VP1 = scholarly merit (unique among social systems), VP2 = competitive admission, VP3 = collegial deliberation (faculty senate), VP4 = mixed (elected rector + appointed deans). The tension between VP1 (merit) and external pressure to bind VP3 to hierarchical management is the central conflict in modern higher education.

18.3 The Structural Pathology of VP1 = Scholarly Merit

"Scholarly merit" is the only variation-point binding in the social-systems catalogue that requires the controller to be assessable on the same standard it controls under. Faculty senates decide on faculty, peer review decides on peer-reviewed work, tenure committees decide on tenure-track scholarship. The system is structurally circular by design — and the circularity works only when the merit signal is not corrupted.

Two corruptions are now well documented. The first is the **proxy-metric substitution**: scholarly merit is hard to measure, but citation counts and journal impact factors are easy. The career structure has migrated to the easy metric, and so the system's most consequential decisions are now made on the proxy rather than the goal. The classic Goodhart pattern: "when a measure becomes a target, it ceases to be a good measure."

The second is **publication bias** — the literature itself is a biased sample, because journals select for positive, novel, surprising findings. The cumulative scholarship the system uses to validate further scholarship has therefore been selected by a mechanism that no longer tracks evidence strength. The 2015 Open Science Collaboration found that only 36% of psychology findings replicated. The collegial-merit decision the structure depends on is being made on data the structure has corrupted.

These are not failures of bad scholars. They are structural consequences of the VP1 binding combined with a measurement choice (impact factor) that the system was not architecturally protected against.

18.4 STPA Highlights

UCA-U1: Journals publish selectively (publication bias), so the literature accumulates a biased sample.

Causal factor: the decision to publish is made after the result is known, by editors and reviewers who reward novelty and significance. A non-replication, a null result, or a careful but unsurprising study faces structural pressure against publication. The cumulative-evidence rung (rung 4) is supposed to correct rung-3 single-study errors; here the rung-4 corpus is itself a rung-3 sample biased at the source.

UCA-U2: Tenure and promotion committees use bibliometrics as proxies for research quality.

Causal factor: committees are overloaded, citation counts are countable, and quality is contested. The path of least administrative resistance is to use h-index, impact factor, and grant size as decision-makers for what is supposed to be a peer-review judgement. Once the proxy is the decision rule, scholars rationally optimise for the proxy.

UCA-U3: Research budgets concentrate at institutions that already hold grant track records.

Causal factor: grant applications are evaluated partly on prior funding success — a Matthew-effect feedback loop that makes the future distribution of resources a function of the past distribution, regardless of merit shifts. The intended rung-3 selection mechanism (peer review of the proposal) is overridden by a rung-1 institutional-reputation signal.

18.5 Fixes That Worked

Pre-registration and Registered Reports — addressing UCA-U1.

Pre-registration commits researchers to hypotheses, methods, and analysis plans before data collection; Registered Reports take this further, with peer review of the protocol before any data exists, and an acceptance-in-principle that survives whatever the result turns out to be. Evidence: the Center for Open Science hosts over 100,000 pre-registrations; more than 300 journals offer Registered Reports. Pre-registered studies replicate at substantially higher rates than unconstrained ones. The structural change: the publication decision is decoupled from the result.

DORA (Declaration on Research Assessment, 2012) and the UK REF panel review — addressing UCA-U2.

DORA, signed by more than 25,000 individuals and 3,000 organisations, commits assessors to evaluate research on its content rather than the journal in which it appeared. The UK Research Excellence Framework moved to article-level panel assessment after the 2014 cycle, explicitly disconnecting reward from journal-level proxies. Evidence: the UK REF 2021 panels read submitted outputs directly; impact factor was excluded as an assessment input. The structural change: the proxy is severed from the decision.

Open access mandates (Plan S, NIH 2008 onward) — partial fix for UCA-U1 and UCA-U3.

Mandating open publication of publicly funded research removes the journal-publisher gatekeeping that selected the published sample. Evidence: Plan S coalition includes Wellcome Trust, ERC, UK Research Councils; compliance rates on funded grants exceed 70% by 2024. The structural change is partial — selection has migrated to preprint visibility — but the original gatekeeping is no longer load-bearing for funded research.



The Design Principle

The fixes do not ask scholars to be more honest. They change the structure so that the proxy metric no longer determines the consequence. Pre-registration moves the publication decision before the result is known. Article-level assessment moves the reward decision off the journal-level proxy. Open access moves the visibility decision off the publisher. Each is a structural intervention against a specific causal pathway STPA identifies.

18.6 Platform Mapping


This system fills all ten universal functional slots identified in the [Ten Social Systems Compared](#):

Functional Slot	How This System Fills It
Authority & Decision-Making	Faculty senate (academic matters); rector/president (executive); board of trustees (strategic oversight)
Membership & Belonging	Matriculated students; academic and administrative staff; alumni (honorary standing)
Resource Allocation	Rector/president budget; competitive grant allocation; tuition and third-party research funding
Norm Setting & Enforcement	Academic regulations; research ethics boards; HR policy; accreditation standards
Dispute Resolution	Academic appeals committees; disciplinary panels; employment tribunals
Legitimation	Knowledge creation and transmission; accreditation; academic freedom as constitutive value
Succession & Continuity	Rector election or appointment; academic tenure; departmental and programme continuity
External Representation	Rector; research partnerships; international agreements; alumni relations
Socialisation	Academic induction; doctoral mentorship; collegiate culture; peer learning communities
Activity Delivery	Teaching; research; knowledge transfer; community engagement and public scholarship

Navigate to the [interactive visualization](#) for the full graph and table.

19. Example: Military

The military is the most hierarchically pure social system — command authority flows strictly top-down, with minimal upward feedback by design. Its SE decomposition reveals why civilian oversight (an external control loop) is architecturally essential: without it, the system has no self-correcting mechanism.



military — Graph

LEVEL	ID	DESCRIPTION	↑	↓
⊙ Goals	G1	Defend the state against external threats	—	R1 R2
	G2	Maintain internal order when directed	—	R1 R3
	G3	Project power and deter aggression	—	R2 R4
	G4	Disaster relief and humanitarian assistance	—	R2 R5
◆ Requirements	R1	Strict command chain from top to soldier	G1 G2	F1
	R2	Trained, equipped, deployable force	G1 G3 G4	F2 F3
	R3	Subordination to civilian authority	G2	F1
	R4	Intelligence and situational awareness	G3	F4
	R5	Logistics — supply, transport, medical	G4	F5
	R6	Military justice framework	—	F6
◇ Functions	F1	Command — orders, coordination, strategy	R1 R3	L1
	F2	Combat — engage adversaries, defend	R2	L2
	F3	Training — drill, exercises, qualification	R2	L3
	F4	Intelligence — collect, analyze, disseminate	R4	L4
	F5	Logistics — ammo, fuel, food, medical	R5	L5
	F6	Military justice — discipline, courts martial	R6	L6
○ Logical Architecture	L1	Command — general staff, HQ, C4ISR	F1	P1
	L2	Combat — brigades, vessels, air wings	F2	P2
	L3	Training — academies, schools, ranges	F3	P3
	L4	Intelligence — SIGINT, HUMINT, cyber	F4	P4
	L5	Logistics — depots, transport, medical corps	F5	P5
	L6	Military Justice — courts martial, JAG, MP	F6	P6
■ Physical Implementation	P1	General staff HQ, situation rooms, comms	L1	—
	P2	Barracks, bases, warships, aircraft, armor	L2	—
	P3	Academies (West Point, Sandhurst, Führungsakademie)	L3	—
	P4	Intelligence agencies, SIGINT stations	L4	—
	P5	Supply depots, field hospitals, transport	L5	—
	P6	Military courts, provost marshal	L6	—

military — Table

19.1 SE Decomposition

The full five-level decomposition is shown in the interactive visualization linked above. Click any node to see its description, parent links, and child links. Use the Table view for the complete traceability matrix.

19.2 Variation Point Bindings

VP1 = command hierarchy (commission from the state), VP2 = enlistment/conscription, VP3 = command order (superior to subordinate), VP4 = promotion by merit/seniority. The binding at VP3 is the most rigid of any social system — orders are obeyed, not debated. This produces decisive action but also enables atrocities when the command chain is corrupt.

19.3 The Structural Pathology of VP3 = Command Order

The strict downward command channel is the military's defining strength. The matching weakness is the absence of any upward channel of equivalent authority. STPA classifies this as an **accountability void**: the controller (the commanding officer) and the interested party (the unit whose conduct is in question) are the same entity. There is no internal mechanism by which a soldier's report of a war crime, a logistics fraud, or a sexual assault can reach a decision-maker outside the chain of command — because by construction, the chain of command is the only authorised channel.

A second consequence follows: the strategic information that civilian leaders need to decide whether to commit forces flows up the same chain that has the strongest incentive to present its readiness favourably. Vietnam-era US body counts and the Soviet General Staff's 1979 Afghanistan readiness reports are the canonical examples — not because the officers were unusually dishonest, but because the structure made honest reporting career-costly without any matching reward.

These are not failures of military discipline. They are structural consequences of the VP3 binding, and they can only be addressed by parallel channels that route around the chain of command.

19.4 STPA Highlights

UCA-M1: The commanding officer decides whether to prosecute felonies committed by the officer's own troops.

Causal factor: the prosecution decision in most militaries historically sat with the convening authority — the same officer whose command culture, training programme, and selection record the case would reflect upon. Reporting rates collapsed; cases that were reported faced an arbiter with a documented institutional incentive against finding fault. The structural conflict of interest is identical to the one Sarbanes-Oxley addressed for corporate auditors.

UCA-M2: Civilian leadership commits forces to combat without parliamentary deliberation.

Causal factor: where the executive can deploy unilaterally, the decision is made by a small group operating at justificatory rung 0/1 (national security authority, classified intelligence). The corrective channels — parliamentary debate, public deliberation, allied consultation — operate at rung 6 and arrive after the commitment is irreversible.

UCA-M3: Strategic readiness is reported through the chain it would indict.

Causal factor: readiness assessments, casualty estimates, and territorial control reports are produced by the same officers whose advancement depends on them being favourable. The civilian decision-maker's process model is populated entirely from this filtered source.

19.5 Fixes That Worked

Wehrbeauftragter (Bundeswehr, Germany, 1957) — addressing UCA-M3 and a class of UCA-M1.

The Parliamentary Commissioner for the Armed Forces has independent access to any unit, any file, and any individual soldier without going through the chain of command, and reports directly to the Bundestag rather than the Defence Ministry. The Commissioner's annual report has repeatedly surfaced structural problems — equipment shortfalls, training gaps, extremism cases, harassment patterns — that command had institutional incentive to conceal. Evidence: sixty-eight years of continuous civilian-controlled defence reform, including the discovery and dismantling of right-wing extremist networks inside KSK (2020) that command had failed to address through internal investigations.

US NDAA 2022 — Special Trial Counsel — addressing UCA-M1.

The 2022 National Defense Authorization Act transferred jurisdiction over sexual assault and a list of other serious felonies from commanding officers to independent Special Trial Counsel, removing the commander from the prosecution decision entirely. Evidence: reporting rates and prosecution rates for referred cases both rose in the first eighteen months; the structural change was identified through sustained advocacy using exactly the conflict-of-interest argument STPA names.

German parliamentary mandate (Article 87a GG and the 1994 Constitutional Court ruling) — addressing UCA-M2.

The Bundeswehr cannot be deployed outside German territory without a Bundestag vote on each deployment, including subsequent renewals. The deliberation operates at rung 6 (parliamentary debate) and produces a public record. Evidence: every German out-of-

area deployment since Kosovo has been mandated and renewed publicly; mission scope has been contested and amended at renewal points (Afghanistan combat role, EUFOR Sahel withdrawal). Whatever one thinks of the outcomes, the process of deliberative commitment is the structural circuit breaker.

**The Design Principle**

None of these remedies weakens the chain of command — the system needs VP3 to remain rigid for its operational logic to work. They install parallel channels at higher rungs that route around the chain of command for specific decisions where the conflict of interest is structural. The Wehrbeauftragter does not give orders. The Special Trial Counsel does not promote. The Bundestag does not command. Each is a circuit breaker on a specific failure path.

19.6 Platform Mapping

This system fills all ten universal functional slots identified in the [Ten Social Systems Compared](#):

Functional Slot	How This System Fills It
Authority & Decision-Making	Chain of command from civilian minister to field commander; joint chiefs as strategic advisory body
Membership & Belonging	Voluntary enlistment or compulsory conscription; rank and unit assignment define standing
Resource Allocation	Parliamentary defence budget; theatre commanders; logistics chain under operational command
Norm Setting & Enforcement	Military law (UCMJ / national equivalent); rules of engagement; code of conduct; Geneva Conventions
Dispute Resolution	Courts-martial; military ombudsman; administrative review boards; civilian oversight tribunal
Legitimation	State mandate for territorial defence; professional military ethos; oath of service
Succession & Continuity	Promotion boards; officer commissioning pipeline; doctrinal continuity through staff colleges
External Representation	Joint chiefs; defence attachés; alliance structures (NATO, bilateral treaties)
Socialisation	Basic training and indoctrination; unit cohesion rituals; rank culture and esprit de corps
Activity Delivery	Combat operations; deterrence; logistics; intelligence; peacekeeping; training and readiness

Navigate to the [interactive visualization](#) for the full graph and table.

20. Example: The Family

The family is the smallest and oldest social system — and the only one where the membership boundary (VP2) is primarily biological rather than contractual. Its SE decomposition reveals a system that must solve all ten platform problems with typically 2–5 people and no formal governance documents.



family — Graph

LEVEL	ID	DESCRIPTION	↑	↓
⊙ Goals	G1	Ensure reproduction and survival of offspring	—	R1 R2
	G2	Provide emotional security, love, belonging	—	R3
	G3	Socialize children into norms and competences	—	R4 R6
	G4	Pool resources and manage economic risk	—	R2 R5
	G5	Transmit identity, heritage, property across generations	—	R5
⬢ Requirements	R1	Pair bond or co-parenting agreement	G1	F1 F3
	R2	Shelter and material provision	G1 G4	F1 F4
	R3	Emotional availability and care	G2	F1 F2
	R4	Authority structure for child-rearing	G3	F3
	R5	Inheritance and succession norms	G4 G5	F4 F5
	R6	Interface with schools, healthcare, state	G3	F6
◇ Functions	F1	Nurture — feed, shelter, protect, care	R1 R2 R3	L2
	F2	Socialize — teach language, norms, values	R3	L4
	F3	Govern — discipline, education, health decisions	R1 R4	L1
	F4	Economic management — earn, budget, save	R2 R5	L3
	F5	Ritual and identity — traditions, storytelling	R5	L4
	F6	External representation — school, state interface	R6	L1
○ Logical Architecture	L1	Parental Authority — decisions, discipline	F3 F6	P1 P2
	L2	Care — nurturing, health, emotional support	F1	P1
	L3	Economic — income, budgeting, property	F4	P3
	L4	Socialization — education, culture, values	F2 F5	P4 P5
▣ Physical Implementation	P1	Home — house or apartment, shared space	L1 L2	—
	P2	Parents/guardians, marriage/partnership	L1	—
	P3	Family income, bank accounts, property, wills	L3	—
	P4	Daily routines, rituals, photo albums, heirlooms	L4	—
	P5	School enrollment, pediatrician, community	L4	—

family — Table

20.1 SE Decomposition

The full five-level decomposition is shown in the interactive visualization linked above. Click any node to see its description, parent links, and child links. Use the Table view for the complete traceability matrix.

20.2 Variation Point Bindings

VP1 = biological/legal bond (unique), VP2 = birth/adoption (non-voluntary for children), VP3 = consensual negotiation (between adults) + parental authority (over children), VP4 = self-selection (founding a new family). The key structural insight: the family is the only system where VP6 (norm enforcement) relies primarily on love rather than coercion.

20.3 The Structural Pathology of Parental Authority over Non-Consenting Members

The family does almost everything right structurally: VP6 (norm enforcement) relies on love rather than coercion, VP3 (decision-making between adults) is consensual negotiation, VP1 (legitimation) is biological bond. The pathology lives in the one part of the

system where consent is not available: the relationship between parental authority and children who did not enter the family by their own choice.

Children are full members of the family system but their consent to its rules is presumed, not given. The system's design assumes the parental authority operates in the child's interest — and structurally, it usually does. Where it does not, the system is unusually closed: the family is a single legal entity in dealings with the state, so external rung-3 channels (a teacher's observation, a doctor's record, a neighbour's report) struggle to reach the inside. The same architectural feature that protects the family from external interference also protects internal harm from external observation.

A second pattern appears in well-functioning families that nonetheless fail structurally. The **triangulation pathology**: when adults stop communicating directly and start communicating through the children, the system has lost its rung-6 deliberative channel between adults and replaced it with a rung-1 authority channel running through a structurally junior member. The fix is not to instruct the children differently; it is to restore the direct adult channel.

These are not failures of love. They are structural consequences of the unique VP2 binding (non-voluntary membership for children) interacting with the unique closure of the family boundary.

20.4 STPA Highlights

UCA-F1: A parent uses rung-0 coercion on a child whose distress signals would be addressed at rung-3 by external attention.

Causal factor: corporal punishment, prolonged shaming, and isolation are rung-0 control actions used in a system whose explicit goal is the child's healthy development — a goal whose evidence (rung-3 developmental health) the rung-0 control action makes harder to read. The parent's process model that the child is "misbehaving" is often the only model in the system, because there is no other adult present to maintain a competing model.

UCA-F2: The family communicates only through the children — no direct adult channel exists.

Causal factor: in family systems where adult communication has degraded, messages flow parent $\leftrightarrow$ child $\leftrightarrow$ parent. The child becomes a controller's instrument (rung-1 authority) at an age where the child cannot refuse the role. The therapy literature identifies this as the canonical structural sign of an unwell family — not because individuals are unwell, but because the connection pattern has reorganised around a structural failure.

UCA-F3: Domestic abuse is invisible to external observers because the family is a single legal entity to the state.

Causal factor: the same architectural feature that gives the family privacy from state intrusion also makes it the social system in which serious harm can persist longest without external rung-3 observation. The internal feedback channels — children's protests, the abused partner's distress — operate at justificatory rungs the abuser has structural authority to filter out.

20.5 Fixes That Worked

Sweden's corporal punishment ban (1979) — addressing UCA-F1.

The first national prohibition of corporal punishment of children, accompanied by a public-information campaign and supportive social services. Evidence: attitudes shifted substantially across one generation (1965: 53% of Swedes supported corporal punishment; 2010: 8%); rates of physical punishment fell, severe abuse incidents fell, child mortality from maltreatment fell. The state operated as an external rung-3 channel that the family could not filter, and changed the cost calculus of rung-0 control. By 2024, more than 60 countries had followed.

Mandatory reporting laws + multi-agency safeguarding — addressing UCA-F3.

Teachers, doctors, social workers, and (in some jurisdictions) clergy and lawyers are required by law to report suspected abuse to a state child-protection agency. Evidence: introduction in successive Anglophone jurisdictions through the 1970s–80s; subsequent UK Children Act 1989 and 2004 reforms after the Jasmine Beckford and Victoria Climbié inquiries; consistent multi-decade rise in identified cases per capita as the reporting channel reaches saturation. The structural change: a rung-3 channel exists that the family cannot filter.

Parental leave policy (Germany Elternzeit, Sweden föräldraförsäkring) — addressing UCA-F2.

Extended paid parental leave for both parents, with a portion non-transferable, structurally increases the time both adults are co-present in the early-childhood phase. Evidence: both countries report measurably higher rates of involved fatherhood, lower rates of post-divorce custody conflict, and (in Sweden's longest-running data) modestly improved parental relationship outcomes. The fix is not relationship counselling; it is time — the structural precondition for an adult-to-adult direct channel to exist in the first place.

Family therapy as institutional rung-2/6 mediation — partial fix for UCA-F2.

Family therapy is structurally an external party that holds the rung-6 deliberative space the family cannot maintain itself, and applies rung-2 consistency checks to the family's own narrative ("you said X, your partner says Y, who is the child reporting to?"). Evidence: well-replicated effect on triangulation patterns specifically; weaker effect when the underlying structural problem is rung-0 violence (where the indicated fix is exit, not therapy).

**The Design Principle**

The family's pathologies are the price of its strengths: the privacy that protects love also hides harm; the parental authority that nurtures can coerce; the closed adult dyad that anchors the system can collapse into triangulation. The fixes that work do not abolish the strengths; they install external rung-3 channels (state observation, mandatory reporting) and structural preconditions (parental leave) that let the system's own corrective loops reach the cases the love-based mechanism does not catch.

20.6 Platform Mapping


This system fills all ten universal functional slots identified in the [Ten Social Systems Compared](#):

Functional Slot	How This System Fills It
Authority & Decision-Making	Parental authority over children; negotiated consensus between adult partners
Membership & Belonging	Biological kinship or legal adoption; marriage or partnership; non-voluntary for children
Resource Allocation	Household income pooling; parental discretion over expenditure; children's allowances
Norm Setting & Enforcement	Parental rules; shared household norms; enforced primarily through relationship and love rather than coercion
Dispute Resolution	Direct negotiation between adults; parental arbitration for child conflicts; family mediation as last resort
Legitimation	Biological bond; legal recognition (marriage, guardianship); emotional attachment and mutual care
Succession & Continuity	Children reaching independence and founding their own families — the system reproduces itself
External Representation	Parents as legal guardians; household as the legal and fiscal unit in dealings with the state
Socialisation	Primary socialisation — language, values, emotional regulation, identity, basic social norms
Activity Delivery	Care, nutrition, housing, educational support, emotional support — the full welfare of members

Navigate to the [interactive visualization](#) for the full graph and table.

21. Example: One-Party State

The one-party state fuses party and government into a single hierarchy, creating the most tightly controlled governance variant. Its SE decomposition reveals a system designed to prevent any alternative power center — which also means it has no structural mechanism for correcting its own errors.



one-party-state — Graph

LEVEL	ID	DESCRIPTION	↑	↓
🎯 Goals	G1	Implement party ideology as organizing principle	—	R1 R3
	G2	Maintain party monopoly on political power	—	R1 R2 R4
	G3	Ensure order, security, and territory	—	R4 R5
	G4	Deliver economic development as legitimacy proof	—	R5 R6
⬢ Requirements	R1	Constitutional enshrinement of party's leading role	G1 G2	F1
	R2	Party-state fusion — parallel hierarchies everywhere	G2	F1 F2
	R3	Ideological conformity — censorship, propaganda	G1	F3
	R4	Security apparatus loyal to the party	G2 G3	F4
	R5	Centralized economic direction	G3 G4	F5
	R6	Cadre system — party membership for advancement	G4	F2
◇ Functions	F1	Party-directed legislation	R1 R2	L1
	F2	Cadre management — recruit, evaluate, promote	R2 R6	L2
	F3	Ideological control — propaganda, censorship	R3	L3
	F4	Security and surveillance	R4	L4
	F5	Economic direction — plans, SOEs, investment	R5	L5
○ Logical Architecture	L1	Party-State Legislative — politburo, congress	F1	P1
	L2	Cadre & Personnel — org department	F2	P2
	L3	Propaganda & Control — media, censorship	F3	P3
	L4	Security — secret police, surveillance	F4	P4
	L5	Economic Planning — commission, SOEs	F5	P5
🏢 Physical Implementation	P1	Politburo / Central Committee	L1	—
	P2	Organization Department, party schools, cadre files	L2	—
	P3	State media (Xinhua, CCTV), Great Firewall	L3	—
	P4	Secret police (Stasi, KGB, MSS), surveillance	L4	—
	P5	State Planning Commission, SOEs, 5-year plans	L5	—

one-party-state — Table

21.1 SE Decomposition

The full five-level decomposition is shown in the interactive visualization linked above. Click any node to see its description, parent links, and child links. Use the Table view for the complete traceability matrix.

21.2 Variation Point Bindings

VP1 = ideological mandate (party), VP2 = territorial citizenship (compulsory), VP3 = party committee decision (autocratic), VP4 = cadre promotion (internal, party-controlled). The binding at VP3 produces the characteristic 'rubber-stamp parliament' — a legislative body that exists structurally but cannot exercise independent legislative function.

21.3 The Structural Pathology of VP3 = Party Committee Decision

The one-party state's defining commitment is that no alternative power center may exist. Every feedback channel that competitive politics, free press, independent judiciary, and civil society provide in a republic is, by VP3 binding, eliminated by design. The system

has no architectural mechanism for error correction: the controller cannot be wrong because the controller cannot be questioned, and there is no body authorised to find it wrong.

This produces a recurrent pattern STPA calls the **self-sealing process model**. The apex receives reports from cadres whose promotion depends on the reports being favourable; the cadres receive instructions from the apex whose authority depends on the instructions being unquestioned. Each level filters out information that the next level above does not want to hear, because each level operates at justificatory rung 1 (party authority) and treats rung-3 contradictory evidence as rung-1 disloyalty. The process model the apex acts on diverges from reality, and there is no internal channel to correct the divergence.

The historical record is severe: the Great Leap Forward famine (30+ million dead, 1959–1961) was driven in significant part by inflated grain-yield reports that the central planning apparatus could not distinguish from reality; the 1986 Chernobyl response began with the local controller's process model that the reactor could not have exploded because the design said so; the December 2019 suppression of Wuhan clinicians' early warnings about a novel coronavirus delayed the response by a critical few weeks. These are not failures of bad individuals — they are structural consequences of a control architecture with no rung-3 corrective channel.

21.4 STPA Highlights

UCA-OP1: Cadre reporting filters bad news upward, so the apex process model diverges from reality.

Causal factor: career advancement is decided by the level above, which has both authority and incentive to reject reports of bad news. The rational cadre transmits good news at face value and bad news with euphemism, delay, or omission. The aggregation over many levels produces an apex process model in which every region is meeting targets, every plan is on schedule, and every dissent is fringe.

UCA-OP2: Censorship suppresses early-warning signals from the population.

Causal factor: a free press exists as a class of rung-3 feedback channel; an unfree press exists as an extension of the apex's rung-1 messaging. The controllers that decide what may be published are part of the same hierarchy whose performance the press would otherwise scrutinise. The result is that real anomalies — a disease outbreak, an industrial accident, a famine — are pre-filtered before they enter the public information environment.

UCA-OP3: Term limits are removed, eliminating the last in-system circuit breaker on incumbent power.

Causal factor: in a system without competitive elections, periodic rotation of the supreme office is the only mechanism that forces a fresh process model into the apex. Removing the constraint (PRC 2018) closes one of the very few remaining structural correction channels, leaving the system maximally dependent on a single individual's judgement remaining accurate.

21.5 Fixes That Worked

The structural fixes for a one-party state are unusually narrow: by construction, they would require relaxing the VP3 binding that defines the system. Where partial fixes have worked, they have done so by **opening rung-3 feedback channels in narrow domains without touching VP1**.

China post-1978 economic reform — addressing a narrowed UCA-OP1.

The reform-era PRC deliberately opened rung-3 feedback in the economic domain — market prices as a rapid signal of supply-demand mismatches, township-village enterprise experimentation, special economic zones as policy laboratories. Evidence: poverty rates fell from above 80% to below 1% over four decades; GDP per capita rose by an order of magnitude. The reform did not touch VP1 (party leadership remained absolute); it opened rung-3 channels in a domain where the apex needed accurate information and the rung-3 evidence (prices, profit/loss) could not easily be falsified upward.

Vietnamese Đổi Mới (1986) — same pattern, different scale.

Same architectural move as the Chinese reform: market-mediated rung-3 feedback in production and trade, party leadership preserved. Evidence: poverty rate fell from 60% to under 5% over 35 years.

Internal audit functions in modern one-party states — narrow defence against UCA-OP1.

The PRC's Central Commission for Discipline Inspection, the SOE audit office, and provincial inspection tours are within-system attempts to insert rung-3 evidence about cadre performance into a process model that would otherwise be entirely rung-1. Evidence:

the mechanism has surfaced and prosecuted significant cases — at the same time, the inspectorate is part of the same hierarchy whose performance it audits, so its independence is structurally limited.

**The Design Principle**

The fixes that work in a one-party state are partial by structural necessity. The VP3 binding is what makes the system the system; relaxing it would change the system's identity. The available reforms are narrow channel openings — economic, technical, anti-corruption — that admit rung-3 evidence in specific domains without admitting a competing power centre. The structural limits of this strategy are themselves a finding: any failure path that depends on rung-3 political evidence reaching the apex (early-warning, dissent, structural critique) is not addressed by the channel openings, because admitting it would dissolve VP3.

21.6 Platform Mapping

This system fills all ten universal functional slots identified in the [Ten Social Systems Compared](#):

Functional Slot	How This System Fills It
Authority & Decision-Making	Party standing committee / Politburo; general secretary; parliament exists but cannot exercise independent legislative function
Membership & Belonging	Compulsory territorial citizenship; party membership is the pathway to political power and state resources
Resource Allocation	Central planning or party-directed state capitalism; no independent budget process; allocation serves party priorities
Norm Setting & Enforcement	Party line; security apparatus; censorship and information control; surveillance of population
Dispute Resolution	Party-controlled courts; no independent judiciary; internal party discipline for cadres
Legitimation	Ideological narrative (socialism, nationalism, historical necessity); economic performance where available
Succession & Continuity	Internal cadre promotion; party congress ratification; consolidation of power by dominant faction
External Representation	General secretary / paramount leader; state diplomatic corps; party-to-party international relations
Socialisation	State education system; propaganda apparatus; party youth organisations; patriotic education campaigns
Activity Delivery	State enterprises; party-directed civil service; public services as instrument of social control

Navigate to the [interactive visualization](#) for the full graph and table.

22. Example: The School (Secondary Education)

The school is the social system the introduction to this book used as its second motivating example. It is also the system most readers have spent more time inside than any other. Its SE decomposition exposes the canonical proxy-metric pathology in its purest form: a system whose declared goal (developing students' capacity to think, reason, create, and contribute) is structurally close to its measurable proxies (exam scores, league-table ranks) without being the same thing.



school — Graph

LEVEL	ID	DESCRIPTION	↑	↓
⊙ Goals	G1	Develop students' capacity to think, reason, create, contribute	—	R1 R2 R3
	G2	Transmit shared knowledge and civic capacity to the next generation	—	R1 R2
	G3	Select and signal student capability to subsequent stages	—	R3 R5
	G4	Provide safe structured care for school-age children	—	R4 R7
	G5	Equalise opportunity across socioeconomic background	—	R5 R6
⬢ Requirements	R1	Qualified teaching staff with pedagogical training	G1 G2	F1
	R2	Defined curriculum aligned to developmental milestones	G1 G2	F1
	R3	Assessment instruments measuring R2 attainment	G1 G3	F2 F7
	R4	Pastoral care structures for safeguarding and welfare	G4	F3 F4
	R5	Statutory and inspectorate framework	G3 G5	F6
	R6	Financial sustainability — state funding and/or fees	G5	F5
	R7	Facilities appropriate to the curriculum	G4	F5
	R8	Interface with parents as co-responsible adults	—	F6
◇ Functions	F1	Curriculum delivery — lessons, projects, practice	R1 R2	L1
	F2	Assessment — measure individual pupil attainment	R3	L2
	F3	Pastoral care — safeguard welfare, mental and physical health	R4	L3
	F4	Behavioural management — norms, disciplinary cases	R4	L3
	F5	Resource management — staffing, budget, timetable	R6 R7	L4 L5
	F6	External liaison — parents, inspectorate, ministry	R5 R8	L6
	F7	Selection signalling — produce qualifications	R3	L2
	F8	Co-curricular activity — sport, arts, civic engagement	—	L1 L3
○ Logical Architecture	L1	Pedagogical — curriculum, teaching teams, materials	F1 F8	P1 P5
	L2	Assessment — internal tests, national exams, transcripts	F2 F7	P6
	L3	Pastoral — form tutors, year heads, safeguarding lead	F3 F4 F8	P1
	L4	Management — head, SLT, department heads, board	F5	P2 P3
	L5	Administration — registrar, timetable, finance, facilities	F5	P4 P7
	L6	External liaison — parents, inspectorate, university advisors	F6	P8 P10
■ Physical Implementation	P1	Teachers — subject specialists, tutors, support staff	L1 L3	—
	P2	Headteacher and senior leadership team	L4	—
	P3	Governance board / Schulleitung	L4	—
	P4	Buildings — classrooms, labs, library, sports	L5	—
	P5	Curriculum documents, textbooks, digital resources	L1	—
	P6	Assessment instruments — GCSE, Abitur, baccalauréat	L2	—
	P7	Information systems — SIMS, ASV, MIS	L5	—
	P8	Inspectorate — Ofsted, Schulaufsicht	L6	—
	P9	Pupils as registered controlled members	—	—
	P10	Parents / guardians as co-responsible adults	L6	—

school — Table

22.1 SE Decomposition

The full five-level decomposition is shown in the interactive visualization linked above. The canonical reference is at knowledge/system-catalogues/social-systems/school/se-decomposition.md.

22.2 Variation Point Bindings

Variation Point	Binding
VP1: Source of Authority	Statutory mandate (state schools) or contract (independent schools); pedagogical authority of qualified teachers
VP2: Membership Boundary	Compulsory for school-age children within catchment; voluntary for staff under contract
VP3: Decision-Making	Hierarchical: ministry/board → headteacher → department head → teacher; pedagogical autonomy at the classroom level
VP4: Succession	Headteacher appointment by board; teacher promotion by tenure/qualification; pupil progression by examination

The key constraint: VP3 vests strategic decision-making (curriculum, assessment, resource allocation) outside the school in the ministry, board, or chain of inspectorates — but the people with the most accurate operational process model of what is actually happening to a given pupil are the teachers, whose voice in the decision-making structure is structurally limited.

22.3 The Structural Pathology of Proxy-Metric Substitution

The school's declared goal (G1 — develop students' capacity for understanding, creation, and contribution) is unmeasurable in any operationally useful timescale: whether a 14-year-old will turn into a flourishing 30-year-old is the actual question, and no school can wait sixteen years to find out whether last year's pedagogy worked.

The system therefore measures proxies — exam scores, attendance, attainment-eight, league-table rank. The proxies were chosen because they were measurable, not because they were the goal. The moment school funding, teacher career progression, and university admission are bound to the proxies, the entire system reorients toward producing them. This is **Pattern 3 — the Proxy Metric** from the cross-system catalogue, in its archetypal form.

A second pathology stacks on top: **selection by stakes**. When a single high-stakes exam controls the next stage (university admission, vocational entry), the system optimises around it more aggressively than the proxy would otherwise warrant. The proxy that was already only correlated with the goal becomes the goal's effective replacement, and a substantial fraction of the curriculum is rebuilt around exam technique rather than the discipline the exam was supposed to assess.

These are not failures of bad teachers or cynical pupils. They are structural consequences of the measurement choice combined with the binding of consequences to the measurement.

22.4 STPA Highlights

UCA-S1: Teacher evaluation is tied to class exam averages.

Causal factor: exam scores are the school's most legible, comparable, and externally accepted output. Teacher evaluation that has to justify its judgement to external auditors gravitates to the same numbers, because alternative judgements (pedagogical quality, individual pupil development) are harder to defend. The teacher rationally responds by teaching to the test — not because the teacher is cynical, but because the controller's process model classifies that response as success.

UCA-S2: School ranking uses only test results, with no parallel signal.

Causal factor: league tables are produced from a single channel (national assessment data). A school whose pedagogical quality is excellent on dimensions the league table does not measure (creativity, civic capacity, longitudinal pupil welfare) appears worse than a school that has reorganised its curriculum around the measured channel. Resource flow follows the league-table signal.

UCA-S3: A single high-stakes exam controls university entry.

Causal factor: selection signals at the end of secondary school carry the entire weight of the next-stage decision. The exam was originally a check on prior school assessment; it has become the decision, and the rest of the system has reorganised around it.

UCA-S4: Pupils with developmental, neurodivergent, or family stressors are graded on the same standardised channel as their unstressed peers.

Causal factor: the rung-3 measurement channel (a standardised exam) does not natively distinguish "did not learn" from "was unable to demonstrate learning under exam conditions". The structural fix is parallel assessment, not better exams.

22.5 Fixes That Worked

Finland's post-1980s reform — addressing UCA-S1 and UCA-S2.

Finland removed national standardised testing before age 16, removed published school league tables, and shifted teacher evaluation from output metrics to professional peer review with substantial pedagogical autonomy. Teacher selection became highly competitive (master's degree, paid pedagogical training, sustained professional development). Evidence: PISA performance over 2000–2018 sustained at or near the top of OECD countries despite — or because of — the removal of the high-stakes measurement channel. The reform was not "teach better"; it was a structural change to which channels carried consequences.

Estonia's national digital infrastructure for assessment — addressing UCA-S2.

Estonia maintained national assessments but disaggregated the channels: longitudinal pupil progress, school value-added, teacher self-evaluation, and parental feedback are reported in parallel rather than collapsed into a single league rank. Evidence: PISA performance has converged on Finland's; teacher status and retention have stabilised.

Germany's Abitur — partial, mixed-evidence fix for UCA-S3.

The German Abitur is partly continuous (the last two years of school grades count substantially) and partly examination (the Abitur exams themselves). The structural intent is to distribute the selection signal across multiple independent measures. Evidence: the system has produced stable graduate quality and an unusually broad vocational pathway alongside the academic one; the proxy pathology still appears at the level of the gymnasium admission decision in primary school, which now carries the load that the Abitur was designed to share.

Pre-registered standardised testing with portfolio supplements — emergent fix for UCA-S4.

Several jurisdictions now require both a standardised exam and a portfolio of teacher-assessed work for university entry. Evidence (early): the parallel channels reduce the selection error against pupils whose abilities are systematically under-measured by exam-format assessment, including many neurodivergent and ESL pupils.

**The Design Principle**

The school's problem is not that it lacks dedication. The Kodak pattern from the introduction applies directly: the people inside the system respond to the operative goals encoded in the measurement and reward structure, and the measurement and reward structure does not connect to the declared goal. The fixes that work are not exhortations to teachers or pupils. They are changes to what is measured, how it feeds back, and what it controls — Finland removed the high-stakes channel, Estonia disaggregated it, Germany distributed it. Each is a structural intervention against a specific causal pathway.

22.6 Platform Mapping

Functional Slot	How This System Fills It
Authority & Decision-Making	Ministry / school board → headteacher → department head → teacher; pedagogical autonomy at classroom level
Membership & Belonging	Compulsory attendance for school-age children in catchment; staff employment contracts; alumni continuity
Resource Allocation	State funding by pupil-count formula; discretionary headteacher budget; parental contributions for non-statutory activities
Norm Setting & Enforcement	School policies; pastoral system; behaviour code; pupil disciplinary procedures; teacher conduct framework
Dispute Resolution	School complaints procedure; pupil disciplinary appeals; staff grievance procedures; external inspectorate as last resort
Legitimation	Statutory mandate; pedagogical qualification; outcomes (degree, employment, civic capacity) — to the extent they can be measured
Succession & Continuity	Headteacher appointment by board; teacher promotion through pedagogical career structure; pupil progression through year groups
External Representation	Headteacher as institutional voice; parent-teacher associations; alumni relations; inspectorate as external observer
Socialisation	Curriculum delivery; school culture; pastoral relationships; peer learning; co-curricular programmes
Activity Delivery	Teaching; learning support; pastoral care; assessment; co-curricular programmes; school meals and welfare

For the full SE decomposition: [interactive visualization](#). For the comparative analysis across all systems: [Ten Social Systems Compared](#). For the cross-system control-structure analysis: [Control Structures](#). The introduction's worked example of the school's structural pathology is at [Introduction §A Closer Example](#).

23. Example: The Hospital (Acute Healthcare)

The hospital is the worked example that bridges the social and technical systems books — the only system in this catalogue with a triadic membership boundary (patient, professional, payer) and the only one in which the cost of a control failure is immediate, physical, and individually attributable. Its SE decomposition exposes a recurrent professional-institution pathology: the people with the most accurate process model of what is happening to a given patient (the bedside clinicians) are not the people whose decisions allocate the resources that determine what they can do.



hospital — Graph

LEVEL	ID	DESCRIPTION	↑	↓
⊙ Goals	G1	Restore, preserve, or extend patients' health	—	R1 R2
	G2	Relieve suffering where cure is impossible	—	R1 R5
	G3	Maintain professional standards and the medical knowledge base	—	R1 R6
	G4	Operate financially sustainably under payer constraints	—	R4 R8
	G5	Train the next generation of clinical staff	—	R1 R6
◆ Requirements	R1	Qualified, licensed clinical workforce	G1 G2 G3 G5	F1 F2 F3
	R2	Diagnostic and therapeutic capability	G1	F2 F3
	R3	Triage and demand management for unscheduled care	—	F1
	R4	Documentation, coding, billing infrastructure	G4	F6
	R5	Safety, infection control, incident reporting	G2	F4 F7
	R6	Governance separating clinical from operational authority	G3 G5	F7 F9
	R7	Interfaces with primary care, social care, post-acute	—	F5
	R8	Capital and operating budget sufficient for R1–R7	G4	F9
◇ Functions	F1	Triage and assessment — receive and route patients	R1 R3	L1
	F2	Diagnosis — investigate, interpret, plan	R1 R2	L1 L2
	F3	Treatment delivery — medical, surgical, nursing	R1 R2	L1 L3
	F4	Monitoring and escalation — observe, react to deterioration	R5	L1
	F5	Discharge and transition of care	R7	L4 L7
	F6	Documentation and billing	R4	L6
	F7	Clinical governance — guidelines, audit, M&M	R5 R6	L5
	F8	Training and supervision	—	L1 L5
	F9	Operational management — staffing, beds, supplies	R6 R8	L4 L6
○ Logical Architecture	L1	Clinical — wards, theatres, ICU, ED, outpatients	F1 F2 F3 F4 F8	P1 P2 P4 P5
	L2	Diagnostic — laboratory, radiology, pathology	F2	P3 P5
	L3	Pharmacy — medicines management, dispensing	F3	P3
	L4	Patient flow — admissions, beds, discharge	F5 F9	P6
	L5	Governance — medical director, safety committee	F7 F8	P1 P8
	L6	Operational — CEO, finance, HR, IT, facilities	F6 F9	P7 P9
	L7	External interface — primary care, social care, payer, inspectorate	F5	P8 P9
■ Physical Implementation	P1	Doctors — consultants, registrars, juniors	L1 L5	—
	P2	Nurses — ward, ICU, theatre, advanced practice	L1	—
	P3	Allied health — physio, pharmacy, technicians	L2 L3	—
	P4	Hospital building — wards, ICU, theatres, ED	L1	—
	P5	Medical devices — imaging, monitors, ventilators	L1 L2	—
	P6	EHR — Epic, Cerner, SAP IS-H	L4	—
	P7	Billing systems — ICD-10, DRG, OPS	L6	—
	P8	Inspectorate / regulator — CQC, Joint Commission, IQTIG	L5 L7	—
	P9	Payer — insurer, Krankenkasse, single-payer body	L6 L7	—
	P10	Patients as the controlled process	—	—

hospital — Table

23.1 SE Decomposition

The full five-level decomposition is shown in the interactive visualization linked above. The canonical reference is at [knowledge/system-catalogues/social-systems/hospital/se-decomposition.md](https://knowledge.system-catalogues/social-systems/hospital/se-decomposition.md).

23.2 Variation Point Bindings

Variation Point	Binding
VP1: Source of Authority	Professional qualification + statutory licence; managerial authority below the clinical decision
VP2: Membership Boundary	Triadic: patient (service recipient), clinician (qualified professional), payer (insurer/state)
VP3: Decision-Making	Clinical decisions belong to the qualified clinician; resource and operational decisions belong to management; the relationship between the two is contested
VP4: Succession	Clinical posts by qualification + interview; management posts by board appointment; medical director as the bridge role

The key constraint: the hospital is the only social system where VP2 (membership) cannot be reduced to a dyadic relationship without losing the system. A purely patient-clinician relationship is medicine; a patient-payer relationship is insurance; a clinician-payer relationship is a contract. The hospital is the institutional form that holds all three together. The pathologies appear where the triangle distorts.

23.3 The Structural Pathology of Triadic Membership

When the payer holds the dominant signalling channel, **operations are optimised for billing categories** — what gets reimbursed gets done, regardless of whether it is what the patient needed. The US healthcare system's well-documented over-procedure rate (unnecessary stents, unnecessary scans, unnecessary admissions) is the direct consequence: the rung-3 medical evidence about marginal benefit is filtered through a rung-1 billing structure that pays per procedure rather than per outcome.

When the clinician holds the dominant signalling channel without a payer constraint, **costs and access diverge**: technically excellent care for the patients who reach it, and a queue of patients who do not because the system has no mechanism to allocate scarce capacity. Single-payer systems with weak demand management (NHS in some periods, Canadian provinces in others) show this pattern.

When the manager holds the dominant signalling channel, **proxy metrics replace clinical outcomes**. The Mid-Staffordshire Hospital scandal (UK, 2005–2009) is the canonical worked example. The trust optimised for the metrics the regulator measured (Foundation Trust status, financial targets, four-hour A&E waits), and the rung-3 clinical evidence about what was happening on the wards — falls, dehydration, neglect, hundreds of preventable deaths — was structurally invisible at the level the regulator could see. The clinicians knew. The whistleblowers tried. The reporting channel went through the manager.

These are not failures of bad people. They are structural consequences of the triadic VP2 binding when one corner of the triangle dominates the others.

23.4 STPA Highlights

UCA-H1: A bedside clinician's observation that a deteriorating patient needs urgent escalation does not reach the decision-maker in time.

Causal factor: the upward escalation path runs through the hierarchy that is structurally biased to deny that the patient is deteriorating (because acknowledging it costs resources, costs reputation, and risks litigation). The structural fix is parallel escalation paths — early-warning scores, rapid-response teams — that bypass the hierarchy for specific clinical states.

UCA-H2: A patient is admitted, treated, and discharged on a clinical pathway optimised for the billing code rather than the clinical reality.

Causal factor: the payer's process model is populated by billing codes; the codes were never designed to be a complete description of the patient. The hospital's operations gravitate to the description the payer can see. Where the description omits relevant clinical reality (complex comorbidity, social factors, palliative trajectory), the operations omit the matching activities.

UCA-H3: A manager-imposed performance target is met by reorganising clinical work in ways that worsen outcomes the target does not measure.

Causal factor: the Mid-Staffordshire pattern: four-hour A&E targets met by relocating patients to other parts of the hospital; financial targets met by reducing nursing staffing. The metric is met; the goal it stood in for is not. This is **Pattern 3 — the Proxy Metric** from the cross-system catalogue, in a system where the proxy substitution kills people.

UCA-H4: An adverse event is reported through the chain it would indict.

Causal factor: the canonical hospital incident-reporting system reports through the same management chain whose performance would be reflected in the incident rate. The structural fix is independent investigation — a coroner who is not employed by the hospital, an inspectorate that does not depend on the hospital's cooperation, a patient ombudsman whose budget is not part of the hospital's.

23.5 Fixes That Worked

Independent inspectorate with statutory access — addressing UCA-H3.

The UK's Care Quality Commission (CQC) was reorganised post-Francis Report (2013) with substantially strengthened powers of unannounced inspection, direct staff and patient interview, and ratings publication. Evidence: CQC inspections have since identified and rated as Inadequate trusts that internal indicators had not surfaced; ratings publication created a public accountability channel the hospital management could not filter. The fix did not remove the proxy metrics — it installed a parallel rung-3 channel (independent inspection) that the proxy-metric channel could not capture.

Bundled payment / value-based purchasing — addressing UCA-H2.

The US Centers for Medicare & Medicaid Services (CMS) bundled-payment programmes (Bundled Payments for Care Improvement, 2013 onward) and value-based purchasing (Hospital Value-Based Purchasing Program, 2010 onward) bind reimbursement to outcomes over an episode of care rather than per-procedure billing codes. Evidence: documented reductions in 30-day readmission rates for bundled DRGs (joint replacement, cardiac); modest cost reductions without quality loss. The structural change: the payer's process model is partially reoriented from billing codes to outcomes.

Early Warning Scores and rapid-response teams — addressing UCA-H1.

The National Early Warning Score (NEWS, NICE 2017) and the rapid-response team model (originally Pittsburgh, late 1990s) install a parallel escalation path keyed to standardised physiological thresholds that bypasses the normal hierarchy. Evidence: meta-analyses report consistent reductions in in-hospital cardiac arrest and unplanned ICU admission where the intervention is fully implemented. The fix is structural: a rung-3 channel (physiological data) and a rung-3 actuator (rapid-response team) operate independently of the rung-1 hierarchy.

Statutory mandatory reporting of adverse events to coroner / patient-safety body — addressing UCA-H4.

UK serious incident reporting to the NHS Patient Safety Incident Response Framework (2022), coroners' inquests with hospital staff as compelled witnesses, and (in selected jurisdictions) statutory duty of candour for clinicians. Evidence: the duty-of-candour move (UK, 2014) shifted significant patient-safety-incident reporting volumes upward; the structural intent — independent rung-3 channel for adverse-event evidence — is in place even where the cultural implementation is still in progress.



The Design Principle

The hospital cannot abolish its triadic structure — patient, clinician, payer is the system. The fixes that work install parallel channels that prevent any one corner of the triangle from monopolising the signalling: an inspectorate the hospital does not employ, an early-warning score the hierarchy does not filter, a payment mechanism that does not collapse the clinical reality to a billing code. Each is a structural intervention against a specific causal pathway the STPA analysis identifies.

23.6 Platform Mapping

Functional Slot	How This System Fills It
Authority & Decision-Making	Medical director / chief of staff (clinical); chief executive / managing director (operational); board (strategic); split responsibility is structurally required
Membership & Belonging	Triadic: patients (episode-bounded), clinical staff (qualification + employment), payer relationship (insurer or state contract)
Resource Allocation	Tariff/contract revenue + state subsidy; budget set by management with clinical input; capacity allocation by clinical demand
Norm Setting & Enforcement	Professional codes (GMC, NMC, equivalents); regulatory standards (CQC, JCI); internal clinical guidelines
Dispute Resolution	Complaints procedure; clinical incident investigation; coroner / medical examiner; litigation as last resort
Legitimation	Professional qualification; statutory licence; demonstrable outcomes; patient experience
Succession & Continuity	Clinical posts by qualification + competitive appointment; management posts by board; institutional continuity through training pipelines
External Representation	Chief executive; medical director; communications; supplier and partner-institution agreements
Socialisation	Clinical training and supervision; ward-level mentoring; professional development; staff induction
Activity Delivery	Diagnosis, treatment, surgery, nursing, rehabilitation, palliative care, emergency response, outpatient services

For the full SE decomposition: [interactive visualization](#). For the comparative analysis across all systems: [Ten Social Systems Compared](#). For the cross-system control-structure analysis: [Control Structures](#).

24. Example: Wikipedia / Open-Source Project

Wikipedia and the canonical open-source project (Linux kernel, Debian, Python) are the test case for whether the book's framework generalises to systems with **no formal authority**. Their SE decomposition reveals that the framework does generalise — but only because their apparent absence of authority is actually a substitution: rung-3 evidence (cited sources, working code, test results) replaces rung-1 institutional authority at the centre of the control loop. When that substitution holds, the system is stable; when it slips, the same failure modes appear that the book documents for institutionally authorised systems.



wikipedia — Graph

LEVEL	ID	DESCRIPTION	↑	↓
🎯 Goals	G1	Produce and maintain a freely accessible body of work	—	R1 R2 R6
	G2	Ensure the work meets a contestable quality standard	—	R2 R3
	G3	Sustain a community of contributors	—	R4 R5
	G4	Remain free in the constitutive sense	—	R7
⬢ Requirements	R1	Open submission mechanism for new contributions	G1	F1
	R2	Review process that accepts, rejects, or amends	G1 G2	F2 F3
	R3	Quality standard evaluable against rung-3 evidence	G2	F2
	R4	Trust hierarchy admitting experienced contributors to rights	G3	F5
	R5	Dispute-resolution mechanism for contested submissions	G3	F4
	R6	Infrastructure — servers, version control, communications	G1	F7
	R7	Legal framework — license, foundation, governance	G4	F8
	R8	Backstop authority for cases consensus cannot resolve	—	F4 F8
◇ Functions	F1	Contribution intake — accept edits / patches	R1	L1
	F2	Review — evaluate against the rung-3 standard	R2 R3	L2
	F3	Merge / publish — commit to the canonical work	R2	L1 L2
	F4	Dispute resolution — handle contested submissions	R5 R8	L4
	F5	Trust management — grant and revoke expanded rights	R4	L3
	F6	Community moderation — conduct, harassment, bad faith	—	L5
	F7	Infrastructure operation — servers, repos, build systems	R6	L6
	F8	Governance — strategy, foundation, escalation	R7 R8	L7
○ Logical Architecture	L1	Contribution — edit interface / pull request / patch	F1 F3	P1 P4
	L2	Review — reviewers, tooling, review norms	F2 F3	P1 P5
	L3	Trust — rights ladder reader → editor → admin / maintainer	F5	P1 P3
	L4	Dispute resolution — talk pages, ArbCom / Steering Council	F4	P3 P6
	L5	Community — code of conduct, moderation, onboarding	F6	P6
	L6	Infrastructure — servers, repositories, mailing lists	F7	P4
	L7	Governance — foundation, BDFL/Steering Council	F8	P2 P3 P7
▣ Physical Implementation	P1	Contributors — anonymous to maintainer	L1 L2 L3	—
	P2	Foundation — Wikimedia, Linux Foundation, PSF	L7	—
	P3	BDFL / Steering Council / ArbCom	L3 L4 L7	—
	P4	Servers, MediaWiki / Git, mailing lists	L1 L6	—
	P5	Sources / test suites / reproducibility tooling	L2	—
	P6	Governance documents — guidelines, conduct code, PEPs	L4 L5	—
	P7	License — CC BY-SA, GPL, MIT, BSD, Apache	L7	—
	P8	External users — readers, downstream developers	—	—

wikipedia — Table

24.1 SE Decomposition

The full five-level decomposition is shown in the interactive visualization linked above. The canonical reference is at knowledge/system-catalogues/social-systems/wikipedia/se-decomposition.md.

24.2 Variation Point Bindings

Variation Point	Binding
VP1: Source of Authority	Tenure + demonstrable contribution — established editors / committers carry weight in proportion to their visible track record
VP2: Membership Boundary	Self-selecting and effectively open; participation gradient from anonymous reader to administrator/maintainer
VP3: Decision-Making	"Rough consensus", supplemented by a "Benevolent Dictator For Life" (BDFL) or arbitration committee for unbreakable deadlocks
VP4: Succession	Emergent: contributors who sustain quality over time accumulate trust and rights; formal roles ratify what has already happened informally

The key structural feature: authority is granted by demonstrated contribution measurable against rung-3 standards (a Wikipedia edit cites sources; a Linux patch must compile and pass tests). The system's distinctiveness in the catalogue is that the rung-3 standard exists before any institutional authority is needed to enforce it — the source citation, the unit test, the compilation result, the reproducibility check.

24.3 The Structural Pathology of Rough Consensus

The system's strength is that VP1 cannot be captured by patronage, inheritance, or capital — because the only currency that buys influence is contribution that survives external evaluation. The matching weaknesses are three.

The tyranny of structurelessness (Jo Freeman's classic critique of the early women's-liberation movement, applied here): when no formal authority exists, informal authority concentrates around individuals whose social-network position, time commitment, and emotional stamina exceeds others'. The early Wikipedia editor with 60,000 edits has accumulated, by tenure alone, an effective veto on consensus calls. The system's claim to operate at rung 6 (open deliberative consensus) is structurally compromised by a rung-1 informal-tenure stratification that consensus calls cannot disentangle from substantive argument.

Edit wars and the dominance of stamina. Where rough consensus fails to converge — typically on contested political, religious, or biographical articles — the editor who can outlast the others controls the outcome. Stamina is the rung at which the contest is decided; it has no truth-tracking property. The arbitration committee was created to resolve exactly these cases, but the cases that reach it are the ones where the underlying disagreement is not actually about facts.

Hostile entry by coordinated minorities. Open membership combined with consensus-without-counting decision-making is exploitable by coordinated factions whose individual contributions look organic. The pattern is documented across Wikipedia (state-sponsored editing campaigns, paid PR editing), open-source projects (governance takeover attempts), and the early Bitcoin community (governance fork dynamics). The defence — community moderation by experienced editors — collides with the informal-authority pathology above.

These are not failures of community spirit. They are structural consequences of the VP1 binding when the rung-3 standard (verifiable evidence) is contested or absent.

24.4 STPA Highlights

UCA-W1: Long-tenured editors dominate consensus calls regardless of the substantive merit of contested edits.

Causal factor: the system has no formal mechanism to weight contributions by current relevance vs accumulated tenure. The rung-1 deference to "established editor" interacts with the rough-consensus mechanism such that the consensus call effectively counts the established editor's opinion above newcomers' regardless of the rung-3 evidence either presents.

UCA-W2: A contested article is controlled by the editor with the most stamina, not the editor with the strongest sources.

Causal factor: the rough-consensus mechanism resolves contested edits by edit-war attrition; the attrition is governed by editor time and stamina, which have no truth-tracking property. The matching failure mode in open-source: maintainer burnout, where the project is captured by whoever is willing to sustain the unpaid review workload.

UCA-W3: A coordinated minority of editors with hidden coordination wins a consensus call against an uncoordinated majority.

Causal factor: the system's mechanisms cannot reliably distinguish independent editors with similar views from coordinated editors presenting similar views. Where coordination is real and hidden, the consensus call is corrupted; where it is suspected but not proven, the suspicion itself corrupts the call.

UCA-W4: The "Benevolent Dictator" resigns or burns out, and no succession mechanism is structurally ready to substitute.

Causal factor: VP4 (succession) is emergent rather than statutory. When the BDFL exits suddenly (Python's Guido van Rossum, 2018), the system enters a structurally ambiguous interregnum until a new authority is recognised. The interregnum is a failure window in which contested decisions cannot be settled.

24.5 Fixes That Worked

Wikipedia's Arbitration Committee — partial fix for UCA-W2 and UCA-W3.

The Wikipedia ArbCom is an elected body of senior editors empowered to decide cases the consensus mechanism cannot. Evidence: the committee has consistently handled the highest-conflict cases (BLP disputes, coordinated-editing investigations, sustained editor conduct) at a rate of dozens per year over two decades. The structural change: a rung-2/6 formal-rule mechanism is installed as a backstop where the rung-6 consensus mechanism fails.

Mandatory disclosure of paid editing — addressing UCA-W3.

Wikipedia's 2014 Terms of Use amendment requires disclosure of paid editing affiliation. Evidence: the disclosure regime is imperfectly enforced but has shifted significant volumes of paid editing from undisclosed to disclosed, where the community can evaluate it explicitly; high-visibility undisclosed-paid-editing cases now reliably trigger blocks.

Python PEP 8016 governance transition (2018) — addressing UCA-W4.

After Guido van Rossum's resignation as BDFL, the Python community executed a structured succession: PEP 8016 established an elected Steering Council, codifying the previously informal authority structure into a statutory one with a fixed term and rotation. Evidence: the transition was completed without forking the language; subsequent governance has surfaced and resolved several controversial decisions (typing syntax, packaging, performance) through the new mechanism. The structural change: VP4 was moved from emergent to statutory before the next succession crisis hit.

Linux kernel maintainer hierarchy — emergent structural defence against UCA-W1 and UCA-W2.

The kernel project distributes authority across hundreds of subsystem maintainers, each with rung-3 demonstrable competence in a specific component. Linus Torvalds's authority is real but exercised only at the top of a deep delegation tree; no single contested area falls entirely under any one person's exclusive control. Evidence: the kernel has scaled to 30+ years and millions of lines without governance crisis, despite many individual flashpoints. The structural feature: VP1 (authority by demonstrated contribution) is finely distributed by domain, so the failure modes that appear when authority concentrates in one place cannot accumulate.



The Design Principle

The wiki / open-source pattern shows that rung-3 evidence (cited sources, working code, reproducible tests) can substitute for rung-1 institutional authority as long as the substitution holds. Where the rung-3 standard is contested, absent, or relegated to stamina contests, the same failure modes that the book documents for institutionally authorised systems reappear — and the same fix applies: install a parallel channel (ArbCom, statutory governance, disclosure regime) that catches the cases the rung-3 mechanism cannot. The conclusion is not that "Wikipedia is different"; it is that every social system in the catalogue is held together by the rung at which its claims are actually evaluated, and Wikipedia just made the rung explicit.

24.6 Platform Mapping

Functional Slot	How This System Fills It
Authority & Decision-Making	Rough consensus + BDFL / Steering Council / ArbCom backstop; subsystem maintainers; experienced contributors as informal authorities
Membership & Belonging	Self-selecting and effectively open; participation gradient from reader/lurker to admin/maintainer; identity by stable handle
Resource Allocation	Volunteer time (primary); foundation budget (Wikimedia, Linux Foundation, PSF) for infrastructure and a thin paid-staff layer
Norm Setting & Enforcement	Verifiability/notability rules; code review and commit policy; code of conduct; community moderation; ArbCom
Dispute Resolution	Talk-page discussion; informal mediation; ArbCom / Steering Council; in extremis, project fork
Legitimation	Demonstrable contribution surviving rung-3 evaluation; cumulative reputation; project usefulness to external users
Succession & Continuity	Emergent rights accretion (Wikipedia "rollback", "admin"); statutory transitions (Python PEP 8016); maintainer delegation chains
External Representation	Foundation (Wikimedia, Linux Foundation, PSF) as legal entity; project spokespersons; press contacts
Socialisation	Onboarding via documentation; mentorship of new contributors; project-specific norms (e.g. "be bold", "rough consensus")
Activity Delivery	Article writing and curation; code authoring and review; bug triage; release management; community support

For the full SE decomposition: [interactive visualization](#). For the comparative analysis across all systems: [Ten Social Systems Compared](#). For the cross-system control-structure analysis: [Control Structures](#).

25. Part III — Development Frameworks

This part applies the same product line analysis to development and project management frameworks, revealing that the "methodology wars" are really about different binding decisions at four well-defined variation points.

25.1 The Frameworks

Framework	Orientation
Waterfall	Plan-driven, sequential
V-Model	Plan-driven, verification-focused
PRINCE2	Governance-focused PM
Scrum	Adaptive, iterative
Kanban	Adaptive, flow-based
Design Thinking	Human-centered, exploratory
DevOps	Continuous delivery
SAFe	Scaled adaptive

25.2 Chapters

- **Eight Frameworks Decomposed** — Full SE hierarchy for each framework
- **The Framework Platform** — Ten universal functional slots
- **Hybrid Architectures** — Three principled merger proposals

26. Product Line Thinking for Development Frameworks

26.1 How Waterfall, Scrum, Kanban, V-Model, PRINCE2, Design Thinking, DevOps, and SAFe Share a Common Platform — and How to Merge Them Better

26.2 Abstract

The "methodology wars" of the past three decades — Waterfall vs. Agile, Scrum vs. Kanban, plan-driven vs. adaptive — have generated more heat than light. This chapter applies the same systems engineering decomposition and product line analysis previously used for social institutions to the domain of development and project management frameworks. By decomposing eight frameworks through a five-level SE hierarchy (Goals → Requirements → Functions → Logical Architecture → Physical Implementation), we discover that these seemingly opposed approaches share a remarkably large common platform and differ primarily at four well-defined variation points. We use this insight to propose three concrete hybrid architectures that merge the best of plan-driven and adaptive approaches — not as vague "be pragmatic" advice, but as structurally principled product-line configurations with explicit binding decisions.

Canonical reference

The normative versions of the eight framework decompositions, the cross-framework platform, the variation points, the hybrid architectures, and the merging principles are maintained as the project's knowledge base under [knowledge/system-catalogues/dev-frameworks/](#). This chapter is the narrative overview; the knowledge base is the reference.

26.3 1. Introduction

1.1 The False Dichotomy

Ask any software team whether they "do Agile or Waterfall" and you'll get a strong opinion. Ask them what exactly they mean by either term and the clarity evaporates. In practice, almost no team does pure Waterfall (sequential phases with no feedback) or pure Scrum (no upfront planning, no documentation, no governance). Most teams operate some hybrid — often an accidental, incoherent one — because the real work demands elements of both.

The problem is not that people lack pragmatism. The problem is that they lack a structural vocabulary for describing which elements of which framework can be combined, and which combinations are architecturally incoherent. Product line thinking provides that vocabulary.

1.2 Scope: Eight Frameworks

We analyze eight widely-used frameworks, chosen to span the full spectrum from plan-driven to adaptive, and from project-focused to flow-focused:

Framework	Orientation	Core Idea
Waterfall	Plan-driven, sequential	Complete each phase before the next begins
V-Model	Plan-driven, verification-focused	Each design phase has a corresponding test phase
PRINCE2	Governance-focused PM	Projects in controlled environments with staged authorization
Scrum	Adaptive, iterative	Self-organizing teams deliver value in fixed-length sprints
Kanban	Adaptive, flow-based	Visualize work, limit WIP, optimize flow — no prescribed iterations
Design Thinking	Human-centered, exploratory	Empathize → Define → Ideate → Prototype → Test
DevOps	Continuous delivery	Merge development and operations into an automated pipeline
SAFe	Scaled adaptive	Agile practices at portfolio, program, and team levels

26.4 2. The Decompositions

Each framework is decomposed through the five-level SE hierarchy with full traceability. Goals name what the framework exists to achieve; requirements translate those goals into must-statements; functions name the activities; logical architecture groups the functions into subsystems; physical implementation names the concrete artefacts, roles, and practices that realise each subsystem.

The full tables are maintained in the knowledge base:

Framework	Decomposition file
Waterfall	knowledge/system-catalogues/dev-frameworks/waterfall/se-decomposition.md
V-Model	knowledge/system-catalogues/dev-frameworks/v-model/se-decomposition.md
PRINCE2	knowledge/system-catalogues/dev-frameworks/prince2/se-decomposition.md
Scrum	knowledge/system-catalogues/dev-frameworks/scrum/se-decomposition.md
Kanban	knowledge/system-catalogues/dev-frameworks/kanban/se-decomposition.md
Design Thinking	knowledge/system-catalogues/dev-frameworks/design-thinking/se-decomposition.md
DevOps	knowledge/system-catalogues/dev-frameworks/devops/se-decomposition.md
SAFe	knowledge/system-catalogues/dev-frameworks/safe/se-decomposition.md

26.5 3. The Development Framework Platform

Once the eight decompositions are aligned level by level, ten universal functional slots emerge that every framework fills in some way. These are **Work Discovery**, **Prioritisation**, **Execution**, **Quality Assurance**, **Integration**, **Delivery**, **Feedback**, **Governance**, **Improvement**, and the framework-specific activity layer that carries the work through them.

The full platform description — what the shared goals, requirements, functions, logical, and physical elements actually are — is in <knowledge/system-catalogues/dev-frameworks/cross-framework/platform.md>.

What the platform means in practice: most framework disputes are not about whether to do a thing, but about where in the lifecycle to do it, how often, and who decides. Those are exactly the questions that variation points make explicit.

26.6 4. The Variation Points

Four primary variation points produce the differences between frameworks:

- **VP1 — Temporal Structure of Work** (phases · sprints · flow · continuous · stages).
- **VP2 — Requirements Stability Assumption** (frozen upfront · evolving throughout · discovered during work).
- **VP3 — Authority & Decision Structure** (hierarchical project manager · self-organising team · product owner · steering committee).
- **VP4 — Quality Assurance Timing** (gated at end · continuous · after each sprint · before each phase transition).

As with social systems, variants are not freely combinable: some binding combinations are coherent, others are structurally contradictory (e.g. "phase-gated temporal structure" with "requirements evolve throughout" is incoherent without a workaround).

The full variation-point catalogue, observed variants, and interaction constraints are in <knowledge/system-catalogues/dev-frameworks/cross-framework/variation-points.md>.

26.7 5. Reuse Analysis Across Frameworks

Elements of one framework can sometimes be adopted by another, sometimes only with adaptation, and sometimes not at all without changing the framework's identity. The full reuse analysis — fully reusable, partially reusable with adaptation, and structurally incompatible — is in <knowledge/system-catalogues/dev-frameworks/cross-framework/reuse-analysis.md>.

26.8 6. Three Proposed Hybrid Architectures

Three concrete hybrid framework configurations are proposed, each with explicit binding decisions at all four variation points:

- **Hybrid A — "Governed Agile"** — for enterprises with regulatory or contractual obligations. Combines PRINCE2-style staged authorisation with Scrum-style execution inside each stage.
- **Hybrid B — "Discovery-Driven Flow"** — for product teams optimising for innovation and speed. Combines Design Thinking discovery with Kanban delivery and DevOps continuous release.
- **Hybrid C — "Scaled Governed Flow"** — for large organisations managing portfolios of diverse work. Combines SAFe portfolio structure with flow-based team execution and phase-gated compliance checkpoints.

Full architectural sketches, binding tables, and trade-offs for all three hybrids are in <knowledge/system-catalogues/dev-frameworks/cross-framework/hybrids.md>.

26.9 7. Principles for Framework Merging

Five principles govern the design of coherent framework hybrids:

- **Compatibility Rule** — binding decisions across variation points must be mutually coherent.
- **Granularity Principle** — variation points bind at the smallest cohesive unit, not the whole framework.
- **Start from the Platform** — build on the shared functional slots first, then add variation.
- **Automate the Right Arm** — whatever part of the lifecycle is repetitive should be automated before it is standardised.
- **Separate Discovery from Delivery** — requirements discovery and execution must not share the same temporal structure.

Full statements of all five principles are in <knowledge/system-catalogues/dev-frameworks/cross-framework/merging-principles.md>.

26.10 8. Conclusion

The methodology wars were never about fundamentally different approaches to building things. They were about different binding decisions at four variation points — temporal structure, requirements stability, authority structure, and QA timing — applied to a shared platform of ten universal functional slots.

Waterfall is not wrong. It is a valid configuration for contexts where requirements are genuinely stable, regulatory traceability is required, and the cost of change is high. Scrum is not wrong. It is a valid configuration for contexts where requirements are volatile, time-to-feedback is critical, and teams are small enough to self-organize. The error — the engineering error, in product-line terms — is applying one configuration to a context that demands the other, or merging elements from both without checking VP compatibility.

Product line thinking offers a way out. By identifying the platform, mapping the variation points, and making binding decisions explicit, teams and organizations can design hybrid frameworks that are principled rather than accidental, coherent rather than contradictory, and tailored rather than dogmatic. The goal is not to find the one true methodology. It is to become fluent in the vocabulary of variation, so that every team can configure the framework that fits their actual context — and explain why.

This analysis applies the product line methodology developed in "Product Line Thinking for Social Systems" to the domain of development and project management frameworks. The SE hierarchy (Goals → Requirements → Functions → Logical Architecture → Physical Implementation) provides the decomposition structure; the product line concepts (platform, variation point, binding, reuse classification) provide the comparative and synthesis methodology.

27. The Framework Platform

27.1 Ten Universal Functional Slots in Development

Every development framework — regardless of philosophy — fills the same functional slots:

Functional Slot	Description
Work Discovery	Determine what to build
Work Prioritization	Decide what first
Work Decomposition	Break large efforts into manageable units
Work Execution	Build the deliverable
Quality Assurance	Verify expectations are met
Integration	Combine outputs into a coherent whole
Delivery	Get the product to users
Feedback Collection	Learn whether the product solves the problem
Process Governance	Authorization, tracking, risk management
Process Improvement	Reflect and evolve the process

27.2 Four Variation Points

The frameworks differ at four points:

1. **VP1: Temporal Structure** — sequential phases · fixed iterations · continuous flow · exploratory loops
2. **VP2: Requirements Stability Assumption** — knowable upfront · emergent · hypotheses
3. **VP3: Authority & Decision Structure** — hierarchical PM · PO + team · shared · automated pipeline
4. **VP4: QA Timing** — end-phase · per-iteration · continuous automated · user validation

The full decomposition and variation point analysis is in [Eight Frameworks Decomposed](#).

28. Hybrid Architectures

28.1 Three Principled Mergers

Hybrid A: "Governed Agile"

For: Regulated industries, contractual delivery (medical devices, automotive, maritime OT, defense)

PRINCE2 stages containing Scrum Sprints. V-Model test levels automated as CI/CD pipeline. Design Thinking discovery upfront.

Hybrid B: "Discovery-Driven Flow"

For: Product teams, SaaS, startups

Design Thinking discovery track running parallel to Kanban delivery track. DevOps CI/CD. No Sprints — continuous pull.

Retrospectives decoupled from delivery cadence.

Hybrid C: "Scaled Governed Flow"

For: Enterprises managing mixed portfolios

Portfolio level on Kanban with PRINCE2 Business Case governance. Program level on SAFe PI cadence. Team level: each team chooses Scrum or Kanban. Shared DevOps platform underneath.

28.2 Five Design Principles

1. **The Compatibility Rule** — check VP binding compatibility before merging
2. **The Granularity Principle** — different VPs can bind at different organizational levels
3. **Start from the Platform** — don't modify one framework; compose from the functional slots
4. **Automate the Right Arm** — V-Model quality structure + DevOps automation = best of both
5. **Separate Discovery from Delivery** — different activities need different rhythms

The full analysis with binding tables is in [Eight Frameworks Decomposed](#).

29. Part IV — STPA Applied

This part applies System-Theoretic Process Analysis to social systems, using the SE decompositions from Part II as input. STPA reveals that the most harmful outcomes of social systems are not random malfunctions but predictable consequences of their control architecture.

29.1 Chapters

- **STPA on Religion** — Full four-step STPA analysis: losses, control structure, 25 UCAs, 6 loss scenarios
- **STPA on Frontier AI** — Worked sociotechnical example: 14 controllers, 65 UCAs, 9 loss scenarios, 22 sequenced remedies, applied to the political economy of frontier AI
- **Control Structures in Social Systems** — Comparing control architectures across the ten systems
- **Architectural Remedies** — Design principles for safer social systems

30. STPA Analysis of Religion as a Social System

30.1 System-Theoretic Process Analysis Applied to the Systems Engineering Decomposition of Religion

At a glance

What this chapter does: runs the four STPA steps on religion as a social system — 8 hazards, a four-level control structure (central doctrinal authority → clergy → congregation → individual), 25 unsafe control actions, 6 loss scenarios connecting structural features to specific historical harms, 5 architectural remedies.

What you need to read it: the SE decomposition of religion ([Part II, Religion](#)), the four STPA steps ([Part I, STPA and STAMP](#)), and the rung apparatus only for §6.4 ([Part I, Justificatory Rungs](#)).

Skip to §6 if you want the cross-cutting structural findings without the full UCA enumeration. **Skip to §7** if you came for the architectural remedies and trust the analysis. **Skip §6.4** if rungs are unfamiliar — the surrounding sections work without them.

30.2 Abstract

System-Theoretic Process Analysis (STPA) is a hazard analysis method developed by Nancy Leveson at MIT, based on the STAMP (Systems-Theoretic Accident Model and Processes) framework. STPA was designed for safety-critical engineering systems — aircraft, nuclear plants, medical devices — but its core logic applies wherever a complex system can produce unintended harmful outcomes through inadequate control rather than mere component failure. This paper applies STPA to religion, treated as a social system previously decomposed through a five-level systems engineering hierarchy. We define the system's losses, model its control structure, identify unsafe control actions across all four STPA categories, and trace loss scenarios to their causal factors. The result is a structural explanation of how religion — a system designed to provide meaning, morality, and community — can produce outcomes that are the exact opposite of its stated goals: meaninglessness, moral corruption, and social destruction.

30.3 1. Background: The SE Decomposition of Religion

1.1 The Five-Level Hierarchy

The following decomposition was established in prior work. It forms the input to the STPA analysis.

Level	ID	Description
Goal	G1	Provide existential meaning and purpose to human life
Goal	G2	Establish moral and ethical orientation for individuals and communities
Goal	G3	Create social cohesion and collective identity
Goal	G4	Offer psychological comfort in the face of suffering and mortality
Req.	R1	Coherent cosmological narrative (creation, afterlife, transcendence)
Req.	R2	Framework for interpreting personal experience as meaningful
Req.	R3	Codified behavioral norms (commandments, precepts, dharma)
Req.	R4	Authority structure to interpret and enforce norms
Req.	R5	Shared symbols, narratives, and practices for group bonding
Req.	R6	Mechanism for distinguishing in-group from out-group
Req.	R7	Rituals addressing grief, loss, and life transitions
Func.	F1	Mythological narration — transmit origin and eschatological stories
Func.	F2	Pastoral care — counsel individuals through crisis and life stages
Func.	F3	Moral legislation — define permissible and forbidden behavior
Func.	F4	Doctrinal governance — interpret scripture, settle disputes
Func.	F5	Ritual performance — enact shared ceremonies
Func.	F6	Missionary expansion — propagate belief system
Log.	L1	Liturgical Subsystem — scripture, preaching, ceremonial calendar
Log.	L2	Pastoral Subsystem — counseling, rites of passage, community support
Log.	L3	Doctrinal/Juridical Subsystem — theology, canon law, ethics boards
Log.	L4	Outreach Subsystem — evangelization, education, charity, media
Phys.	P1	Churches, mosques, temples, synagogues, gurdwaras
Phys.	P2	Sacred texts — Bible, Quran, Torah, Vedas, Tripitaka
Phys.	P3	Clergy — priests, imams, rabbis, pastors, monks, gurus
Phys.	P4	Parish networks, sanghas, congregations, umma structures
Phys.	P5	Central authorities — Vatican, Al-Azhar, ecumenical councils
Phys.	P6	Legal systems — canon law, fatwa apparatus, halakhic courts
Phys.	P7	Missionary organizations, madrasas, Sunday schools
Phys.	P8	Religious media — TV networks, podcasts, publishing houses

1.2 Why STPA?

Traditional risk analyses ask: "What component can fail?" STPA asks a different question: "What control actions, if inadequate, can lead to losses — even when every component is working as designed?" This is precisely the right question for religion, where the most devastating outcomes (inquisitions, abuse cover-ups, radicalization, exclusion of vulnerable people) typically occur not because the system has "broken" but because its control mechanisms are operating exactly as designed — just in a way that produces harm.

30.4 2. STPA Step 1: Define Purpose of the Analysis

2.1 System-Level Losses

STPA begins by identifying **losses** — outcomes that the system's stakeholders would consider unacceptable. For religion as a social system, we define:

Loss ID	Loss Description	Affected Goal
L1	Individuals suffer existential despair, alienation, or loss of meaning because of the religious system (not despite it)	Negates G1
L2	Moral harm — the system produces, enables, or conceals immoral behavior (abuse, violence, corruption, exploitation)	Negates G2
L3	Social destruction — the system causes division, persecution, hatred, or war between groups	Negates G3
L4	Psychological damage — the system inflicts trauma, guilt, fear, or mental illness on individuals	Negates G4
L5	Loss of institutional legitimacy — the system loses the trust and authority it needs to function	Negates all
L6	Suppression of truth — the system suppresses scientific knowledge, critical thinking, or individual conscience	Negates G1, G2

Note: each loss is the negation of a system goal. This is the defining feature of STPA applied to social systems — the hazards are not external threats but self-inflicted contradictions where the system's own control actions produce outcomes opposite to its stated purpose.

2.2 System-Level Hazards

Hazards are system states that, combined with worst-case environmental conditions, lead to losses. We derive them from the SE requirements:

Hazard ID	Hazard Description	Related Losses
H1	Doctrine becomes unfalsifiable and self-sealing — no internal mechanism can correct doctrinal errors	L1, L2, L6
H2	Authority structure becomes unaccountable — no effective oversight of those who interpret and enforce norms	L2, L4, L5
H3	In-group/out-group mechanism escalates to dehumanization of outsiders	L3
H4	Moral legislation becomes absolutist — no allowance for context, conscience, or evolving understanding	L1, L2, L4
H5	Pastoral care relationship is exploited — power asymmetry enables abuse	L2, L4, L5
H6	Missionary function prioritizes expansion over welfare of converts or target populations	L2, L3
H7	Financial resources are extracted from vulnerable populations without accountability	L2, L5
H8	Ritual performance becomes coercive — participation is compelled, not voluntary	L1, L4

2.3 System-Level Constraints

Each hazard implies a constraint — a condition the system must maintain to prevent the hazard:

Constraint ID	Constraint	Prevents Hazard
SC1	Doctrine must include mechanisms for self-correction, reinterpretation, and engagement with external knowledge	H1
SC2	Authority figures must be subject to transparent oversight, accountability, and removal procedures	H2
SC3	In-group/out-group distinctions must not extend to denial of outsiders' human dignity or rights	H3
SC4	Moral legislation must preserve space for individual conscience and contextual judgment	H4
SC5	Pastoral relationships must have safeguards: codes of conduct, mandatory reporting, external recourse	H5
SC6	Missionary activity must respect the autonomy and existing culture of target populations	H6
SC7	Financial management must be transparent, audited, and separated from spiritual authority	H7
SC8	Ritual participation must be voluntary, with no material or social penalty for non-participation	H8

30.5 3. STPA Step 2: Model the Control Structure

3.1 The STAMP Control Structure

STAMP models systems as hierarchical control structures where controllers issue control actions to controlled processes, and receive feedback. For religion:





3.2 Why these control loops — and why only these?

The test for a control loop

STPA imposes three necessary conditions on every arrow in a control structure. All three must hold simultaneously; failure of any one disqualifies the relationship from being modelled as a control loop:

- 1. **Controller (authority condition):** the upstream element issues observable, authoritative control actions — directives the downstream element is structurally expected to obey, not merely influence it may choose to consider.
- 2. **Controlled process (response condition):** the downstream element's state is expected to change in response to those control actions as part of normal system operation.
- 3. **Feedback channel (closure condition):** some signal — however weak, delayed, or filtered — travels back from the controlled process to the controller and can in principle update the controller's process model.

A relationship that passes all three tests is a control loop and must be modelled. A relationship that fails any one test is modelled differently: as peer coordination (aggregated into an existing feedback channel), as an environmental disturbance, or as a belief inside a controller's process model.

Applying the test to the SE decomposition

The four loops follow directly from applying these three conditions to the physical elements named in §1.1. The authority source that makes relationships authoritative rather than merely influential in every case is R4 ("authority structure to interpret and enforce norms") — a requirement the system itself specifies.

Loop	Controller	Controlled Process	Condition 1: authority	Condition 2: state change	Condition 3: feedback
1	Central Authority (P5, P6)	Clergy (P3)	R4 + claimed divine mandate; clergy bound under canon law / fatwa / denominational discipline	Clergy doctrine, appointments, removal — clergy behaviour demonstrably changes in response	Reports from clergy, synod deliberations, petitions (filtered but structurally present)
2	Clergy (P3)	Congregation (P4)	R4 delegated downward; authority to preach, administer sacraments, impose social discipline	Congregation attendance, expressed belief, ritual participation change in response	Confession, counselling, donation levels, attendance — reach clergy as indirect behavioural signals
3	Congregation (P4)	Individual	R5, R6 (shared identity, in/out-group enforcement); peer authority over social standing and belonging	Individual's public behaviour; expressed belief, ritual participation change to maintain belonging	Gossip, reputation signals, observation of conformity or deviance feed back to community process model
4	Individual (self)	Own behaviour	Internalised R1–R3; conscience acts as an internal authority that the person is expected — by the system's own design — to obey	Own behaviour, choices, emotional state, and spiritual self-assessment change	Emotional feedback: guilt, consolation, doubt — conscience functions as internal sensor closing the loop

These are the only four pairs in §1.1 where one element holds designated authority to issue control actions the other is structurally expected to obey. Remove any loop and the transmission path from doctrine to individual behaviour breaks — no alternative channel carries an authoritative control action all the way to the individual.

Why every other plausible relationship fails the test

The table below applies the three conditions to every candidate relationship commonly raised when this control structure is presented. In every case the candidate fails at least Condition 1.

Candidate relationship	Fails condition	Reason
Clergy ↔ Clergy (lateral, same rank)	1 — authority	Clergy of equal rank do not issue authoritative control actions to each other. Synods, councils, and theological deliberations are coordination inputs that aggregate into Loop 1's feedback channel, not independent loops. Where genuine rank differences exist (bishop over priest, abbot over monk), the relationship resolves to Loop 1 or Loop 2 at finer granularity — not a new, fifth loop.
Congregation ↔ Congregation (between parishes)	1 — authority	Parish networks are parallel peers. No congregation holds designated authority over another. Binding coordination between congregations routes through Loop 1 — the central authority that owns both. Direct inter-congregational social pressure carries no structural control action.
Individual or congregation → Authority (upward)	1 — authority	In the hierarchical institutional form this analysis targets, laity have no structural authority to issue control actions upward. Their upward channel is feedback, already represented in every loop. Note: congregationalist polities (Quaker, many Baptist, Reform Jewish) genuinely invert Loop 1 — authority flows from the congregation upward. Those constitute a genuinely different control structure under the same family name and require their own analysis.
God → Authority	1 — authority	The system claims this loop, but divine control actions are not observably issued or measurably received. What actually operates is the Authority's belief that it is divinely controlled — a feature of the Authority's process model (shown inside the top box of the diagram), not a separate causal loop. STPA models what exists in the causal structure, not what is asserted.
Civil state / science / secular media	1 — authority	These elements lie in the system's environment. They can disturb the system or supply external corrective signals, but they do not issue control actions the religion is architecturally bound to obey. They appear in §7 as externally imposed circuit breakers, not as internal control loops.
Other religions	1 — authority	No cross-denominational structural authority exists by definition of denomination. Influence between traditions is cultural and bilateral; control is absent in both directions.
Schism / sect formation	Not a loop at all	This is a phase transition, not a control relationship. When Loop 1 fails to correct accumulating error, a fragment of the controlled process exits the hierarchy entirely and becomes the apex of a new control structure. Schism is a symptom of loop failure (analysed in §5), not a sixth relationship within the existing hierarchy.
The upward feedback arrows on each loop	Not a new loop	Every loop already includes a feedback channel (the upward arrows in the diagram). Those arrows satisfy Condition 3 of their respective loop. They do not satisfy Condition 1 of a new loop — feedback is not control. Treating them as control would mean the congregation governs the clergy and the clergy governs the central authority, which is structurally false in the institutional form under analysis.

A different decomposition — congregational polity, decentralised charismatic movement, household-religion — would yield a different control structure with different loops. Within the hierarchical institutional decomposition analysed here, no fifth loop exists, because no fifth pair of elements satisfies all three conditions simultaneously.

3.3 Key Observations about the Control Structure

Feedback poverty at the top. The Central Doctrinal Authority receives feedback primarily through its own appointees (clergy), who have incentives to report compliance rather than problems. External feedback (scientific findings, secular criticism, abuse reports) is structurally classified as hostile input, not corrective signal.

Asymmetric power at every level. Each controller has significantly more power over the controlled process than the controlled process has to send corrective feedback upward. This is by design (R4: authority to interpret and enforce), but it means that control errors propagate downward unchecked.

Self-confirming process models. At every level, the controller's mental model of the controlled process contains built-in resistance to disconfirmation. If a believer doubts, the model says "doubt is temptation" — the model explains away the very feedback that would correct it.

30.6 4. STPA Step 3: Identify Unsafe Control Actions

STPA categorizes unsafe control actions (UCAs) into four types:

- 1. **Not providing** a control action that is needed
- 2. **Providing** a control action that causes a hazard
- 3. **Providing** a control action **too early, too late, or out of sequence**
- 4. **Providing** a control action that is **applied too long or stopped too soon**

4.1 UCAs for Central Doctrinal Authority → Clergy

Eight UCAs at this loop, dominated by Not provided (doctrine and discipline that should be issued and is not) and Provided (doctrine that should not be issued and is).

All eight UCAs — Central Authority → Clergy			
UCA ID	Type	Unsafe Control Action	Hazard
UCA-1	Not provided	Authority does NOT update doctrine when scientific evidence invalidates cosmological claims (e.g., geocentrism, age of earth, evolution)	H1
UCA-2	Not provided	Authority does NOT remove or discipline clergy credibly accused of abuse	H2, H5
UCA-3	Provided	Authority ISSUES doctrine that dehumanizes outsiders (infidels, heretics, apostates as subhuman or damned)	H3
UCA-4	Provided	Authority EXCOMMUNICATES or PUNISHES members who raise legitimate doctrinal questions	H1, H4
UCA-5	Provided	Authority APPOINTS clergy based on doctrinal loyalty rather than pastoral competence or moral character	H2, H5
UCA-6	Too late	Authority CORRECTS a doctrinal error only after centuries of harm (e.g., Galileo rehabilitation in 1992, 359 years late)	H1, H6
UCA-7	Too long	Authority MAINTAINS absolutist moral prohibition long after ethical understanding has evolved (e.g., blanket prohibitions that cause measurable suffering)	H4
UCA-8	Not provided	Authority does NOT establish transparent financial auditing of religious institutions	H7

4.2 UCAs for Clergy → Congregation

Seven UCAs at this loop, including the highest-stakes UCA-9 (pastoral relationship used for exploitation) and the cover-up pattern at UCA-15.

All seven UCAs — Clergy → Congregation

UCA ID	Type	Unsafe Control Action	Hazard
UCA-9	Provided	Clergy USES pastoral relationship to sexually, emotionally, or financially exploit vulnerable individuals	H5
UCA-10	Provided	Clergy PREACHES hatred or contempt toward out-groups (other religions, LGBTQ+ persons, ethnic minorities, apostates)	H3
UCA-11	Provided	Clergy WITHHOLDS sacraments or pastoral care as punishment for questioning doctrine	H1, H4, H8
UCA-12	Not provided	Clergy does NOT provide pastoral support during genuine crisis (grief, mental illness, abuse) because doctrinal framework has no model for the problem	H4
UCA-13	Provided	Clergy DEMANDS financial contributions from impoverished believers through spiritual coercion (prosperity gospel, indulgences, obligatory tithing)	H7
UCA-14	Too long	Clergy MAINTAINS shunning of excommunicated member long after the person has suffered disproportionate social death	H4, L3
UCA-15	Not provided	Clergy does NOT report abuse to civil authorities, instead handling it "internally"	H2, H5

4.3 UCAs for Congregation → Individual

Five UCAs at the peer-enforcement loop, dominated by social-pressure mechanisms operating at justificatory rung 1 (community authority) where rung-3 evidence (developmental harm) would warrant restraint.

All five UCAs — Congregation → Individual

UCA ID	Type	Unsafe Control Action	Hazard
UCA-16	Provided	Congregation SHUNS member who expresses doubt, questions doctrine, or leaves the faith	H1, H4, H8
UCA-17	Provided	Congregation ENFORCES social conformity on matters of personal conscience (dress, diet, relationships, career choices) through gossip and exclusion	H4, H8
UCA-18	Not provided	Congregation does NOT intervene when they observe clergy abusing a member	H2, H5
UCA-19	Provided	Congregation PRESSURES children into public faith commitments (baptism, confirmation, bar mitzvah) before they can meaningfully consent	H8
UCA-20	Too early	Congregation LABELS children as sinful, damned, or spiritually deficient at a developmentally inappropriate age	H4

4.4 UCAs for Individual (Self-Control / Internalized Beliefs)

Five UCAs at the self-control loop: the individual's internalised model of the system acts as a controller that filters their own reasoning and choices.

All five UCAs — Individual self-control

UCA ID	Type	Unsafe Control Action	Hazard
UCA-21	Provided	Individual SUPPRESSES legitimate doubt, curiosity, or moral intuition because they have internalized "doubt = sin"	H1, H6
UCA-22	Provided	Individual REMAINS in an abusive religious environment because they believe leaving = damnation	H4, H5
UCA-23	Not provided	Individual does NOT seek professional mental health support because the religious framework teaches that faith alone should suffice	H4
UCA-24	Provided	Individual REJECTS medical treatment for self or dependents based on doctrinal prohibition	H4, L1
UCA-25	Too long	Individual CONTINUES religious practices that cause measurable psychological harm (extreme fasting, self-flagellation, sleep deprivation) beyond any devotional purpose	H4

30.7 5. STPA Step 4: Identify Loss Scenarios

Each UCA has **causal factors** — the reasons why the unsafe control action occurs. STPA organizes these into two categories:

A. Why would the controller issue the UCA? (Flawed process model, incorrect feedback, conflicting goals)

B. Why would the control action not be executed properly? (Communication failure, implementation breakdown)

5.1 Loss Scenario LS-1: Doctrine Fails to Self-Correct (UCA-1, UCA-6 → H1 → L1, L6)

The scenario: Central authority maintains a cosmological or moral doctrine that contradicts established scientific evidence or evolved ethical understanding. Believers experience cognitive dissonance between their faith and their knowledge, leading to either intellectual suppression (L6) or loss of faith and meaning (L1).

Why the controller issues this UCA:

- **Flawed process model:** The authority's model holds that doctrine is divinely revealed and therefore cannot be wrong. This makes self-correction logically impossible within the model — any correction would imply the divine source was wrong.
- **Conflicting goals:** Correcting doctrine threatens the authority's legitimacy (SC1 conflicts with R4). If the authority was wrong about X, what else might it be wrong about? The authority perceives self-correction as existentially threatening.
- **Inadequate feedback:** Believers who experience the dissonance often leave silently rather than escalating. The authority receives feedback only from those who stay — a survivorship bias that confirms the model.
- **Temporal delay:** By the time evidence is overwhelming, the authority has centuries of institutional commitment to the doctrine. Reversal costs increase with time (sunk cost fallacy at institutional scale).

Structural fix (from SC1): Build doctrinal revision mechanisms into the system. Historical examples: Vatican II (1962–65) as a partial self-correction mechanism. Judaism's tradition of Talmudic disputation, which structurally preserves minority opinions. Islam's concept of *ijtihad* (independent reasoning), when not suppressed. The key architectural principle: doctrine should be versionable, with explicit revision processes, not immutable.

5.2 Loss Scenario LS-2: Abuse Cover-Up (UCA-2, UCA-15 → H2, H5 → L2, L4, L5)

The scenario: A clergy member abuses a congregant (sexually, emotionally, financially). Other clergy and the authority become aware but do not remove the abuser or report to civil authorities. Instead, the abuser is transferred, the victim is silenced, and the institution protects its reputation. This is the pattern documented in the Catholic Church, the Southern Baptist Convention, Jehovah's Witnesses, and numerous other contexts.

Why the controller issues this UCA:

- **Flawed process model:** The authority's model holds that (a) clergy are called by God and therefore inherently trustworthy, (b) the institution's reputation is essential to its mission, and (c) internal resolution is preferable to secular intervention. All three beliefs are structurally embedded in R4 and the control structure.
- **Conflicting goals:** Protecting the institution (G3: social cohesion, R4: authority structure) conflicts with protecting the individual (G2: moral orientation, G4: comfort). The control structure systematically prioritizes the institution because the controllers are the institution.
- **Feedback suppression:** The victim's feedback (complaint, report) enters the system through the very structure that has an interest in suppressing it. There is no independent feedback channel that bypasses the authority hierarchy.
- **Accountability void:** The controller who should discipline the abuser (bishop, denominational leader) is also the person whose career and institution are damaged by disclosure. The judge and the interested party are the same entity.

Structural fix (from SC2, SC5): (a) Mandatory external reporting to civil authorities — removing the internal-only handling option. (b) Independent ombudsman or safeguarding board with authority to investigate and sanction, staffed by people outside the clerical hierarchy. (c) Separation of pastoral authority from disciplinary authority — the person who counsels the victim must not be the person who decides whether to report. (d) Statute of limitations reform within canon law. These are architectural changes to the control structure itself, not mere policy adjustments.

5.3 Loss Scenario LS-3: Radicalization through In-Group/Out-Group Escalation (UCA-3, UCA-10 → H3 → L3)

The scenario: The mechanism for creating social cohesion (R5, R6: shared identity and in-group/out-group distinction) escalates from "we are different from them" to "they are inferior" to "they are dangerous" to "they must be destroyed." This progression has driven religious wars, pogroms, crusades, inquisitions, and modern terrorism.

Why the controller issues this UCA:

- **Flawed process model:** The in-group/out-group mechanism (R6) has no built-in limit. The same mechanism that creates belonging ("we share this faith") also creates hostility ("they threaten our faith"). The model does not distinguish between "different" and "dangerous."
- **Positive feedback loop:** Preaching against out-groups increases in-group solidarity (measurable by attendance, donations, emotional intensity). This positive feedback rewards escalation. The more extreme the rhetoric, the stronger the group cohesion — until the rhetoric produces violence.
- **External trigger:** The escalation is often catalyzed by external threat (real or perceived). Economic stress, political marginalization, or military conflict activates the in-group/out-group mechanism and pushes it toward its extreme.
- **Missing control action:** There is no structural "circuit breaker" that de-escalates when rhetoric approaches dehumanization. The authority may agree with the escalation, or may be too distant to intervene in local radicalization.

Structural fix (from SC3): (a) Explicit doctrinal commitment to universal human dignity that overrides in-group/out-group distinctions — structurally, this means SC3 must be higher-priority than R6 in the requirements hierarchy. (b) Interfaith engagement as a standing function, not an exceptional one — regular contact with out-groups reduces dehumanization (contact hypothesis). (c) Internal monitoring for escalatory rhetoric, with de-escalation authority assigned to a specific role.

5.4 Loss Scenario LS-4: Spiritual Coercion and Psychological Damage (UCA-11, UCA-16, UCA-17, UCA-20 → H4, H8 → L1, L4)

The scenario: An individual (often a child or young person) is subjected to religious teaching that instills deep fear (hellfire, damnation, divine punishment), shame (sinfulness, unworthiness), or guilt (for normal developmental experiences like doubt, sexual feelings, or intellectual curiosity). Social enforcement (shunning, withdrawal of love, public shaming) reinforces the psychological impact. The individual develops religious trauma, anxiety disorders, or complex PTSD.

Why the controller issues this UCA:

- **Flawed process model:** The clergy and congregation model holds that (a) fear of divine punishment is a legitimate motivator, (b) shame is a corrective emotion, and (c) social exclusion of deviants protects the community. These beliefs are embedded in R3 (codified norms), R6 (in-group/out-group), and the enforcement functions (F3, F4).
- **Missing feedback:** Children and psychologically dependent adults cannot provide effective corrective feedback to the system. They lack the power, the vocabulary, and often the awareness that their experience is harmful. By the time they can articulate the damage, they may have left the system entirely — at which point their feedback is classified as "apostate hostility."
- **Normalization:** When an entire community shares the same beliefs and practices, harmful behaviors are invisible as harmful. Telling a child they are "born in sin" feels like education, not psychological abuse, to a community that genuinely believes it.

Structural fix (from SC4, SC8): (a) Age-appropriate religious education that does not use fear or shame as motivational tools — this requires explicit pedagogical standards within the religious education subsystem. (b) Voluntary participation norms: no social penalty for non-attendance, non-belief, or questioning. (c) Mental health literacy within pastoral training — clergy should be trained to recognize when religious practice is causing psychological harm and to refer to professional support. (d) External child safeguarding standards applied to religious education, just as they are applied to secular schools.

5.5 Loss Scenario LS-5: Financial Exploitation (UCA-8, UCA-13 → H7 → L2, L5)

The scenario: Religious authority extracts financial resources from believers through spiritual coercion — "God wants you to give," "tithing is obligatory," "sowing seeds of faith will bring financial blessings" (prosperity gospel). Funds are not transparently accounted for. Leaders live in luxury while congregants struggle.

Why the controller issues this UCA:

- **Conflicting goals:** The system needs financial sustainability (R7 in the SE decomposition) but has no structural separation between spiritual authority and financial management. The person who tells you what God wants is also the person who tells you how much to give.
- **Missing oversight:** Unlike secular nonprofits, many religious organizations are exempt from financial reporting requirements in numerous jurisdictions. The control structure has no external audit function.
- **Doctrinal justification:** The prosperity gospel and similar doctrines create a self-sealing system: if you give and prosper, it proves the doctrine; if you give and don't prosper, your faith was insufficient — give more.

Structural fix (from SC7): (a) Mandatory financial transparency — published annual reports, independent external audit, separation of spiritual authority from financial authority. (b) Congregation-elected financial oversight body with real power (a "Kassenprüfung" model, borrowed from the German Verein). (c) Doctrinal prohibition on linking financial giving to spiritual outcomes.

5.6 Loss Scenario LS-6: Suppression of Individual Conscience (UCA-4, UCA-7, UCA-21 → H1, H4 → L1, L6)

The scenario: A believer's individual moral conscience or intellectual understanding conflicts with official doctrine. The system's response is to demand submission of conscience to authority. The individual either suppresses their authentic moral judgment (psychological harm, L4) or is punished for expressing it (exclusion, L1). In either case, the system has sacrificed one of its own goals (G2: moral orientation) to preserve another (R4: authority structure).

Why the controller issues this UCA:

- **Architectural conflict:** R4 (authority to interpret and enforce) and G2 (moral orientation) are in structural tension. R4 was designed to serve G2 (authority enables moral guidance), but R4 can become an end in itself — authority for authority's sake — at which point it overrides G2.
- **Process model deficiency:** The authority's model does not contain the possibility that the individual's conscience could be more morally correct than the doctrine. The model assumes doctrinal infallibility, which means any dissent is by definition error.

Structural fix (from SC1, SC4): (a) Explicit doctrinal recognition of conscience as a valid moral source — Catholicism actually has this (Dignitatis Humanae, Vatican II) but struggles to operationalize it when conscience contradicts doctrine. (b) Structured dissent mechanisms: the ability to formally register disagreement without excommunication, analogous to a "dissenting opinion" in a court ruling. (c) Periodic doctrinal review processes that engage lay members and external scholars, not just clergy.

30.8 6. Cross-Cutting Structural Findings

6.1 The Feedback Problem Is Universal

Every loss scenario shares a common structural deficiency: **inadequate upward feedback**. The control structure of religion is designed for top-down transmission (doctrine flows from authority to believer) but has weak, filtered, or suppressed bottom-up feedback. This is not a bug but a feature of the design (R4: authority to interpret and enforce). However, it means that control errors are systematically uncorrected.

In control systems engineering, this is called **open-loop control** — the controller acts without adequate measurement of the actual system state. Open-loop control works only when the process is perfectly predictable. Human societies are never perfectly predictable. Therefore, open-loop control of human behavior through religion will always produce drift toward loss conditions.

The architectural remedy is to **close the loop**: create genuine, unfiltered feedback channels from the controlled process (believers, congregations) to the controller (authority), with structural protections ensuring the feedback cannot be suppressed.

6.2 The Self-Sealing Model Problem

Multiple UCAs are enabled by **self-sealing process models** — mental models held by controllers that explain away any evidence that would contradict them:

Self-Sealing Pattern	How It Works
"Doubt is sin"	Eliminates the primary corrective signal (questioning)
"Suffering is divine test"	Reframes harm as benefit — the worse it gets, the more it "proves" the model
"External criticism is persecution"	Classifies corrective feedback from outside the system as hostile action
"Apostates were never true believers"	Discounts the testimony of the people with the most relevant experience
"God's ways are beyond human understanding"	Makes the model unfalsifiable by definition

In control engineering, a process model that cannot be updated by feedback is called **unobservable**. An unobservable system cannot be controlled safely. The architectural remedy is to ensure that every process model in the control structure includes explicit **update mechanisms** — conditions under which the model itself is revised.

6.3 The Conflicting-Goals Problem

The SE decomposition reveals that several system goals are in structural tension:

Goal Pair	Tension
G2 (moral orientation) vs. R4 (authority)	Authority enables morality but can also override conscience
G3 (social cohesion) vs. R6 (in-group/out-group)	Cohesion requires boundaries, but boundaries enable exclusion and hatred
G4 (comfort) vs. R3 (norms)	Comfort requires unconditional acceptance, but norms require conditional compliance
G1 (meaning) vs. R1 (cosmological narrative)	Meaning requires honest engagement with reality; a fixed narrative may conflict with reality

These tensions are not design flaws — they are inherent in any system that simultaneously serves individual and collective needs. The design flaw is when the system has **no mechanism for managing the tension**: no process for deciding which goal takes priority when they conflict, and no authority empowered to make that trade-off transparently.

6.4 The rung mismatch behind the four findings

The previous chapter [Justificatory Rungs](#) introduced a seven-rung ladder of standards under which control actions and feedback signals are accepted. Reading §§6.1–6.3 through that lens, **rung mismatch subsumes them all** for this system:

- **"Feedback poverty at the top"** is the rung-1 controller refusing rung-3 input — Pattern A (asymmetric loop). The rung-1 hierarchy classifies abuse data, scientific evidence, and demographic harm as rung-1 hostility and filters them out.
- **"Self-confirming process models"** is the controller's acceptance-rung filter operationalised. Each entry in the table in §6.2 ("doubt is sin," "suffering is divine test," "external criticism is persecution," "apostates were never true believers," "God's ways are beyond human understanding") is a specific rule for re-classifying rung-3+ feedback as rung-1 noise.
- **"Conflicting goals"** are tractable when the system can deliberate at rung 6 about how to balance them; in a rung-1 system, whichever side has more authority simply wins.

The system simultaneously runs **Pattern B — Claimed-Rung Inflation**: it claims rung 6 (sacralised ultimate legitimacy) while operating at rung 1 (clerical authority). The gap is enforced doctrinally rather than closed, which is why the failures are predictable and recurrent.

The architectural fixes in §7 are all attempts to insert **rung-3 channels that route around the rung-1 filter** without forcing the controller to abandon its rung-1 self-understanding. Mandatory civil reporting, independent ombudspersons, lay safeguarding boards, doctrinal version control with external scholarship, and pre-registered scientific-consultation processes are all implementations of the same architectural pattern.

This is the single most important structural finding of the analysis: religion's harmful outcomes are not random malfunctions but the **predictable consequence of operating a sacralised rung-1 control architecture while claiming the authority of rung 6**. The canonical reference, including the rung-tagged control structure and the rung-acceptance filter as a recurring causal factor in every loss scenario, is in [knowledge/system-catalogues/social-systems/religion/applied-se-analysis.md](#).

30.9 7. Recommendations: Architectural Improvements

Based on the STPA analysis, we propose the following changes to religion's control structure. These are stated as architectural principles, applicable across traditions:

7.1 Close the Feedback Loop

- Establish **independent feedback channels** that bypass the authority hierarchy: ombudsman, external review boards, anonymous reporting systems.
- Require **periodic surveys of congregational wellbeing** — measuring psychological health, not just attendance and donations.
- Create **formal exit interviews** for members who leave, treated as system feedback rather than apostasy.

7.2 Make Process Models Updatable

- Institute **doctrinal version control** — every doctrine should have a stated scope, a review date, and a revision process.
- Require **engagement with external knowledge** (science, psychology, ethics) as a standing input to doctrinal governance, not an exceptional event.
- Preserve **dissenting opinions** in the doctrinal record, as courts preserve minority opinions.

7.3 Separate Control Functions

- Separate **spiritual authority** (preaching, sacraments, counseling) from **administrative authority** (finances, personnel, discipline).
- Separate **pastoral care** from **norm enforcement** — the person who counsels you should not be the person who punishes you.
- Establish **independent safeguarding** with authority to investigate and override clergy decisions on abuse cases.

7.4 Install Circuit Breakers

- Define **escalation limits** on in-group/out-group rhetoric: explicit doctrinal commitment to universal human dignity that cannot be overridden by any other teaching.
- Establish **proportionality review** for disciplinary actions: shunning, excommunication, and other exclusion mechanisms should require documented justification, appeal process, and time limits.
- Require **informed consent** for participation in intensive religious practices (extended fasting, retreats, exorcism rituals) with right of withdrawal.

7.5 Protect Vulnerable Populations

- Apply **secular child safeguarding standards** to all religious education and youth activities.
 - Prohibit **fear-based instruction** (hellfire, damnation) for children below a developmentally appropriate age.
 - Require **mandatory external reporting** of all abuse allegations to civil authorities, with no exceptions for internal handling.
-

30.10 8. Conclusion

STPA reveals that religion's most harmful outcomes are not random malfunctions but **predictable consequences of its control architecture**. A system designed for top-down doctrinal transmission with weak upward feedback, self-sealing process models, unaccountable authority, and no structural mechanism for managing goal conflicts will, over sufficient time, produce every one of the losses we identified — regardless of the sincerity or goodness of the people within it.

This is not an indictment of religion's purpose. The goals — meaning, morality, community, comfort — are among the most important things humans seek. It is an indictment of religion's control architecture as historically implemented. The same goals could be served by a system with closed feedback loops, updatable doctrines, separated authority, independent oversight, and protected individual conscience. Some religious traditions have moved in this direction; none has fully arrived.

The value of STPA is that it moves the conversation from "religion is good" vs. "religion is harmful" — a debate that generates more heat than light — to a structural question: **which control actions, under which conditions, produce which outcomes?** That question has specific, actionable answers. The fixes are architectural, not theological. They do not require abandoning faith. They require redesigning the system that carries it.

This analysis applies STPA (Leveson, 2012) to the SE decomposition of religion developed in prior work. The four STPA steps — define losses, model control structure, identify UCAs, trace loss scenarios — are followed systematically. All hazard and constraint IDs trace back to the original SE decomposition's goals and requirements.

References

- Leveson, N.G. (2012). Engineering a Safer World: Systems Thinking Applied to Safety. MIT Press.
- Leveson, N.G. (2004). "A New Accident Model for Engineering Safer Systems." Safety Science, 42(4), 237–270.
- Leveson, N.G. & Thomas, J.P. (2018). STPA Handbook. MIT Partnership for Systems Approaches to Safety and Security.

31. STPA Analysis of Frontier AI as a Sociotechnical System

31.1 What Will AI Do to Us? A Systems-Engineering Reading

At a glance

What this chapter does: the longest worked example in the book. Treats the global political economy of frontier AI as a sociotechnical system, draws its 14 controllers, identifies five structural mismatches (§3), traces them to nine failure scenarios (§4), derives a four-phase sequence of 22 remedies (§5), and lists what the analysis cannot tell you (§6).

What you need to read it: Part I in full — and especially [Justificatory Rungs](#), because the chapter uses rung tags as routine shorthand. STPA's four steps are folded into §§3–5 rather than labelled (see the heading map below).

If you only have 30 minutes: read §1 (questions), §2 (method), §3.6 (what the five mismatches buy us), §4.10 (what the nine scenarios buy us), §5.1 (the four-phase sequence), and §7 (coda). Everything else is enumeration the conclusions rest on.

Heading map onto STPA Steps: §1–2 = framing · §3 (system + mismatches) = Step 2 (Control Structure) · §3.1–3.5 = Step 3 (Unsafe Control Actions) compressed by mismatch type · §4 = Step 4 (Loss Scenarios) · §5 = Architectural Remedies · §6 = scope limits · §7 = coda.

31.2 Abstract

The diffusion of frontier artificial intelligence is reshaping the global economy and political order, but most public analysis of this fact takes the form of strategic essays — narratives that identify forces and rank scenarios without pointing at specific levers. This chapter applies System-Theoretic Process Analysis (STPA), with justification-rung tagging from the social-systems extension developed in [Part I](#), to the global political economy of frontier AI. We identify fourteen distinct controllers, twenty-six control actions, sixty-five unsafe control actions across the four standard categories, and nine loss scenarios connecting specific structural mismatches to specific harms. We then derive twenty-two concrete remedies organised into a four-phase sequence with explicit decision points. The result is a structural reformulation of the questions about AI's effects on jobs, power, and geopolitics — turning forecasting questions into action-shaped questions of which feedback channels are missing and which institutions are needed to install them. The full analytical artefacts are filed under [projects/problems/ai-impact-analysis/](#); this chapter distils them into prose for a non-specialist reader.

31.3 1. The questions, and why the usual answers do not satisfy

Three questions about artificial intelligence circulate in every serious conversation about the next decade.

The first is about jobs. If software writes itself, what becomes of the people who wrote it? If diagnoses, contracts, lesson plans, research summaries, and software designs are produced in seconds by machines that cost a fraction of a human salary to operate, how many human roles survive? Will the result be a new wave of unemployment, or a new kind of leisure?

The second is about power. The systems that now demonstrate near- human reasoning at language tasks were built by perhaps a dozen laboratories, financed by perhaps a hundred investors, and trained on perhaps a few hundred specialised computers in a few jurisdictions. If the technology becomes the central productive force of the next era, what stops the firms that own the substrate from owning, by extension, much of what the substrate produces?

The third is about geopolitics. The same technology is being built in two large economic blocs that mistrust one another. The American bloc's frontier laboratories sit at the top of the global capability ranking; the Chinese bloc has matched them at lower rungs and is catching up at the top. Export controls, sanctions, and chip-supply dependencies have already produced the first skirmishes of an infrastructure war. Where does this end?

Each question has been answered many times. The answers are sophisticated, but they share a common shape and a common limitation. The shape is the strategic essay: a narrative that identifies the main forces, ranks them by importance, and concludes with

a list of probable scenarios or a cautious warning. The limitation is that the essay's conclusions, however confidently stated, are free-floating. They are not connected to specific mechanisms in the world that one could intervene on. The essay tells the reader what is likely to happen; it rarely tells the reader where the system can be moved.

This chapter does something different. It treats the AI economy and its political surroundings as a system in the engineer's sense: a collection of parts connected by flows and feedback loops, each of which can be examined and, if it is part of a failure pathway, changed. The discipline that does this is called systems engineering, and a particular variant of it called system-theoretic process analysis (STPA) was originally developed for safety-critical machines like aircraft and medical devices. We use it here for a different class of artefact: not a machine, but the global political economy of frontier AI.

The reader gains, by the end, a structural picture of the system, five specific identified failure mechanisms inside it, nine specific scenarios in which those mechanisms produce identifiable harms, and twenty-two concrete interventions arranged into a four-phase implementation plan. This is more than an essay can deliver, and the reason it is more is the method. So we begin with the method.

31.4 2. Two ways of thinking about the system

A strategic essay about AI says things that sound like this. Power will concentrate, but open-weight models from China will keep prices honest. Workers will be displaced, but new categories of work will emerge as they always have. The geopolitical risk is significant but probably manageable. The most important uncertainty is whether redistribution mechanisms catch up to the productivity shock. These sentences are useful for orientation. They are also, on closer examination, the kind of observations that one could make about almost any technology in almost any decade. They do not tell you which lever to pull. They do not even tell you which levers exist.

A systems-engineering analysis, by contrast, begins by asking who is actually in the system, what each of them controls, and what information each of them receives. It records this not as a narrative but as a small set of explicit lists: the controllers, the things they control, the control actions they issue, the feedback channels they receive, and — crucially — the internal model each controller carries of the thing it is controlling. It then asks, for each control action, what it would mean for that control action to be issued unsafely (provided when it should not be, withheld when it should not be, given too late, given for too long). Each unsafe possibility is recorded. Then it asks, for each unsafe possibility, why the unsafe action could occur — the controller's model could be wrong, the feedback could be missing or filtered, the action could be misexecuted, the controlled process could have changed underneath. Each cause suggests a remedy: fix the model, fix the feedback, fix the actuator, fix the structure.

This sounds like ordinary strategic thinking with extra paperwork, and at first glance the difference is just that there are tables where the essayist has paragraphs. The deeper difference shows up at three points. First, the act of writing down every control action and every feedback channel forces the analyst to confront parts of the system that the essayist could quietly skip past; the omissions become visible. Second, recording each unsafe control action against the four standard categories produces a checklist of things-that-could-go-wrong that is exhaustive in a way that strategic intuition is not. The analyst is no longer free to remember only the failure modes that fit the narrative. Third, the loss-scenario step demands that the analyst not just assert that a failure could occur but trace the causal chain through the structure. This is what produces remedies. A narrative observation that "regulators are slow" is not actionable; a loss scenario showing that labour ministries operate on a years cycle while capability deployment happens on a months cycle, so the corrective signal arrives after the cohort it was meant to help is already displaced points directly at automatic stabilisers indexed to a fast-moving benchmark.

There is a further refinement, particular to systems engineering applied to social and political systems: every flow in the structure can be tagged with the justificatory rung on which it operates. The rungs are a seven-step ladder of how a claim or a control action gets its force. Rung 0 is naked coercion: do it or else. Rung 1 is authority: do it because I say so, or because the tradition says so, or because our group says so. Rung 2 is formal consistency: this follows from rules we both accept. Rung 3 is empirical evidence: this works against rivals when tested. Rung 4 is replicated cumulative evidence. Rung 5 is meta-rational integration of multiple methods under uncertainty. Rung 6 is deliberative legitimacy — the standard by which questions about what should be done are settled in open discourse among those affected.

Each of these rungs can be claimed. Each can be operated. The two are not always the same, and where they differ, the system typically produces specific kinds of damage. A laboratory that claims to operate at rung 6 (its safety framework is "for the benefit of humanity") but actually operates at rung 1 (the framework binds only the laboratory itself, by its own choice, amendable by its own governance) is in a claim/operation gap. A funding flow that goes downward at rung 3 (financial-return diligence) but receives capability evidence upward at rung 1 (the lab's own benchmarks, in the lab's own framing) is in an asymmetric loop — the loop appears closed because there are arrows in both directions, but the upward arrow is filtered by the downward arrow's framing, and corrective information cannot get through.

Once you start labelling every arrow with its rung, three patterns keep appearing. The first is the asymmetric loop just described. The second is the claim/operation gap. The third, which I will explain when we encounter an instance, is cross-loop rung imposition — when an institution tries to settle a question that belongs at one rung using the authority of a different rung. Each of these patterns has a generic remedy that the engineering literature catalogues. The remedies are structural — they change the architecture, not the personalities — and that is what makes them robust.

We now have everything we need to look at the AI system itself. We will list its parts, draw its loops, mark the mismatches that the rung-tagging method surfaces, walk through the failures these mismatches produce, and end with the remedies they imply. The reader who finds the technical vocabulary unfamiliar can lean on the prose: we will introduce each new term in plain English at the moment it enters the chapter.

31.5 3. The system

To draw the AI system's control structure honestly we have to be willing to write down a longer list of actors than a strategic essay would normally bother with. The longer list matters because the failures we are looking for happen in the connections between actors, and a connection is invisible if either of its endpoints is missing from the picture.

At the apex sits something we will call the global polity: diffuse, slow, but the source of legitimacy on which most controllers ultimately depend. Below it, three large state groupings — the United States, the People's Republic of China, and the European Union together with other states — each with several distinct organs. The United States in this analysis is not one actor but six: the Congress that passes legislation, the Bureau of Industry and Security that issues export controls, the Treasury that runs sanctions and tax, the Federal Trade Commission that runs antitrust, the AI Safety Institute that runs technical evaluation, and the national security apparatus that classifies and designates. The People's Republic of China is three: the party centre that sets strategic direction, the Ministry of Industry and Information Technology that runs industrial policy, and the Cyberspace Administration that regulates content and approves models. The European Union is two: the Commission with its AI Office, and the member states. Other states — the United Kingdom, Japan, South Korea, India, the Netherlands, Taiwan, the Gulf monarchies — appear as a coalition layer whose compliance or defection determines whether the bigger states' levers actually move anything.

Beneath the state layer sit the capital markets in their various forms — public equities, private capital, sovereign wealth, debt. These are the controllers that decide which laboratories can sustain frontier training runs and which compute providers can sustain capital-expenditure cycles. Beneath them sit the frontier laboratories themselves, perhaps a dozen worldwide, each with its internal mission narrative, its internal capability assessments, and its internal safety framework. Beneath the laboratories sit the compute infrastructure providers — the chip designers, the foundries, the hyperscale cloud operators, the energy and grid companies. To one side of this hierarchy, neither below nor above the labs but functionally parallel to them, sits the open-weights ecosystem: the laboratories, mostly Chinese but including some Western publishers, that release model weights free for download. On the other side sit the application-layer companies that turn model capability into deployed software for paying customers. Below those, the information platforms that distribute everything. Below them the labour ministries and education systems that interface with the people whom all of this affects, and the information regulators that look after the integrity of the public information environment. At the bottom — though "bottom" is wrong because they are the controlled process this whole apparatus ultimately produces effects on — are the workers, consumers, and citizens of the cognitive-knowledge economy.

This is fourteen distinct kinds of controller, and a quarter of a hundred when we count the sub-units. The point of writing them all down is not exhaustiveness for its own sake. The point is that each of them issues control actions — directives, allocations, investments, releases, mandates, statutes, lawsuits, salaries — and each of them receives feedback — financial returns, benchmarks, statistics, news, lawsuits, elections, market prices. The shape and the rung of those flows is what determines whether the system as a whole works. A strategic essay that talks about "the AI industry" or "the labour market" without pulling these controllers apart cannot see, for example, that the antitrust authority's process model treats concentration as something that happens in mergers, while concentration in this market is happening in long-duration compute-allocation contracts that no merger review touches.

We can summarise the structure in a single picture. Money flows downward — capital allocates to laboratories and compute providers; laboratories pay compute providers for training; compute providers pay foundries for chips; everyone hires people. Capability flows outward — laboratories release models, application companies wrap those models into products, products reach workers and consumers. Information flows in two directions — capability claims and incident reports flow upward from labs to capital, regulators, and the public; benchmarks and prices flow horizontally between labs and the open-weights ecosystem; employment and price statistics flow upward from the labour and consumer base to ministries and to the polity. State controllers issue rules and enforcement actions downward; backlash flows upward when other channels fail. The picture is dense, but it is not infinite. In the analysis companion to this chapter, every arrow is enumerated; here we only need the shape.

The diagram below shows the principal controllers and the channels between them, with rung tags on the arrows and the five mismatches (M1–M5) marked at the boundaries where they occur. Solid arrows are control actions; dashed arrows are feedback.

```

flowchart TB
    POLITY["Global polity<br/>(r6 claimed / r1 operating)"]
    STATES["State controllers<br/>US · PRC · EU · others"]
    CAPITAL["Capital markets<br/>(operating r3 → r1)"]
    LABS["Frontier AI labs<br/>(claimed r6 / operating r1)"]
    INFRA["Compute infrastructure<br/>chips · fabs · hyperscale · energy"]
    OPEN["Open-weights ecosystem<br/>(price-setter)"]
    APPS["Application layer"]
    PLATS["Information platforms"]
    LABMIN["Labour ministries / education"]
    WORK["Workers / consumers<br/>(controlled process)"]

    POLITY -- "norms (r1/r6)" --> STATES
    STATES -- "rules / sanctions (r2 → r0)" --> INFRA
    STATES -- "AI Act / evals (r2/r4)" --> LABS
    CAPITAL -- "funding (r3 → r1) M2" --> LABS
    CAPITAL -- "capex (r3)" --> INFRA
    LABS -- "model release (r3)" --> APPS
    LABS -- "RSPs (claimed r6 / op r1) M1" --> POLITY
    INFRA -- "compute allocation (r3) M5" --> LABS
    OPEN -- "weight release (r3)" --> APPS
    APPS -- "deployment (r3)" --> WORK
    PLATS -- "amplification (r3 → r0/r1)" --> WORK
    LABMIN -- "retraining (r2/r4) M3" --> WORK

    WORK -. "wages, employment (r3/r4 + r0)" .-> LABMIN
    WORK -. "backlash (r0)" .-> POLITY
    LABS -. "incidents (r3 inside / r1 disclosed) M4" .-> STATES
    LABS -. "capability evidence (r1-wrapped)" .-> CAPITAL
    OPEN -. "open benchmarks (r3)" .-> POLITY

    classDef apex fill:#fef3c7,stroke:#b45309,color:#1f2937;
    classDef state fill:#d1ecf1,stroke:#1d4ed8,color:#1f2937;
    classDef capital fill:#f1f3f4,stroke:#4b5562,color:#1f2937;
    classDef tech fill:#d1ecf1,stroke:#1d4ed8,color:#1f2937;
    classDef ground fill:#f9f9f9,stroke:#4b5562,color:#1f2937;

    class POLITY apex;
    class STATES,LABMIN state;
    class CAPITAL capital;
    class LABS,INFRA,OPEN,APPS,PLATS tech;
    class WORK ground;
  
```

When we tag the arrows with their justificatory rungs, five specific mismatches stand out. Each of them is a particular case of one of the dangerous patterns we introduced in the previous section. Each is a structural reason why specific failures occur.

3.1 The first mismatch: lab safety frameworks claim more legitimacy than they operate

Frontier laboratories publish documents — responsible scaling policies, deployment policies, constitutional AI charters — that claim deliberative legitimacy. Their tone is the tone of public ethics: the lab is binding itself, on behalf of humanity, to thresholds that cannot be crossed without specified cause. The language sits at rung 6, the rung at which questions about what should be done are settled.

The operating mechanism, however, is rung 1. The lab is bound because the lab has chosen to be bound. The lab amends its framework when its own governance decides to amend it. There is no external party with standing to enforce, no court to which violations can be brought, no parliament that can refuse a budget on a non-compliant lab. This is the claim/operation gap. The gap is not a sign of bad faith — many of these documents are written in good faith by people who believe in them. The gap is structural: a self-binding commitment, however sincere, is operationally a rung-1 commitment, regardless of how it sounds.

The damage is that public trust calibrates to the higher claim and gets eroded at the rate of the gap. Every release that rolls back a previous threshold, every quiet amendment to an evaluation protocol, every framework that survives commercial pressure only in form, contributes to a slow loss of legitimacy that the labs cannot themselves staunch — because the legitimacy was never theirs to issue.

3.2 The second mismatch: capital cannot evaluate the laboratories it funds

Frontier laboratories are funded by capital allocators — venture funds, growth equity, public markets, sovereign wealth. The funds flow downward at rung 3: financial diligence, term-sheet evaluation, comparable-round analysis. The capability evidence flows back upward through the laboratories' own benchmark selection, the laboratories' own model cards, the laboratories' own mission narrative. This is rung 1 in disguise — capability claims wrapped in the lab's chosen frame.

The result is what the engineering literature calls an asymmetric loop. The arrows exist in both directions, but the upward arrow is filtered by the framing of the downward arrow. Capital cannot ask, with the rigour its own diligence claims, whether a lab's mission claims are falsifiable: the very benchmarks against which the question would be asked were chosen by the entity making the claims. The fund can either underwrite the mission narrative as given or refuse to fund — and at the scale of capital required for frontier training runs, refusing means watching a competitor fund the same lab on more pliable terms.

The consequence is that capital structure exerts a quiet but relentless pressure on operating practice. Funding rounds occur on a quarterly cadence; safety frameworks are amended on a yearly cadence; capability releases happen between rounds. When the cadences disagree, the cadence with the money wins.

3.3 The third mismatch: the labour-policy cycle cannot keep up with the capability-deployment cycle

Labour ministries and education systems track unemployment, sectoral employment, vacancies, and wage statistics. These indicators arrive at rung 3 or rung 4 — empirical aggregates, sometimes cumulative — but they arrive on a quarterly or annual cycle. Retraining programmes and curricular updates respond on a multi-year cycle. Capability deployment, by contrast, moves on a months cycle: a new model is released, application companies integrate it, automation is in production within weeks.

The mismatch in cycle times means that the corrective signal — a labour-market indicator showing that a class of workers is being displaced — arrives after the cohort it would have informed is already displaced. The signal cannot land in time to affect the decisions that produced the displacement. When this happens repeatedly, the rung-3/4 channel loses its informational value, and the only channel left is rung 0: political backlash, which arrives at the speed of news cycles and elections rather than at the speed of statistics. The corrective intent of the statistical channel is overrun by the rung-0 escape valve.

This is not a story about regulators being slow in some abstract sense. It is a specific identification of two cycles whose periods no longer line up — and the structural fix is not to be smarter, but to install instruments whose triggers fire automatically on a fast index, removing the requirement that a slow channel confirm.

3.4 The fourth mismatch: incidents inside the labs do not reach outside observers

When something goes wrong inside a frontier lab — a misuse case caught in red-team testing, an unexpected capability surfaced late in evaluation, a deployment incident in the wild — the information first exists at rung 3, inside the lab, in the internal logs. Outside observers — the AI Safety Institute, the EU AI Office, the press, the public — depend on what the lab chooses to disclose, which arrives wrapped in rung-1 framing (the lab's authority on the meaning of its own incidents).

The loop appears closed. Labs do publish model cards and post-mortems. But the closure is in form only. The rung-3 evidence that exists inside is filtered through rung-1 disclosure before reaching the controllers whose process models it would correct. This is the same asymmetric pattern as the capital/lab loop, in a different boundary. It is the structural reason why the press keeps discovering things about labs that the labs already knew.

3.5 The fifth mismatch: antitrust looks at mergers; concentration happens in contracts

Antitrust authorities — the FTC and DoJ in the United States, DG-COMP in the European Union, SAMR in China — evaluate market concentration through merger review and conduct cases. Their process models track mergers and acquisitions, dominant-firm behaviour, and pricing patterns. The cycle on which they operate is years. The threshold above which a transaction triggers review is set in statute.

Concentration in this market, however, does not flow primarily through mergers. It flows through long-duration exclusive compute-allocation contracts between hyperscale providers and frontier laboratories, through bundled tying arrangements between chip designers and their largest buyers, through grandfathered grid-interconnection rights at data-centre sites. None of these trigger merger review, because none of them are mergers. They happen at the contract level, on a quarters cycle, and they foreclose entry by smaller actors before any antitrust authority has cause to look.

This is cross-loop rung imposition in the third dangerous-patterns sense. The rung-2 forum (statutory antitrust review) is being asked to adjudicate a dispute generated below its threshold and beneath its cadence, and it cannot. The forum is the wrong shape for the adjudication. The structural fix is not to make antitrust faster; it is to give the forum a different statutory mandate that captures contract-level conduct.

3.6 What having five mismatches in hand changes

A strategic essay can speak generally about "the AI industry's governance problem." Naming five specific mismatches changes the conversation in three ways.

It changes what is observable. Once we know that the lab/capital loop is asymmetric, we can look for the symptom: capability claims that survive funding rounds but not independent evaluation. We can know what would falsify the diagnosis, and what would confirm it.

It changes what counts as a remedy. A general call for "better AI governance" is not a remedy. A specific structural intervention that turns an asymmetric loop into a symmetric one — say, by inserting a statutory pre-deployment evaluation channel that does not pass through the lab's own benchmark selection — is a remedy. The intervention does not depend on anyone changing their mind; it changes the architecture.

It changes what survives political turbulence. Personal commitments and voluntary frameworks last as long as the people and the incentives that sustain them. Structural changes — new statutes, new institutions, new contractual rules — last as long as the political settlement that created them, which is typically much longer. The remedies the engineering literature catalogues are the structural kind.

 Key Finding — the five rung mismatches

- Tagging every flow in the AI control structure with its justificatory rung surfaces five specific structural mismatches:
- **M1.** Lab safety frameworks claim rung-6 deliberative legitimacy but operate at rung-1 self-binding (Pattern B — claimed-rung inflation).
 - **M2.** Capital flows downward at rung-3 due-diligence; capability evidence flows back at rung-1 lab-selected benchmarks (Pattern A — asymmetric loop).
 - **M3.** Labour-policy cycle (years) is mismatched to capability-deployment cycle (months); the rung-3/4 corrective channel arrives after the cohort it would have helped is already displaced.
 - **M4.** Incident evidence exists at rung-3 inside labs but reaches outside observers wrapped in rung-1 disclosure (Pattern A again, at a different boundary).
 - **M5.** Antitrust operates at rung-2 with a years cycle and a merger-review threshold; concentration happens at rung-3 in long-duration contracts on a quarters cycle and below the threshold (Pattern C — cross-loop rung imposition).
- These five mismatches account for most of the loss scenarios developed in §4 and motivate most of the remedies in §5.

We turn now to what specifically goes wrong when these mismatches are not addressed.

31.6 4. The nine ways the system can fail

A loss scenario, in the engineering vocabulary, is a description of how a particular bad outcome arises from particular failures inside the structure. It names the controllers whose process models go wrong, the feedback channels that go missing or get filtered, the controlled processes that drift into states from which exit is costly, and the external disturbances that finish the job. A loss scenario is, in effect, a small story about how the system produces harm, told in the language of its own architecture.

The full analysis identifies nine such scenarios for the AI system. They are not predictions. Each of them describes a pathway by which a specific identified harm could be produced, traceable to specific mechanisms one could intervene on. Some pathways are already partly running; others are latent. None is inevitable. The analytic purpose of writing them down is to attach each remedy to a specific failure pathway, so that the policy debate has a structural rather than a rhetorical anchor.

4.1 Concentration entrenches at the infrastructure layer

The first scenario is about how AI's economic gains end up concentrated in the hands of a small number of firms — but not through the obvious channel one would expect.

Antitrust authorities, as we saw in section 3.5, look at mergers. They look in the wrong place. The actual mechanism of concentration is contractual. A hyperscale cloud provider signs a multi-year exclusive capacity agreement with a frontier laboratory. A chip designer commits its leading-edge supply, a year ahead, to its three largest buyers. A foundry pre-allocates capacity windows. A data centre

receives a grid-interconnection right while smaller applicants sit in a queue. Each of these is a single transaction that, on its own, looks like ordinary commerce. None of them crosses a merger threshold. None of them appears in any one controller's process model as concentration.

The cumulative effect is foreclosure. Smaller laboratories, open-weights publishers, application companies that depend on non-frontier compute, all find that the affordable supply has been contracted away in advance. Entry becomes a financial impossibility. The controlled process — the structure of the AI infrastructure market — moves to a state from which exit is costly: the contracts have already been signed; the grid interconnections have already been allocated; the leading-edge fab capacity is forward-booked for years.

The corrective channel exists, in principle. Concentration ratios, prices, and entry rates are measured. But they are measured at the slow rate at which antitrust evidence accumulates — over years — while the foreclosing transactions happen quarterly. By the time the evidence is strong enough to act on, the architecture is already locked in. The remedy class is cadence-matched controls: rules that fire on the transaction cycle, not on the evidence cycle. The specific remedies in the catalogue are reservation requirements on leading-edge supply (a defined fraction must be sold to non-frontier customers), an expedited review window on long-duration contracts during defined capability-jump moments, and a merger-review reform that captures contract-level foreclosure in the same statutory net as ordinary M&A.

4.2 Lab safety frameworks become ceremonial

The second scenario follows directly from the first and second mismatches, working together. Capital pressures laboratories to release on a fast cadence to defend market share. The laboratories' safety frameworks claim deliberative legitimacy. The frameworks are amended between releases, in response to commercial pressure, to allow the next release. The amendment process is internal; the legitimacy claim is external. The gap widens each cycle.

The damage here is not principally that any particular release is unsafe. It is that the legitimacy claim erodes. Each visible amendment to a safety threshold, each evaluation that produces a result the lab finds inconvenient and reframes, each model that ships with caveats no external party could have written, drains the rung-6 reservoir the labs implicitly drew on when they published their first safety frameworks. At some point the public notices that the frameworks are advisory, the amendments are unilateral, and the claims are aspirational. When that happens, the labs can no longer issue legitimate safety assurances of any kind.

The remedy class here works in two directions at once. It lowers the claim: the lab admits, structurally, that its frameworks are self-binding rung-1 commitments — useful, but limited. It raises the operation: an external evaluator with statutory authority enters the loop. The specific remedies in the catalogue are mandatory pre-deployment evaluation by an independent body (the AI Safety Institutes are early examples, operating today as voluntary partners but candidates for statutory upgrade), structural separation of safety governance from commercial governance inside labs, and a capital-structure mission-constraint mechanism that gives a public agent — a state attorney general, a regulator — standing to enforce the mission constraint, rather than relying on shareholder litigation that will never come.

4.3 Workers cannot keep up with capability deployment

The third scenario is about the labour market and runs through the third mismatch. Capability deploys on a months cycle — model release, application integration, deployment in customer workflow. Labour ministries' instruments — retraining programmes, curriculum updates, unemployment-insurance schemes — operate on a years cycle. Cohort-specific damage to early-career cognitive workers — the hardest-hit group because their tasks are most substitutable — is invisible in the headline employment statistics that ministries track.

The corrective channel is structurally late. By the time the employment data shows a problem, the cohort the data described has already entered the labour force, found that the entry-level roles they trained for are gone, and either dropped out, accepted underemployment, or organised politically. The political channel operates at rung 0 (electoral backlash) on a fast timescale that the rung-3 statistical channel cannot match. The rung-0 channel arrives, in other words, exactly when the rung-3 channel was supposed to have prevented its arrival.

The remedy class is automatic stabilisers indexed to a fast benchmark of task displacement, not to headline employment. Indexed unemployment insurance, indexed earned-income tax credit, indexed retraining credits: instruments whose triggers fire without requiring legislative confirmation per cycle. To these add cohort-specific transition programmes for the early-career group — recognising that the affected cohort changes year by year as substitution sweeps through different occupations — and portable benefits that decouple healthcare and pension continuity from continuous employment with a single employer. None of these require

new theory; they require the index, the statutory trigger, and the political decision to commit ahead of the displacement rather than after.

4.4 The productivity surplus accrues to capital and stays there

The fourth scenario is the distributional one. Suppose the productivity gain from AI is real — every careful estimate suggests it is, on the order of fractions of a percentage point per year of GDP growth, possibly more. The question is who receives it.

The mechanisms by which a productivity gain reaches workers, as wages, depend on labour bargaining power. The mechanisms by which it reaches consumers, as lower prices, depend on competitive markets. The mechanisms by which it reaches the public, as taxation, depend on a fiscal architecture calibrated to the new factor shares. None of these is automatic. If labour's bargaining power falls (because AI substitutes for substitutable cognitive tasks), the wage channel narrows. If markets concentrate (scenario 4.1), the price channel narrows. If tax instruments calibrated to corporate income do not adjust for the shift toward AI-as-capital that flows to less-taxed bases, the public channel narrows.

The structural diagnosis is that no controller in the existing system is positioned to act on cross-sectoral productivity-gain capture as such. Tax authorities track corporate income, capital gains, and payroll — three categories that together miss the factor-share shift. There is, in the engineering vocabulary, a controller absent from the loop. The remedy class is to add one: a transmission mechanism, of which the catalogue lists four candidates — a productivity-indexed dividend (analogue: Alaska Permanent Fund, but indexed to AI productivity rather than oil), a value-added tax on AI inference with hypothecated dividend, an expanded earned-income tax credit indexed to the displacement benchmark, and sovereign-wealth participation in AI infrastructure with public dividend.

These are not interchangeable. The productivity-indexed dividend is conceptually clean but politically the hardest to land. The inference VAT is operationally simpler because it taxes the visible economic flow. The EITC expansion is the most familiar because the instrument already exists. Sovereign-wealth participation requires a fund and a governance mechanism; some jurisdictions have one (Norway, Singapore, the Gulf), the United States does not.

The point is not which one is correct; it is that without at least one of them, the productivity surplus has no transmission channel to labour, consumers, or the public. It will accrue to capital not because capital is rapacious but because capital is the only controller in the loop with a defined claim on it.

4.5 The open-weights price-setting mechanism collapses

The fifth scenario concerns a controller most strategic essays do not name as a controller at all: the open-weights ecosystem. Chinese laboratories and a handful of Western publishers release model weights free for download. The releases serve many purposes, but at the system level they perform one specific function: they set a price ceiling on the model layer. Frontier laboratories cannot price closed-weight inference far above what an adequately-equipped competitor can do with a recent open-weight release. The ceiling exists.

If the ceiling collapses, the model layer becomes a place where a small number of closed-weight providers can extract whatever the market will bear. The collapse can happen in three ways. Western governments could regulate open-weights releases out of existence — well-intentioned, on misuse-risk grounds, without recognising that the regulation also removes the price-setting mechanism. Capital could flee the few Western open-weights labs whose business is harder than their closed-weights competitors', leaving Chinese labs as the only credible open-weights suppliers — at which point the political acceptability of using their weights becomes a separate concern. Or Chinese labs could shift strategy and stop releasing open weights, for reasons of their own, and no one outside has the leverage to insist.

The diagnostic point is that the open-weights ecosystem performs a function that no controller is assigned to protect. Regulators treat open weights primarily as a misuse risk; market participants treat them as a competitive threat or competitive crutch. No body publishes a measure of the frontier-versus-open-weights gap as a system-level state variable that regulators are obliged to consider. The remedy class is what the dangerous-patterns catalogue calls domain-appropriate rung containment: route open-weights questions to a forum that holds both misuse risk and price-setting in superposition, and require any rule affecting open weights to document its impact on the gap. Concretely: publish the gap, calibrate misuse-risk rules narrowly enough to preserve the ecosystem below those thresholds, and treat the ≤ 12 -month gap as an explicit policy goal rather than an emergent outcome.

4.6 Compute competition escalates to kinetic conflict

The sixth scenario is the most extreme. Compute is geographically concentrated. Leading-edge logic chips come overwhelmingly from a single foundry on a single island. Hyperscale capacity is concentrated across a small number of operators in a small number of regions. Energy and grid interconnections are concentrated where electricity is cheap. The geography of AI infrastructure is fragile in ways the geography of, say, the global oil supply once was — a small number of physical chokepoints whose disruption would stop the flow.

National security apparatuses on both sides of the US–PRC relationship watch each other's compute build-out. Each side's process model treats the other's build-out as preparation for adversarial use. The defensive interpretation — each side is building because each side fears being shut out — is filtered out on both sides by the rung-1 strategic identity that does not admit the other's defensive logic. There is no confidence-building channel at the AI-infrastructure level. The only available channels are coercive (export controls, sanctions, designations) and rhetorical (state speeches).

A scenario in which a regional crisis, perhaps unrelated to AI itself, transmits directly into the AI-infrastructure layer is not implausible. A blockade or a strike that disrupts foundry output. A sanctions package that pulls one side into commercial escalation. A reciprocal designation that classifies a category of research on both sides simultaneously. Each step is small; the cumulative path is from rung-2 rule-making through rung-0 coercion toward kinetic action.

The remedy class here is the hardest in the catalogue and the most important. An independent inspection or observation regime for AI-infrastructure deployments — analogous to the IAEA for nuclear material, the OPCW for chemical weapons, or the now-defunct Open Skies regime — would provide a rung-3 verification channel that bypasses the rung-1 filter on both sides simultaneously. It would not solve the conflict; it would create a forum in which evidence about the system's actual state could enter both sides' decision-making, reducing the chance that either side acts on a process-model error. The argument for such a regime, in the engineering vocabulary, is structural rather than rhetorical: the conflict is being driven by mutual process-model error, and only a r3 channel external to both sides can correct mutual error symmetrically.

This will not be built quickly. It depends on a precondition — shared technical vocabulary about what AI-infrastructure verification could even mean — that current track-1.5 dialogues are only beginning to produce. It depends, further, on geographic diversification of fab capacity, so that verification has somewhere to look that is not already monopolised. The sequencing is therefore: track-1.5 dialogue first (months), fab diversification in parallel (years), inspection regime later (decade).

4.7 The Chinese mirror

The seventh scenario is a specific consequence of the second, mirrored across the bloc boundary. Inside the Chinese state-lab system, the same Pattern A asymmetric loop operates that we identified in section 3.2 between US capital and US labs. The party centre receives capability evidence from PRC labs filtered through r1 mission-conformant disclosure; PRC labs cannot, in their own internal process, fully separate r3 evidence from r1 authority because the r1 structure is constitutive rather than contingent.

The damage is symmetric to what happens in the US case, with one critical difference. The US side is open enough that external r3 channels (independent benchmarks, AISI evaluations, academic critique) can in principle reach the apex. Whether they do depends on whether structural separation is enforceable. The Chinese side has no equivalent external channel: the structure that produces the asymmetry is the same structure that excludes external correction.

The implication is uncomfortable but operationally important. Unilateral remedies on the US side can produce evidence that PRC analysts read; capability claims that survive Western r3 evaluation also survive their own internal r1 wrapping. An externalised r3 channel becomes a correction channel for both sides at once. But the only remedy that works directly on the PRC side is a bilateral one — an inspection regime of the kind described in scenario 4.6, which exists outside both national hierarchies. This is one of several reasons that the bilateral remedies are higher-leverage than they appear.

4.8 The information environment degrades faster than authentication scales

The eighth scenario is about epistemics. The volume of AI-generated content already exceeds the capacity of human readers to authenticate it manually and is approaching the capacity of automated detection systems to track it. Provenance signals exist — cryptographic content credentials, watermarking — but their adoption is voluntary, and adversarial actors strip them at low cost. Information platforms whose business model is engagement-driven amplification have no incentive to favour verified content over high-engagement AI content; moderation, if it happens, imposes a unilateral cost on a single platform in a multi-platform competition for attention.

The structural diagnosis is that the controlled process — the information environment — is moving to a state in which trust in any signal degrades trust in all signals, and from that state recovery is difficult. The corrective feedback channel (measurement of provenance failure rates) does not exist at the aggregate level; only individual incidents are tracked. The controller responsible (information regulators) tracks takedowns, not aggregate state.

The remedy class is to move provenance from the platform layer (where it can be stripped) to the capture layer (camera, OS) and the transmission layer (model output API, where it can be cryptographically attested). To this add a statutory regime for aggregate measurement of provenance-failure rate, so that the question "is the information environment getting worse, by how much" can be answered. Watermarking alone is insufficient; the analogy is with seatbelt laws — voluntary mechanisms have a ceiling; mandatory mechanisms with measurement do not.

4.9 Rules calibrated to capability tiers already obsolete

The ninth scenario is the regulatory cadence one, mirrored. Legislatures pass AI rules calibrated to a defined capability threshold — typically a specific compute or capability level named in the statute. Algorithmic-efficiency improvements move the threshold under the rule by a factor of three or four every year. The regulator's process model treats the threshold as a fixed boundary; in reality it is a moving target. Five years after the rule passes, the rule applies to capabilities the intended scope did not anticipate, while capabilities that match the intended scope have moved into a less-regulated regime.

The corrective mechanism — recalibration — depends on re-passage of primary legislation, which requires legislative attention that is itself rationed. By the time the rule is re-passed, the threshold has moved again. The remedy class is to write rules whose thresholds move with a published benchmark, and to grant a standing body the authority to recalibrate without primary legislation, subject to a notice-and-comment process. This is not novel — securities regulators in many jurisdictions recalibrate margin and capital requirements without primary legislation — but it requires a deliberate political decision to delegate, and the political appetite for delegation varies by jurisdiction.

4.10 What having nine scenarios gives us

Nine scenarios is more than a strategic essay would provide and fewer than the underlying analysis admits. The number is roughly right because each scenario corresponds to one identifiable failure mechanism. Together they cover all the unsafe control actions catalogued in the analysis; each remedy class points at specific UCAs whose failure pathway it interrupts.

The nine scenarios are not independent. Scenario 4.1 (concentration) makes scenario 4.4 (surplus capture) worse. Scenario 4.2 (ceremonial safety) makes scenario 4.8 (epistemic degradation) worse. Scenarios 4.6 and 4.7 (the geopolitical pair) interact symmetrically with each other. The remedies, in turn, are not independent: as we will see in the next section, several of them share preconditions, several form sequencing chains, and the structurally hardest remedies are exactly those whose absence would cost the most to live with.

31.7 5. What can be done

Twenty-two specific remedies follow from the analysis. They divide into five families, each addressed at one or more of the nine loss scenarios. The family-level summary first; then the sequencing.

The first family is measurement and feedback channels. Several of the most damaging mismatches occur because the system lacks an honest measurement of its own state. Regulators do not know the frontier-versus-open-weights gap; nobody publishes the aggregate provenance-failure rate; labour ministries track headline employment but not task-by-task displacement. Four remedies fill these channels: a quarterly assessment of the open-weights gap, a quarterly aggregate provenance-failure measurement, a task-displacement benchmark indexed by occupation cohort, and a track-1.5 dialogue process that builds shared technical vocabulary between the US and PRC technical communities about what AI-infrastructure verification could mean. Each of these is operationally close to existing instruments and politically modest. Each of them, however, is a precondition for remedies in later families.

The second family is rules built on the new measurements. Once the open-weights gap is published, a targeted misuse-risk regime can be calibrated narrowly enough to preserve the ecosystem below those thresholds. Once the provenance-failure rate is measured, mandatory provenance becomes specific rather than gestural. Once the task-displacement benchmark exists, automatic stabilisers — unemployment insurance, retraining credits, EITC top-ups — can be indexed to it, so that triggers fire on the fast cycle that the disturbance arrives on. Reservation requirements on leading-edge supply, capability-tier rules indexed to a moving benchmark with sunset clauses, and a standing recalibration authority for AI rules each address a specific identified pathway. A mandatory pre-

deployment evaluation regime — turning the AI Safety Institutes from voluntary partners into statutory ones — closes the asymmetric loop between labs and capital.

The third family is structural remedies. These are the interventions that change the architecture rather than the operating parameters. Merger-review reform that captures contract-level foreclosure. Structural separation of safety governance from commercial governance inside frontier laboratories. A statutory public-benefit-corporation class with mission constraints enforceable by a public agent, not by shareholder litigation. Cohort-specific transition programmes for early-career cognitive workers. Portable benefits that decouple healthcare and pension continuity from continuous employment. A value-added tax on inference, hypothecated to a public dividend. Sovereign-wealth participation in AI infrastructure, where the jurisdiction has a fund. These are harder to land politically and slower to take effect, but each of them addresses a structural failure that no operating- parameter adjustment can fix.

The fourth family is the slow geopolitical remedies. Three of the most consequential remedies operate on a decade horizon because their preconditions are physical or diplomatic. Geographic diversification of fab capacity — bringing leading- edge logic-chip production to a strategically meaningful share outside the present chokepoint — is funded but slow; the constraint is sustained political commitment across electoral cycles. The independent inspection or observation regime for AI infrastructure depends on shared technical vocabulary, which the track-1.5 process slowly produces, and on geographic diversification, so that verification has somewhere to look that is not already monopolised. A productivity-indexed dividend, of the Alaska-Permanent-Fund kind but indexed to AI productivity, operates on a decade horizon because it requires both a fiscal instrument (which the inference VAT could provide) and a governance mechanism (which sovereign wealth could provide). Each of these is the long-cycle remedy whose absence the others cannot make up for.

The fifth family is the politically novel ones. Two remedies do not fit cleanly into any of the four prior categories. An expedited capability-jump merger-review window, granting antitrust authorities a temporary hold on long-duration compute contracts during defined capability-jump moments, is a small intervention with a precise function but no real precedent. An employer-independent benefits regime is operationally familiar but politically contested in jurisdictions where employer-tied benefits define the post-war social contract. These are listed separately not because they are less important, but because they require political adoption on terms the system does not yet provide.

5.1 The four-phase sequence

The 22 remedies do not all enter the world at once. They enter in a specific order constrained by the dependency graph among them, and the analysis derives a four-phase implementation plan.

The first phase, months one through eighteen, is measurement. The four feedback-channel remedies — the open-weights gap, the provenance-failure rate, the task-displacement benchmark, the track-1.5 dialogue — can begin without statutory change. An AI Safety Institute could initiate the open-weights gap publication tomorrow, as a non-binding research output. A national statistics office could establish a provenance-measurement programme on existing authority. A labour-statistics office could publish a task-displacement index using existing labour data and AI task mapping. The track-1.5 dialogue is already running in several forums; sustaining and broadening it requires nothing more than philanthropic funding and continued attention. By month twelve, each of these channels will have produced a first batch of evidence; by month eighteen, the question of whether to formalise each one statutorily will have a real answer.

The second phase, months twelve through forty-eight, builds rules on the new feedback channels. This is where most of the operationally familiar remedies sit: indexed automatic stabilisers, indexed EITC, mandatory provenance, reservation requirements on leading-edge supply, capability-tier rules with sunset clauses, a recalibration authority, mandatory pre- deployment evaluation, a targeted misuse-risk regime for open weights. Each of these is a piece of legislation or a regulation written against a measurement that the first phase produced. Each is operationally close to existing instruments. Each requires political work but not novel theory.

The third phase, months thirty-six through eighty-four, is the structural remedy phase. Merger-review reform, structural separation of lab safety and commercial governance, the statutory PBC class with public-agent enforcement, cohort-specific transitions, portable benefits, the inference VAT with hypothecated dividend, sovereign-wealth participation. These require political capital that the first two phases will have generated, and an empirical record that shows whether the operating-parameter remedies were sufficient on their own (the analysis suggests they will not be, but the evidence by month thirty-six will be specific). Each of these is structurally novel and politically harder than the second-phase remedies; each addresses a pathway the second-phase remedies cannot close.

The fourth phase, months sixty through one hundred and forty-four, is the long-cycle remedy phase. Fab diversification will already be running in parallel; what completes in this phase is the inspection regime, the productivity-indexed dividend, and the expedited capability-jump review window if it has not already been adopted. These are the remedies whose absence imposes the largest tail cost

— concentration unbroken, geopolitical escalation unverified, surplus capture untransmitted — and whose adoption is hardest because their political horizon exceeds electoral cycles.

The decision-memo companion to this chapter sets out the dependency graph in full and identifies, at each phase boundary, the question that the prior phase must have answered before the next phase becomes feasible. By month twelve: did anyone publish the open-weights gap? By month eighteen: is the displacement benchmark stable? By month thirty-six: have at least three of the eight phase-two rules been adopted? At each decision point, the analysis describes both what success looks like and what failure implies for downstream sequencing. This is what distinguishes a planning artefact from a wish list.

**The four-phase sequence at a glance**

Phase	Months	Theme	Representative remedies
1	1-18	Measurement and feedback channels	Open-weights gap publication · provenance-failure measurement · task-displacement benchmark · track-1.5 dialogue
2	12-48	Rules built on phase-1 measurements	Mandatory pre-deployment evaluation · indexed automatic stabilisers · indexed EITC · mandatory provenance · reservation requirements · capability-tier rules with sunset · recalibration authority · targeted misuse-risk regime
3	36-84	Structural remedies	Merger-review reform for contract foreclosure · structural separation of safety governance · statutory PBC class · cohort-specific transitions · portable benefits · inference VAT · sovereign-wealth participation
4	60-144	Long-cycle remedies	Geographic fab diversification · independent inspection regime · productivity-indexed dividend · expedited capability-jump review

Each phase's gating decision determines whether the next phase is structurally feasible. Phase 1 is reversible and operationally familiar; phase 4 is largely irreversible once started.

5.2 The non-state actor's leverage

The remedies are not, as a class, things that one private actor can do alone. They are political and statutory. But several of them have entry points where a non-state actor — a foundation, a research institute, a coalition of application companies, a group of labour organisations — can move first and, by moving first, make the statutory work tractable.

Phase-one measurement is the cleanest example. Each of the four channels can be initiated as a non-binding publication; once the data exists, the statutory mandate becomes specific rather than gestural. Drafting model-bill language for phase-two rules is work that legal-policy institutes do well, and adoption rates of model-bill language are unusually high in jurisdictions where the legislature is overstretched. Coalition-building among open-weights-dependent companies is work that market participants can do without waiting for political alignment. Track-1.5 dialogue is the canonical philanthropic-funded geopolitical intervention; it is not a substitute for state action, but it produces the vocabulary state action requires.

The leverage point for any non-state actor, then, is to ask which of the 22 remedies in which phase has an entry point that matches the actor's position, and to commit to the entry-point work for at least one full decision cycle. The dependency graph makes the leverage points explicit; it also makes clear which work, however well-meaning, is not on the critical path.

5.3 Reversibility

A final consideration constrains the remedy set. Most of the proposals are reversible at moderate cost within their first cycle. Statistical publications can be discontinued. Indexed triggers can be recalibrated. Pre-deployment evaluation requirements can be relaxed. Reservation requirements can be adjusted. This is structurally important because it means a risk-averse stakeholder can support phase-one and most of phase-two with limited downside. The risk of getting the rule wrong is not the risk of locking in an error.

The exceptions are the structural and slow-cycle remedies. Once established, an inspection regime is hard to withdraw because withdrawal would itself signal an escalation posture. Sovereign-wealth participation in AI infrastructure cannot be unwound without selling at unfavourable timing. Fab diversification is sunk capital. A productivity-indexed dividend, once it becomes a household income source, is politically irreversible in practice. A loss-averse stakeholder considering tail risks — the kinetic-escalation risk in scenario 4.6, the cohort-displacement risk in scenario 4.3 — would weight these irreversible remedies more heavily despite their cost. A risk-averse stakeholder weights phase-one and phase-two preferentially.

The point of putting reversibility in the analysis explicitly is that different stakeholders, with different risk tolerances, will rank the same remedies differently. The analysis does not adjudicate the ranking. It surfaces the trade-off so that the adjudication can happen on terms a deliberative process can accept.

31.8 6. What the analysis does not do

A method that produces twenty-two specific remedies is also a method that risks being read as a forecast or a manifesto. It is neither, and it is worth being explicit about what the limits are before closing.

It does not predict which of the nine loss scenarios will occur, or in what combination, or on what timescale. The scenarios are possibility statements: each one describes a pathway from the current structure to a specific bad outcome, and the structure does in fact have those pathways. Whether the outcomes occur depends on which control actions are issued, by which controllers, in which sequence, and with what feedback. The analysis maps the terrain; it does not tell the reader how the journey ends.

It does not settle value disputes. The remedies are derived from the loss scenarios, and the loss scenarios are derived from stakeholder-named losses. Different stakeholders would name different losses. A reader who does not consider concentrated ownership of the AI substrate to be a loss — perhaps because they believe the firms in question are well-governed and well-intentioned — will not find scenario 4.1 a problem and will not weight remedies R-1 through R-3 highly. The analysis cannot adjudicate that disagreement; it can only show that if the loss is taken to be a loss, then certain remedies follow with specific properties. The reader's own values determine which "ifs" are live.

It does not tell any particular government or firm what to do. Decisions belong to controllers; the analysis describes what each controller could in principle do and what consequences each choice would have for the controlled processes. A decision-memo addressed to a specific stakeholder — say, a foundation allocating its grant budget over five years, or a labour union prioritising its lobbying agenda, or a cabinet office choosing between policy packages — would be a different document, derived from this analysis but committing to that stakeholder's goals and constraints.

It does not, finally, exhaust the possible failure modes of the AI system. The analysis is a first cycle in the engineering sense: every loss scenario is one that the controller / control- action / feedback / process-model decomposition can surface. There may well be failure modes that this decomposition cannot see — failures that emerge from interactions across multiple loops simultaneously, failures that depend on non-controllers the analysis treats as environment, failures whose causes lie below the threshold of the current decomposition. A second cycle, with refined boundaries and a tighter decomposition, would find some of these. The analysis we have is bounded in the way every engineering analysis is bounded: by the granularity at which the system was modelled.

31.9 7. Coda: the question reformulated

We started with three questions: about jobs, about power, about geopolitics. The structured analysis allows us to reformulate each of them in a way that points at action.

The jobs question — will AI cause mass joblessness or freer prosperity? — is, structurally, the question of whether the labour-policy cycle (years) will be brought into alignment with the capability-deployment cycle (months) before the cohort displacement that the mismatch produces becomes irreversible. The technology produces the surplus; the institutions decide whether the surplus reaches workers, consumers, the public, or only capital. The analysis identifies the specific instruments that would close the cycle gap (indexed automatic stabilisers, a displacement benchmark, cohort-specific transitions, portable benefits) and the specific instruments that would transmit the surplus (the inference VAT, the productivity-indexed dividend, the EITC top-up, sovereign-wealth participation). Whether any of these is adopted is a political choice made under known trade-offs — not, as the strategic essay would have it, an open question with a wide range of plausible answers.

The power question — will a few firms own everything? — becomes, structurally, the question of whether the contract-level foreclosure mechanism currently moving compute infrastructure into a small number of hands is captured by an antitrust regime

calibrated to the right cadence and the right threshold. The analysis identifies merger-review reform, reservation requirements on leading-edge supply, and an expedited capability- jump review window as the specific instruments. It also identifies the open-weights ecosystem as a price-setting mechanism whose suppression would be the worst single thing that could happen at the model layer, and identifies the specific regulatory pathway (well-meaning misuse-risk regulation that fails to recognise the price-setting function) by which the suppression could occur.

The geopolitics question — will the US-PRC AI competition escalate? — becomes, structurally, the question of whether a rung-3 verification channel can be installed between the two blocs before mutual process-model error converts a series of rung-2 coercive moves into a kinetic exchange. The analysis identifies track-1.5 dialogue as the precondition, fab diversification as the parallel constraint, and an infrastructure-inspection regime as the long-cycle remedy. It also identifies the symmetric Pattern-A failure on both sides that makes mutual error structurally likely, and explains why unilateral US-side remedies (especially mandatory pre-deployment evaluation) produce evidence that bypasses the rung-1 filter on the PRC side as well, making them more valuable than they appear when each is considered in isolation.

None of these reformulations resolves the questions. Each of them turns a forecasting question into a structural question, and a structural question is one a deliberative process can act on. A political community that refuses to take action on a forecasting question — because the future is uncertain, because experts disagree, because the right thing to do is unclear — can take action on a structural question, because the structural question is not "what will happen" but "what feedback channels are missing and what institutions are needed to install them." The first question admits indefinite postponement; the second does not.

That, finally, is what an engineering reading offers that an essay does not. The engineering reading produces a list of specific things that could be done, in a specific order, by specific actors, with specific costs and specific reversibility profiles. It does not promise that any of them will be done. It does promise that, if one of them is done, the system will respond in the way the analysis predicts — because the prediction is not a guess about behaviour but a derivation from structure.

Whether the structures change is, in the end, the same question as whether the political communities affected by AI choose to change them. The technology will not make the choice. Capital will not make the choice. The labs will not make the choice. The choice has to be made by the controllers whose process models the structural analysis maps and whose feedback channels the remedies repair. And the choice can only be made if the controllers see the structure clearly enough to know which levers exist.

This chapter has tried to make the structure visible. The rest is up to whoever is reading.

32. Control Structures in Social Systems

32.1 The Universal Control Pattern

Every social system has a hierarchical control structure with the same basic shape. Authority flows downward; feedback — ideally — flows upward. STPA analyses this structure by asking five questions of every system:

1. **Feedback richness** — how many independent channels carry information from the base to the top, and how accurately do they represent reality?
2. **Self-sealing tendency** — can the process model at the top be corrected by information from below, or does it filter out contradictory signals?
3. **Accountability voids** — where is the controller also the interested party, structurally preventing independent oversight?
4. **Circuit breakers** — what mechanisms exist to interrupt a harmful escalating loop before it causes irreversible damage?
5. **Rung match** (social-systems extension; see [Justificatory Rungs](#)) — at what justificatory rungs do the loop's control action and feedback operate, and is the loop symmetric? Does the system's claimed rung match its operating rung?

The four levels of the universal pattern are:

Level	Role
Supreme Authority	Issues doctrine, law, strategy, or policy
Intermediate Controllers	Interpret and transmit downward; aggregate and filter upward
Local Community	Enforces norms through social mechanisms; primary source of ground-truth feedback
Individual	Internalises norms; self-regulates; ultimate recipient of system outcomes

What varies across systems is: who occupies each level, what mechanisms connect them, and how honest the upward flow is.

Canonical reference

The technique definitions — the five diagnostic questions and the four dangerous patterns — live in [knowledge/se-techniques/control-structures/](#) and (for the rung extension) [knowledge/se-techniques/justification-rungs/](#). The per-system profiles for all ten social systems, with claimed-rung and operating-rung tags, live in [knowledge/system-catalogues/social-systems/cross-system/control-structure-profiles.md](#). The cross-system rung comparison is in [knowledge/system-catalogues/social-systems/cross-system/justification-rungs-by-system.md](#).

32.2 The Ten Systems: Control Structure Profiles

For each of the ten social systems, the same four-level pattern is filled in with the specific controllers, intermediate actors, local communities, and individuals that occupy each role, and each system is rated on the four diagnostic questions.

The complete profiles — Kingdom, Republic, Theocracy, One-Party State, Corporation, University, Military, Religion (Church), Family, Verein — and the cross-system comparison table are in [knowledge/system-catalogues/social-systems/cross-system/control-structure-profiles.md](#). A short summary of the pattern of feedback richness, self-sealing tendency, accountability voids, and circuit breakers across the ten systems appears there as a single table.

The key observation the profiles support is that the same control-structure weakness recurs in structurally identical forms across very different systems: the monarch judging the Crown's own claims, the bishop investigating his own clergy, the commanding officer prosecuting his own troops, the board setting the pay of executives who chose the board, the peer reviewing his own competitor's work. The institutional vocabulary differs; the structural pattern is the same.

32.3 The Four Most Dangerous Structural Patterns

Across all systems, four patterns recur wherever the most serious harm occurs. The first three are domain-general; the fourth is the social-systems extension introduced in [Justificatory Rungs](#).

Pattern 1: The Accountability Void. The controller and the interested party are the same entity. No external actor has authority to review the decision. The remedy in every case is the same: structurally separate the controller from the interested party, and give the separating body genuine authority.

Pattern 2: The Self-Sealing Process Model. The top of the control structure populates its own process model — its internal representation of reality — from sources it controls. Contradictory information is excluded by design (doctrine), by incentive (career risk for bearers of bad news), or by structure (classification, privacy norms). The model diverges from reality while the authority acts with increasing confidence. The remedy: establish independent information channels with direct access to the apex, reporting to a body outside the hierarchy.

Pattern 3: The Proxy Metric. The system measures something that is correlated with its actual goal, attaches resources and careers to that measure, and then watches as the system reorients toward producing the measure rather than the goal. Exam scores replace education. Quarterly earnings replace sustainable value creation. Publication counts replace scientific knowledge. Compliance reports replace operational readiness. The remedy: continuously examine what is being measured and whether the feedback loop connects to the goal or to a proxy, and restructure the measurement accordingly.

Pattern 4: Rung Asymmetry. The downward control action and the upward feedback on the same loop operate at different justificatory rungs. The controller transmits commands at rung 1 (authority, tradition) but the corrective signals it would need to register arrive — or are expected to arrive — at rung 3 or higher (empirical, audited, replicated). The controller's process model classifies the higher-rung feedback as out-of-band and filters it out. Instances: doctrinal authorities receiving abuse reports as rung-1 hostility rather than rung-3 evidence; autocracies receiving independent press reports as treason rather than information; family systems treating children's developmental signals as defiance rather than health data. The remedy: insert an independent feedback channel that operates at rung 3 (or higher), routes around the rung-1 filter, and reaches the controller's process model directly. Most of the canonical architectural remedies in the next chapter are rung-elevation moves of this kind.

These four patterns are the recurring targets of the architectural remedies in the next chapter. The technique-level statements — with the failure mode each corresponds to in a generic control loop, and the remedy preconditions — are in [knowledge/se-techniques/control-structures/dangerous-patterns.md](#) (Patterns 1–4) and [knowledge/se-techniques/justification-rungs/dangerous-mismatches.md](#) (Pattern 4 in detail, plus two further rung-specific patterns: Claimed-Rung Inflation and Cross-Loop Rung Imposition).

33. Architectural Remedies

33.1 The Analysis Generates the Fix

The introduction to this book used a secondary school to illustrate how structural analysis works: once you have traced which feedback loop is misconfigured and why, you know exactly which connection to change. Finland changed those connections and the outcomes changed with them.

This chapter applies the same logic across five social systems. For each, STPA identifies the unsafe control actions, traces their causal factors, and derives the architectural remedy. Where a remedy has been implemented and its effects measured, that evidence is included. The aim is to show that STPA is a design method — not merely a diagnostic — and that its prescriptions are testable.



Canonical reference

The full case studies — unsafe control actions, causal factors, remedies, and the real-world evidence for each — are in [knowledge/system-catalogues/social-systems/cross-system/remedies-case-studies.md](#). This chapter gives a narrative walk-through of the five worked cases; the knowledge file is the reference.

33.2 The Five Worked Cases

1. Democracy (Republic) — the Weimar → Basic Law transition

The Weimar Republic (1919–1933) was formally democratic, constitutionally sophisticated, and destroyed from within in under fifteen years. Its failures — no judicial review of constitutionality, an emergency-powers path from parliamentary deadlock to authoritarian rule, no defence against anti-constitutional parties, and a capturable information environment — map onto four specific unsafe control actions.

The Federal Republic of Germany's Basic Law (1949) addressed each of them: the Bundesverfassungsgericht with power to strike down laws; the constructive vote of no confidence, which lets a Chancellor be removed only by simultaneously electing a successor; Article 21 GG allowing the Constitutional Court to ban parties seeking to undermine the free democratic basic order; and public broadcasting governed by civil-society councils rather than government officials. The result is a democracy that has operated without interruption for 75 years — not because Germans are different, but because the control structure is different.

2. Corporation — Sarbanes-Oxley, Mitbestimmung, and Cadbury

The corporation's defining structural feature is that authority derives from capital ownership. The three characteristic failures are: auditor dependence on the audited client, short-term compensation structures that outlive their consequences, and employee exclusion from governance despite holding the most accurate operational process models.

Three remedies, each tested in practice: Sarbanes-Oxley severed the auditor's consulting revenue from the audited client and made financial statement accuracy a personal criminal-liability matter for CEOs and CFOs. German co-determination (Mitbestimmung) placed worker representatives on supervisory boards of companies with more than 2,000 employees, bringing a different time horizon into the decision room. The UK Cadbury Code separated CEO and Chairman roles and required independent non-executive directors on key committees. None of the three asks executives to be more virtuous; each changes the structure so that the incentives point in a different direction.

3. Religion — mandatory reporting and Vatican II

Where bishops controlled both pastoral care and external communication, the authority that received abuse reports also decided whether to report externally. Mandatory reporting laws — enacted in Australia after the 2013–2017 Royal Commission, in Ireland after the 2009 Murphy Report, and elsewhere — bypass the institutional hierarchy by legal channel, changing the cost calculus of concealment from "this protects us" to "this destroys us."

The self-sealing property of doctrinal authority is harder to fix. The Second Vatican Council (1962–1965) was a structural opening to lay participation, ecumenical engagement, and modern scholarship. Its reforms were partially reversed in subsequent decades — itself a data point: the opening worked where it was maintained, showing the structural analysis was correct, and showing that reform requires sustained commitment at the goals level, not only the connections level.

4. University — pre-registration, Registered Reports, and DORA

Two failures dominate: selective reporting of positive results (the replication crisis: only 36% of psychology findings replicated in the 2015 Open Science Collaboration study), and impact-factor substitution for research quality (the classic proxy-metric failure).

The structural fixes: pre-registration and Registered Reports decouple the publication decision from the result, so that researchers commit to hypotheses, methods, and analysis plans before data collection. The Center for Open Science now hosts over 100,000 pre-registrations, and more than 300 journals offer Registered Reports. Pre-registered studies replicate at substantially higher rates. DORA and the UK REF's move toward article-level panel assessment disconnect the reward structure from the specific proxy metric. Both fixes change the connection between "result" and "consequence for the researcher" — and behaviour changes accordingly.

5. Military — parliamentary authorisation and independent prosecution

German armed forces cannot be deployed outside German territory without a Bundestag vote. The Wehrbeauftragter — the Parliamentary Commissioner for the Armed Forces — has independent access to any unit, any file, and any individual soldier; and reports directly to the Bundestag rather than the Defence Ministry. Germany has maintained continuous civilian control of its military since 1955; the Commissioner has regularly surfaced structural problems that command had institutional incentive to conceal.

For military justice, the US National Defense Authorization Act 2022 transferred jurisdiction over sexual assault and other serious felonies from commanding officers to independent Special Trial Counsel, removing the commanding officer from the prosecution decision entirely. Reporting rates rose; prosecution rates for referred cases rose. The structural conflict of interest — commanding officer deciding whether to prosecute conduct that reflects on the command — was identified through sustained advocacy using exactly the structural argument this book makes, and directly addressed.

33.3 Cross-System Pattern

Across all five systems, the same three structural failures recur: an accountability void, a self-sealing process model, and a proxy metric that has replaced the goal. The full cross-system table — failure type by system, with each proven remedy mapped to the corresponding control-loop fix — is in <knowledge/system-catalogues/social-systems/cross-system/remedies-case-studies.md>.

The remedies are not generic. Each closes a specific causal pathway identified by the structural analysis — and each has been tested in practice. That precision is what distinguishes architectural reform from exhortation.

33.4 Reading the Remedies as Rung-Elevation Moves

The previous chapter introduced the seven [Justificatory Rungs](#) and the dangerous **Pattern A — Asymmetric Loop**: rung-1 control downward with rung-3 evidence trying, and failing, to feed back upward. Re-read through that lens, **all thirteen catalogued architectural remedies share a common move**: the controller's authority remains at rung 1 (legal, institutional, or sacred); a

parallel feedback channel is installed that operates at rung 3 (independent audit, statutory investigation, replicated study, criminal-liability disclosure); and the new channel routes around the rung-1 hierarchy that previously filtered upward signal.

Remedy	Pre-remedy feedback rung	Post-remedy feedback rung
Bundesverfassungsgericht (Basic Law)	1 (parliamentary self-judgment)	2/6 (legal-rational + constitutional principle)
Sarbanes-Oxley external audit	1 (management self-reports)	3 (audited financials with criminal liability)
Mitbestimmung worker board representation	1 (management self-reviews)	6 (deliberative — multiple stakeholder perspectives)
Mandatory abuse reporting (Religion)	1 (clergy report internally)	3 (statutory report to civil authority)
Vatican II lay councils	1 (clerical hierarchy)	1+2 (lay participation; partially reversed)
Pre-registration / Registered Reports	3 (publication-biased trial sample)	4 (cumulative evidence without publication-bias filter)
DORA / article-level assessment	1 (journal prestige proxy)	3 (article-level merit)
Parliamentary mandate (Bundeswehr)	0/1 (executive decision)	6 (parliamentary deliberation)
Wehrbeauftragter (independent military prosecutor)	1 (commander investigates own troops)	3 (independent investigation)

The full thirteen-remedy table with detailed deltas is in [knowledge/system-catalogues/social-systems/cross-system/remedies-case-studies.md](#) §6.

The remedies do not try to convert rung-1 controllers to higher rungs — that would dissolve the institution they protect. They parallel the rung-1 channel with a rung-3 (or rung-2/6) channel routed to a separate oversight body whose existence does not depend on the controller. Where remedies have been incomplete (Vatican II's partial post-Council reversal) it is because the rung-elevated channel was made revocable by the same rung-1 hierarchy it was supposed to correct — which is the precondition this book's structural fixes already require.

This is the canonical architectural pattern for repairing **Pattern A — Asymmetric Loop** in the social-systems catalogue. It is the unifying idea behind the apparently disparate reforms above.

34. About This Book

34.1 Origin

This book grew from a series of conversations exploring what happens when you take systems engineering seriously — not just for aircraft and software, but for the institutions that shape human life.

The starting point was a simple question: if we try to build something in life — whether in family, in politics, or in engineering — we need to distinguish between goals, requirements, functions, logical architecture, and physical implementation. What happens when you apply this discipline to religion, to political systems, to development methodologies?

The answer turned out to be more interesting than expected.

34.2 Approach

Every analysis in this book follows the same method:

1. **Define** what a system is — establishing a precise working definition before any decomposition begins
2. **Decompose** the system through the five-level SE hierarchy
3. **Compare** systems by identifying their shared platform and variation points (product line thinking)
4. **Analyse** the control structure using STPA to identify how the system can produce outcomes opposite to its goals
5. **Propose** architectural remedies — structural changes, not just policy adjustments

34.3 Technology

This book is built with [MkDocs](#) and the [Material](#) theme, hosted on [GitHub Pages](#), and deployed automatically via GitHub Actions.

Source: github.com/ccabos/systems

34.4 License

This work is a personal research and documentation project.