

Systems Thinking

From SE to STPA

*Applying Systems Engineering Decomposition,
Product Line Thinking, and STPA
to Social and Technical Systems*

C. Cabos

<https://ccabos.github.io/systems/>

Table of Contents

Home	3
Introduction	5
Part I -- Foundations	9
What Is a System?	10
The Five-Level SE Hierarchy	13
Product Line Thinking	15
STPA and STAMP	17
Justificatory Rungs	19
Part II -- Social Systems	24
Religion	25
Kingdom & Republic	29
Ten Social Systems Compared	32
Verein (e.V.)	52
Democracy	56
Theocracy	62
Corporation	66
University	72
Military	76
Family	80
One-Party State	84
Part III -- Development Frameworks	88
Eight Frameworks Decomposed	89
The Framework Platform	93
Hybrid Architectures	94
Part IV -- STPA Applied	95
STPA on Religion	96
STPA on Frontier AI	111
Control Structures in Social Systems	124
Architectural Remedies	126
About	129

Systems Thinking — From SE to STPA

Applying Systems Engineering Decomposition, Product Line Thinking, and STPA to Social and Technical Systems

A place to start

In 1975, Kodak built the world's first digital camera — and buried it. Twenty years later, digital photography destroyed the company that invented it. The engineers were excellent. The technology was real. The failure was structural: the organisation was wired to protect film revenues, and it reliably smothered anything that threatened them. No personnel change could have fixed it. The design had to change. Most institutional failures look like this. **The Introduction** develops this argument through several cases before any framework or equation appears.

This book explores a single, powerful idea: **the tools we use to engineer safe aircraft and reliable software apply with surprising precision to human institutions** — political regimes, religions, corporations, families, and development methodologies.

We use three frameworks, layered on top of each other:

The Five-Level SE Hierarchy — decompose any system into Goals → Requirements → Functions → Logical Architecture → Physical Implementation, with full traceability.

Product Line Thinking — compare systems by identifying their shared platform (what's common) and their variation points (where they diverge), then study which elements can be reused across system types.

STPA (System-Theoretic Process Analysis) — model the system as a hierarchical control structure and systematically identify unsafe control actions that produce outcomes opposite to the system's stated goals.

Structure

The book is organized in four parts:

	Focus
Introduction	Why structural analysis beats moral exhortation — argued through concrete cases
Part I — Foundations	What is a system? The three frameworks: SE hierarchy, product lines, STPA
Part II — Social Systems	Ten social systems decomposed and compared
Part III — Development Frameworks	Waterfall, Scrum, Kanban, DevOps, and others — analysed as a product family
Part IV — STPA Applied	STPA applied to religion and other social systems

The Core Insight

Social systems look more different from each other than they actually are. A kingdom and a republic, a church and a corporation, a family and a Verein — these seem like categorically different things. But when decomposed through a rigorous systems engineering hierarchy, they reveal a **shared platform** of functional slots, and they differ primarily at a small number of well-defined **variation points**.

Product line thinking does not reduce human institutions to interchangeable parts. It does the opposite: it makes the *genuine* differences visible by stripping away the illusion that everything is different.

And STPA reveals that the most harmful outcomes of social systems are not random malfunctions but **predictable consequences of their control architecture** — consequences that can be addressed through architectural redesign, not just moral exhortation.

Introduction — What This Book Does and Why It Works

Start with a Puzzle

In the 1970s, Kodak employed more than 60,000 people in Rochester, New York. The company was profitable, technically excellent, and entirely dominant in its market. It also, in 1975, built the world's first digital camera. The engineer who built it, Steve Sasson, presented the prototype to management. The response was cautious appreciation and a quiet burial. The product was not released. Twenty years later, digital photography destroyed Kodak's business, and the company filed for bankruptcy in 2012.

What happened? The simple story is that Kodak's managers were short-sighted or afraid. But that explanation predicts the wrong fix. If you replace short-sighted managers with visionary ones, and nothing else changes, you get the same outcome — because the problem was not the people. Kodak had a different problem, one that no personnel change could solve.

Kodak's film business generated enormous profit margins. Every internal reporting line, every budget process, every promotion criterion was calibrated to protect and extend that revenue stream. The digital camera project reported into the film division. Its budget competed against film products. Its success was measured against metrics designed for film. The *structure* of the organisation — the connections between its parts and the goals those connections served — guaranteed that any technology threatening film would be smothered, regardless of who was in charge. Kodak did not fail because of bad people. It failed because its organisational structure contained an architectural flaw that turned a threat into an impossibility.

This is a systems problem. And systems problems have systems solutions.

A Closer Example

Consider something more familiar: a secondary school.

The school's declared goal is to educate students — to develop their capacity to think, reason, create, and contribute. Nobody disputes this. Teachers enter the profession because they believe in it. Administrators genuinely want students to succeed. Parents care about their children's development. The people are good and their intentions are aligned.

And yet, in many such schools, the following happens: students who understand nothing about history can pass a history exam. Teachers who are brilliant at inspiring curiosity get poor performance reviews. Students optimise for grades rather than understanding, teachers optimise for exam results rather than education, and school administrators optimise for league table rankings rather than learning. Everybody does their job correctly. The system produces the wrong outcomes.

Why? Because the *connections* in the system — the feedback loops that link behaviour to reward — do not connect to the declared goal. They connect to a measurable proxy for the goal. Exam scores are measurable; understanding is not. The moment you attach resources, careers, and prestige to exam scores, the whole system reorients toward producing exam scores. This is not cynicism — it is the mechanical consequence of the control structure.

The fix is not to motivate teachers more, threaten students more, or exhort administrators to care about real learning. None of those interventions touch the structure. The fix requires changing *what is measured, how it feeds back, and what it controls* — which means changing the connections. That is an architectural intervention, not a moral one.

But what does that architectural fix look like, concretely? This is where the engineering approach earns its keep. Once you have traced the causal chain — teacher evaluation derives from class exam scores, school ranking derives from published test data, university admission derives from GPA — each link in that chain becomes a candidate for redesign:

Misconfigured connection	Architectural remedy
Teacher appraisal tied only to exam averages	Add peer review of lesson quality and longitudinal tracking of how former students actually perform
School ranking uses only test results	Require portfolio assessment, inspection reports, and employer or university feedback as parallel signals

Misconfigured connection**Architectural remedy**

Single high-stakes exam controls university entry	Reduce the stakes of any one instrument; distribute the selection signal across multiple independent measures
---	---

The remedies are not exotic. Several school systems have implemented versions of them. Finland removed published school rankings and standardised national testing before age 16; teacher evaluation shifted to professional peer review; the feedback loops changed. The outcomes changed with them — not because Finnish children or teachers are different, but because the control structure is different.

This is the pattern the methods in this book are designed to produce: not merely a diagnosis of what went wrong, but a precise identification of *which connection to change* — and confidence that changing it will change the outcome. The analysis generates the remedy. The two are the same work.

The Pattern Appears Everywhere

Once you see this pattern, you find it in every kind of institution.

Religions often declare a goal of compassion, care for the vulnerable, and guidance toward transcendence. The institutional structure — hierarchy, doctrine, territory, resources — has its own logic. When institutional survival conflicts with the declared goal, the structure tends to win. The result is not hypocrisy by individuals but an architectural conflict between declared goals and operative feedback loops. The Catholic Church's response to the sexual abuse crisis is inexplicable as individual moral failure at scale; it is entirely explicable as the output of a control structure in which the preservation of institutional reputation was connected, through many feedback loops, to the decision-making of bishops — while the protection of victims was not.

Democracies are designed on the premise that distributing power across competing institutions prevents any single actor from accumulating enough control to subvert the system. When democracies fail, the failure is rarely a sudden coup. It is typically a sequence of individually legal changes — each one defensible in isolation — that progressively weakens the connections between institutions, concentrates decision-making, and removes the feedback mechanisms (free press, independent courts, competitive elections) that generate corrective signals. The structure changes before anyone notices the system has changed.

Corporations fail not because they hire bad people but because they build control structures that reward the wrong behaviour. The 2008 financial crisis did not require criminal intent at every level. It required a mortgage originator whose bonus was tied to volume, not quality; a securitiser whose fee was tied to deal flow, not risk; a rating agency whose revenue depended on the issuers it rated; and a regulator whose authority did not extend to the instruments being created. Each actor responded rationally to the incentives their position created. The system produced catastrophe.

Families can be analysed the same way. A family that communicates only through the children, where no direct channel exists between parents, has a structural problem. Replacing the children with different children — or telling the existing children to communicate more clearly — does not address it. The connection pattern needs to change.

In each case, the same diagnosis applies: the *declared goals* of the system diverge from the *operative goals* encoded in its feedback structure, and the people inside it respond to the operative goals. The fix requires architectural change — changing connections, changing what the system measures, changing what triggers what — not just changing the people or the rhetoric.

What Engineering Knows

Engineers have been building complex systems — aircraft, nuclear plants, air traffic control, spacecraft — for decades, and those systems must work. The cost of failure is immediate and physical. This pressure has produced a set of analytical disciplines that are ruthlessly practical: they do not ask whether the people involved are good; they assume people will respond to the structure they are in, and they ask whether the structure is safe.

Three of those disciplines are central to this book.

Systems Engineering decomposition takes any system — a jet engine, a hospital, a political party — and breaks it into five layers: the goals it pursues, the requirements those goals generate, the functions it must perform, the logical

architecture that organises those functions, and the physical components that realise that architecture. This decomposition makes the structure explicit and legible. You cannot analyse what you cannot see, and most social systems are analysed at the level of individuals and events rather than at the level of structure. The SE hierarchy changes that.

Product Line Thinking observes that many systems that look very different are actually variations on a common underlying architecture. Boeing builds multiple aircraft types from a shared engineering platform. IKEA produces thousands of products from a shared manufacturing and distribution system. This book argues that kingdoms and republics, corporations and universities, religions and militaries, are variations on a shared social system platform — and that understanding which elements are shared and which are genuinely different is more useful than treating each system as entirely unique.

STPA — System-Theoretic Process Analysis — was developed by Nancy Leveson at MIT after studying catastrophic failures in complex socio-technical systems. Its central insight is that most catastrophic failures in modern systems do not result from component breakdown; they result from *control failures* — situations where the commands, feedback, and authority relationships between parts of the system produce unsafe actions even when each part is functioning normally. STPA does not ask "what broke?" It asks "what could the control structure do that it should not, and what is it not doing that it should?" Applied to social systems, it produces a systematic map of the ways an institution can harm the people it was built to serve — not through individual malice, but through structural design.

What This Book Does

This book takes those three engineering tools and applies them to the social systems most of us live inside.

We decompose ten social systems — the kingdom, the republic, the theocracy, the one-party state, the corporation, the university, the military, the family, the church, and the Verein — through the five-level SE hierarchy. Each decomposition produces a table: goals at the top, physical implementation at the bottom, with traceable connections between every level. The decomposition makes visible what is usually invisible: what the system is actually *for*, how it is actually *organised*, and where the critical connections are.

We then apply Product Line Thinking to compare these decompositions. The result is striking: beneath the obvious differences in flag, creed, and custom, these ten systems share a common architectural platform of ten functional slots that every human institution must fill. They differ at a small number of variation points — source of authority, membership boundary, decision-making mechanism, succession mechanism, legitimation narrative, and norm enforcement. Understanding the variation points tells you what can be changed and what cannot: institutional reform that crosses variation-point constraints will fail, not because of bad intentions, but because the structural constraints are real.

Finally, we apply STPA to trace how these systems can turn against their own stated goals. We show how the control structure of a religion can produce exactly the outcomes the religion forbids. We show how the control structure of a democracy can produce authoritarian outcomes if the feedback mechanisms are progressively dismantled. And we identify, for each type of system, the architectural features — the design choices — that make harmful outcomes more or less likely.

What This Book Does Not Do

It does not argue that all institutions are secretly the same, or that cultural and historical particularity does not matter. The variation points are real and consequential. A theocracy and a republic are genuinely different systems, and those differences explain a great deal about how they behave.

It does not argue that structural analysis replaces moral and political judgement. It is a precondition for it. You cannot judge a system's design wisely if you cannot see the design.

And it does not promise that structural reform is easy. The Kodak story illustrates the problem: the people inside a system are rewarded by the system's operative goals. Those who propose changing the structure threaten the operative goals of the people the structure rewards. Structural reform is resisted not because people are perverse but because the structure makes resistance rational.

What this book does promise is that structural analysis makes the conversation more honest. When we understand that harmful outcomes are *predictable consequences of identifiable architectural features*, we can argue about design instead of character. That is a more productive argument — and a fairer one.

How to Read This Book

If you want the examples first — continue reading here and move directly to Part II, which applies all three frameworks to ten social systems. The theoretical foundations are in Part I, but the examples in Part II are self-contained.

If you want the theory first — Part I develops the three frameworks systematically, starting with a precise definition of what a system is and ending with the full STPA methodology.

If you want the most unsettling chapter — go directly to the STPA analysis of religion in Part IV. It is the worked example that brings all three frameworks together and shows, in formal detail, how a system designed to serve the transcendent can be structurally configured to serve itself instead.

The introduction you just read is the argument in compressed form. The rest of the book is the evidence.

Part I — Foundations

This part introduces the foundational definition of a system and the three analytical frameworks that form the backbone of the book. Each framework is well-established in engineering practice; our contribution is applying them systematically to social and organizational systems.

- **What Is a System?** — the foundational definition
- **The Five-Level SE Hierarchy** — the decomposition instrument
- **Product Line Thinking** — the comparative methodology
- **STPA and STAMP** — the safety and control analysis
- **Justificatory Rungs** — the epistemic-standard extension to STPA for social systems

What Is a System?

Canonical reference

The normative statement of the foundational concepts in this chapter — the working definition, the replaceability hierarchy, emergence, boundaries, and the distinction between technical and non-technical systems — is maintained as the project's knowledge base under `knowledge/foundations/`. This chapter is the narrative discussion; the knowledge base is the reference.

Before We Decompose, We Must Define

Every analytical framework rests on an ontology — a theory of what kinds of things exist. Before applying the SE hierarchy, product line thinking, or STPA, we need a precise answer to the question that precedes all of them: *what makes something a system in the first place?*

Two families of answers dominate the literature. They are not mutually exclusive; this chapter argues they are complementary and that understanding the tension between them clarifies what our analytical tools are actually doing.

Theory I — Structure: Parts, Connections, and Goals

The first and most widely held view holds that a system is constituted by three elements:

1. **Parts** — the distinguishable components that can, in principle, be identified and listed
2. **Connections** — the relationships, flows, and dependencies between parts
3. **Goals** — the purpose or purposes the system pursues, whether declared or operative

The Replaceability Hierarchy

A significant observation follows from this ontology: the three elements are not equally essential to system identity, and they differ in how easily they can be replaced without the system ceasing to be *the same system*.

Element	Replaceability	Consequence of Change
Parts	High — most easily swapped	System continues; this is routine maintenance
Connections	Moderate — harder to change	System is restructured; behaviour shifts significantly
Goals	Low — very difficult to change	System is transformed; what was there before is effectively a different system

A human body replaces most of its cells over years and remains the same organism. A corporation can change its entire workforce and remain recognisably the same corporation. But a corporation whose goal shifts from profit to public service is, in any meaningful sense, a different institution wearing the same legal clothes. The Roman Catholic Church that abandoned the salvation of souls as its operative goal would not be a reformed Catholic Church — it would be something else entirely occupying the same physical infrastructure.

This hierarchy maps directly onto the SE decomposition used throughout this book. Physical implementations (P) correspond to parts. Logical architecture and functions (L, F) encode the connection patterns. Goals and requirements (G, R) encode purpose. Institutional reform that touches only P-level elements is maintenance. Reform that reaches G-level is transformation — and the historical record shows it is both rare and destabilising.

Critical Assessment

The structural theory is productive and largely correct, but it has four significant gaps.

1. Goals are often contested, latent, or split. The framework assumes a system has goals. Real social systems typically have *declared* goals and *operative* goals that diverge substantially. A regulatory agency declared to protect the public may operate primarily to protect the regulated industry. A religion declared to save souls may operate primarily to perpetuate its own institutional authority. The structural theory needs to distinguish these two types of goal, or it will mistake the label for the contents. This book addresses the divergence directly in the STPA analysis (Part IV), where unsafe control actions typically arise from exactly this split.

2. Connections may outrank goals as the determinant of system identity. Donella Meadows argued that the *feedback structure* of a system — its loops, delays, and reinforcing dynamics — determines its long-term behaviour more than any stated goal. A system whose goals are replaced but whose feedback loops remain intact will tend to revert to its former behaviour over time. The Soviet Union changed its declared goals repeatedly; the feedback structure of centralised authority, information suppression, and reward for loyalty produced remarkably consistent outcomes across six decades and several ideological pivots. The replaceability ordering may need revision: connections as structure may be *more* fundamental to identity than goals as declaration.

3. Every part is itself a system. The framework is valid but incomplete: it defers the definition rather than grounding it. A part is a system with its own parts, connections, and goals. This recursion is not a defect — it reflects the genuine nested structure of complex systems — but it means the theory requires a boundary condition: a decision about where to stop decomposing. That decision is not given by the theory itself.

4. The boundary question is missing. The most important definitional act for any system is drawing its boundary: what is inside and what is outside. For social systems this is particularly fraught. Who is "in" a state? Citizens only? Residents? Who counts as a member of a religion — the registered faithful, the active practitioner, the nominal cultural adherent? The boundary choice determines what counts as a part, what counts as an internal connection, and what counts as an external relationship. A theory of systems that omits the boundary is like a theory of cells that omits the membrane.

Theory II — Behaviour: Emergence

The second view defines systems not by what they are made of but by what they *do* that their parts cannot. A system is characterised by **emergent properties** — properties that are present in the organised whole but absent from any of the individual parts considered in isolation.

No individual neuron is conscious. No individual water molecule is wet. No individual trader sets market prices. Consciousness, wetness, and market prices are real; they exist; but they exist only at the level of the organised system, not at the level of its parts. If emergence is present, something is a system. If it is absent — if the whole is simply the sum of its parts — what we have is an aggregate, not a system.

This provides an observational test: can the behaviour of the whole be predicted by summing the behaviours of the parts? If yes: aggregate. If no: system.

Critical Assessment

The emergence criterion captures something real and important. But it has three significant weaknesses that limit its usefulness as a *definition* of system.

1. Emergence is descriptive, not constitutive. The emergence criterion tells you what to observe once you have already decided you are looking at a system; it does not tell you how to identify or engineer one. You cannot look for emergence before you have decided what to treat as the parts — and that decision already presupposes a structural theory. Emergence is the *output* of a system's structure, not the criterion that defines it. For an SE-based approach, this is a practical limitation: you design structure and observe emergence as a consequence; you cannot design emergence directly.

2. The criterion is trivially satisfied by aggregates. A pile of sand has a height and an angle of repose that no individual grain possesses. A crowd has a density that no individual person has. These are properties present in the aggregate but absent from the parts — but no one seriously claims a pile of sand is a system in the relevant sense. The emergence criterion must be qualified: we need to specify *which* emergent properties matter. The relevant ones are goal-directed or self-maintaining behaviours — and as soon as we say this, we have smuggled goal back in through the back door, returning to Theory I.

3. The strong/weak emergence distinction matters. Philosophers distinguish *weak emergence* (in principle reducible to lower-level description, but practically too complex to compute) from *strong emergence* (genuinely irreducible, a new causal power not explainable even in principle from the parts). Almost all emergence in social systems is weak:

complex, non-linear, and surprising, but not metaphysically mysterious. Treating social emergence as strong emergence — as if "the system" were a force beyond analysis — is an error that forecloses the very analysis this book attempts.

Synthesis: The Working Definition

The two theories are not rivals. They are the **inside view** and the **outside view** of the same reality.

- **From the inside** (Theory I): a system is a *structure* — organised parts, patterned connections, directing goals, within a defined boundary.
- **From the outside** (Theory II): a system is an *actor* — something whose boundary-level behaviour cannot be read off from its parts in isolation.

Neither is sufficient alone. A structure with no emergent behaviour is a machine, not a system in the interesting sense. An emergent phenomenon with no describable structure is magic, not analysis.

Adding the boundary gives the working definition this book operates with throughout:

Working Definition

A system is a **bounded structure of parts and connections organised by goals**, whose behaviour at the boundary **cannot be predicted from its parts in isolation**.

Four corollaries follow from this definition and recur throughout the analysis:

1. **Changing parts is maintenance.** It preserves the system.
2. **Changing connections restructures the system.** It changes what the system does, often dramatically.
3. **Changing goals transforms the system.** The result is a different system with inherited infrastructure.
4. **Changing the boundary redefines the system.** What counts as internal, external, and relational changes entirely.

The SE hierarchy in the next chapter is a formal instrument for making this structure explicit and tractable. Goals become the G-level. Connections between functional components become the L-level. Parts become the P-level. The boundary is the decision to start the decomposition at all.

Why This Matters for Social Systems

Technical systems — aircraft, software, power grids — are designed with explicit goals, documented connections, and identifiable parts. Their structure is legible because engineers made it so.

Social systems — states, religions, corporations, families — were not designed in the same deliberate sense. Their goals are often implicit, their connections informal, and their boundaries contested. The structural theory predicts that this makes them harder to reform: if the connection patterns are not documented, you cannot know what you are changing. If operative goals are not distinguished from declared goals, reform efforts target the label and leave the engine untouched.

The emergent-behaviour theory adds a further warning: social systems produce outcomes that no participant intended and no individual can fully understand. This is not mysticism — it is the practical consequence of non-linear feedback in complex structures. It means that moral responsibility is always distributed across the structure, never fully located in any single part.

Both warnings are taken seriously in what follows.

The Five-Level SE Hierarchy

Canonical reference

The normative definition of the five-level hierarchy, its traceability rules, ID conventions, decomposition order, and failure pathologies is maintained as the project's knowledge base under `knowledge/se-techniques/goals-requirements-hierarchy/`. This chapter is the narrative introduction; the knowledge base is the reference.

A Universal Decomposition Instrument

Systems engineering decomposes any system — technical or social — into five levels of abstraction. Each level answers a different question, and each item at one level traces to items at the next level down, forming a directed acyclic graph with full traceability.

The Five Levels

Level	Prefix	Question Answered
Goals	G	Why does this system exist? What overarching objectives does it pursue?
Requirements	R	What must be true? What constraints, criteria, and conditions are derived from the goals?
Functions	F	What does the system do? What operations does it perform to meet requirements?
Logical Architecture	L	How is it organized? How are functions grouped into interacting subsystems?
Physical Implementation	P	What is it made of? What concrete people, buildings, documents, tools, and processes realize the architecture?

Why Five Levels?

The five levels are not arbitrary. They correspond to the fundamental decisions that any system designer — whether an engineer or a lawmaker — must make, in order:

- 1. Purpose before structure.** You cannot design a system without knowing what it is for (Goals).
- 2. Constraints before mechanisms.** You cannot choose mechanisms without knowing the constraints they must satisfy (Requirements).
- 3. Functions before grouping.** You cannot group things into subsystems without knowing what functions need to exist (Functions).
- 4. Logical before physical.** You should decide how subsystems interact before choosing which specific technologies, people, or buildings implement them (Logical → Physical).

Skipping a level — for example, jumping from goals directly to physical implementation — produces systems that "work" but are fragile, incoherent, and resistant to change. This is as true of a parliament as it is of a software architecture.

Traceability Rules

Every item in the hierarchy must be connected:

- Every Goal links to at least one Requirement

- Every Requirement links to at least one Function
- Every Function links to at least one Logical Architecture element
- Every Logical Architecture element links to at least one Physical Implementation
- A child can have multiple parents (many-to-many is valid)
- No orphan nodes — every item must be traceable upward to a goal and downward to an implementation

These rules ensure **completeness** (every goal is implemented) and **necessity** (every implementation serves a goal). When applied to social systems, they reveal a common pathology: physical implementations that have lost their connection to any goal ("we do this because we've always done it") and goals that have no implementation ("we believe in this but have no mechanism to achieve it").

Example: The German Verein

A brief example to make the abstraction concrete:

Goal

G1 — Enable citizens to pursue a shared non-commercial purpose collectively

Req

R1 — Written Satzung defining the purpose

R3 — Democratic decision-making with equal voting rights

Func

F3 — Democratic assembly — Mitgliederversammlungen, elect Vorstand

Logical

L1 — Governance Subsystem — Satzung, Wahlordnung, Geschäftsordnung

Physical

P2 — Mitgliederversammlung (annual meeting), Protokolle, Kassenprüfung

The traceability chain: G1 → R3 → F3 → L1 → P2. Every physical artifact (the annual meeting, the minutes, the financial audit) traces back to a goal. If you removed the annual meeting, you would break the trace to G1 — the system would lose its mechanism for collective purpose.

This is the instrument. In the following chapters, we apply it to systems that are not usually analyzed this way.

Product Line Thinking

Canonical reference

The normative definitions of platform, variation points, binding decisions, the reuse gradient, the coherence rule, the workaround principle, and the product-line architect's toolkit are maintained as the project's knowledge base under `knowledge/se-techniques/product-line-engineering/`. This chapter is the narrative introduction; the knowledge base is the reference.

Platform, Variation Points, and Reuse

A **product line** (or product family) is a set of systems that share a **common platform** but differ at defined **variation points**. The engineering discipline has three core ideas:

The Platform

The platform is the set of architectural elements — goals, requirements, functions, subsystems, and physical components — that are shared across all (or most) members of the family. In automotive terms, this is the chassis. In social system terms, this is the governance function, the membership mechanism, the resource allocation subsystem.

Variation Points

Variation points are the locations in the architecture where the platform deliberately leaves a decision open, because different family members will make different choices. A variation point has a defined *interface* (what connects to it) and a set of *variants* (what can be plugged in).

Automotive Analogy

Interface: "Here, a powertrain connects to the drivetrain."

Variants: 2.0L diesel · electric motor · plug-in hybrid

Everything below the interface (suspension, brakes, steering) works with any variant.

Political Analogy

Interface: "Here, executive authority is sourced and legitimated."

Variants: hereditary succession · popular election · party appointment · clerical selection

Everything below the interface (courts, civil service, tax authority) works with any variant.

Binding Decisions

Binding decisions are the specific choices made at each variation point for a given product. Once a binding is made, it constrains which other bindings are compatible. Not all combinations are coherent.

The Reuse Gradient

Our analysis reveals a general principle: **reuse feasibility decreases as you move upward in the SE hierarchy.**

SE Level	Reuse Potential
Physical Implementation	Highest — a tax authority works under any regime
Logical Architecture	High — subsystems transfer with adaptation
Functions	Medium — mechanisms may need interface adjustment
Requirements	Low — must be checked for goal compatibility
Goals	Lowest — goals define the system type; changing them means changing the system

This gradient has a practical corollary: **start reform at the bottom**. Modernize physical implementations first (professional civil service, independent courts, transparent budgets), then work upward. Top-down reform — imposing new goals before building the implementing infrastructure — is the recipe for failed states and institutional collapse.

The Coherence Rule

A system is coherent when its binding decisions at each variation point are mutually compatible, and when every element at every SE level traces consistently to the goals above it and the implementations below it. Incoherent systems — those with contradictory bindings — are unstable.

The Workaround Principle

When a desirable module from another system type conflicts with a binding decision, the most successful reforms create a **workaround** — a modified version that preserves the function while respecting the constraint. The constitutional monarchy is the paradigm case.

These concepts are applied extensively in Part II (social systems) and Part III (development frameworks).

STPA and STAMP

Canonical reference

The normative statements of STAMP (the accident model) and STPA (the four-step hazard analysis procedure, the four UCA types, and the rules for translating an SE decomposition into STPA inputs) are maintained as the project's knowledge base under `knowledge/se-techniques/stamp/` and `knowledge/se-techniques/stpa/`. This chapter is the narrative introduction; the knowledge base is the reference.

System-Theoretic Process Analysis

The STAMP Accident Model

STAMP (Systems-Theoretic Accident Model and Processes), developed by Nancy Leveson at MIT, treats accidents not as the result of component failures but as **emergent properties of complex systems** — caused by inadequate control, including failures between components that may each be functioning exactly as designed.

This is precisely the right model for social systems, where the most devastating outcomes (inquisitions, abuse cover-ups, financial crises, failed reforms) typically occur not because the system has "broken" but because its control mechanisms are operating exactly as designed — just in a way that produces harm.

The Four Steps of STPA

Step	What It Produces
Step 1: Define Purpose	Losses (unacceptable outcomes), Hazards (system states leading to losses), System-level Constraints
Step 2: Model Control Structure	Hierarchical diagram of controllers, control actions, feedback channels, and process models
Step 3: Identify Unsafe Control Actions	Systematic enumeration of UCAs across four types
Step 4: Identify Loss Scenarios	Causal analysis: why each UCA occurs, what contextual factors enable it

The Four UCA Types

For every control action in the system, STPA asks four questions:

1. **Not providing** — What happens if this control action is *not* issued when it should be?
2. **Providing** — What happens if this control action *is* issued when it shouldn't be, or is the wrong action?
3. **Too early / too late / wrong order** — What happens if the timing is wrong?
4. **Applied too long / stopped too soon** — What happens if the duration is wrong?

From SE Decomposition to STPA

The SE hierarchy feeds directly into STPA:

- **Goals** define what the system must achieve → STPA **Losses** are the negation of goals
- **Requirements** define constraints → STPA **System-level Constraints** are safety requirements
- **Logical Architecture** defines subsystems → STPA **Control Structure** maps controllers and controlled processes
- **Physical Implementation** identifies actors → STPA identifies who issues which control actions

This integration is applied in detail in Part IV.

Rung tagging — the social-systems extension

For social systems, Step 2 acquires one more move: tag every control action and every feedback channel with the **justificatory rung** at which it operates (rung 0 coercion through rung 6 deliberative legitimacy), and tag each controller with both its *claimed* and *operating* rung.

Rung information is not derivable from the SE decomposition — it is its own elicitation activity, performed during Step 2 — and it is required by Step 3 to identify rung-mismatch UCAs (over-rung command, under-rung command, asymmetric loop). The previous chapter Justificatory Rungs introduces the ladder and the three dangerous mismatch patterns; the canonical integration with STPA Steps 1–4 is at [knowledge/se-techniques/justification-rungs/application-to-stpa.md](#).

For engineering systems where every loop operates at rung 3–4 throughout, rung tagging adds nothing and is omitted. Use this extension where the system is social, institutional, or policy-shaped — where the question being asked is *whether the loop's rung fits its function*.

References

- Leveson, N.G. (2012). *Engineering a Safer World: Systems Thinking Applied to Safety*. MIT Press.
- Leveson, N.G. & Thomas, J.P. (2018). *STPA Handbook*. MIT Partnership for Systems Approaches to Safety and Security.
- Leveson, N.G. (2004). "A New Accident Model for Engineering Safer Systems." *Safety Science*, 42(4), 237–270.

Justificatory Rungs

Canonical reference

The normative statement of the seven-rung ladder, its three rung-mismatch patterns, and the seven hooks into STPA Steps 1–4 are maintained as the project's knowledge base under [knowledge/se-techniques/justification-rungs/](#). This chapter is the narrative introduction; the knowledge base is the reference.

What this chapter adds to the book

Parts I and IV introduce STPA — the safety-and-control analysis at the heart of this book. STPA tells us *which control actions are unsafe*. It does not tell us *why a control action that should land does not*. That second question is the everyday experience of social systems: the bishop's sermon does not sway the empirically-minded congregant, the scientist's data does not move the believer, the parliament's resolution does not constrain the autocrat. In each case the control action is issued, transmitted, and received — but the receiver does not accept it as a *reason* to change behaviour.

The reason is that human beings act on different **standards** for what counts as a decisive reason. There are seven of them in this book's analysis. They form a ladder, and most of the catastrophic failures of social systems are traceable to mismatches on that ladder: between the two ends of a control loop, or between what a controller *claims* and what it *operates on*.

The seven rungs at a glance

Rung	Standard	Where you meet it
0	Power / coercion ("do it or else")	Coups, monopolies, threats
1	Authority / rhetoric / identity	Slogans, scripture, "as a parent..."
2	Formal consistency	Math proofs, court syllogisms
3	Empirical testability	Randomised trials, falsification
4	Cumulative evidence and consilience	Meta-analyses, IPCC reports
5	Meta-rational integration	Risk engineering, AI-safety, complex-systems policy
6	Normative legitimacy / deliberative ethics	Constitutional courts, citizens' assemblies

Each later rung adds a *check* that earlier rungs lack:

- **Rung 0** → **Rung 1**. The receiver moves voluntarily, without a threat, because they recognise the source as legitimate. Rung 1 has no check on truth or consistency — but it has stopped using raw force.
- **Rung 1** → **Rung 2**. A consistency check appears: even an authority cannot make an inconsistent argument win. Contradictions lose. This is the first *internal* check.
- **Rung 2** → **Rung 3**. An *external* check appears: data can sink a consistent theory if its premises do not match the world.
- **Rung 3** → **Rung 4**. The check is broadened against single-study flukes, publication bias, and method artefacts. Findings have to hold up across replications and methods.
- **Rung 4** → **Rung 5**. The receiver becomes able to step *outside* any single method, recognise its limits, and combine methods under uncertainty. Rung 5 is rare, fragile, and cognitively costly.

• **Rung 5 → Rung 6.** A second kind of question appears that no amount of rung-5 sharpening can answer: *what should we do?* Rung 6 is the check that surviving open inclusive discourse imposes on claims about *oughts*.

The full rung definitions and examples are in `knowledge/se-techniques/justification-rungs/rungs.md`.

The seven rungs in detail

Each rung admits four questions: what is the standard, where do you meet it, what does the rung *add* over the previous one, and what does it still *lack* — and one extra question that matters for STPA: how does the rung *show up* in a control structure when something goes wrong.

Rung 0 — Power / coercion

The standard is "do it or else." Compliance is enforced by threat of sanction, violence, or exclusion; there is no shared criterion for what is true, valid, or fair. Two parties at rung 0 cannot disagree productively — they can only enforce or submit. Examples range from coups and occupation regimes to organised crime and parental "because I said so."

What rung 0 *lacks*: any common standard at all. In an STPA control structure this shows up as control actions that land without persuasion and feedback that either flatters the controller or is suppressed. Loss scenarios concentrate around the absence of corrective feedback — the controller cannot be wrong because the controller cannot be questioned.

Rung 1 — Authority / rhetoric / identity

A claim is decisive because of *who* makes it, *what tradition* it comes from, or *what group* the listener identifies with. The listener moves voluntarily, but without an obligation to check truth or consistency. Examples: campaign slogans, scripture, "as a parent...", brand loyalty, expert testimony where the listener cannot evaluate the substance, charismatic leadership.

What rung 1 *adds* over rung 0: voluntary acceptance — the listener moves without a sanction being threatened. What rung 1 *lacks*: any internal check. Two contradictory rung-1 authorities cannot adjudicate the contradiction without escalating to rung 0 (force) or rung 2+ (some standard outside personal authority). In STPA terms, rung-1 controllers treat *who said it* as the deciding factor; feedback that comes from the wrong source is dismissed regardless of content. The self-sealing process models catalogued in the next chapter are most often rung-1 process models.

Rung 2 — Formal consistency

A conclusion follows validly from premises both parties already accept. The argument can be checked step by step; contradictions lose. Examples: mathematical proofs, court syllogisms, formal logic, rule-based bureaucratic decisions, legal reasoning within an established code.

What rung 2 *adds* over rung 1: an *internal* check. Even an authority cannot make an inconsistent argument win — a rung-2 challenge can defeat a rung-1 claim by exhibiting contradiction within the claim's own premises. What rung 2 *lacks*: an *external* check. Premises themselves are not tested against the world. Loops that operate at rung 2 produce internally coherent decisions that may diverge from external reality — the characteristic failure is the system that "follows the rules" straight into harm.

Rung 3 — Empirical testability

A hypothesis is decisive only if it survives tests against rivals — controlled experiments, falsification, prediction against unseen data. Examples: randomised trials, falsification in physical sciences, A/B tests, engineering acceptance tests, audited financial statements.

What rung 3 *adds*: an *external* check. Data can sink a neat theory; premises themselves become accountable to observation. What rung 3 *lacks*: protection against one-off flukes, publication bias, p-hacking, and the small-N problem. A single positive trial does not yet make a robust claim. Rung-3 feedback channels are powerful but fragile — a single biased sample can defeat them — so loops that operate at rung 3 require active maintenance of the channel's integrity (independence of the auditor, pre-registration of trials, replication budgets).

Rung 4 — Cumulative evidence and consilience

A claim is decisive when results survive replication and line up across multiple methods, populations, and disciplines. Meta-analyses, systematic reviews, and cross-method triangulation beat single studies. Examples: evidence-based medicine, IPCC assessment reports, well-replicated psychology effects, cross-discipline syntheses, regulatory science.

What rung 4 *adds*: robustness against flukes, biases, and method artefacts. A finding that holds up across independent methods is harder to explain away. What rung 4 *lacks*: a way to handle situations in which the methods themselves are in dispute, or where no single method applies, or where the question is partly normative. Rung-4 feedback is slow but durable; systems that depend on it (climate policy, drug regulation) have characteristic failure modes around delay — the rung-4 signal arrives only after irreversible action.

Rung 5 — Meta-rational integration

A claim is decisive when the speaker can show *why* a method works, *when it breaks*, and *how to combine multiple methods* under uncertainty. The speaker steps outside any single formalism to choose the right one for the question. Examples: risk engineering, AI-safety research, complex-systems policy labs, integrated assessment in climate policy, intelligence analysis at its best.

What rung 5 *adds*: an acknowledgement that no single method covers all messy domains, and a discipline for choosing among methods rather than defaulting to one. Rung 5 is where one becomes able to talk sensibly about model failure, unknown unknowns, and the limits of one's own evidence. What rung 5 *lacks*: any answer to *what should be done* once the empirics are settled. Rung 5 sharpens facts; it does not deliver oughts. Rung-5 controllers can hold multiple feedback channels in superposition — a rung-3 RCT, a rung-1 stakeholder intuition, a rung-2 legal constraint — and weight them appropriately. This is rare, fragile, and cognitively costly.

Rung 6 — Normative legitimacy / deliberative ethics

A claim is decisive when it survives open, inclusive discourse and meets normative criteria — fairness, reversibility, universalisability, respect for the autonomy of the affected. Even perfect facts and models do not tell us *oughts*; society still has to justify rules to all affected by them. Examples: constitutional courts, citizens' assemblies, Habermasian deliberative democracy, well-functioning parliamentary deliberation, bioethics committees.

What rung 6 *adds*: a way to settle questions about what should be done that cannot be reduced to fact — distribution, risk acceptance, value tradeoffs, intergenerational duties. Rung 6 is the only rung that handles *value pluralism* without smuggling one party's values in as facts. What rung 6 *lacks*: speed and reach. Rung-6 deliberation is slow, expensive, and does not scale to many decisions per unit time, so most operational control stays at lower rungs. The characteristic failure mode is *displacement*: when the rung-6 channel is suppressed, the unresolved normative question accumulates pressure that eventually escapes through rung 0 (revolution).

What every controller carries on every loop

Drawing these rungs onto an STPA control structure means every arrow gets a tag, and every controller acquires two extra attributes: the rung at which it accepts *commands* from above and the rung at which it accepts *feedback* from below. The two are often different on the same controller, and that asymmetry is where most rung-related failures live. The next section catalogues the three patterns of mismatch.

Why the ladder matters for STPA

STPA Step 2 produces a control structure: controllers issuing control actions to controlled processes, with feedback returning. For social systems, every arrow in that structure operates at *some rung*. The bishop-to-priest arrow is rung 1 (sacred authority). The scientist-to-policymaker arrow is rung 3 (empirical testability). The court-to-parliament arrow is rung 2 (formal consistency) backed by rung 6 (constitutional principle).

Tagging the rungs makes three failure modes visible:

Same rung is necessary for transmission

If the controller speaks at a rung the receiver does not operate on, the control action does not land. A scientist briefing a charismatic authority at rung 3 ("the data shows...") to a rung-1 audience does not produce behavioural change — the audience cannot process *data* as a reason. This failure mode is silent: the controller may believe the action was issued, transmitted, and obeyed.

Same rung is not sufficient — and at low rungs it is dangerous

Two actors both at rung 1 transmit perfectly but have *no internal check*. The system efficiently propagates whatever the upper controller wants, and the same low-rung filter applies to feedback. The corrective signal that would expose error is never registered as evidence — it is registered as disloyalty, hostility, or sin.

This is the dominant failure mode of *Religion* (the Part IV worked example): rung-1 control downward is honest and

stable, but rung-3 reality (abuse evidence, scientific findings, demographic harm) tries to enter the upward channel and is reclassified as rung-1 noise.

Different rungs on the two arrows of one loop

The most subtle pattern: the *same loop* operates at different rungs in each direction. The controller issues commands at rung 1 but expects feedback at rung 3. Or it issues commands at rung 3 but receives feedback only at rung 1. Either way the loop is structurally broken: information flows but does not correct.

This is what the cross-system catalogue identifies in religion, the unreformed military justice system, board-captured corporations, and family systems where the privacy norm absorbs all upward signal into a rung-1 filter.

Three dangerous mismatches

The seven rungs admit many possible per-arrow combinations. Three recur across the catalogued social systems and account for most of the rung-related failures.

Pattern A — Asymmetric loop

Rung-1 control downward, rung-3 (or higher) feedback expected upward. The controller's process model rejects the higher-rung feedback unless it arrives wrapped in rung-1 packaging — which strips the empirical content.

Generic remedy: insert an independent rung-3 channel that parallels (does not replace) the rung-1 channel and reaches the apex of control directly. Mandatory external audit, independent ombudspersons, pre-registered studies, and constitutional courts hearing rung-3 evidence against rung-1 sovereign claims are all implementations.

Pattern B — Claimed-rung inflation

The system publicly claims a high rung (rung 6 deliberation, rung 4 cumulative evidence) but operates at rung 0/1. The catastrophic systems of the twentieth century — one-party states claiming "people's democracy," theocracies claiming sacralised ultimate legitimacy — are inflations of this kind.

Generic remedy: close the gap. Either lower the claim to match operation, or raise the operation to match the claim. The first is rare because high claims confer near-term legitimacy; the second is what most catalogued architectural remedies attempt.

Pattern C — Cross-loop rung imposition

A controller legitimate at one rung in its own loop tries to extend that rung's authority into adjacent loops where a different rung is appropriate. Rung-1 religious authority claiming jurisdiction over rung-3 scientific questions; rung-3 scientific authority ruling on rung-6 normative questions; markets pricing questions whose decisive considerations sit at rung 6.

Generic remedy: domain-appropriate rung containment via subsidiarity rules and standing forums for the rung the question requires.

The full catalogue is in `knowledge/se-techniques/justification-rungs/dangerous-mismatches.md`.

Claimed rung vs operating rung

Every controller has *two* rungs. The **claimed rung** is the standard the controller's public legitimacy appeals to — what the system tells participants and outsiders it operates on. The **operating rung** is the standard at which control actions and feedback actually run.

The gap between the two is itself a diagnostic. The widest gaps produce the most catastrophic failures, because participants experience the gap empirically and the system's legitimacy collapses faster than the system can correct.

Across the ten social systems in this book's catalogue, the gaps distribute as follows:

System	Claimed rung	Operating rung	Gap
Kingdom	1	0/1	Small
Republic	6	3–6 mixed	Variable; opens with capture

System	Claimed rung	Operating rung	Gap
Theocracy	6 (sacralised)	1	Large
One-Party State	6	0/1	Extreme
Corporation	3	1–3 mixed	Moderate
University	3–4	1–3 mixed	Moderate
Military	1+2+6 (modern)	0/1	Small in tradition; widens under stress
Religion	1+6 (sacralised)	1	Large
Family	Variable	Variable	The gap is the diagnostic
Verein	6	6+2	Smallest

The full per-system analysis is in [knowledge/system-catalogues/social-systems/cross-system/justification-rungs-by-system.md](https://knowledge.system-catalogues/social-systems/cross-system/justification-rungs-by-system.md).

A normative caveat

STPA in its original form is deliberately neutral about which controllers are "better" — it describes control loops without ranking them. The rung ladder is openly normative: rung 6 sits above rung 0 in a way that rung 0 cannot reciprocate. We flag this explicitly rather than smuggle it in.

The ladder is a useful add-on *for social systems where epistemic standards matter*. It is not a generic STPA tool. Engineering systems mostly operate at rung 3–4 throughout, the rungs do not mismatch, and the ladder adds nothing. Use this technique where the question being asked is *whether the loop's rung fits its function* — not where the question is purely informational.

Forward references

- The technique is applied to the religion case study in STPA on Religion. Religion is the cleanest example of Pattern A asymmetric-loop failure.
- The dangerous patterns and their remedies generalise across all ten social systems — see Control Structures in Social Systems and Architectural Remedies.
- The full canonical reference is at knowledge/se-techniques/justification-rungs/.

Part II — Social Systems

This part applies the SE hierarchy and product line framework to ten social systems across three families.

The Systems

Governance Family

Systems that rule a territory and its population:

- Kingdom (Traditional Monarchy)
- Republic (Modern Democratic)
- Theocracy
- One-Party State

Organization Family

Systems that coordinate collective effort toward a declared purpose:

- Corporation (For-Profit)
- University
- Military
- Religion (Church / Religious Organisation)

Intimate Systems Family

Systems at the scale of direct human relationships:

- The Family
- The Verein (Voluntary Association)

Chapters

- **Religion** — The founding decomposition that started this analysis
- **Kingdom & Republic** — Side-by-side comparison with reuse analysis
- **Ten Social Systems Compared** — Full decompositions, platform analysis, variation points, cross-family reuse, and interactive explorers

Religion — SE Decomposition

The Founding Example

Religion was the first social system we decomposed through the five-level SE hierarchy. It serves as the worked example throughout this book and as the input to the STPA analysis in Part IV.

The Decomposition

Goal

- G1** — Provide existential meaning and purpose to human life
- G2** — Establish moral and ethical orientation for individuals and communities
- G3** — Create social cohesion and collective identity
- G4** — Offer psychological comfort in the face of suffering and mortality

Req

- R1** — Coherent cosmological narrative (creation, afterlife, transcendence)
- R2** — Framework for interpreting personal experience as meaningful
- R3** — Codified behavioral norms (commandments, precepts, dharma)
- R4** — Authority structure to interpret and enforce norms
- R5** — Shared symbols, narratives, and practices for group bonding
- R6** — Mechanism for distinguishing in-group from out-group
- R7** — Rituals addressing grief, loss, and life transitions

Func

- F1** — Mythological narration — transmit origin and eschatological stories
- F2** — Pastoral care — counsel individuals through crisis and life stages
- F3** — Moral legislation — define permissible and forbidden behavior
- F4** — Doctrinal governance — interpret scripture, settle disputes
- F5** — Ritual performance — enact shared ceremonies (prayer, sacraments, festivals)
- F6** — Missionary expansion — propagate belief system to new populations

Log

- L1** — Liturgical Subsystem — scripture reading, preaching, ceremonial calendar
- L2** — Pastoral Subsystem — counseling, rites of passage, community support
- L3** — Doctrinal/Juridical Subsystem — theology, canon law, ethics boards
- L4** — Outreach Subsystem — evangelization, education, charity, media

Phys

- P1** — Churches, mosques, temples, synagogues, gurdwaras
- P2** — Sacred texts — Bible, Quran, Torah, Vedas, Tripitaka
- P3** — Clergy — priests, imams, rabbis, pastors, monks, gurus
- P4** — Parish networks, sanghas, congregations, umma structures
- P5** — Central authorities — Vatican, Al-Azhar, ecumenical councils
- P6** — Legal systems — canon law, fatwa apparatus, halakhic courts, Sharia boards
- P7** — Missionary organizations, madrasas, Sunday schools, yeshivas
- P8** — Religious media — TV networks, podcasts, publishing houses

Key Structural Observations

R4 (authority) is the most consequential requirement. It enables F3 (moral legislation) and F4 (doctrinal governance), which together determine how the system interprets and enforces its own norms. When R4 becomes an end in itself — authority for authority's sake — the system begins to produce outcomes opposite to its goals. This dynamic is analyzed in detail in the STPA analysis.

R6 (in-group/out-group) is a double-edged mechanism. It serves G3 (social cohesion) but contains no built-in limit. The same mechanism that creates belonging can escalate to dehumanization. The absence of a structural "circuit breaker" is a key hazard identified in the STPA.

The Liturgical and Pastoral subsystems (L1, L2) serve different goals that can conflict. L1 (liturgy) serves G1 (meaning) and G3 (cohesion) through shared ritual. L2 (pastoral care) serves G4 (comfort) through individual attention. When institutional resources are scarce, L1 tends to win — communal performance is more visible than individual care.

See Also

The political system that fuses religious authority with state governance shares many structural elements with this decomposition. Explore it in the



Theocracy Interactive Explorer — graph

LEVEL	ID	DESCRIPTION	↑	↓
⊙ Goals	G1	Implement divine law as foundation of governance	—	R1 R2
	G2	Ensure spiritual salvation and moral purity	—	R3 R6
	G3	Maintain order, security, and territory	—	R4 R5
	G4	Preserve clerical authority as sole divine interpreter	—	R2 R3
◆ Requirements	R1	Sacred text as supreme legal source	G1	F1
	R2	Clerical qualification for governance	G1 G4	F1 F4
	R3	Religious conformity enforcement	G2 G4	F3
	R4	Military loyal to the theocratic order	G3	F4
	R5	Revenue system (religious taxes, endowments)	G3	F5
	R6	Education transmitting doctrine as civic education	G2	F6
◇ Functions	F1	Divine legislation — derive laws from scripture	R1 R2	L1
	F2	Religious adjudication — disputes per sacred law	—	L2
	F3	Moral enforcement — police virtue, suppress dissent	R3	L3
	F4	Military command and territorial defense	R2 R4	L4
	F5	Revenue collection and treasury	R5	L5
	F6	Doctrinal education	R6	L6
○ Logical Architecture	L1	Clerical-Legislative — supreme authority, councils	F1	P1 P2
	L2	Religious-Judicial — Sharia courts, tribunals	F2	P3
	L3	Enforcement — morality police, censorship	F3	P4
	L4	Military — army, revolutionary guard	F4	P5
	L5	Fiscal — treasury, religious taxes	F5	P6
	L6	Education — seminaries, madrasas	F6	P7
■ Physical Implementation	P1	Supreme Leader / Council of Guardians	L1	—
	P2	Sacred text as constitutional document	L1	—
	P3	Religious courts, Sharia judges	L2	—
	P4	Morality police, censorship boards	L3	—
	P5	Revolutionary Guard / military forces	L4	—
	P6	Religious treasury, zakat, endowments (waqf)	L5	—
	P7	Seminaries (Hawza), madrasas	L6	—

Theocracy Interactive Explorer — table

Kingdom & Republic — A Reuse Analysis

The Constitutional Monarchy as a Product-Line Variant



→ Interactive Explorer: Republic/Democracy — graph

LEVEL	ID	DESCRIPTION	↑	↓
⊙ Goals	G1	Derive governing authority from consent of the governed	—	R1 R2
	G2	Protect individual rights and prevent tyranny	—	R2 R3 R4
	G3	Ensure order, security, and territorial integrity	—	R3 R5
	G4	Enable peaceful, periodic transfer of executive power	—	R1 R6
◆ Requirements	R1	Free, periodic elections with universal suffrage	G1 G4	F1 F4
	R2	Codified constitution with enforceable rights	G1 G2	F1 F3
	R3	Separation of powers — legislative, executive, judicial	G2 G3	F1 F2 F3
	R4	Independent judiciary with constitutional review	G2	F3
	R5	Professional civil service	G3	F2
	R6	Free press and public accountability	G4	F4
◇ Functions	F1	Legislative process — enact laws through parliament	R1 R2 R3	L1
	F2	Executive governance — implement policy, manage ministries	R3 R5	L2
	F3	Constitutional adjudication — protect rights	R2 R3 R4	L3
	F4	Electoral & accountability administration	R1 R6	L4
○ Logical Architecture	L1	Legislative Subsystem — parliament, committees	F1	P1 P2
	L2	Executive Subsystem — head of government, ministries	F2	P3
	L3	Judicial Subsystem — constitutional court, courts	F3	P4
	L4	Electoral & Transparency Subsystem	F4	P5
■ Physical Implementation	P1	Parliament (Bundestag, Congress)	L1	—
	P2	Constitution, bill of rights	L1	—
	P3	President/Chancellor, cabinet, civil servants	L2	—
	P4	Constitutional court (BVerfG, Supreme Court)	L3	—
	P5	Election commission, polling stations, free press	L4	—

→ Interactive Explorer: Republic/Democracy — table

Key Finding

A constitutional monarchy reuses approximately **80% of the republican physical implementation** and **60% of its logical architecture**, while preserving the one goal-level requirement (dynastic continuity) that defines it as a distinct system variant.

Reuse Classification

Fully Reusable — Adopt Directly

Republican Element	Kingdom Equivalent	Historical Example
Independent judiciary	Royal courts → constitutional courts	UK Constitutional Reform Act 2005
Professional service	civil Patronage officials → merit-based Beamte	Stein-Hardenberg reforms, Prussia 1807
Modern fiscal system	Royal treasury → transparent tax authority + audit	Prussian fiscal reforms 1807–1815
Free press	No equivalent → introduce press freedom	Sweden's Freedom of the Press Act 1766

Partially Reusable — Adapt with Modifications

Republican Element	Adaptation Required	Historical Example
Parliament	Co-legislative body alongside crown, not replacing it	UK Parliament Act 1911
Written constitution	Limits royal power to enumerated prerogatives	Belgian Constitution 1831
Elections	For legislature only — head of state remains hereditary	Every constitutional monarchy
Separation of powers	Partial — monarch bridges branches as constitutional organ	Norwegian Grunnloven 1814

Incompatible — Cannot Reuse Without Abolishing the Monarchy

Republican Element	Why Incompatible
Popular sovereignty as sole source of authority	Eliminates hereditary authority by definition
Periodic election of head of state	Replaces hereditary succession
Electoral administration for choosing the executive	Structurally replaces dynastic succession

The full decomposition of both systems and the detailed cross-system comparison are in Ten Social Systems Compared.

Product Line Thinking for Social Systems

Applying Systems Engineering Decomposition and Platform Reuse to Human Institutions

Abstract

Product line engineering — the discipline of designing a family of systems that share a common platform while differing at defined variation points — is a mature practice in automotive, aerospace, and software engineering. This chapter argues that the same framework applies with surprising precision to social systems: political regimes, organizations, religious institutions, and even families. By decomposing ten social systems through a five-level systems engineering hierarchy (Goals → Requirements → Functions → Logical Architecture → Physical Implementation), we reveal that human institutions share far more structural DNA than their surface differences suggest. We identify a "social system platform" — a common architecture present across nearly all institutions — and map the variation points that produce distinct system types. The result is a conceptual toolkit for understanding institutional reform, cross-cultural comparison, and the structural reasons why some institutions modernize successfully while others collapse under attempted change.

Canonical reference

The normative versions of the ten decompositions, the social-system platform, and the cross-system analyses are maintained as the project's knowledge base under [knowledge/system-catalogues/social-systems/](#). This chapter is the narrative overview; the knowledge base is the reference.

1. Introduction: Why Product Lines?

1.1 The Problem of Institutional Comparison

When we compare a kingdom to a republic, a corporation to a university, or a church to a military, we tend to focus on what makes them *different*. A king inherits power; a president is elected. A corporation maximizes shareholder value; a university pursues knowledge. These differences are real and important. But they obscure a deeper structural truth: these systems are far more alike than they are different.

Consider: every one of these systems needs to make collective decisions, allocate resources, resolve disputes, recruit and manage members, maintain legitimacy, and adapt to external change. The *mechanisms* differ, but the *functional slots* are nearly identical. This is precisely the situation that product line engineering was invented to exploit.

1.2 The Product Line Concept in Systems Engineering Terms

A product line is a set of systems that share a **common platform** but differ at defined **variation points**. The platform is the set of architectural elements — goals, requirements, functions, subsystems, and physical components — that are shared across all (or most) members of the family. In automotive terms, this is the chassis. In social-system terms, this is the governance function, the membership mechanism, the resource allocation subsystem.

Variation points are the locations in the architecture where the platform deliberately leaves a decision open. A variation point has a defined interface and a set of variants. In political terms: "here, executive authority is sourced and legitimated" is the interface; "hereditary succession" or "popular election" are the variants.

Binding decisions are the specific choices made at each variation point for a given product. Once a binding is made, it constrains which other bindings are compatible.

The full technique — platform, variation points, binding decisions, reuse gradient, coherence rule, workaround principle — is defined in [knowledge/se-techniques/product-line-engineering/](#).

1.3 The Five-Level SE Hierarchy

Throughout this chapter, we decompose every system using a five-level hierarchy: Goals → Requirements → Functions

→ Logical Architecture → Physical Implementation. Each item at one level traces to items at the next level down. This traceability is what makes the product line analysis possible: we can identify which elements are shared and which are variant, at every level of abstraction.

The level definitions and traceability rules are in <knowledge/se-techniques/goals-requirements-hierarchy/>.

2. The Decompositions: Ten Social Systems

We decompose ten social systems, grouped into three families:

- **Governance family** — Kingdom (Traditional Monarchy), Republic (Modern Democracy), Theocracy, One-Party State.
- **Organization family** — Corporation (For-Profit), University, Military, Church / Religious Organization.
- **Intimate systems family** — Family, Verein (Voluntary Association).

Each decomposition names the system's goals, requirements, functions, logical architecture, and physical implementation, with traceability across all five levels. The full tables are maintained in the knowledge base:

System	Decomposition file
Kingdom	knowledge/system-catalogues/social-systems/kingdom/se-decomposition.md
Democracy / Republic	knowledge/system-catalogues/social-systems/democracy/se-decomposition.md
Theocracy	knowledge/system-catalogues/social-systems/theocracy/se-decomposition.md
One-Party State	knowledge/system-catalogues/social-systems/one-party-state/se-decomposition.md
Corporation	knowledge/system-catalogues/social-systems/corporation/se-decomposition.md
University	knowledge/system-catalogues/social-systems/university/se-decomposition.md
Military	knowledge/system-catalogues/social-systems/military/se-decomposition.md
Religion	knowledge/system-catalogues/social-systems/religion/se-decomposition.md
Family	knowledge/system-catalogues/social-systems/family/se-decomposition.md
Verein	knowledge/system-catalogues/social-systems/verein/se-decomposition.md

Each file presents the full five-level hierarchy with IDs, descriptions, and cross-level traceability. The interactive explorers linked at the bottom of this chapter render the same decompositions as browsable graphs.

3. The Social System Platform

Once you align the ten decompositions level by level, a common core emerges: a set of functional slots that every system fills, however differently.

The slots are: **Authority**, **Membership**, **Resource Allocation**, **Norms**, **Dispute Resolution**, **Legitimation**, **Succession**, **External Representation**, **Socialisation**, and **Activity Delivery**. Every social system — whether a monarchy, a corporation, or a family — contains all ten slots, because they correspond to problems every collective must solve.

The full platform description (goals, requirements, functions, logical, and physical elements shared across all ten systems) is in <knowledge/system-catalogues/social-systems/cross-system/platform.md>.

The platform is not an abstract ideal — it is the specific structural residue that appears when you decompose enough systems the same way and look at what is left in common.

4. The Variation Points

What varies across the family is concentrated at a small number of well-defined variation points. Six have been identified:

- **VP1 — Source of Authority** (hereditary · popular · party · divine · capital · voluntary · statutory · ...)
- **VP2 — Membership Boundary** (birth · contract · profession · faith · election · ...)
- **VP3 — Decision-Making Mechanism** (command · vote · consensus · ritual · market · ...)
- **VP4 — Succession Mechanism** (inheritance · election · appointment · nomination · ...)
- **VP5 — Legitimation Narrative** (tradition · law · results · doctrine · contract · ...)
- **VP6 — Norm Enforcement Mechanism** (coercion · social sanction · ritual · employment · ...)

Variants are not freely combinable: some combinations are coherent, others are structurally contradictory. The full set of variation points, observed variants, and interaction constraints is in <knowledge/system-catalogues/social-systems/cross-system/variation-points.md>.

5. Cross-Family Reuse: Case Studies

The binding decisions at each variation point constrain which elements can be transplanted between systems. Five worked cases illustrate this:

1. **Kingdom adopting republican structures** — the constitutional monarchy as a workaround that preserves hereditary legitimation while grafting on parliamentary legislation.
2. **Church adopting corporate structures** — institutional professionalization as partial reuse of corporate management under a divine-mandate top.
3. **University adopting corporate structures** — New Public Management as an amber-zone reuse with documented side effects.
4. **One-party state adopting republican fiscal structures** — China's reform model as selective transplantation below the authority variation point.
5. **Family adopting Verein structures** — democratic parenting as controlled opening of parental authority.

All five case studies are in <knowledge/system-catalogues/social-systems/cross-system/reuse-case-studies.md>.

6. Implications and Principles

Four principles recur across the case studies:

- **The Reuse Gradient.** Reuse feasibility decreases as you move upward in the SE hierarchy. Physical-level reuse is easy; goal-level reuse is almost impossible without transforming the system.
- **The Coherence Rule.** A system's binding decisions must be mutually compatible. Incoherent systems either resolve the contradiction (reform or revolution) or persist in institutional dysfunction.
- **The Workaround Principle.** When a desirable element conflicts with a binding decision, the most successful reforms create a modified version that preserves the function while respecting the constraint.
- **The Product Line Architect's Toolkit.** An eight-step procedure for designing or reforming a system: decompose both source and target, identify the platform, map the variation points, check binding constraints, classify reuse potential, start at the bottom, design workarounds, accept red items as system-defining.

Full statements of all four principles are in

7. Conclusion

Social systems look more different from each other than they actually are. A kingdom and a republic, a church and a corporation, a family and a Verein – these seem like categorically different things. But when decomposed through a rigorous systems engineering hierarchy, they reveal a shared platform of functional slots, and they differ primarily at a small number of well-defined variation points.

Product line thinking does not reduce human institutions to interchangeable parts. It does the opposite: it makes the *genuine* differences visible by stripping away the illusion that everything is different. Once you see that a kingdom's judicial system and a republic's judicial system fill the same functional slot with similar physical implementations, you can focus your analytical energy on the *actual* variation point – the source of executive authority – and reason clearly about which reforms are compatible with that binding and which are not.

This framework is not merely academic. Every successful institutional modernization in history – the Stein-Hardenberg reforms, the Meiji Restoration, the post-WWII German constitution, Vatican II, China's economic opening – can be understood as a product-line reuse exercise: identify what can be transplanted, adapt what must be adapted, and respect what cannot be changed without changing the system's identity.

The five-level SE decomposition is the analytical instrument. The variation point map is the strategic guide. And the coherence rule is the warning system. Together, they form a toolkit for anyone who wants to reform an institution – or simply to understand why institutions are the way they are.

Interactive Explorers

Explore each social system's SE decomposition interactively:

System	Explorer
Corporation	



→ Open — graph

LEVEL	ID	DESCRIPTION	↑	↓
⊙ Goals	G1	Generate sustainable profit and shareholder value	—	R1 R3
	G2	Deliver products/services satisfying market demand	—	R3 R5
	G3	Provide livelihood and purpose to employees	—	R4 R5
	G4	Ensure organizational continuity	—	R1 R2
◆ Requirements	R1	Legal incorporation with limited liability	G1 G4	F1 F5
	R2	Board accountable to shareholders	G4	F1
	R3	Revenue exceeding costs	G1 G2	F2 F3
	R4	Hierarchical management with clear authority	G3	F1 F4
	R5	Employment contracts	G2 G3	F4
◇ Functions	F1	Strategic direction — vision, budgets, appointments	R1 R2 R4	L1
	F2	Product/service delivery	R3	L2
	F3	Financial management — accounting, treasury	R3	L3
	F4	HR management — recruit, develop, compensate	R4 R5	L4
	F5	Legal and compliance	R1	L5
○ Logical Architecture	L1	Governance — board, shareholders meeting, audit	F1	P1 P2
	L2	Operations — product dev, manufacturing, supply chain	F2	P3
	L3	Financial — accounting, treasury, controlling	F3	P4
	L4	HR — recruiting, payroll, training	F4	P5
	L5	Legal & Compliance — counsel, regulatory	F5	P1
■ Physical Implementation	P1	Articles of incorporation, board resolutions	L1 L5	—
	P2	CEO, C-suite, management hierarchy	L1	—
	P3	Factories, offices, IT, supply chain	L2	—
	P4	ERP (SAP/Oracle), bank accounts, annual report	L3	—
	P5	Employment contracts, HR software, training	L4	—

→ Open — table

| | Democracy (Republic) |




→ Open — graph

LEVEL	ID	DESCRIPTION	↑	↓
⊙ Goals	G1	Derive governing authority from consent of the governed	—	R1 R2
	G2	Protect individual rights and prevent tyranny	—	R2 R3 R4
	G3	Ensure order, security, and territorial integrity	—	R3 R5
	G4	Enable peaceful, periodic transfer of executive power	—	R1 R6
◆ Requirements	R1	Free, periodic elections with universal suffrage	G1 G4	F1 F4
	R2	Codified constitution with enforceable rights	G1 G2	F1 F3
	R3	Separation of powers — legislative, executive, judicial	G2 G3	F1 F2 F3
	R4	Independent judiciary with constitutional review	G2	F3
	R5	Professional civil service	G3	F2
	R6	Free press and public accountability	G4	F4
◇ Functions	F1	Legislative process — enact laws through parliament	R1 R2 R3	L1
	F2	Executive governance — implement policy, manage ministries	R3 R5	L2
	F3	Constitutional adjudication — protect rights	R2 R3 R4	L3
	F4	Electoral & accountability administration	R1 R6	L4
○ Logical Architecture	L1	Legislative Subsystem — parliament, committees	F1	P1 P2
	L2	Executive Subsystem — head of government, ministries	F2	P3
	L3	Judicial Subsystem — constitutional court, courts	F3	P4
	L4	Electoral & Transparency Subsystem	F4	P5
■ Physical Implementation	P1	Parliament (Bundestag, Congress)	L1	—
	P2	Constitution, bill of rights	L1	—
	P3	President/Chancellor, cabinet, civil servants	L2	—
	P4	Constitutional court (BVerfG, Supreme Court)	L3	—
	P5	Election commission, polling stations, free press	L4	—

→ *Open — table*

|| Family |




→ Open — graph

LEVEL	ID	DESCRIPTION	↑	↓
⊙ Goals	G1	Ensure reproduction and survival of offspring	—	R1 R2
	G2	Provide emotional security, love, belonging	—	R3
	G3	Socialize children into norms and competences	—	R4 R6
	G4	Pool resources and manage economic risk	—	R2 R5
	G5	Transmit identity, heritage, property across generations	—	R5
◆ Requirements	R1	Pair bond or co-parenting agreement	G1	F1 F3
	R2	Shelter and material provision	G1 G4	F1 F4
	R3	Emotional availability and care	G2	F1 F2
	R4	Authority structure for child-rearing	G3	F3
	R5	Inheritance and succession norms	G4 G5	F4 F5
	R6	Interface with schools, healthcare, state	G3	F6
◇ Functions	F1	Nurture — feed, shelter, protect, care	R1 R2 R3	L2
	F2	Socialize — teach language, norms, values	R3	L4
	F3	Govern — discipline, education, health decisions	R1 R4	L1
	F4	Economic management — earn, budget, save	R2 R5	L3
	F5	Ritual and identity — traditions, storytelling	R5	L4
	F6	External representation — school, state interface	R6	L1
○ Logical Architecture	L1	Parental Authority — decisions, discipline	F3 F6	P1 P2
	L2	Care — nurturing, health, emotional support	F1	P1
	L3	Economic — income, budgeting, property	F4	P3
	L4	Socialization — education, culture, values	F2 F5	P4 P5
■ Physical Implementation	P1	Home — house or apartment, shared space	L1 L2	—
	P2	Parents/guardians, marriage/partnership	L1	—
	P3	Family income, bank accounts, property, wills	L3	—
	P4	Daily routines, rituals, photo albums, heirlooms	L4	—
	P5	School enrollment, pediatrician, community	L4	—

→ Open — table

| | Military |



→ Open — graph

LEVEL	ID	DESCRIPTION	↑	↓
⊙ Goals	G1	Defend the state against external threats	—	R1 R2
	G2	Maintain internal order when directed	—	R1 R3
	G3	Project power and deter aggression	—	R2 R4
	G4	Disaster relief and humanitarian assistance	—	R2 R5
◆ Requirements	R1	Strict command chain from top to soldier	G1 G2	F1
	R2	Trained, equipped, deployable force	G1 G3 G4	F2 F3
	R3	Subordination to civilian authority	G2	F1
	R4	Intelligence and situational awareness	G3	F4
	R5	Logistics — supply, transport, medical	G4	F5
	R6	Military justice framework	—	F6
◇ Functions	F1	Command — orders, coordination, strategy	R1 R3	L1
	F2	Combat — engage adversaries, defend	R2	L2
	F3	Training — drill, exercises, qualification	R2	L3
	F4	Intelligence — collect, analyze, disseminate	R4	L4
	F5	Logistics — ammo, fuel, food, medical	R5	L5
	F6	Military justice — discipline, courts martial	R6	L6
○ Logical Architecture	L1	Command — general staff, HQ, C4ISR	F1	P1
	L2	Combat — brigades, vessels, air wings	F2	P2
	L3	Training — academies, schools, ranges	F3	P3
	L4	Intelligence — SIGINT, HUMINT, cyber	F4	P4
	L5	Logistics — depots, transport, medical corps	F5	P5
	L6	Military Justice — courts martial, JAG, MP	F6	P6
■ Physical Implementation	P1	General staff HQ, situation rooms, comms	L1	—
	P2	Barracks, bases, warships, aircraft, armor	L2	—
	P3	Academies (West Point, Sandhurst, Führungsakademie)	L3	—
	P4	Intelligence agencies, SIGINT stations	L4	—
	P5	Supply depots, field hospitals, transport	L5	—
	P6	Military courts, provost marshal	L6	—

→ Open — table

| | One-Party State |




→ Open — graph

LEVEL	ID	DESCRIPTION	↑	↓
⊙ Goals	G1	Implement party ideology as organizing principle	—	R1 R3
	G2	Maintain party monopoly on political power	—	R1 R2 R4
	G3	Ensure order, security, and territory	—	R4 R5
	G4	Deliver economic development as legitimacy proof	—	R5 R6
◆ Requirements	R1	Constitutional enshrinement of party's leading role	G1 G2	F1
	R2	Party-state fusion — parallel hierarchies everywhere	G2	F1 F2
	R3	Ideological conformity — censorship, propaganda	G1	F3
	R4	Security apparatus loyal to the party	G2 G3	F4
	R5	Centralized economic direction	G3 G4	F5
	R6	Cadre system — party membership for advancement	G4	F2
◇ Functions	F1	Party-directed legislation	R1 R2	L1
	F2	Cadre management — recruit, evaluate, promote	R2 R6	L2
	F3	Ideological control — propaganda, censorship	R3	L3
	F4	Security and surveillance	R4	L4
	F5	Economic direction — plans, SOEs, investment	R5	L5
○ Logical Architecture	L1	Party-State Legislative — politburo, congress	F1	P1
	L2	Cadre & Personnel — org department	F2	P2
	L3	Propaganda & Control — media, censorship	F3	P3
	L4	Security — secret police, surveillance	F4	P4
	L5	Economic Planning — commission, SOEs	F5	P5
■ Physical Implementation	P1	Politburo / Central Committee	L1	—
	P2	Organization Department, party schools, cadre files	L2	—
	P3	State media (Xinhua, CCTV), Great Firewall	L3	—
	P4	Secret police (Stasi, KGB, MSS), surveillance	L4	—
	P5	State Planning Commission, SOEs, 5-year plans	L5	—

→ *Open — table*

|| Theocracy |



→ Open — graph

LEVEL	ID	DESCRIPTION	↑	↓
⊙ Goals	G1	Implement divine law as foundation of governance	—	R1 R2
	G2	Ensure spiritual salvation and moral purity	—	R3 R6
	G3	Maintain order, security, and territory	—	R4 R5
	G4	Preserve clerical authority as sole divine interpreter	—	R2 R3
◆ Requirements	R1	Sacred text as supreme legal source	G1	F1
	R2	Clerical qualification for governance	G1 G4	F1 F4
	R3	Religious conformity enforcement	G2 G4	F3
	R4	Military loyal to the theocratic order	G3	F4
	R5	Revenue system (religious taxes, endowments)	G3	F5
	R6	Education transmitting doctrine as civic education	G2	F6
◇ Functions	F1	Divine legislation — derive laws from scripture	R1 R2	L1
	F2	Religious adjudication — disputes per sacred law	—	L2
	F3	Moral enforcement — police virtue, suppress dissent	R3	L3
	F4	Military command and territorial defense	R2 R4	L4
	F5	Revenue collection and treasury	R5	L5
	F6	Doctrinal education	R6	L6
○ Logical Architecture	L1	Clerical-Legislative — supreme authority, councils	F1	P1 P2
	L2	Religious-Judicial — Sharia courts, tribunals	F2	P3
	L3	Enforcement — morality police, censorship	F3	P4
	L4	Military — army, revolutionary guard	F4	P5
	L5	Fiscal — treasury, religious taxes	F5	P6
	L6	Education — seminaries, madrasas	F6	P7
■ Physical Implementation	P1	Supreme Leader / Council of Guardians	L1	—
	P2	Sacred text as constitutional document	L1	—
	P3	Religious courts, Sharia judges	L2	—
	P4	Morality police, censorship boards	L3	—
	P5	Revolutionary Guard / military forces	L4	—
	P6	Religious treasury, zakat, endowments (waqf)	L5	—
	P7	Seminaries (Hawza), madrasas	L6	—

→ Open — table



→ Open — graph

LEVEL	ID	DESCRIPTION	↑	↓
⊙ Goals	G1	Advance knowledge through research	—	R1 R5
	G2	Educate and credential the next generation	—	R2 R6
	G3	Serve society through knowledge transfer	—	R1 R2
	G4	Preserve academic freedom and autonomy	—	R3 R5
◆ Requirements	R1	Peer review as knowledge validation standard	G1 G3	F1
	R2	Structured curricula leading to degrees	G2 G3	F2
	R3	Academic self-governance — senates, elected rectors	G4	F3
	R4	Sustainable funding — tuition, grants, state	—	F5
	R5	Tenure protecting freedom of inquiry	G1 G4	F1 F3
	R6	Accreditation by external bodies	G2	F4
◇ Functions	F1	Research — conduct and publish scholarship	R1 R5	L2
	F2	Teaching — lectures, seminars, examinations	R2	L3
	F3	Academic governance — senate, appointments	R3 R5	L1
	F4	Student administration — admissions, graduation	R6	L4
	F5	Financial management — budgets, grants	R4	L4
○ Logical Architecture	L1	Governance — senate, rector, faculty councils	F3	P1 P2
	L2	Research — departments, institutes, labs, libraries	F1	P3
	L3	Teaching — curricula, lecture halls, exams	F2	P3 P4
	L4	Administration — student office, finance, HR	F4 F5	P4 P5
■ Physical Implementation	P1	University charter, Grundordnung	L1	—
	P2	Rector, deans, senate, faculty councils	L1	—
	P3	Labs, libraries, lecture halls, institutes	L2 L3	—
	P4	Student registry (HISinOne), exam office	L3 L4	—
	P5	Endowment, state funding, DFG/ERC grants	L4	—

→ Open — table



→ Open — graph

LEVEL	ID	DESCRIPTION	↑	↓
⊙ Goals	G1	Enable citizens to pursue a shared non-commercial purpose collectively	—	R1 R2
	G2	Provide a democratic, self-governing legal structure	—	R3 R4
	G3	Create social cohesion and community identity	—	R5 R6
	G4	Ensure institutional continuity independent of individuals	—	R2 R4 R7
◆ Requirements	R1	Written Satzung defining the non-commercial purpose	G1	F1
	R2	Legal personality through Vereinsregister (e.V.)	G1 G4	F1 F2
	R3	Democratic decision-making with equal voting rights	G2	F3
	R4	Elected Vorstand accountable to membership	G2 G4	F3 F4
	R5	Regular activities bringing members together	G3	F5
	R6	Open, voluntary membership	G3	F2
	R7	Financial sustainability through dues, donations, grants	G4	F6
◇ Functions	F1	Constitutional governance — draft and enforce Satzung	R1 R2	L1
	F2	Membership administration — admit, track, manage	R2 R6	L2
	F3	Democratic assembly — Mitgliederversammlungen, elections	R3 R4	L1
	F4	Executive management — represent Verein externally	R4	L3
	F5	Activity delivery — events, training, competitions	R5	L4
	F6	Financial management — dues, bookkeeping, Gemeinnützigkeit	R7	L2 L3
○ Logical Architecture	L1	Governance — Satzung, Versammlung, Wahlordnung	F1 F3	P1 P2
	L2	Administration — member database, dues, Vereinsregister	F2 F6	P3 P4
	L3	Executive — Vorstand, contracts, bank account	F4 F6	P5 P6
	L4	Operations — Abteilungen, coaches, facilities	F5	P7 P8
■ Physical Implementation	P1	Satzung, Geschäftsordnung, Beitragsordnung	L1	—
	P2	Mitgliederversammlung, Protokolle, Kassenprüfung	L1	—
	P3	Membership software (SEWOBE, Vereinsflieger)	L2	—
	P4	Vereinsregister entry, Finanzamt registration	L2	—
	P5	Vorstand (Vorsitzender, Schatzmeister, Schriftführer)	L3	—
	P6	Vereinskonto, DATEV, Zuwendungsbestätigungen	L3	—
	P7	Vereinsheim, sports fields, rehearsal rooms	L4	—
	P8	Abteilungen (Fußball, Turnen, Chor), Übungsleiter	L4	—

→ Open — table

Example: Der Deutsche Verein (e.V.)

The Verein is Germany's universal civic institution — over 600,000 registered associations covering everything from football to beekeeping. It is the simplest social system that instantiates all ten platform functional slots, making it an ideal worked example.



** → Interactive D3 Visualization** — graph

LEVEL	ID	DESCRIPTION	↑	↓
⊙ Goals	G1	Enable citizens to pursue a shared non-commercial purpose collectively	—	R1 R2
	G2	Provide a democratic, self-governing legal structure	—	R3 R4
	G3	Create social cohesion and community identity	—	R5 R6
	G4	Ensure institutional continuity independent of individuals	—	R2 R4 R7
◆ Requirements	R1	Written Satzung defining the non-commercial purpose	G1	F1
	R2	Legal personality through Vereinsregister (e.V.)	G1 G4	F1 F2
	R3	Democratic decision-making with equal voting rights	G2	F3
	R4	Elected Vorstand accountable to membership	G2 G4	F3 F4
	R5	Regular activities bringing members together	G3	F5
	R6	Open, voluntary membership	G3	F2
	R7	Financial sustainability through dues, donations, grants	G4	F6
◇ Functions	F1	Constitutional governance — draft and enforce Satzung	R1 R2	L1
	F2	Membership administration — admit, track, manage	R2 R6	L2
	F3	Democratic assembly — Mitgliederversammlungen, elections	R3 R4	L1
	F4	Executive management — represent Verein externally	R4	L3
	F5	Activity delivery — events, training, competitions	R5	L4
	F6	Financial management — dues, bookkeeping, Gemeinnützigkeit	R7	L2 L3
○ Logical Architecture	L1	Governance — Satzung, Versammlung, Wahlordnung	F1 F3	P1 P2
	L2	Administration — member database, dues, Vereinsregister	F2 F6	P3 P4
	L3	Executive — Vorstand, contracts, bank account	F4 F6	P5 P6
	L4	Operations — Abteilungen, coaches, facilities	F5	P7 P8
■ Physical Implementation	P1	Satzung, Geschäftsordnung, Beitragsordnung	L1	—
	P2	Mitgliederversammlung, Protokolle, Kassenprüfung	L1	—
	P3	Membership software (SEWOBE, Vereinsflieger)	L2	—
	P4	Vereinsregister entry, Finanzamt registration	L2	—
	P5	Vorstand (Vorsitzender, Schatzmeister, Schriftführer)	L3	—
	P6	Vereinskonto, DATEV, Zuwendungsbestätigungen	L3	—
	P7	Vereinsheim, sports fields, rehearsal rooms	L4	—
	P8	Abteilungen (Fußball, Turnen, Chor), Übungsleiter	L4	—

** → Interactive D3 Visualization** — table

SE Decomposition

The full five-level decomposition is shown in the interactive visualization linked above. Click any node to see its description, parent links, and child links. Use the Table view for the complete traceability matrix.

Variation Point Bindings

The Verein binds VP1 (authority) to *voluntary democratic participation*, VP2 (membership) to *open and voluntary*, VP3 (decisions) to *one-person-one-vote*, and VP4 (succession) to *periodic election*. This makes it structurally closest to a republic, but at micro-scale and with a non-commercial purpose constraint.

Platform Mapping

This system fills all ten universal functional slots identified in the Ten Social Systems Compared:

Functional Slot	How This System Fills It
Authority & Decision-Making	Mitgliederversammlung (general assembly) as sovereign body; elected Vorstand handles day-to-day execution

Functional Slot	How This System Fills It
Membership & Belonging	Voluntary application; admission approved by Vorstand; conditions and fees defined in the Satzung
Resource Allocation	Membership fees (Beiträge); donations; assembly-approved annual budget; independent Kassenprüfer audit
Norm Setting & Enforcement	Satzung (articles of association); Geschäftsordnung (procedural rules); group social norms
Dispute Resolution	Ehrengericht (honour court) or Vorstand mediation; expulsion requires assembly vote by absolute majority
Legitimation	Shared non-commercial purpose; democratic participation; voluntary membership as constitutive feature
Succession & Continuity	Annual or biennial election of Vorstand at Mitgliederversammlung; written Satzung ensures institutional memory
External Representation	1. Vorsitzender (chair) as sole legal representative (Vertretungsberechtigung); umbrella organisation membership
Socialisation	Club culture and shared practices; activity rituals; mentoring of new members into group norms
Activity Delivery	Regular activities (training sessions, meetings, events, competitions) serving the stated purpose

*Navigate to the



interactive visualization — graph

LEVEL	ID	DESCRIPTION	↑	↓
⊙ Goals	G1	Enable citizens to pursue a shared non-commercial purpose collectively	–	R1 R2
	G2	Provide a democratic, self-governing legal structure	–	R3 R4
	G3	Create social cohesion and community identity	–	R5 R6
	G4	Ensure institutional continuity independent of individuals	–	R2 R4 R7
◆ Requirements	R1	Written Satzung defining the non-commercial purpose	G1	F1
	R2	Legal personality through Vereinsregister (e.V.)	G1 G4	F1 F2
	R3	Democratic decision-making with equal voting rights	G2	F3
	R4	Elected Vorstand accountable to membership	G2 G4	F3 F4
	R5	Regular activities bringing members together	G3	F5
	R6	Open, voluntary membership	G3	F2
	R7	Financial sustainability through dues, donations, grants	G4	F6
◇ Functions	F1	Constitutional governance — draft and enforce Satzung	R1 R2	L1
	F2	Membership administration — admit, track, manage	R2 R6	L2
	F3	Democratic assembly — Mitgliederversammlungen, elections	R3 R4	L1
	F4	Executive management — represent Verein externally	R4	L3
	F5	Activity delivery — events, training, competitions	R5	L4
	F6	Financial management — dues, bookkeeping, Gemeinnützigkeit	R7	L2 L3
○ Logical Architecture	L1	Governance — Satzung, Versammlung, Wahlordnung	F1 F3	P1 P2
	L2	Administration — member database, dues, Vereinsregister	F2 F6	P3 P4
	L3	Executive — Vorstand, contracts, bank account	F4 F6	P5 P6
	L4	Operations — Abteilungen, coaches, facilities	F5	P7 P8
■ Physical Implementation	P1	Satzung, Geschäftsordnung, Beitragsordnung	L1	–
	P2	Mitgliederversammlung, Protokolle, Kassenprüfung	L1	–
	P3	Membership software (SEWOBE, Vereinsflieger)	L2	–
	P4	Vereinsregister entry, Finanzamt registration	L2	–
	P5	Vorstand (Vorsitzender, Schatzmeister, Schriftführer)	L3	–
	P6	Vereinskonto, DATEV, Zuwendungsbestätigungen	L3	–
	P7	Vereinsheim, sports fields, rehearsal rooms	L4	–
	P8	Abteilungen (Fußball, Turnen, Chor), Übungsleiter	L4	–

interactive visualization — table

for the full graph and table.*

Example: Democracy (Modern Republic)

The democratic republic is the governance family's most complex variant, with the richest set of checks, balances, and feedback mechanisms of all ten systems analysed in this book. Its SE decomposition reveals why it is also the most structurally resilient — it has more redundant control paths, more independent feedback channels, and more explicit separation of authority than any other social system architecture.



** → Interactive D3 Visualization** — graph

LEVEL	ID	DESCRIPTION	↑	↓
⊙ Goals	G1	Derive governing authority from consent of the governed	—	R1 R2
	G2	Protect individual rights and prevent tyranny	—	R2 R3 R4
	G3	Ensure order, security, and territorial integrity	—	R3 R5
	G4	Enable peaceful, periodic transfer of executive power	—	R1 R6
◆ Requirements	R1	Free, periodic elections with universal suffrage	G1 G4	F1 F4
	R2	Codified constitution with enforceable rights	G1 G2	F1 F3
	R3	Separation of powers — legislative, executive, judicial	G2 G3	F1 F2 F3
	R4	Independent judiciary with constitutional review	G2	F3
	R5	Professional civil service	G3	F2
	R6	Free press and public accountability	G4	F4
◇ Functions	F1	Legislative process — enact laws through parliament	R1 R2 R3	L1
	F2	Executive governance — implement policy, manage ministries	R3 R5	L2
	F3	Constitutional adjudication — protect rights	R2 R3 R4	L3
	F4	Electoral & accountability administration	R1 R6	L4
○ Logical Architecture	L1	Legislative Subsystem — parliament, committees	F1	P1 P2
	L2	Executive Subsystem — head of government, ministries	F2	P3
	L3	Judicial Subsystem — constitutional court, courts	F3	P4
	L4	Electoral & Transparency Subsystem	F4	P5
■ Physical Implementation	P1	Parliament (Bundestag, Congress)	L1	—
	P2	Constitution, bill of rights	L1	—
	P3	President/Chancellor, cabinet, civil servants	L2	—
	P4	Constitutional court (BVerfG, Supreme Court)	L3	—
	P5	Election commission, polling stations, free press	L4	—

** → Interactive D3 Visualization** — table

SE Decomposition

The full five-level decomposition is shown in the interactive visualization above. Click any node to see its description and connections. Use the Table view for the complete traceability matrix.

Variation Point Bindings

Variation Point	Binding
VP1: Source of Authority	Popular sovereignty — the only legitimate source of governing power is the consent of the governed
VP2: Membership Boundary	Citizenship by territory, birth, or naturalisation
VP3: Decision-Making	Representative democratic vote — competitive elections with universal suffrage
VP4: Succession	Periodic election — no hereditary or self-perpetuating leadership

The key constraint: VP1 (popular sovereignty) structurally forces VP3 and VP4 to include genuine competitive elections. A system that calls itself democratic but controls elections, eliminates competition, or makes succession non-periodic has changed VP1 — it is no longer a democracy in the product-line sense, whatever its name.

Why This System Is Structurally Unusual

Every other system in this book concentrates authority at the top and manages the risk of that concentration through informal norms, cultural expectations, or individual virtue. The republic does something structurally different: it

distributes authority across institutions that are designed to check each other, and it builds the checking mechanism into the constitution rather than leaving it to goodwill.

The result is the highest feedback richness of any system: elections connect government behaviour to electoral consequences; independent courts connect government action to constitutional standards; a free press connects official narrative to observable reality; civil society connects individual grievance to collective voice. These channels operate in parallel and are deliberately independent of one another — so that capturing one does not disable the others.

This redundancy is not accidental. It is the product of historical learning from democratic failures — most clearly documented in the engineering of Germany's Basic Law after the Weimar Republic's collapse.

STPA Highlights

STPA identifies the conditions under which democracy's redundant architecture degrades. The most important unsafe control actions are:

UCA-D1: Legislature enacts legislation that dismantles constitutional democracy, with no body able to invalidate it.

The Weimar Republic had no mechanism for judicial review of constitutionality. The Enabling Act of 1933 was passed by a formally valid parliamentary vote. There was no circuit breaker. The causal factor is structural, not moral: even a legislature of good people cannot prevent constitutional erosion if no external authority can constrain it.

UCA-D2: Government collapses without a functioning successor, creating conditions for emergency rule.

Weimar's Article 48 combined easily dismissed coalition governments with presidential emergency powers. The path from parliamentary deadlock to authoritarian rule was structurally open. Fourteen chancellors in fourteen years exhausted the system's legitimacy before it was overthrown.

UCA-D3: An anti-constitutional party uses democratic procedures to eliminate democracy.

When a system's rules contain no defence against participants who intend to abolish the rules, the system can be closed by the actors it most needs to exclude. The structural openness that is democracy's greatest strength becomes a vulnerability without a corresponding structural defence.

UCA-D4: The information environment is captured, severing corrective feedback between government behaviour and electoral accountability.

When citizens cannot see what the government is doing, they cannot vote against it. Press freedom and media independence are not decorative features of democracy — they are the feedback channel that makes elections meaningful.

Fixes That Worked: Germany's Basic Law (1949)

The Federal Republic of Germany is the most thoroughly documented case of constitutional engineering explicitly designed to prevent identified structural failure modes. Its designers worked from what was in effect a structural post-mortem of Weimar — identifying each failure pattern and designing a structural countermeasure.

UCA	Weimar structural failure	Basic Law remedy
UCA-D1	No judicial review of constitutionality	Bundesverfassungsgericht — independent constitutional court with power to strike down unconstitutional legislation. Has struck down hundreds of laws since 1951.
UCA-D2	Easy government collapse + emergency powers	Constructive vote of no confidence — the Bundestag can only remove a Chancellor by simultaneously electing a successor. No more leadership vacuum. Invoked twice in 75 years, both times appropriately.
UCA-D3	No defence against anti-constitutional parties	Article 21 GG — parties seeking to undermine the free democratic basic order can be banned by the Constitutional Court. SRP banned 1952; KPD 1956. The mechanism functions as a structural deterrent.
UCA-D4	Centralised or capturable press	Public broadcasting (ARD, ZDF) governed by Rundfunkräte — councils of civil society representatives, not government officials. Financially independent.

The Design Principle

Germany's post-Weimar architects did not trust good intentions. They trusted structure. The Basic Law has supported a functioning democracy for 75 years through reunification, financial crises, and significant polarisation — not because Germans are inherently different, but because the control structure was deliberately engineered to be resilient to the failure modes that destroyed Weimar.

Platform Mapping

Functional Slot	How This System Fills It
Authority & Decision-Making	Popular sovereignty expressed through elected representative institutions
Membership & Belonging	Citizenship by birth, territory, or naturalisation; universal suffrage
Resource Allocation	Parliamentary budget process; taxation by elected legislature
Norm Setting & Enforcement	Legislative process with constitutional review; independent judiciary
Dispute Resolution	Independent courts; constitutional court for fundamental rights
Legitimation	Consent of the governed; periodic competitive elections
Succession & Continuity	Periodic elections; constitutional transfer of power
External Representation	Elected executive; treaty ratification by legislature
Socialisation	Public education; civic education; free press
Activity Delivery	Civil service; public services; rule of law

*For the full SE decomposition:



interactive visualization — graph

LEVEL	ID	DESCRIPTION	↑	↓
⊙ Goals	G1	Derive governing authority from consent of the governed	—	R1 R2
	G2	Protect individual rights and prevent tyranny	—	R2 R3 R4
	G3	Ensure order, security, and territorial integrity	—	R3 R5
	G4	Enable peaceful, periodic transfer of executive power	—	R1 R6
◆ Requirements	R1	Free, periodic elections with universal suffrage	G1 G4	F1 F4
	R2	Codified constitution with enforceable rights	G1 G2	F1 F3
	R3	Separation of powers — legislative, executive, judicial	G2 G3	F1 F2 F3
	R4	Independent judiciary with constitutional review	G2	F3
	R5	Professional civil service	G3	F2
	R6	Free press and public accountability	G4	F4
◇ Functions	F1	Legislative process — enact laws through parliament	R1 R2 R3	L1
	F2	Executive governance — implement policy, manage ministries	R3 R5	L2
	F3	Constitutional adjudication — protect rights	R2 R3 R4	L3
	F4	Electoral & accountability administration	R1 R6	L4
○ Logical Architecture	L1	Legislative Subsystem — parliament, committees	F1	P1 P2
	L2	Executive Subsystem — head of government, ministries	F2	P3
	L3	Judicial Subsystem — constitutional court, courts	F3	P4
	L4	Electoral & Transparency Subsystem	F4	P5
■ Physical Implementation	P1	Parliament (Bundestag, Congress)	L1	—
	P2	Constitution, bill of rights	L1	—
	P3	President/Chancellor, cabinet, civil servants	L2	—
	P4	Constitutional court (BVerfG, Supreme Court)	L3	—
	P5	Election commission, polling stations, free press	L4	—

interactive visualization — table

. For the comparative analysis across all ten systems: Ten Social Systems Compared. For the cross-system control structure analysis: Control Structures.*

Example: Theocracy

The theocracy derives all authority from a sacred source, making it the governance type with the strongest legitimization narrative — and the weakest feedback mechanisms. Its SE decomposition reveals a structurally closed system where doctrine, law, education, and enforcement form a self-reinforcing loop.



** → Interactive D3 Visualization** — graph

LEVEL	ID	DESCRIPTION	↑	↓
⊙ Goals	G1	Implement divine law as foundation of governance	—	R1 R2
	G2	Ensure spiritual salvation and moral purity	—	R3 R6
	G3	Maintain order, security, and territory	—	R4 R5
	G4	Preserve clerical authority as sole divine interpreter	—	R2 R3
◆ Requirements	R1	Sacred text as supreme legal source	G1	F1
	R2	Clerical qualification for governance	G1 G4	F1 F4
	R3	Religious conformity enforcement	G2 G4	F3
	R4	Military loyal to the theocratic order	G3	F4
	R5	Revenue system (religious taxes, endowments)	G3	F5
	R6	Education transmitting doctrine as civic education	G2	F6
◇ Functions	F1	Divine legislation — derive laws from scripture	R1 R2	L1
	F2	Religious adjudication — disputes per sacred law	—	L2
	F3	Moral enforcement — police virtue, suppress dissent	R3	L3
	F4	Military command and territorial defense	R2 R4	L4
	F5	Revenue collection and treasury	R5	L5
	F6	Doctrinal education	R6	L6
○ Logical Architecture	L1	Clerical-Legislative — supreme authority, councils	F1	P1 P2
	L2	Religious-Judicial — Sharia courts, tribunals	F2	P3
	L3	Enforcement — morality police, censorship	F3	P4
	L4	Military — army, revolutionary guard	F4	P5
	L5	Fiscal — treasury, religious taxes	F5	P6
	L6	Education — seminaries, madrasas	F6	P7
■ Physical Implementation	P1	Supreme Leader / Council of Guardians	L1	—
	P2	Sacred text as constitutional document	L1	—
	P3	Religious courts, Sharia judges	L2	—
	P4	Morality police, censorship boards	L3	—
	P5	Revolutionary Guard / military forces	L4	—
	P6	Religious treasury, zakat, endowments (waqf)	L5	—
	P7	Seminaries (Hawza), madrasas	L6	—

** → Interactive D3 Visualization** — table

SE Decomposition

The full five-level decomposition is shown in the interactive visualization linked above. Click any node to see its description, parent links, and child links. Use the Table view for the complete traceability matrix.

Variation Point Bindings

VP1 = divine mandate, VP2 = territorial + religious conformity, VP3 = doctrinal interpretation by clergy, VP4 = clerical selection (conclave, council, or dynastic). The binding at VP1 constrains VP3 so severely that popular input to governance is structurally impossible.

Platform Mapping

This system fills all ten universal functional slots identified in the Ten Social Systems Compared:

Functional Slot	How This System Fills It
Authority & Decision-Making	Supreme religious authority (Pope, Supreme Leader, Caliph); clerical council; no independent secular check

Functional Slot	How This System Fills It
Membership & Belonging	Territorial citizenship combined with required religious conformity; apostasy carries legal or social penalty
Resource Allocation	Religious endowments (waqf, church estates); state budget under clerical control; no separation of sacred and fiscal authority
Norm Setting & Enforcement	Religious law (canon law, sharia, halakha); clerical courts; religious police where applicable
Dispute Resolution	Religious courts; clerical arbitration; all dispute resolution channels pass through the religious hierarchy
Legitimation	Divine mandate; sacred texts as constitutional authority; clerical monopoly on authoritative interpretation
Succession & Continuity	Clerical selection: conclave, council deliberation, or dynastic clerical succession
External Representation	Supreme religious leader conducts state diplomacy; religious and foreign policy are indistinguishable
Socialisation	Mandatory religious education (madrasas, seminaries, religious schools); ritual participation as civic obligation
Activity Delivery	Religious services; welfare distribution (zakat, tithing); education; governance fused with religious practice

*Navigate to the



interactive visualization — graph


LEVEL	ID	DESCRIPTION	↑	↓
⊙ Goals	G1	Implement divine law as foundation of governance	—	R1 R2
	G2	Ensure spiritual salvation and moral purity	—	R3 R6
	G3	Maintain order, security, and territory	—	R4 R5
	G4	Preserve clerical authority as sole divine interpreter	—	R2 R3
◆ Requirements	R1	Sacred text as supreme legal source	G1	F1
	R2	Clerical qualification for governance	G1 G4	F1 F4
	R3	Religious conformity enforcement	G2 G4	F3
	R4	Military loyal to the theocratic order	G3	F4
	R5	Revenue system (religious taxes, endowments)	G3	F5
	R6	Education transmitting doctrine as civic education	G2	F6
◇ Functions	F1	Divine legislation — derive laws from scripture	R1 R2	L1
	F2	Religious adjudication — disputes per sacred law	—	L2
	F3	Moral enforcement — police virtue, suppress dissent	R3	L3
	F4	Military command and territorial defense	R2 R4	L4
	F5	Revenue collection and treasury	R5	L5
	F6	Doctrinal education	R6	L6
○ Logical Architecture	L1	Clerical-Legislative — supreme authority, councils	F1	P1 P2
	L2	Religious-Judicial — Sharia courts, tribunals	F2	P3
	L3	Enforcement — morality police, censorship	F3	P4
	L4	Military — army, revolutionary guard	F4	P5
	L5	Fiscal — treasury, religious taxes	F5	P6
	L6	Education — seminaries, madrasas	F6	P7
■ Physical Implementation	P1	Supreme Leader / Council of Guardians	L1	—
	P2	Sacred text as constitutional document	L1	—
	P3	Religious courts, Sharia judges	L2	—
	P4	Morality police, censorship boards	L3	—
	P5	Revolutionary Guard / military forces	L4	—
	P6	Religious treasury, zakat, endowments (waqf)	L5	—
	P7	Seminaries (Hawza), madrasas	L6	—

interactive visualization — table

for the full graph and table.*

Example: Corporation (For-Profit)

The for-profit corporation is the organizational family's most studied variant. Its SE decomposition reveals a system optimised for a single metric — profit — with governance designed to serve capital holders. This binding produces both extraordinary efficiency in capital allocation and a set of well-documented structural pathologies that follow directly from the architecture, not from individual moral failure.



** → Interactive D3 Visualization** — graph

LEVEL	ID	DESCRIPTION	↑	↓
⊙ Goals	G1	Generate sustainable profit and shareholder value	—	R1 R3
	G2	Deliver products/services satisfying market demand	—	R3 R5
	G3	Provide livelihood and purpose to employees	—	R4 R5
	G4	Ensure organizational continuity	—	R1 R2
◆ Requirements	R1	Legal incorporation with limited liability	G1 G4	F1 F5
	R2	Board accountable to shareholders	G4	F1
	R3	Revenue exceeding costs	G1 G2	F2 F3
	R4	Hierarchical management with clear authority	G3	F1 F4
	R5	Employment contracts	G2 G3	F4
◇ Functions	F1	Strategic direction — vision, budgets, appointments	R1 R2 R4	L1
	F2	Product/service delivery	R3	L2
	F3	Financial management — accounting, treasury	R3	L3
	F4	HR management — recruit, develop, compensate	R4 R5	L4
	F5	Legal and compliance	R1	L5
○ Logical Architecture	L1	Governance — board, shareholders meeting, audit	F1	P1 P2
	L2	Operations — product dev, manufacturing, supply chain	F2	P3
	L3	Financial — accounting, treasury, controlling	F3	P4
	L4	HR — recruiting, payroll, training	F4	P5
	L5	Legal & Compliance — counsel, regulatory	F5	P1
■ Physical Implementation	P1	Articles of incorporation, board resolutions	L1 L5	—
	P2	CEO, C-suite, management hierarchy	L1	—
	P3	Factories, offices, IT, supply chain	L2	—
	P4	ERP (SAP/Oracle), bank accounts, annual report	L3	—
	P5	Employment contracts, HR software, training	L4	—

→ Interactive D3 Visualization — table

SE Decomposition

The full five-level decomposition is shown in the interactive visualization above. Click any node to see its description and connections. Use the Table view for the complete traceability matrix.

Variation Point Bindings

Variation Point	Binding
VP1: Source of Authority	Capital ownership — authority is proportional to shareholding
VP2: Membership Boundary	Employment contract (voluntary, conditional, terminable)
VP3: Decision-Making	Board vote weighted by share ownership
VP4: Succession	Board appointment — not election, not merit review by peers

The key structural feature: employees are the system's primary operators — they hold the most accurate process models of what the system actually does — but they have no formal governance authority. The people who know most about the corporation's real functioning are structurally excluded from its decision-making.

The Structural Pathology of VP1 = Capital Ownership

Every variation-point binding creates trade-offs. VP1 = capital ownership has two well-documented consequences.

Time horizon misalignment. Shareholders can exit in milliseconds (liquid markets). Employees, suppliers, and

communities cannot exit as quickly. The governance structure therefore weights the interests of the fastest-exit stakeholder most heavily — which creates systematic pressure toward decisions that serve 90-day earnings cycles over multi-year value creation.

Operational knowledge exclusion. The people with the most accurate information about product quality, customer relationships, safety risks, and long-term capabilities — the workers — have no formal channel through which that information reaches the decision-making authority. Executives receive information that has been filtered upward through layers of management with their own incentives. The board receives information filtered through the executives. The process model at the governance level is systematically less accurate than the process model held by the operators.

These are not failures of bad management. They are structural consequences of the VP1 binding.

STPA Highlights

UCA-C1: The external auditor certifies materially false financial statements.

Causal factor: Arthur Andersen earned \$25 million in audit fees and \$27 million in consulting fees from Enron in 2000. The auditor's business revenue depended on client satisfaction. Accurate qualification of the accounts would end the relationship; passing the accounts continued it. The control structure made honest auditing financially irrational for the auditor — not because Arthur Andersen's people were unusually corrupt, but because the incentive structure made the corrupted outcome individually optimal.

UCA-C2: The board approves compensation that rewards short-term stock price over long-term value creation.

Causal factor: Options vesting over 3–4 years, combined with quarterly earnings reports and average executive tenure of 5 years, encodes a time horizon into the compensation structure that is shorter than the time horizon of the decisions being made. Leveraged acquisitions, capacity-depleting cost cuts, and aggressive accounting are individually rational under this incentive structure and collectively destructive at scale.

UCA-C3: Workers have no voice in decisions that directly affect them.

Causal factor: VP1 = capital ownership structurally excludes the operational knowledge held by workers from the governance process. Decisions about working conditions, investment priorities, and cost structures are made by people whose information about operational reality is systematically less complete than the people executing those decisions.

Fixes That Worked

Sarbanes-Oxley Act 2002 — addressing UCA-C1

After Enron and WorldCom, the US Congress made these structural changes:

- Prohibited audit firms from providing most non-audit services (consulting, internal audit outsourcing) to their audit clients — severing the revenue dependency that created the conflict.
- Required CEO and CFO to personally certify financial statement accuracy under criminal penalty, removing deniability.
- Created the Public Company Accounting Oversight Board (PCAOB) as an independent inspectorate, replacing self-regulation by the audit profession.

Evidence: No Enron-scale accounting fraud has recurred in US-listed companies at comparable magnitude in the 22 years since SOX. The structural conflict of interest was directly addressed.

German co-determination — Mitbestimmung (1976) — addressing UCA-C2 and UCA-C3

Germany's Mitbestimmungsgesetz requires companies with more than 2,000 employees to give workers equal representation on the supervisory board (Aufsichtsrat). In Germany's two-tier structure, the Aufsichtsrat oversees management strategy without running day-to-day operations. Worker representatives bring:

- A different time horizon (employment stability, long-term skill investment, not quarterly share price)
- Operational knowledge that capital holders structurally lack
- Incentive to flag risks visible to the workforce before they surface in financial results

Evidence: German DAX companies maintained significantly more stable employment during the 2008–2009 financial crisis and the 2020 COVID recession than comparable US firms. German industrial companies (Volkswagen, Siemens, BASF, Bosch) have sustained R&D investment cycles measured in decades and maintained global technological leadership through multiple economic downturns. The Mitbestimmung structure is a contributing — not the sole — causal factor; it works in combination with Germany's long-term banking relationships and patient capital model.

UK Corporate Governance Code (post-Cadbury 1992) — addressing UCA-C2

Following UK corporate scandals (Maxwell, BCCI, Polly Peck), the Cadbury Report mandated:

- Separation of CEO and Chairman roles — preventing a single person from both running the company and overseeing themselves.
- Majority independent non-executive directors on audit, remuneration, and nomination committees.
- Annual shareholder vote on executive compensation (say-on-pay).

Evidence: UK listed companies report >90% compliance. The structural concentration of unchecked personal power that characterised the Maxwell case has been architecturally eliminated as a governance norm for listed companies.


The Design Principle

None of these remedies asks executives or auditors to be more virtuous. They change the structure so that the incentive points in a different direction: severing the auditor's financial dependence on the auditee; changing the time horizon encoded in the compensation structure; adding voices with different interests and different knowledge to the decision room. Structural change, not moral exhortation.

Platform Mapping

Functional Slot	How This System Fills It
<i>Authority & Decision-Making</i>	<i>Board of Directors; weighted shareholder vote</i>
<i>Membership & Belonging</i>	<i>Employment contract; shareholder register</i>
<i>Resource Allocation</i>	<i>Capital budgeting; executive discretion within board-approved limits</i>
<i>Norm Setting & Enforcement</i>	<i>HR policies; compliance function; employment law</i>
<i>Dispute Resolution</i>	<i>Employment tribunals; shareholder litigation; regulatory enforcement</i>
<i>Legitimation</i>	<i>Profit generation; shareholder returns; employment provision</i>
<i>Succession & Continuity</i>	<i>Board appointment; executive succession planning</i>
<i>External Representation</i>	<i>CEO; investor relations; legal and regulatory affairs</i>
<i>Socialisation</i>	<i>Onboarding; training; corporate culture</i>
<i>Activity Delivery</i>	<i>Operations; product/service development and delivery</i>

*For the full SE decomposition:



interactive visualization — graph

LEVEL	ID	DESCRIPTION	↑	↓
⊙ Goals	G1	Generate sustainable profit and shareholder value	—	R1 R3
	G2	Deliver products/services satisfying market demand	—	R3 R5
	G3	Provide livelihood and purpose to employees	—	R4 R5
	G4	Ensure organizational continuity	—	R1 R2
◆ Requirements	R1	Legal incorporation with limited liability	G1 G4	F1 F5
	R2	Board accountable to shareholders	G4	F1
	R3	Revenue exceeding costs	G1 G2	F2 F3
	R4	Hierarchical management with clear authority	G3	F1 F4
	R5	Employment contracts	G2 G3	F4
◇ Functions	F1	Strategic direction — vision, budgets, appointments	R1 R2 R4	L1
	F2	Product/service delivery	R3	L2
	F3	Financial management — accounting, treasury	R3	L3
	F4	HR management — recruit, develop, compensate	R4 R5	L4
	F5	Legal and compliance	R1	L5
○ Logical Architecture	L1	Governance — board, shareholders meeting, audit	F1	P1 P2
	L2	Operations — product dev, manufacturing, supply chain	F2	P3
	L3	Financial — accounting, treasury, controlling	F3	P4
	L4	HR — recruiting, payroll, training	F4	P5
	L5	Legal & Compliance — counsel, regulatory	F5	P1
■ Physical Implementation	P1	Articles of incorporation, board resolutions	L1 L5	—
	P2	CEO, C-suite, management hierarchy	L1	—
	P3	Factories, offices, IT, supply chain	L2	—
	P4	ERP (SAP/Oracle), bank accounts, annual report	L3	—
	P5	Employment contracts, HR software, training	L4	—

interactive visualization — table

. For the comparative analysis across all ten systems: Ten Social Systems Compared. For the cross-system control structure analysis: Control Structures.*

Example: University

The university is the organization family's most internally contested variant — it must simultaneously serve truth (research), education (teaching), and self-governance (academic freedom), goals that regularly conflict. Its SE decomposition exposes why corporatization threatens the system's core logic.



** → Interactive D3 Visualization** — graph

LEVEL	ID	DESCRIPTION	↑	↓
⊙ Goals	G1	Advance knowledge through research	—	R1 R5
	G2	Educate and credential the next generation	—	R2 R6
	G3	Serve society through knowledge transfer	—	R1 R2
	G4	Preserve academic freedom and autonomy	—	R3 R5
◆ Requirements	R1	Peer review as knowledge validation standard	G1 G3	F1
	R2	Structured curricula leading to degrees	G2 G3	F2
	R3	Academic self-governance — senates, elected rectors	G4	F3
	R4	Sustainable funding — tuition, grants, state	—	F5
	R5	Tenure protecting freedom of inquiry	G1 G4	F1 F3
	R6	Accreditation by external bodies	G2	F4
◇ Functions	F1	Research — conduct and publish scholarship	R1 R5	L2
	F2	Teaching — lectures, seminars, examinations	R2	L3
	F3	Academic governance — senate, appointments	R3 R5	L1
	F4	Student administration — admissions, graduation	R6	L4
	F5	Financial management — budgets, grants	R4	L4
○ Logical Architecture	L1	Governance — senate, rector, faculty councils	F3	P1 P2
	L2	Research — departments, institutes, labs, libraries	F1	P3
	L3	Teaching — curricula, lecture halls, exams	F2	P3 P4
	L4	Administration — student office, finance, HR	F4 F5	P4 P5
■ Physical Implementation	P1	University charter, Grundordnung	L1	—
	P2	Rector, deans, senate, faculty councils	L1	—
	P3	Labs, libraries, lecture halls, institutes	L2 L3	—
	P4	Student registry (HISinOne), exam office	L3 L4	—
	P5	Endowment, state funding, DFG/ERC grants	L4	—

→ Interactive D3 Visualization — table

SE Decomposition

The full five-level decomposition is shown in the interactive visualization linked above. Click any node to see its description, parent links, and child links. Use the Table view for the complete traceability matrix.

Variation Point Bindings

VP1 = scholarly merit (unique among social systems), VP2 = competitive admission, VP3 = collegial deliberation (faculty senate), VP4 = mixed (elected rector + appointed deans). The tension between VP1 (merit) and external pressure to bind VP3 to hierarchical management is the central conflict in modern higher education.

Platform Mapping

This system fills all ten universal functional slots identified in the Ten Social Systems Compared:

Functional Slot	How This System Fills It
Authority & Decision-Making	Faculty senate (academic matters); rector/president (executive); board of trustees (strategic oversight)
Membership & Belonging	Matriculated students; academic and administrative staff; alumni (honorary standing)
Resource Allocation	Rector/president budget; competitive grant allocation; tuition and third-party research funding
Norm Setting & Enforcement	Academic regulations; research ethics boards; HR policy; accreditation standards

Functional Slot	How This System Fills It
Dispute Resolution	Academic appeals committees; disciplinary panels; employment tribunals
Legitimation	Knowledge creation and transmission; accreditation; academic freedom as constitutive value
Succession & Continuity	Rector election or appointment; academic tenure; departmental and programme continuity
External Representation	Rector; research partnerships; international agreements; alumni relations
Socialisation	Academic induction; doctoral mentorship; collegiate culture; peer learning communities
Activity Delivery	Teaching; research; knowledge transfer; community engagement and public scholarship

*Navigate to the



interactive visualization — graph


LEVEL	ID	DESCRIPTION	↑	↓
⊙ Goals	G1	Advance knowledge through research	—	R1 R5
	G2	Educate and credential the next generation	—	R2 R6
	G3	Serve society through knowledge transfer	—	R1 R2
	G4	Preserve academic freedom and autonomy	—	R3 R5
◆ Requirements	R1	Peer review as knowledge validation standard	G1 G3	F1
	R2	Structured curricula leading to degrees	G2 G3	F2
	R3	Academic self-governance — senates, elected rectors	G4	F3
	R4	Sustainable funding — tuition, grants, state	—	F5
	R5	Tenure protecting freedom of inquiry	G1 G4	F1 F3
	R6	Accreditation by external bodies	G2	F4
◇ Functions	F1	Research — conduct and publish scholarship	R1 R5	L2
	F2	Teaching — lectures, seminars, examinations	R2	L3
	F3	Academic governance — senate, appointments	R3 R5	L1
	F4	Student administration — admissions, graduation	R6	L4
	F5	Financial management — budgets, grants	R4	L4
○ Logical Architecture	L1	Governance — senate, rector, faculty councils	F3	P1 P2
	L2	Research — departments, institutes, labs, libraries	F1	P3
	L3	Teaching — curricula, lecture halls, exams	F2	P3 P4
	L4	Administration — student office, finance, HR	F4 F5	P4 P5
■ Physical Implementation	P1	University charter, Grundordnung	L1	—
	P2	Rector, deans, senate, faculty councils	L1	—
	P3	Labs, libraries, lecture halls, institutes	L2 L3	—
	P4	Student registry (HISinOne), exam office	L3 L4	—
	P5	Endowment, state funding, DFG/ERC grants	L4	—

interactive visualization — table

for the full graph and table.*

Example: Military

The military is the most hierarchically pure social system — command authority flows strictly top-down, with minimal upward feedback by design. Its SE decomposition reveals why civilian oversight (an external control loop) is architecturally essential: without it, the system has no self-correcting mechanism.



** → Interactive D3 Visualization** — graph

LEVEL	ID	DESCRIPTION	↑	↓
⊙ Goals	G1	Defend the state against external threats	—	R1 R2
	G2	Maintain internal order when directed	—	R1 R3
	G3	Project power and deter aggression	—	R2 R4
	G4	Disaster relief and humanitarian assistance	—	R2 R5
◆ Requirements	R1	Strict command chain from top to soldier	G1 G2	F1
	R2	Trained, equipped, deployable force	G1 G3 G4	F2 F3
	R3	Subordination to civilian authority	G2	F1
	R4	Intelligence and situational awareness	G3	F4
	R5	Logistics — supply, transport, medical	G4	F5
	R6	Military justice framework	—	F6
◇ Functions	F1	Command — orders, coordination, strategy	R1 R3	L1
	F2	Combat — engage adversaries, defend	R2	L2
	F3	Training — drill, exercises, qualification	R2	L3
	F4	Intelligence — collect, analyze, disseminate	R4	L4
	F5	Logistics — ammo, fuel, food, medical	R5	L5
	F6	Military justice — discipline, courts martial	R6	L6
○ Logical Architecture	L1	Command — general staff, HQ, C4ISR	F1	P1
	L2	Combat — brigades, vessels, air wings	F2	P2
	L3	Training — academies, schools, ranges	F3	P3
	L4	Intelligence — SIGINT, HUMINT, cyber	F4	P4
	L5	Logistics — depots, transport, medical corps	F5	P5
	L6	Military Justice — courts martial, JAG, MP	F6	P6
■ Physical Implementation	P1	General staff HQ, situation rooms, comms	L1	—
	P2	Barracks, bases, warships, aircraft, armor	L2	—
	P3	Academies (West Point, Sandhurst, Führungsakademie)	L3	—
	P4	Intelligence agencies, SIGINT stations	L4	—
	P5	Supply depots, field hospitals, transport	L5	—
	P6	Military courts, provost marshal	L6	—

→ Interactive D3 Visualization — table

SE Decomposition

The full five-level decomposition is shown in the interactive visualization linked above. Click any node to see its description, parent links, and child links. Use the Table view for the complete traceability matrix.

Variation Point Bindings

VP1 = command hierarchy (commission from the state), VP2 = enlistment/conscription, VP3 = command order (superior to subordinate), VP4 = promotion by merit/seniority. The binding at VP3 is the most rigid of any social system — orders are obeyed, not debated. This produces decisive action but also enables atrocities when the command chain is corrupt.

Platform Mapping

This system fills all ten universal functional slots identified in the Ten Social Systems Compared:

Functional Slot	How This System Fills It
Authority & Decision-Making	Chain of command from civilian minister to field commander; joint chiefs as strategic advisory body

Functional Slot	How This System Fills It
Membership & Belonging	Voluntary enlistment or compulsory conscription; rank and unit assignment define standing
Resource Allocation	Parliamentary defence budget; theatre commanders; logistics chain under operational command
Norm Setting & Enforcement	Military law (UCMJ / national equivalent); rules of engagement; code of conduct; Geneva Conventions
Dispute Resolution	Courts-martial; military ombudsman; administrative review boards; civilian oversight tribunal
Legitimation	State mandate for territorial defence; professional military ethos; oath of service
Succession & Continuity	Promotion boards; officer commissioning pipeline; doctrinal continuity through staff colleges
External Representation	Joint chiefs; defence attachés; alliance structures (NATO, bilateral treaties)
Socialisation	Basic training and indoctrination; unit cohesion rituals; rank culture and esprit de corps
Activity Delivery	Combat operations; deterrence; logistics; intelligence; peacekeeping; training and readiness

*Navigate to the



interactive visualization — graph


LEVEL	ID	DESCRIPTION	↑	↓
⊙ Goals	G1	Defend the state against external threats	—	R1 R2
	G2	Maintain internal order when directed	—	R1 R3
	G3	Project power and deter aggression	—	R2 R4
	G4	Disaster relief and humanitarian assistance	—	R2 R5
◆ Requirements	R1	Strict command chain from top to soldier	G1 G2	F1
	R2	Trained, equipped, deployable force	G1 G3 G4	F2 F3
	R3	Subordination to civilian authority	G2	F1
	R4	Intelligence and situational awareness	G3	F4
	R5	Logistics — supply, transport, medical	G4	F5
	R6	Military justice framework	—	F6
◇ Functions	F1	Command — orders, coordination, strategy	R1 R3	L1
	F2	Combat — engage adversaries, defend	R2	L2
	F3	Training — drill, exercises, qualification	R2	L3
	F4	Intelligence — collect, analyze, disseminate	R4	L4
	F5	Logistics — ammo, fuel, food, medical	R5	L5
	F6	Military justice — discipline, courts martial	R6	L6
○ Logical Architecture	L1	Command — general staff, HQ, C4ISR	F1	P1
	L2	Combat — brigades, vessels, air wings	F2	P2
	L3	Training — academies, schools, ranges	F3	P3
	L4	Intelligence — SIGINT, HUMINT, cyber	F4	P4
	L5	Logistics — depots, transport, medical corps	F5	P5
	L6	Military Justice — courts martial, JAG, MP	F6	P6
■ Physical Implementation	P1	General staff HQ, situation rooms, comms	L1	—
	P2	Barracks, bases, warships, aircraft, armor	L2	—
	P3	Academies (West Point, Sandhurst, Führungsakademie)	L3	—
	P4	Intelligence agencies, SIGINT stations	L4	—
	P5	Supply depots, field hospitals, transport	L5	—
	P6	Military courts, provost marshal	L6	—

interactive visualization — table

for the full graph and table.*

Example: The Family

The family is the smallest and oldest social system — and the only one where the membership boundary (VP2) is primarily biological rather than contractual. Its SE decomposition reveals a system that must solve all ten platform problems with typically 2–5 people and no formal governance documents.



** → Interactive D3 Visualization** — graph

LEVEL	ID	DESCRIPTION	↑	↓
⊙ Goals	G1	Ensure reproduction and survival of offspring	—	R1 R2
	G2	Provide emotional security, love, belonging	—	R3
	G3	Socialize children into norms and competences	—	R4 R6
	G4	Pool resources and manage economic risk	—	R2 R5
	G5	Transmit identity, heritage, property across generations	—	R5
◆ Requirements	R1	Pair bond or co-parenting agreement	G1	F1 F3
	R2	Shelter and material provision	G1 G4	F1 F4
	R3	Emotional availability and care	G2	F1 F2
	R4	Authority structure for child-rearing	G3	F3
	R5	Inheritance and succession norms	G4 G5	F4 F5
	R6	Interface with schools, healthcare, state	G3	F6
◇ Functions	F1	Nurture — feed, shelter, protect, care	R1 R2 R3	L2
	F2	Socialize — teach language, norms, values	R3	L4
	F3	Govern — discipline, education, health decisions	R1 R4	L1
	F4	Economic management — earn, budget, save	R2 R5	L3
	F5	Ritual and identity — traditions, storytelling	R5	L4
	F6	External representation — school, state interface	R6	L1
○ Logical Architecture	L1	Parental Authority — decisions, discipline	F3 F6	P1 P2
	L2	Care — nurturing, health, emotional support	F1	P1
	L3	Economic — income, budgeting, property	F4	P3
	L4	Socialization — education, culture, values	F2 F5	P4 P5
■ Physical Implementation	P1	Home — house or apartment, shared space	L1 L2	—
	P2	Parents/guardians, marriage/partnership	L1	—
	P3	Family income, bank accounts, property, wills	L3	—
	P4	Daily routines, rituals, photo albums, heirlooms	L4	—
	P5	School enrollment, pediatrician, community	L4	—

→ Interactive D3 Visualization — table

SE Decomposition

The full five-level decomposition is shown in the interactive visualization linked above. Click any node to see its description, parent links, and child links. Use the Table view for the complete traceability matrix.

Variation Point Bindings

VP1 = biological/legal bond (unique), VP2 = birth/adoption (non-voluntary for children), VP3 = consensual negotiation (between adults) + parental authority (over children), VP4 = self-selection (founding a new family). The key structural insight: the family is the only system where VP6 (norm enforcement) relies primarily on love rather than coercion.


Platform Mapping

This system fills all ten universal functional slots identified in the Ten Social Systems Compared:

Functional Slot	How This System Fills It
Authority & Decision-Making	Parental authority over children; negotiated consensus between adult partners
Membership & Belonging	Biological kinship or legal adoption; marriage or partnership; non-voluntary for children
Resource Allocation	Household income pooling; parental discretion over expenditure; children's allowances

Functional Slot	How This System Fills It
Norm Setting & Enforcement	Parental rules; shared household norms; enforced primarily through relationship and love rather than coercion
Dispute Resolution	Direct negotiation between adults; parental arbitration for child conflicts; family mediation as last resort
Legitimation	Biological bond; legal recognition (marriage, guardianship); emotional attachment and mutual care
Succession & Continuity	Children reaching independence and founding their own families — the system reproduces itself
External Representation	Parents as legal guardians; household as the legal and fiscal unit in dealings with the state
Socialisation	Primary socialisation — language, values, emotional regulation, identity, basic social norms
Activity Delivery	Care, nutrition, housing, educational support, emotional support — the full welfare of members

*Navigate to the



interactive visualization — graph


LEVEL	ID	DESCRIPTION	↑	↓
⊙ Goals	G1	Ensure reproduction and survival of offspring	—	R1 R2
	G2	Provide emotional security, love, belonging	—	R3
	G3	Socialize children into norms and competences	—	R4 R6
	G4	Pool resources and manage economic risk	—	R2 R5
	G5	Transmit identity, heritage, property across generations	—	R5
◆ Requirements	R1	Pair bond or co-parenting agreement	G1	F1 F3
	R2	Shelter and material provision	G1 G4	F1 F4
	R3	Emotional availability and care	G2	F1 F2
	R4	Authority structure for child-rearing	G3	F3
	R5	Inheritance and succession norms	G4 G5	F4 F5
	R6	Interface with schools, healthcare, state	G3	F6
◇ Functions	F1	Nurture — feed, shelter, protect, care	R1 R2 R3	L2
	F2	Socialize — teach language, norms, values	R3	L4
	F3	Govern — discipline, education, health decisions	R1 R4	L1
	F4	Economic management — earn, budget, save	R2 R5	L3
	F5	Ritual and identity — traditions, storytelling	R5	L4
	F6	External representation — school, state interface	R6	L1
○ Logical Architecture	L1	Parental Authority — decisions, discipline	F3 F6	P1 P2
	L2	Care — nurturing, health, emotional support	F1	P1
	L3	Economic — income, budgeting, property	F4	P3
	L4	Socialization — education, culture, values	F2 F5	P4 P5
■ Physical Implementation	P1	Home — house or apartment, shared space	L1 L2	—
	P2	Parents/guardians, marriage/partnership	L1	—
	P3	Family income, bank accounts, property, wills	L3	—
	P4	Daily routines, rituals, photo albums, heirlooms	L4	—
	P5	School enrollment, pediatrician, community	L4	—

interactive visualization — table

for the full graph and table.*

Example: One-Party State

The one-party state fuses party and government into a single hierarchy, creating the most tightly controlled governance variant. Its SE decomposition reveals a system designed to prevent any alternative power center — which also means it has no structural mechanism for correcting its own errors.



** → Interactive D3 Visualization** — graph

LEVEL	ID	DESCRIPTION	↑	↓
⊙ Goals	G1	Implement party ideology as organizing principle	—	R1 R3
	G2	Maintain party monopoly on political power	—	R1 R2 R4
	G3	Ensure order, security, and territory	—	R4 R5
	G4	Deliver economic development as legitimacy proof	—	R5 R6
◆ Requirements	R1	Constitutional enshrinement of party's leading role	G1 G2	F1
	R2	Party-state fusion — parallel hierarchies everywhere	G2	F1 F2
	R3	Ideological conformity — censorship, propaganda	G1	F3
	R4	Security apparatus loyal to the party	G2 G3	F4
	R5	Centralized economic direction	G3 G4	F5
	R6	Cadre system — party membership for advancement	G4	F2
◇ Functions	F1	Party-directed legislation	R1 R2	L1
	F2	Cadre management — recruit, evaluate, promote	R2 R6	L2
	F3	Ideological control — propaganda, censorship	R3	L3
	F4	Security and surveillance	R4	L4
	F5	Economic direction — plans, SOEs, investment	R5	L5
○ Logical Architecture	L1	Party-State Legislative — politburo, congress	F1	P1
	L2	Cadre & Personnel — org department	F2	P2
	L3	Propaganda & Control — media, censorship	F3	P3
	L4	Security — secret police, surveillance	F4	P4
	L5	Economic Planning — commission, SOEs	F5	P5
■ Physical Implementation	P1	Politburo / Central Committee	L1	—
	P2	Organization Department, party schools, cadre files	L2	—
	P3	State media (Xinhua, CCTV), Great Firewall	L3	—
	P4	Secret police (Stasi, KGB, MSS), surveillance	L4	—
	P5	State Planning Commission, SOEs, 5-year plans	L5	—

** → *Interactive D3 Visualization*** — *table*

SE Decomposition

The full five-level decomposition is shown in the interactive visualization linked above. Click any node to see its description, parent links, and child links. Use the Table view for the complete traceability matrix.

Variation Point Bindings

VP1 = ideological mandate (party), VP2 = territorial citizenship (compulsory), VP3 = party committee decision (autocratic), VP4 = cadre promotion (internal, party-controlled). The binding at VP3 produces the characteristic 'rubber-stamp parliament' — a legislative body that exists structurally but cannot exercise independent legislative function.

Platform Mapping

This system fills all ten universal functional slots identified in the Ten Social Systems Compared:

Functional Slot	How This System Fills It
Authority & Decision-Making	Party standing committee / Politburo; general secretary; parliament exists but cannot exercise independent legislative function
Membership & Belonging	Compulsory territorial citizenship; party membership is the pathway to political power and state resources
Resource Allocation	Central planning or party-directed state capitalism; no independent budget process; allocation serves party priorities

Functional Slot	How This System Fills It
Norm Setting & Enforcement	Party line; security apparatus; censorship and information control; surveillance of population
Dispute Resolution	Party-controlled courts; no independent judiciary; internal party discipline for cadres
Legitimation	Ideological narrative (socialism, nationalism, historical necessity); economic performance where available
Succession & Continuity	Internal cadre promotion; party congress ratification; consolidation of power by dominant faction
External Representation	General secretary / paramount leader; state diplomatic corps; party-to-party international relations
Socialisation	State education system; propaganda apparatus; party youth organisations; patriotic education campaigns
Activity Delivery	State enterprises; party-directed civil service; public services as instrument of social control

*Navigate to the



interactive visualization — graph

LEVEL	ID	DESCRIPTION	↑	↓
⊙ Goals	G1	Implement party ideology as organizing principle	—	R1 R3
	G2	Maintain party monopoly on political power	—	R1 R2 R4
	G3	Ensure order, security, and territory	—	R4 R5
	G4	Deliver economic development as legitimacy proof	—	R5 R6
◆ Requirements	R1	Constitutional enshrinement of party's leading role	G1 G2	F1
	R2	Party-state fusion — parallel hierarchies everywhere	G2	F1 F2
	R3	Ideological conformity — censorship, propaganda	G1	F3
	R4	Security apparatus loyal to the party	G2 G3	F4
	R5	Centralized economic direction	G3 G4	F5
	R6	Cadre system — party membership for advancement	G4	F2
◇ Functions	F1	Party-directed legislation	R1 R2	L1
	F2	Cadre management — recruit, evaluate, promote	R2 R6	L2
	F3	Ideological control — propaganda, censorship	R3	L3
	F4	Security and surveillance	R4	L4
	F5	Economic direction — plans, SOEs, investment	R5	L5
○ Logical Architecture	L1	Party-State Legislative — politburo, congress	F1	P1
	L2	Cadre & Personnel — org department	F2	P2
	L3	Propaganda & Control — media, censorship	F3	P3
	L4	Security — secret police, surveillance	F4	P4
	L5	Economic Planning — commission, SOEs	F5	P5
■ Physical Implementation	P1	Politburo / Central Committee	L1	—
	P2	Organization Department, party schools, cadre files	L2	—
	P3	State media (Xinhua, CCTV), Great Firewall	L3	—
	P4	Secret police (Stasi, KGB, MSS), surveillance	L4	—
	P5	State Planning Commission, SOEs, 5-year plans	L5	—

interactive visualization — table

for the full graph and table.*

Part III — Development Frameworks

This part applies the same product line analysis to development and project management frameworks, revealing that the "methodology wars" are really about different binding decisions at four well-defined variation points.

The Frameworks

Framework	Orientation
Waterfall	Plan-driven, sequential
V-Model	Plan-driven, verification-focused
PRINCE2	Governance-focused PM
Scrum	Adaptive, iterative
Kanban	Adaptive, flow-based
Design Thinking	Human-centered, exploratory
DevOps	Continuous delivery
SAFe	Scaled adaptive

Chapters

- **Eight Frameworks Decomposed** — Full SE hierarchy for each framework
- **The Framework Platform** — Ten universal functional slots
- **Hybrid Architectures** — Three principled merger proposals

Product Line Thinking for Development Frameworks

How Waterfall, Scrum, Kanban, V-Model, PRINCE2, Design Thinking, DevOps, and SAFe Share a Common Platform — and How to Merge Them Better

Abstract

The "methodology wars" of the past three decades — Waterfall vs. Agile, Scrum vs. Kanban, plan-driven vs. adaptive — have generated more heat than light. This chapter applies the same systems engineering decomposition and product line analysis previously used for social institutions to the domain of development and project management frameworks. By decomposing eight frameworks through a five-level SE hierarchy (Goals → Requirements → Functions → Logical Architecture → Physical Implementation), we discover that these seemingly opposed approaches share a remarkably large common platform and differ primarily at four well-defined variation points. We use this insight to propose three concrete hybrid architectures that merge the best of plan-driven and adaptive approaches — not as vague "be pragmatic" advice, but as structurally principled product-line configurations with explicit binding decisions.

Canonical reference

The normative versions of the eight framework decompositions, the cross-framework platform, the variation points, the hybrid architectures, and the merging principles are maintained as the project's knowledge base under [knowledge/system-catalogues/dev-frameworks/](#). This chapter is the narrative overview; the knowledge base is the reference.

1. Introduction

1.1 The False Dichotomy

Ask any software team whether they "do Agile or Waterfall" and you'll get a strong opinion. Ask them what exactly they mean by either term and the clarity evaporates. In practice, almost no team does pure Waterfall (sequential phases with no feedback) or pure Scrum (no upfront planning, no documentation, no governance). Most teams operate some hybrid — often an accidental, incoherent one — because the real work demands elements of both.

The problem is not that people lack pragmatism. The problem is that they lack a *structural vocabulary* for describing which elements of which framework can be combined, and which combinations are architecturally incoherent. Product line thinking provides that vocabulary.

1.2 Scope: Eight Frameworks

We analyze eight widely-used frameworks, chosen to span the full spectrum from plan-driven to adaptive, and from project-focused to flow-focused:

Framework	Orientation	Core Idea
Waterfall	Plan-driven, sequential	Complete each phase before the next begins
V-Model	Plan-driven, verification-focused	Each design phase has a corresponding test phase
PRINCE2	Governance-focused PM	Projects in controlled environments with staged authorization
Scrum	Adaptive, iterative	Self-organizing teams deliver value in fixed-length sprints
Kanban	Adaptive, flow-based	Visualize work, limit WIP, optimize flow — no prescribed iterations
Design Thinking	Human-centered, exploratory	Empathize → Define → Ideate → Prototype → Test
DevOps	Continuous delivery	Merge development and operations into an automated pipeline

Framework	Orientation	Core Idea
SAFe	Scaled adaptive	Agile practices at portfolio, program, and team levels

2. The Decompositions

Each framework is decomposed through the five-level SE hierarchy with full traceability. Goals name what the framework exists to achieve; requirements translate those goals into must-statements; functions name the activities; logical architecture groups the functions into subsystems; physical implementation names the concrete artefacts, roles, and practices that realise each subsystem.

The full tables are maintained in the knowledge base:

Framework	Decomposition file
Waterfall	knowledge/system-catalogues/dev-frameworks/waterfall/se-decomposition.md
V-Model	knowledge/system-catalogues/dev-frameworks/v-model/se-decomposition.md
PRINCE2	knowledge/system-catalogues/dev-frameworks/prince2/se-decomposition.md
Scrum	knowledge/system-catalogues/dev-frameworks/scrum/se-decomposition.md
Kanban	knowledge/system-catalogues/dev-frameworks/kanban/se-decomposition.md
Design Thinking	knowledge/system-catalogues/dev-frameworks/design-thinking/se-decomposition.md
DevOps	knowledge/system-catalogues/dev-frameworks/devops/se-decomposition.md
SAFe	knowledge/system-catalogues/dev-frameworks/safe/se-decomposition.md

3. The Development Framework Platform

Once the eight decompositions are aligned level by level, ten universal functional slots emerge that every framework fills in some way. These are **Work Discovery**, **Prioritisation**, **Execution**, **Quality Assurance**, **Integration**, **Delivery**, **Feedback**, **Governance**, **Improvement**, and the framework-specific activity layer that carries the work through them.

The full platform description — what the shared goals, requirements, functions, logical, and physical elements actually are — is in `knowledge/system-catalogues/dev-frameworks/cross-framework/platform.md`.

What the platform means in practice: most framework disputes are not about whether to do a thing, but about *where* in the lifecycle to do it, *how often*, and *who decides*. Those are exactly the questions that variation points make explicit.

4. The Variation Points

Four primary variation points produce the differences between frameworks:

- **VP1 — Temporal Structure of Work** (phases · sprints · flow · continuous · stages).
- **VP2 — Requirements Stability Assumption** (frozen upfront · evolving throughout · discovered during work).

- **VP3 — Authority & Decision Structure** (hierarchical project manager · self-organising team · product owner · steering committee).
- **VP4 — Quality Assurance Timing** (gated at end · continuous · after each sprint · before each phase transition).

As with social systems, variants are not freely combinable: some binding combinations are coherent, others are structurally contradictory (e.g. "phase-gated temporal structure" with "requirements evolve throughout" is incoherent without a workaround).

The full variation-point catalogue, observed variants, and interaction constraints are in <knowledge/system-catalogues/dev-frameworks/cross-framework/variation-points.md>.

5. Reuse Analysis Across Frameworks

Elements of one framework can sometimes be adopted by another, sometimes only with adaptation, and sometimes not at all without changing the framework's identity. The full reuse analysis — fully reusable, partially reusable with adaptation, and structurally incompatible — is in <knowledge/system-catalogues/dev-frameworks/cross-framework/reuse-analysis.md>.

6. Three Proposed Hybrid Architectures

Three concrete hybrid framework configurations are proposed, each with explicit binding decisions at all four variation points:

- **Hybrid A — "Governed Agile"** — for enterprises with regulatory or contractual obligations. Combines PRINCE2-style staged authorisation with Scrum-style execution inside each stage.
- **Hybrid B — "Discovery-Driven Flow"** — for product teams optimising for innovation and speed. Combines Design Thinking discovery with Kanban delivery and DevOps continuous release.
- **Hybrid C — "Scaled Governed Flow"** — for large organisations managing portfolios of diverse work. Combines SAFe portfolio structure with flow-based team execution and phase-gated compliance checkpoints.

Full architectural sketches, binding tables, and trade-offs for all three hybrids are in <knowledge/system-catalogues/dev-frameworks/cross-framework/hybrids.md>.

7. Principles for Framework Merging

Five principles govern the design of coherent framework hybrids:

- **Compatibility Rule** — binding decisions across variation points must be mutually coherent.
- **Granularity Principle** — variation points bind at the smallest cohesive unit, not the whole framework.
- **Start from the Platform** — build on the shared functional slots first, then add variation.
- **Automate the Right Arm** — whatever part of the lifecycle is repetitive should be automated before it is standardised.
- **Separate Discovery from Delivery** — requirements discovery and execution must not share the same temporal structure.

Full statements of all five principles are in <knowledge/system-catalogues/dev-frameworks/cross-framework/merging-principles.md>.

8. Conclusion

The methodology wars were never about fundamentally different approaches to building things. They were about different *binding decisions* at four variation points — temporal structure, requirements stability, authority structure, and

QA timing — applied to a *shared platform* of ten universal functional slots.

Waterfall is not wrong. It is a valid configuration for contexts where requirements are genuinely stable, regulatory traceability is required, and the cost of change is high. Scrum is not wrong. It is a valid configuration for contexts where requirements are volatile, time-to-feedback is critical, and teams are small enough to self-organize. The error — the engineering error, in product-line terms — is applying one configuration to a context that demands the other, or merging elements from both without checking VP compatibility.

Product line thinking offers a way out. By identifying the platform, mapping the variation points, and making binding decisions explicit, teams and organizations can design hybrid frameworks that are *principled* rather than accidental, *coherent* rather than contradictory, and *tailored* rather than dogmatic. The goal is not to find the one true methodology. It is to become fluent in the vocabulary of variation, so that every team can configure the framework that fits their actual context — and explain why.

This analysis applies the product line methodology developed in "Product Line Thinking for Social Systems" to the domain of development and project management frameworks. The SE hierarchy (Goals → Requirements → Functions → Logical Architecture → Physical Implementation) provides the decomposition structure; the product line concepts (platform, variation point, binding, reuse classification) provide the comparative and synthesis methodology.

The Framework Platform

Ten Universal Functional Slots in Development

Every development framework — regardless of philosophy — fills the same functional slots:

Functional Slot	Description
Work Discovery	Determine what to build
Work Prioritization	Decide what first
Work Decomposition	Break large efforts into manageable units
Work Execution	Build the deliverable
Quality Assurance	Verify expectations are met
Integration	Combine outputs into a coherent whole
Delivery	Get the product to users
Feedback Collection	Learn whether the product solves the problem
Process Governance	Authorization, tracking, risk management
Process Improvement	Reflect and evolve the process

Four Variation Points

The frameworks differ at four points:

1. **VP1: Temporal Structure** — sequential phases · fixed iterations · continuous flow · exploratory loops
2. **VP2: Requirements Stability Assumption** — knowable upfront · emergent · hypotheses
3. **VP3: Authority & Decision Structure** — hierarchical PM · PO + team · shared · automated pipeline
4. **VP4: QA Timing** — end-phase · per-iteration · continuous automated · user validation

The full decomposition and variation point analysis is in *Eight Frameworks Decomposed*.

Hybrid Architectures

Three Principled Mergers

Hybrid A: "Governed Agile"

For: Regulated industries, contractual delivery (medical devices, automotive, maritime OT, defense)

PRINCE2 stages containing Scrum Sprints. V-Model test levels automated as CI/CD pipeline. Design Thinking discovery upfront.

Hybrid B: "Discovery-Driven Flow"

For: Product teams, SaaS, startups

Design Thinking discovery track running parallel to Kanban delivery track. DevOps CI/CD. No Sprints — continuous pull. Retrospectives decoupled from delivery cadence.

Hybrid C: "Scaled Governed Flow"

For: Enterprises managing mixed portfolios

Portfolio level on Kanban with PRINCE2 Business Case governance. Program level on SAFe PI cadence. Team level: each team chooses Scrum or Kanban. Shared DevOps platform underneath.

Five Design Principles

1. **The Compatibility Rule** — check VP binding compatibility before merging
2. **The Granularity Principle** — different VPs can bind at different organizational levels
3. **Start from the Platform** — don't modify one framework; compose from the functional slots
4. **Automate the Right Arm** — V-Model quality structure + DevOps automation = best of both
5. **Separate Discovery from Delivery** — different activities need different rhythms

The full analysis with binding tables is in Eight Frameworks Decomposed.

Part IV — STPA Applied

This part applies System-Theoretic Process Analysis to social systems, using the SE decompositions from Part II as input. STPA reveals that the most harmful outcomes of social systems are not random malfunctions but predictable consequences of their control architecture.

Chapters

- **STPA on Religion** — Full four-step STPA analysis: losses, control structure, 25 UCAs, 6 loss scenarios
- **STPA on Frontier AI** — Worked sociotechnical example: 14 controllers, 65 UCAs, 9 loss scenarios, 22 sequenced remedies, applied to the political economy of frontier AI
- **Control Structures in Social Systems** — Comparing control architectures across the ten systems
- **Architectural Remedies** — Design principles for safer social systems

STPA Analysis of Religion as a Social System

System-Theoretic Process Analysis Applied to the Systems Engineering Decomposition of Religion

Abstract

System-Theoretic Process Analysis (STPA) is a hazard analysis method developed by Nancy Leveson at MIT, based on the STAMP (Systems-Theoretic Accident Model and Processes) framework. STPA was designed for safety-critical engineering systems — aircraft, nuclear plants, medical devices — but its core logic applies wherever a complex system can produce unintended harmful outcomes through inadequate control rather than mere component failure. This paper applies STPA to religion, treated as a social system previously decomposed through a five-level systems engineering hierarchy. We define the system's losses, model its control structure, identify unsafe control actions across all four STPA categories, and trace loss scenarios to their causal factors. The result is a structural explanation of how religion — a system designed to provide meaning, morality, and community — can produce outcomes that are the exact opposite of its stated goals: meaninglessness, moral corruption, and social destruction.

1. Background: The SE Decomposition of Religion

1.1 The Five-Level Hierarchy

The following decomposition was established in prior work. It forms the input to the STPA analysis.

Goal

- G1** — Provide existential meaning and purpose to human life
- G2** — Establish moral and ethical orientation for individuals and communities
- G3** — Create social cohesion and collective identity
- G4** — Offer psychological comfort in the face of suffering and mortality

Req

- R1** — Coherent cosmological narrative (creation, afterlife, transcendence)
- R2** — Framework for interpreting personal experience as meaningful
- R3** — Codified behavioral norms (commandments, precepts, dharma)
- R4** — Authority structure to interpret and enforce norms
- R5** — Shared symbols, narratives, and practices for group bonding
- R6** — Mechanism for distinguishing in-group from out-group
- R7** — Rituals addressing grief, loss, and life transitions

Func

- F1** — Mythological narration — transmit origin and eschatological stories
- F2** — Pastoral care — counsel individuals through crisis and life stages
- F3** — Moral legislation — define permissible and forbidden behavior

F4 — Doctrinal governance — interpret scripture, settle disputes

F5 — Ritual performance — enact shared ceremonies

F6 — Missionary expansion — propagate belief system

Log

L1 — Liturgical Subsystem — scripture, preaching, ceremonial calendar

L2 — Pastoral Subsystem — counseling, rites of passage, community support

L3 — Doctrinal/Juridical Subsystem — theology, canon law, ethics boards

L4 — Outreach Subsystem — evangelization, education, charity, media

Phys

P1 — Churches, mosques, temples, synagogues, gurdwaras

P2 — Sacred texts — Bible, Quran, Torah, Vedas, Tripitaka

P3 — Clergy — priests, imams, rabbis, pastors, monks, gurus

P4 — Parish networks, sanghas, congregations, umma structures

P5 — Central authorities — Vatican, Al-Azhar, ecumenical councils

P6 — Legal systems — canon law, fatwa apparatus, halakhic courts

P7 — Missionary organizations, madrasas, Sunday schools

P8 — Religious media — TV networks, podcasts, publishing houses

1.2 Why STPA?

Traditional risk analyses ask: "What component can fail?" STPA asks a different question: "What control actions, if inadequate, can lead to losses — even when every component is working as designed?" This is precisely the right question for religion, where the most devastating outcomes (inquisitions, abuse cover-ups, radicalization, exclusion of vulnerable people) typically occur not because the system has "broken" but because its control mechanisms are operating exactly as designed — just in a way that produces harm.

2. STPA Step 1: Define Purpose of the Analysis

2.1 System-Level Losses

STPA begins by identifying **losses** — outcomes that the system's stakeholders would consider unacceptable. For religion as a social system, we define:

Loss ID	Loss Description	Affected Goal
L1	Individuals suffer existential despair, alienation, or loss of meaning because of the religious system (not despite it)	Negates G1
L2	Moral harm — the system produces, enables, or conceals immoral behavior (abuse, violence, corruption, exploitation)	Negates G2
L3	Social destruction — the system causes division, persecution, hatred, or war between groups	Negates G3
L4	Psychological damage — the system inflicts trauma, guilt, fear, or mental illness on individuals	Negates G4
L5	Loss of institutional legitimacy — the system loses the trust and authority it needs to function	Negates all

Loss ID	Loss Description	Affected Goal
L6	Suppression of truth — the system suppresses scientific knowledge, critical thinking, or individual conscience	Negates G1, G2

Note: each loss is the *negation* of a system goal. This is the defining feature of STPA applied to social systems — the hazards are not external threats but *self-inflicted contradictions* where the system's own control actions produce outcomes opposite to its stated purpose.

2.2 System-Level Hazards

Hazards are system states that, combined with worst-case environmental conditions, lead to losses. We derive them from the SE requirements:

Hazard ID	Hazard Description	Related Losses
H1	Doctrine becomes unfalsifiable and self-sealing — no internal mechanism can correct doctrinal errors	L1, L2, L6
H2	Authority structure becomes unaccountable — no effective oversight of those who interpret and enforce norms	L2, L4, L5
H3	In-group/out-group mechanism escalates to dehumanization of outsiders	L3
H4	Moral legislation becomes absolutist — no allowance for context, conscience, or evolving understanding	L1, L2, L4
H5	Pastoral care relationship is exploited — power asymmetry enables abuse	L2, L4, L5
H6	Missionary function prioritizes expansion over welfare of converts or target populations	L2, L3
H7	Financial resources are extracted from vulnerable populations without accountability	L2, L5
H8	Ritual performance becomes coercive — participation is compelled, not voluntary	L1, L4

2.3 System-Level Constraints


Each hazard implies a constraint — a condition the system must maintain to prevent the hazard:

Constraint ID	Constraint	Prevents Hazard
SC1	Doctrine must include mechanisms for self-correction, reinterpretation, and engagement with external knowledge	H1
SC2	Authority figures must be subject to transparent oversight, accountability, and removal procedures	H2
SC3	In-group/out-group distinctions must not extend to denial of outsiders' human dignity or rights	H3
SC4	Moral legislation must preserve space for individual conscience and contextual judgment	H4
SC5	Pastoral relationships must have safeguards: codes of conduct, mandatory reporting, external recourse	H5
SC6	Missionary activity must respect the autonomy and existing culture of target populations	H6
SC7	Financial management must be transparent, audited, and separated from spiritual authority	H7
SC8	Ritual participation must be voluntary, with no material or social penalty for non-participation	H8

3. STPA Step 2: Model the Control Structure

3.1 The STAMP Control Structure

STAMP models systems as hierarchical control structures where controllers issue control actions to controlled processes, and receive feedback. For religion:



exclusion | | - Participation

INDIVIDUAL BELIEVER

Process Model:

- Faith provides meaning
- Clergy speak for God
- Doubt is sinful
- Suffering has purpose

3.2 Why these control loops — and why only these?

The test for a control loop

STPA imposes three necessary conditions on every arrow in a control structure. All three must hold simultaneously; failure of any one disqualifies the relationship from being modelled as a control loop:

1. **Controller (authority condition):** the upstream element issues *observable, authoritative* control actions – directives the downstream element is structurally expected to obey, not merely influence it may choose to consider.
2. **Controlled process (response condition):** the downstream element's state is expected to change in response to those control actions as part of normal system operation.
3. **Feedback channel (closure condition):** some signal – however weak, delayed, or filtered – travels back from the controlled process to the controller and can in principle update the controller's process model.

A relationship that passes all three tests is a control loop and must be modelled. A relationship that fails any one test is modelled differently: as peer coordination (aggregated into an existing feedback channel), as an environmental disturbance, or as a belief inside a controller's process model.

Applying the test to the SE decomposition

The four loops follow directly from applying these three conditions to the physical elements named in §1.1. The authority source that makes relationships authoritative rather than merely influential in every case is R4 ("authority structure to interpret and enforce norms") – a requirement the system *itself* specifies.

Loop	Controller	Controlled Process	Condition 1: authority	Condition 2: state change	Condition 3: feedback
1	Central Authority (P5, Clergy (P3) P6)		R4 + claimed divine mandate; clergy bound under canon law / fatwa / denominational discipline	Clergy doctrine, appointments, removal of clergy behaviour demonstrably changes in response	Reports from synod deliberations, petitions (filtered but structurally in present)

Loop	Controller	Controlled Process	Condition 1: authority	Condition 2: state change	Condition 3: feedback
2	Clergy (P3)	Congregation (P4)	R4 delegated downward; authority to preach, administer sacraments, impose social discipline	Congregation attendance, expressed belief, ritual participation change in response	Confession, counselling, donation levels, attendance — reach clergy as indirect behavioural signals
3	Congregation (P4)	Individual	R5, R6 (shared identity, in/out-group enforcement); peer authority over social standing and belonging	Individual's public behaviour, expressed belief, ritual participation change to maintain belonging	Gossip, reputation signals, observation of conformity or deviance feed back to community process model
4	Individual (self)	Own behaviour	Internalised R1–R3; conscience acts as an internal authority that the person is expected — by the system's own design — to obey	Own behaviour, choices, emotional state, and spiritual self-assessment change	Emotional feedback: guilt, consolation, doubt — conscience functions as internal sensor closing the loop

These are the only four pairs in §1.1 where one element holds *designated authority to issue control actions* the other is *structurally expected to obey*. Remove any loop and the transmission path from doctrine to individual behaviour breaks — no alternative channel carries an authoritative control action all the way to the individual.

Why every other plausible relationship fails the test

The table below applies the three conditions to every candidate relationship commonly raised when this control structure is presented. In every case the candidate fails at least Condition 1.

Candidate relationship	Fails condition	Reason
Clergy ↔ Clergy (lateral, same rank)	1 — authority	Clergy of equal rank do not issue authoritative control actions to each other. Synods, councils, and theological deliberations are coordination inputs that aggregate into Loop 1's feedback channel, not independent loops. Where genuine rank differences exist (bishop over priest, abbot over monk), the relationship resolves to Loop 1 or Loop 2 at finer granularity — not a new, fifth loop.
Congregation ↔ Congregation (between parishes)	1 — authority	Parish networks are parallel peers. No congregation holds designated authority over another. Binding coordination between congregations routes through Loop 1 — the central authority that owns both. Direct inter-congregational social pressure carries no structural control action.

Candidate relationship	Fails condition	Reason
Individual congregation Authority (upward) or → 1 — authority		In the hierarchical institutional form this analysis targets, laity have no structural authority to issue control actions upward. Their upward channel is feedback, already represented in every loop. Note: congregationalist polities (Quaker, many Baptist, Reform Jewish) genuinely invert Loop 1 — authority flows from the congregation upward. Those constitute a genuinely different control structure under the same family name and require their own analysis.
God → Authority	1 — authority	The system claims this loop, but divine control actions are not observably issued or measurably received. What actually operates is the Authority's belief that it is divinely controlled — a feature of the Authority's process model (shown inside the top box of the diagram), not a separate causal loop. STPA models what exists in the causal structure, not what is asserted.
Civil state / science / secular media	1 — authority	These elements lie in the system's environment. They can disturb the system or supply external corrective signals, but they do not issue control actions the religion is architecturally bound to obey. They appear in §7 as externally imposed circuit breakers, not as internal control loops.
Other religions	1 — authority	No cross-denominational structural authority exists by definition of denomination. Influence between traditions is cultural and bilateral; control is absent in both directions.
Schism / sect formation	Not a loop at all	This is a phase transition, not a control relationship. When Loop 1 fails to correct accumulating error, a fragment of the controlled process exits the hierarchy entirely and becomes the apex of a new control structure. Schism is a symptom of loop failure (analysed in §5), not a sixth relationship within the existing hierarchy.
The upward feedback arrows on each loop	Not a new loop	Every loop already includes a feedback channel (the upward arrows in the diagram). Those arrows satisfy Condition 3 of their respective loop. They do not satisfy Condition 1 of a new loop — feedback is not control. Treating them as control would mean the congregation governs the clergy and the clergy governs the central authority, which is structurally false in the institutional form under analysis.

A different decomposition — congregational polity, decentralised charismatic movement, household-religion — would yield a different control structure with different loops. Within the hierarchical institutional decomposition analysed here, no fifth loop exists, because no fifth pair of elements satisfies all three conditions simultaneously.

3.3 Key Observations about the Control Structure

Feedback poverty at the top. The Central Doctrinal Authority receives feedback primarily through its own appointees (clergy), who have incentives to report compliance rather than problems. External feedback (scientific findings, secular criticism, abuse reports) is structurally classified as hostile input, not corrective signal.

Asymmetric power at every level. Each controller has significantly more power over the controlled process than the controlled process has to send corrective feedback upward. This is by design (R4: authority to interpret and enforce), but it means that control errors propagate downward unchecked.

Self-confirming process models. At every level, the controller's mental model of the controlled process contains built-in resistance to disconfirmation. If a believer doubts, the model says "doubt is temptation" — the model *explains away* the very feedback that would correct it.

4. STPA Step 3: Identify Unsafe Control Actions

STPA categorizes unsafe control actions (UCAs) into four types:

1. **Not providing** a control action that is needed
2. **Providing** a control action that causes a hazard
3. **Providing** a control action **too early, too late, or out of sequence**
4. **Providing** a control action that is **applied too long or stopped too soon**

4.1 UCAs for Central Doctrinal Authority → Clergy

UCA ID	Type	Unsafe Control Action	Hazard
UCA-1	Not provided	Authority does NOT update doctrine when scientific evidence invalidates cosmological claims (e.g., geocentrism, age of earth, evolution)	H1
UCA-2	Not provided	Authority does NOT remove or discipline clergy credibly accused of abuse	H2, H5
UCA-3	Provided	Authority ISSUES doctrine that dehumanizes outsiders (infidels, heretics, apostates as subhuman or damned)	H3
UCA-4	Provided	Authority EXCOMMUNICATES or PUNISHES members who raise legitimate doctrinal questions	H1, H4
UCA-5	Provided	Authority APPOINTS clergy based on doctrinal loyalty rather than pastoral competence or moral character	H2, H5
UCA-6	Too late	Authority CORRECTS a doctrinal error only after centuries of harm (e.g., Galileo rehabilitation in 1992, 359 years late)	H1, H6
UCA-7	Too long	Authority MAINTAINS absolutist moral prohibition long after ethical understanding has evolved (e.g., blanket prohibitions that cause measurable suffering)	H4
UCA-8	Not provided	Authority does NOT establish transparent financial auditing of religious institutions	H7

4.2 UCAs for Clergy → Congregation

UCA ID	Type	Unsafe Control Action	Hazard
UCA-9	Provided	Clergy USES pastoral relationship to sexually, emotionally, or financially exploit vulnerable individuals	H5
UCA-10	Provided	Clergy PREACHES hatred or contempt toward out-groups (other religions, LGBTQ+ persons, ethnic minorities, apostates)	H3
UCA-11	Provided	Clergy WITHHOLDS sacraments or pastoral care as punishment for questioning doctrine	H1, H4, H8
UCA-12	Not provided	Clergy does NOT provide pastoral support during genuine crisis (grief, mental illness, abuse) because doctrinal framework has no model for the problem	H4
UCA-13	Provided	Clergy DEMANDS financial contributions from impoverished believers through spiritual coercion (prosperity gospel, indulgences, obligatory tithing)	H7

UCA ID	Type	Unsafe Control Action	Hazard
UCA-14	Too long	Clergy MAINTAINS shunning of excommunicated member long after the person has suffered disproportionate social death	H4, L3
UCA-15	Not provided	Clergy does NOT report abuse to civil authorities, instead handling it "internally"	H2, H5

4.3 UCAs for Congregation → Individual

UCA ID	Type	Unsafe Control Action	Hazard
UCA-16	Provided	Congregation SHUNS member who expresses doubt, questions doctrine, or leaves the faith	H1, H4, H8
UCA-17	Provided	Congregation ENFORCES social conformity on matters of personal conscience (dress, diet, relationships, career choices) through gossip and exclusion	H4, H8
UCA-18	Not provided	Congregation does NOT intervene when they observe clergy abusing a member	H2, H5
UCA-19	Provided	Congregation PRESSURES children into public faith commitments (baptism, confirmation, bar mitzvah) before they can meaningfully consent	H8
UCA-20	Too early	Congregation LABELS children as sinful, damned, or spiritually deficient at a developmentally inappropriate age	H4

4.4 UCAs for Individual (Self-Control / Internalized Beliefs)

UCA ID	Type	Unsafe Control Action	Hazard
UCA-21	Provided	Individual SUPPRESSES legitimate doubt, curiosity, or moral intuition because they have internalized "doubt = sin"	H1, H6
UCA-22	Provided	Individual REMAINS in an abusive religious environment because they believe leaving = damnation	H4, H5
UCA-23	Not provided	Individual does NOT seek professional mental health support because the religious framework teaches that faith alone should suffice	H4
UCA-24	Provided	Individual REJECTS medical treatment for self or dependents based on doctrinal prohibition	H4, L1
UCA-25	Too long	Individual CONTINUES religious practices that cause measurable psychological harm (extreme fasting, self-flagellation, sleep deprivation) beyond any devotional purpose	H4

5. STPA Step 4: Identify Loss Scenarios

Each UCA has **causal factors** — the reasons why the unsafe control action occurs. STPA organizes these into two categories:

A. Why would the controller issue the UCA? (Flawed process model, incorrect feedback, conflicting goals)

B. Why would the control action not be executed properly? (Communication failure, implementation breakdown)

5.1 Loss Scenario LS-1: Doctrine Fails to Self-Correct (UCA-1, UCA-6 → H1 → L1, L6)

The scenario: Central authority maintains a cosmological or moral doctrine that contradicts established scientific evidence or evolved ethical understanding. Believers experience cognitive dissonance between their faith and their knowledge, leading to either intellectual suppression (L6) or loss of faith and meaning (L1).

Why the controller issues this UCA:

- **Flawed process model:** The authority's model holds that doctrine is divinely revealed and therefore cannot be wrong. This makes self-correction logically impossible within the model — any correction would imply the divine source was wrong.
- **Conflicting goals:** Correcting doctrine threatens the authority's legitimacy (SC1 conflicts with R4). If the authority was wrong about X, what else might it be wrong about? The authority perceives self-correction as existentially threatening.
- **Inadequate feedback:** Believers who experience the dissonance often leave silently rather than escalating. The authority receives feedback only from those who stay — a survivorship bias that confirms the model.
- **Temporal delay:** By the time evidence is overwhelming, the authority has centuries of institutional commitment to the doctrine. Reversal costs increase with time (sunk cost fallacy at institutional scale).

Structural fix (from SC1): Build doctrinal revision mechanisms into the system. Historical examples: Vatican II (1962–65) as a partial self-correction mechanism. Judaism's tradition of Talmudic disputation, which structurally preserves minority opinions. Islam's concept of *ijtihad* (independent reasoning), when not suppressed. The key architectural principle: doctrine should be *versionable*, with explicit revision processes, not *immutable*.

5.2 Loss Scenario LS-2: Abuse Cover-Up (UCA-2, UCA-15 → H2, H5 → L2, L4, L5)

The scenario: A clergy member abuses a congregant (sexually, emotionally, financially). Other clergy and the authority become aware but do not remove the abuser or report to civil authorities. Instead, the abuser is transferred, the victim is silenced, and the institution protects its reputation. This is the pattern documented in the Catholic Church, the Southern Baptist Convention, Jehovah's Witnesses, and numerous other contexts.

Why the controller issues this UCA:

- **Flawed process model:** The authority's model holds that (a) clergy are called by God and therefore inherently trustworthy, (b) the institution's reputation is essential to its mission, and (c) internal resolution is preferable to secular intervention. All three beliefs are structurally embedded in R4 and the control structure.
- **Conflicting goals:** Protecting the institution (G3: social cohesion, R4: authority structure) conflicts with protecting the individual (G2: moral orientation, G4: comfort). The control structure systematically prioritizes the institution because the controllers are the institution.
- **Feedback suppression:** The victim's feedback (complaint, report) enters the system through the very structure that has an interest in suppressing it. There is no independent feedback channel that bypasses the authority hierarchy.
- **Accountability void:** The controller who should discipline the abuser (bishop, denominational leader) is also the person whose career and institution are damaged by disclosure. The judge and the interested party are the same entity.

Structural fix (from SC2, SC5): (a) Mandatory external reporting to civil authorities — removing the internal-only handling option. (b) Independent ombudsman or safeguarding board with authority to investigate and sanction, staffed by people outside the clerical hierarchy. (c) Separation of pastoral authority from disciplinary authority — the person who counsels the victim must not be the person who decides whether to report. (d) Statute of limitations reform within canon law. These are architectural changes to the control structure itself, not mere policy adjustments.

5.3 Loss Scenario LS-3: Radicalization through In-Group/Out-Group Escalation (UCA-3, UCA-10 → H3 → L3)

The scenario: The mechanism for creating social cohesion (R5, R6: shared identity and in-group/out-group distinction) escalates from "we are different from them" to "they are inferior" to "they are dangerous" to "they must be destroyed."

This progression has driven religious wars, pogroms, crusades, inquisitions, and modern terrorism.

Why the controller issues this UCA:

- **Flawed process model:** The in-group/out-group mechanism (R6) has no built-in limit. The same mechanism that creates belonging ("we share this faith") also creates hostility ("they threaten our faith"). The model does not distinguish between "different" and "dangerous."
- **Positive feedback loop:** Preaching against out-groups increases in-group solidarity (measurable by attendance, donations, emotional intensity). This positive feedback rewards escalation. The more extreme the rhetoric, the stronger the group cohesion — until the rhetoric produces violence.
- **External trigger:** The escalation is often catalyzed by external threat (real or perceived). Economic stress, political marginalization, or military conflict activates the in-group/out-group mechanism and pushes it toward its extreme.
- **Missing control action:** There is no structural "circuit breaker" that de-escalates when rhetoric approaches dehumanization. The authority may agree with the escalation, or may be too distant to intervene in local radicalization.

Structural fix (from SC3): (a) Explicit doctrinal commitment to universal human dignity that *overrides* in-group/out-group distinctions — structurally, this means SC3 must be higher-priority than R6 in the requirements hierarchy. (b) Interfaith engagement as a standing function, not an exceptional one — regular contact with out-groups reduces dehumanization (contact hypothesis). (c) Internal monitoring for escalatory rhetoric, with de-escalation authority assigned to a specific role.

5.4 Loss Scenario LS-4: Spiritual Coercion and Psychological Damage (UCA-11, UCA-16, UCA-17, UCA-20 → H4, H8 → L1, L4)

The scenario: An individual (often a child or young person) is subjected to religious teaching that instills deep fear (hellfire, damnation, divine punishment), shame (sinfulness, unworthiness), or guilt (for normal developmental experiences like doubt, sexual feelings, or intellectual curiosity). Social enforcement (shunning, withdrawal of love, public shaming) reinforces the psychological impact. The individual develops religious trauma, anxiety disorders, or complex PTSD.

Why the controller issues this UCA:

- **Flawed process model:** The clergy and congregation model holds that (a) fear of divine punishment is a legitimate motivator, (b) shame is a corrective emotion, and (c) social exclusion of deviants protects the community. These beliefs are embedded in R3 (codified norms), R6 (in-group/out-group), and the enforcement functions (F3, F4).
- **Missing feedback:** Children and psychologically dependent adults cannot provide effective corrective feedback to the system. They lack the power, the vocabulary, and often the awareness that their experience is harmful. By the time they can articulate the damage, they may have left the system entirely — at which point their feedback is classified as "apostate hostility."
- **Normalization:** When an entire community shares the same beliefs and practices, harmful behaviors are invisible as harmful. Telling a child they are "born in sin" feels like education, not psychological abuse, to a community that genuinely believes it.

Structural fix (from SC4, SC8): (a) Age-appropriate religious education that does not use fear or shame as motivational tools — this requires explicit pedagogical standards within the religious education subsystem. (b) Voluntary participation norms: no social penalty for non-attendance, non-belief, or questioning. (c) Mental health literacy within pastoral training — clergy should be trained to recognize when religious practice is causing psychological harm and to refer to professional support. (d) External child safeguarding standards applied to religious education, just as they are applied to secular schools.

5.5 Loss Scenario LS-5: Financial Exploitation (UCA-8, UCA-13 → H7 → L2, L5)

The scenario: Religious authority extracts financial resources from believers through spiritual coercion — "God wants you to give," "tithing is obligatory," "sowing seeds of faith will bring financial blessings" (prosperity gospel). Funds are not transparently accounted for. Leaders live in luxury while congregants struggle.

Why the controller issues this UCA:

- **Conflicting goals:** The system needs financial sustainability (R7 in the SE decomposition) but has no structural separation between spiritual authority and financial management. The person who tells you what God wants is also the person who tells you how much to give.

- **Missing oversight:** Unlike secular nonprofits, many religious organizations are exempt from financial reporting requirements in numerous jurisdictions. The control structure has no external audit function.
- **Doctrinal justification:** The prosperity gospel and similar doctrines create a self-sealing system: if you give and prosper, it proves the doctrine; if you give and don't prosper, your faith was insufficient — give more.

Structural fix (from SC7): (a) Mandatory financial transparency — published annual reports, independent external audit, separation of spiritual authority from financial authority. (b) Congregation-elected financial oversight body with real power (a "Kassenprüfung" model, borrowed from the German Verein). (c) Doctrinal prohibition on linking financial giving to spiritual outcomes.

5.6 Loss Scenario LS-6: Suppression of Individual Conscience (UCA-4, UCA-7, UCA-21 → H1, H4 → L1, L6)

The scenario: A believer's individual moral conscience or intellectual understanding conflicts with official doctrine. The system's response is to demand submission of conscience to authority. The individual either suppresses their authentic moral judgment (psychological harm, L4) or is punished for expressing it (exclusion, L1). In either case, the system has sacrificed one of its own goals (G2: moral orientation) to preserve another (R4: authority structure).

Why the controller issues this UCA:

- **Architectural conflict:** R4 (authority to interpret and enforce) and G2 (moral orientation) are in structural tension. R4 was designed to serve G2 (authority *enables* moral guidance), but R4 can become an end in itself — authority for authority's sake — at which point it *overrides* G2.
- **Process model deficiency:** The authority's model does not contain the possibility that the individual's conscience could be more morally correct than the doctrine. The model assumes doctrinal infallibility, which means any dissent is by definition error.

Structural fix (from SC1, SC4): (a) Explicit doctrinal recognition of conscience as a valid moral source — Catholicism actually has this (Dignitatis Humanae, Vatican II) but struggles to operationalize it when conscience contradicts doctrine. (b) Structured dissent mechanisms: the ability to formally register disagreement without excommunication, analogous to a "dissenting opinion" in a court ruling. (c) Periodic doctrinal review processes that engage lay members and external scholars, not just clergy.

6. Cross-Cutting Structural Findings

6.1 The Feedback Problem Is Universal

Every loss scenario shares a common structural deficiency: **inadequate upward feedback**. The control structure of religion is designed for top-down transmission (doctrine flows from authority to believer) but has weak, filtered, or suppressed bottom-up feedback. This is not a bug but a feature of the design (R4: authority to interpret and enforce). However, it means that control errors are systematically uncorrected.

In control systems engineering, this is called **open-loop control** — the controller acts without adequate measurement of the actual system state. Open-loop control works only when the process is perfectly predictable. Human societies are never perfectly predictable. Therefore, open-loop control of human behavior through religion will always produce drift toward loss conditions.

The architectural remedy is to **close the loop**: create genuine, unfiltered feedback channels from the controlled process (believers, congregations) to the controller (authority), with structural protections ensuring the feedback cannot be suppressed.

6.2 The Self-Sealing Model Problem

Multiple UCAs are enabled by **self-sealing process models** — mental models held by controllers that explain away any evidence that would contradict them:

Self-Sealing Pattern	How It Works
"Doubt is sin"	Eliminates the primary corrective signal (questioning)
"Suffering is divine test"	Reframes harm as benefit — the worse it gets, the more it "proves" the model

Self-Sealing Pattern	How It Works
"External criticism is persecution"	Classifies corrective feedback from outside the system as hostile action
"Apostates were never true believers"	Discounts the testimony of the people with the most relevant experience
"God's ways are beyond human understanding"	Makes the model unfalsifiable by definition

In control engineering, a process model that cannot be updated by feedback is called **unobservable**. An unobservable system cannot be controlled safely. The architectural remedy is to ensure that every process model in the control structure includes explicit **update mechanisms** — conditions under which the model itself is revised.

6.3 The Conflicting-Goals Problem

The SE decomposition reveals that several system goals are in structural tension:

Goal Pair	Tension
G2 (moral orientation) vs. R4 (authority)	Authority enables morality but can also override conscience
G3 (social cohesion) vs. R6 (in-group/out-group)	Cohesion requires boundaries, but boundaries enable exclusion and hatred
G4 (comfort) vs. R3 (norms)	Comfort requires unconditional acceptance, but norms require conditional compliance
G1 (meaning) vs. R1 (cosmological narrative)	Meaning requires honest engagement with reality; a fixed narrative may conflict with reality

These tensions are not design flaws — they are inherent in any system that simultaneously serves individual and collective needs. The design flaw is when the system has **no mechanism for managing the tension**: no process for deciding which goal takes priority when they conflict, and no authority empowered to make that trade-off transparently.

6.4 The rung mismatch behind the four findings

The previous chapter Justificatory Rungs introduced a seven-rung ladder of standards under which control actions and feedback signals are accepted. Reading §§6.1–6.3 through that lens, **rung mismatch subsumes them all** for this system:

- **"Feedback poverty at the top"** is the rung-1 controller refusing rung-3 input — Pattern A (asymmetric loop). The rung-1 hierarchy classifies abuse data, scientific evidence, and demographic harm as rung-1 hostility and filters them out.
- **"Self-confirming process models"** is the controller's *acceptance-rung filter* operationalised. Each entry in the table in §6.2 ("doubt is sin," "suffering is divine test," "external criticism is persecution," "apostates were never true believers," "God's ways are beyond human understanding") is a specific rule for re-classifying rung-3+ feedback as rung-1 noise.
- **"Conflicting goals"** are tractable when the system can deliberate at rung 6 about how to balance them; in a rung-1 system, whichever side has more authority simply wins.

The system simultaneously runs **Pattern B — Claimed-Rung Inflation**: it claims rung 6 (sacralised ultimate legitimacy) while operating at rung 1 (clerical authority). The gap is enforced doctrinally rather than closed, which is why the failures are predictable and recurrent.

The architectural fixes in §7 are all attempts to insert **rung-3 channels that route around the rung-1 filter** without forcing the controller to abandon its rung-1 self-understanding. Mandatory civil reporting, independent ombudspersons, lay safeguarding boards, doctrinal version control with external scholarship, and pre-registered scientific-consultation processes are all implementations of the same architectural pattern.

This is the single most important structural finding of the analysis: religion's harmful outcomes are not random malfunctions but the **predictable consequence of operating a sacralised rung-1 control architecture while claiming the authority of rung 6**. The canonical reference, including the rung-tagged control structure and the rung-acceptance filter as a recurring causal factor in every loss scenario, is in <knowledge/system-catalogues/social-systems/religion/applied-se-analysis.md>.

7. Recommendations: Architectural Improvements

Based on the STPA analysis, we propose the following changes to religion's control structure. These are stated as architectural principles, applicable across traditions:

7.1 Close the Feedback Loop

- Establish **independent feedback channels** that bypass the authority hierarchy: ombudsman, external review boards, anonymous reporting systems.
- Require **periodic surveys of congregational wellbeing** — measuring psychological health, not just attendance and donations.
- Create **formal exit interviews** for members who leave, treated as system feedback rather than apostasy.

7.2 Make Process Models Updatable

- Institute **doctrinal version control** — every doctrine should have a stated scope, a review date, and a revision process.
- Require **engagement with external knowledge** (science, psychology, ethics) as a standing input to doctrinal governance, not an exceptional event.
- Preserve **dissenting opinions** in the doctrinal record, as courts preserve minority opinions.

7.3 Separate Control Functions

- Separate **spiritual authority** (preaching, sacraments, counseling) from **administrative authority** (finances, personnel, discipline).
- Separate **pastoral care** from **norm enforcement** — the person who counsels you should not be the person who punishes you.
- Establish **independent safeguarding** with authority to investigate and override clergy decisions on abuse cases.

7.4 Install Circuit Breakers

- Define **escalation limits** on in-group/out-group rhetoric: explicit doctrinal commitment to universal human dignity that cannot be overridden by any other teaching.
- Establish **proportionality review** for disciplinary actions: shunning, excommunication, and other exclusion mechanisms should require documented justification, appeal process, and time limits.
- Require **informed consent** for participation in intensive religious practices (extended fasting, retreats, exorcism rituals) with right of withdrawal.

7.5 Protect Vulnerable Populations

- Apply **secular child safeguarding standards** to all religious education and youth activities.
 - Prohibit **fear-based instruction** (hellfire, damnation) for children below a developmentally appropriate age.
 - Require **mandatory external reporting** of all abuse allegations to civil authorities, with no exceptions for internal handling.
-

8. Conclusion

STPA reveals that religion's most harmful outcomes are not random malfunctions but **predictable consequences of its control architecture**. A system designed for top-down doctrinal transmission with weak upward feedback, self-sealing process models, unaccountable authority, and no structural mechanism for managing goal conflicts will, over sufficient time, produce every one of the losses we identified — regardless of the sincerity or goodness of the people within it.

This is not an indictment of religion's *purpose*. The goals — meaning, morality, community, comfort — are among the

most important things humans seek. It is an indictment of religion's *control architecture* as historically implemented. The same goals could be served by a system with closed feedback loops, updatable doctrines, separated authority, independent oversight, and protected individual conscience. Some religious traditions have moved in this direction; none has fully arrived.

The value of STPA is that it moves the conversation from "religion is good" vs. "religion is harmful" — a debate that generates more heat than light — to a structural question: **which control actions, under which conditions, produce which outcomes?** That question has specific, actionable answers. The fixes are architectural, not theological. They do not require abandoning faith. They require redesigning the system that carries it.

This analysis applies STPA (Leveson, 2012) to the SE decomposition of religion developed in prior work. The four STPA steps — define losses, model control structure, identify UCAs, trace loss scenarios — are followed systematically. All hazard and constraint IDs trace back to the original SE decomposition's goals and requirements.

References

- Leveson, N.G. (2012). *Engineering a Safer World: Systems Thinking Applied to Safety*. MIT Press.
- Leveson, N.G. (2004). "A New Accident Model for Engineering Safer Systems." *Safety Science*, 42(4), 237–270.
- Leveson, N.G. & Thomas, J.P. (2018). *STPA Handbook*. MIT Partnership for Systems Approaches to Safety and Security.

STPA Analysis of Frontier AI as a Sociotechnical System

What Will AI Do to Us? A Systems-Engineering Reading

Abstract

The diffusion of frontier artificial intelligence is reshaping the global economy and political order, but most public analysis of this fact takes the form of strategic essays — narratives that identify forces and rank scenarios without pointing at specific levers. This chapter applies System-Theoretic Process Analysis (STPA), with justification-rung tagging from the social-systems extension developed in Part I, to the global political economy of frontier AI. We identify fourteen distinct controllers, twenty-six control actions, sixty-five unsafe control actions across the four standard categories, and nine loss scenarios connecting specific structural mismatches to specific harms. We then derive twenty-two concrete remedies organised into a four-phase sequence with explicit decision points. The result is a structural reformulation of the questions about AI's effects on jobs, power, and geopolitics — turning forecasting questions into action-shaped questions of which feedback channels are missing and which institutions are needed to install them. The full analytical artefacts are filed under [projects/problems/ai-impact-analysis/](#); this chapter distils them into prose for a non-specialist reader.

1. The questions, and why the usual answers do not satisfy

Three questions about artificial intelligence circulate in every serious conversation about the next decade.

The first is about jobs. If software writes itself, what becomes of the people who wrote it? If diagnoses, contracts, lesson plans, research summaries, and software designs are produced in seconds by machines that cost a fraction of a human salary to operate, how many human roles survive? Will the result be a new wave of unemployment, or a new kind of leisure?

The second is about power. The systems that now demonstrate near-human reasoning at language tasks were built by perhaps a dozen laboratories, financed by perhaps a hundred investors, and trained on perhaps a few hundred specialised computers in a few jurisdictions. If the technology becomes the central productive force of the next era, what stops the firms that own the substrate from owning, by extension, much of what the substrate produces?

The third is about geopolitics. The same technology is being built in two large economic blocs that mistrust one another. The American bloc's frontier laboratories sit at the top of the global capability ranking; the Chinese bloc has matched them at lower rungs and is catching up at the top. Export controls, sanctions, and chip-supply dependencies have already produced the first skirmishes of an infrastructure war. Where does this end?

Each question has been answered many times. The answers are sophisticated, but they share a common shape and a common limitation. The shape is the strategic essay: a narrative that identifies the main forces, ranks them by importance, and concludes with a list of probable scenarios or a cautious warning. The limitation is that the essay's conclusions, however confidently stated, are *free-floating*. They are not connected to specific mechanisms in the world that one could intervene on. The essay tells the reader what is likely to happen; it rarely tells the reader where the system can be moved.

This chapter does something different. It treats the AI economy and its political surroundings as a system in the engineer's sense: a collection of parts connected by flows and feedback loops, each of which can be examined and, if it is part of a failure pathway, changed. The discipline that does this is called systems engineering, and a particular variant of it called system-theoretic process analysis (STPA) was originally developed for safety-critical machines like aircraft and medical devices. We use it here for a different class of artefact: not a machine, but the global political economy of frontier AI.

The reader gains, by the end, a structural picture of the system, five specific identified failure mechanisms inside it, nine specific scenarios in which those mechanisms produce identifiable harms, and twenty-two concrete interventions arranged into a four-phase implementation plan. This is more than an essay can deliver, and the reason it is more is the method. So we begin with the method.

2. Two ways of thinking about the system

A strategic essay about AI says things that sound like this. *Power will concentrate, but open-weight models from China will keep prices honest. Workers will be displaced, but new categories of work will emerge as they always have. The geopolitical risk is significant but probably manageable. The most important uncertainty is whether redistribution mechanisms catch up to the productivity shock.* These sentences are useful for orientation. They are also, on closer examination, the kind of observations that one could make about almost any technology in almost any decade. They do not tell you which lever to pull. They do not even tell you which levers exist.

A systems-engineering analysis, by contrast, begins by asking who is actually in the system, what each of them controls, and what information each of them receives. It records this not as a narrative but as a small set of explicit lists: the controllers, the things they control, the *control actions* they issue, the *feedback channels* they receive, and — crucially — the *internal model* each controller carries of the thing it is controlling. It then asks, for each control action, what it would mean for that control action to be issued unsafely (provided when it should not be, withheld when it should not be, given too late, given for too long). Each unsafe possibility is recorded. Then it asks, for each unsafe possibility, *why* the unsafe action could occur — the controller's model could be wrong, the feedback could be missing or filtered, the action could be misexecuted, the controlled process could have changed underneath. Each cause suggests a remedy: fix the model, fix the feedback, fix the actuator, fix the structure.

This sounds like ordinary strategic thinking with extra paperwork, and at first glance the difference is just that there are tables where the essayist has paragraphs. The deeper difference shows up at three points. First, the act of writing down every control action and every feedback channel forces the analyst to confront parts of the system that the essayist could quietly skip past; the omissions become visible. Second, recording each unsafe control action against the four standard categories produces a checklist of things-that-could-go-wrong that is exhaustive in a way that strategic intuition is not. The analyst is no longer free to remember only the failure modes that fit the narrative. Third, the loss-scenario step demands that the analyst not just *assert* that a failure could occur but *trace* the causal chain through the structure. This is what produces remedies. A narrative observation that "regulators are slow" is not actionable; a loss scenario showing that *labour ministries operate on a years cycle while capability deployment happens on a months cycle, so the corrective signal arrives after the cohort it was meant to help is already displaced* points directly at automatic stabilisers indexed to a fast-moving benchmark.

There is a further refinement, particular to systems engineering applied to social and political systems: every flow in the structure can be tagged with the *justificatory rung* on which it operates. The rungs are a seven-step ladder of how a claim or a control action gets its force. Rung 0 is naked coercion: do it or else. Rung 1 is authority: do it because *I* say so, or because the tradition says so, or because our group says so. Rung 2 is formal consistency: this follows from rules we both accept. Rung 3 is empirical evidence: this works against rivals when tested. Rung 4 is replicated cumulative evidence. Rung 5 is meta-rational integration of multiple methods under uncertainty. Rung 6 is deliberative legitimacy — the standard by which questions about what should be done are settled in open discourse among those affected.

Each of these rungs can be claimed. Each can be operated. The two are not always the same, and where they differ, the system typically produces specific kinds of damage. A laboratory that claims to operate at rung 6 (its safety framework is "for the benefit of humanity") but actually operates at rung 1 (the framework binds only the laboratory itself, by its own choice, amendable by its own governance) is in a *claim/operation* gap. A funding flow that goes downward at rung 3 (financial-return diligence) but receives capability evidence upward at rung 1 (the lab's own benchmarks, in the lab's own framing) is in an *asymmetric loop* — the loop appears closed because there are arrows in both directions, but the upward arrow is filtered by the downward arrow's framing, and corrective information cannot get through.

Once you start labelling every arrow with its rung, three patterns keep appearing. The first is the asymmetric loop just described. The second is the claim/operation gap. The third, which I will explain when we encounter an instance, is *cross-loop rung imposition* — when an institution tries to settle a question that belongs at one rung using the authority of a different rung. Each of these patterns has a generic remedy that the engineering literature catalogues. The remedies are structural — they change the architecture, not the personalities — and that is what makes them robust.

We now have everything we need to look at the AI system itself. We will list its parts, draw its loops, mark the mismatches that the rung-tagging method surfaces, walk through the failures these mismatches produce, and end with the remedies they imply. The reader who finds the technical vocabulary unfamiliar can lean on the prose: we will introduce each new term in plain English at the moment it enters the chapter.

3. The system

To draw the AI system's control structure honestly we have to be willing to write down a longer list of actors than a strategic essay would normally bother with. The longer list matters because the failures we are looking for happen in the connections between actors, and a connection is invisible if either of its endpoints is missing from the picture.

At the apex sits something we will call *the global polity*: diffuse, slow, but the source of legitimacy on which most controllers ultimately depend. Below it, three large state groupings — the United States, the People's Republic of China, and the European Union together with other states — each with several distinct organs. The United States in this analysis is not one actor but six: the Congress that passes legislation, the Bureau of Industry and Security that issues export controls, the Treasury that runs sanctions and tax, the Federal Trade Commission that runs antitrust, the AI Safety Institute that runs technical evaluation, and the national security apparatus that classifies and designates. The People's Republic of China is three: the party centre that sets strategic direction, the Ministry of Industry and Information Technology that runs industrial policy, and the Cyberspace Administration that regulates content and approves models. The European Union is two: the Commission with its AI Office, and the member states. Other states — the United Kingdom, Japan, South Korea, India, the Netherlands, Taiwan, the Gulf monarchies — appear as a coalition layer whose compliance or defection determines whether the bigger states' levers actually move anything.

Beneath the state layer sit the capital markets in their various forms — public equities, private capital, sovereign wealth, debt. These are the controllers that decide which laboratories can sustain frontier training runs and which compute providers can sustain capital-expenditure cycles. Beneath them sit the frontier laboratories themselves, perhaps a dozen worldwide, each with its internal mission narrative, its internal capability assessments, and its internal safety framework. Beneath the laboratories sit the compute infrastructure providers — the chip designers, the foundries, the hyperscale cloud operators, the energy and grid companies. To one side of this hierarchy, neither below nor above the labs but functionally parallel to them, sits the open-weights ecosystem: the laboratories, mostly Chinese but including some Western publishers, that release model weights free for download. On the other side sit the application-layer companies that turn model capability into deployed software for paying customers. Below those, the information platforms that distribute everything. Below them the labour ministries and education systems that interface with the people whom all of this affects, and the information regulators that look after the integrity of the public information environment. At the bottom — though "bottom" is wrong because they are the controlled process this whole apparatus ultimately produces effects on — are the workers, consumers, and citizens of the cognitive-knowledge economy.

This is fourteen distinct kinds of controller, and a quarter of a hundred when we count the sub-units. The point of writing them all down is not exhaustiveness for its own sake. The point is that each of them issues control actions — directives, allocations, investments, releases, mandates, statutes, lawsuits, salaries — and each of them receives feedback — financial returns, benchmarks, statistics, news, lawsuits, elections, market prices. The shape and the rung of those flows is what determines whether the system as a whole works. A strategic essay that talks about "the AI industry" or "the labour market" without pulling these controllers apart cannot see, for example, that the antitrust authority's process model treats concentration as something that happens in mergers, while concentration in this market is happening in long-duration compute-allocation contracts that no merger review touches.

We can summarise the structure in a single picture. Money flows downward — capital allocates to laboratories and compute providers; laboratories pay compute providers for training; compute providers pay foundries for chips; everyone hires people. Capability flows outward — laboratories release models, application companies wrap those models into products, products reach workers and consumers. Information flows in two directions — capability claims and incident reports flow upward from labs to capital, regulators, and the public; benchmarks and prices flow horizontally between labs and the open-weights ecosystem; employment and price statistics flow upward from the labour and consumer base to ministries and to the polity. State controllers issue rules and enforcement actions downward; backlash flows upward when other channels fail. The picture is dense, but it is not infinite. In the analysis companion to this chapter, every arrow is enumerated; here we only need the shape.

The diagram below shows the principal controllers and the channels between them, with rung tags on the arrows and the five mismatches (M1–M5) marked at the boundaries where they occur. Solid arrows are control actions; dashed arrows are feedback.

[Diagram: see online version]

When we tag the arrows with their justificatory rungs, five specific mismatches stand out. Each of them is a particular case of one of the dangerous patterns we introduced in the previous section. Each is a structural reason why specific failures occur.

3.1 The first mismatch: lab safety frameworks claim more legitimacy than they operate

Frontier laboratories publish documents — *responsible scaling policies*, *deployment policies*, *constitutional AI charters* — that claim deliberative legitimacy. Their tone is the tone of public ethics: the lab is binding itself, on behalf of humanity, to thresholds that cannot be crossed without specified cause. The language sits at rung 6, the rung at which questions about what *should* be done are settled.

The operating mechanism, however, is rung 1. The lab is bound because the lab has chosen to be bound. The lab amends its framework when its own governance decides to amend it. There is no external party with standing to enforce, no court to which violations can be brought, no parliament that can refuse a budget on a non-compliant lab. This is the

claim/operation gap. The gap is not a sign of bad faith — many of these documents are written in good faith by people who believe in them. The gap is structural: a self-binding commitment, however sincere, is operationally a rung-1 commitment, regardless of how it sounds.

The damage is that public trust calibrates to the higher claim and gets eroded at the rate of the gap. Every release that rolls back a previous threshold, every quiet amendment to an evaluation protocol, every framework that survives commercial pressure only in form, contributes to a slow loss of legitimacy that the labs cannot themselves staunch — because the legitimacy was never theirs to issue.

3.2 The second mismatch: capital cannot evaluate the laboratories it funds

Frontier laboratories are funded by capital allocators — venture funds, growth equity, public markets, sovereign wealth. The funds flow downward at rung 3: financial diligence, term-sheet evaluation, comparable-round analysis. The capability evidence flows back upward through the laboratories' own benchmark selection, the laboratories' own model cards, the laboratories' own mission narrative. This is rung 1 in disguise — capability claims wrapped in the lab's chosen frame.

The result is what the engineering literature calls an *asymmetric loop*. The arrows exist in both directions, but the upward arrow is filtered by the framing of the downward arrow. Capital cannot ask, with the rigour its own diligence claims, whether a lab's mission claims are falsifiable: the very benchmarks against which the question would be asked were chosen by the entity making the claims. The fund can either underwrite the mission narrative as given or refuse to fund — and at the scale of capital required for frontier training runs, refusing means watching a competitor fund the same lab on more pliable terms.

The consequence is that capital structure exerts a quiet but relentless pressure on operating practice. Funding rounds occur on a quarterly cadence; safety frameworks are amended on a yearly cadence; capability releases happen between rounds. When the cadences disagree, the cadence with the money wins.

3.3 The third mismatch: the labour-policy cycle cannot keep up with the capability-deployment cycle

Labour ministries and education systems track unemployment, sectoral employment, vacancies, and wage statistics. These indicators arrive at rung 3 or rung 4 — empirical aggregates, sometimes cumulative — but they arrive on a quarterly or annual cycle. Retraining programmes and curricular updates respond on a multi-year cycle. Capability deployment, by contrast, moves on a months cycle: a new model is released, application companies integrate it, automation is in production within weeks.

The mismatch in cycle times means that the corrective signal — a labour-market indicator showing that a class of workers is being displaced — arrives after the cohort it would have informed is already displaced. The signal cannot land in time to affect the decisions that produced the displacement. When this happens repeatedly, the rung-3/4 channel loses its informational value, and the only channel left is rung 0: political backlash, which arrives at the speed of news cycles and elections rather than at the speed of statistics. The corrective intent of the statistical channel is overrun by the rung-0 escape valve.

This is not a story about regulators being slow in some abstract sense. It is a specific identification of two cycles whose periods no longer line up — and the structural fix is not to be smarter, but to install instruments whose triggers fire automatically on a fast index, removing the requirement that a slow channel confirm.

3.4 The fourth mismatch: incidents inside the labs do not reach outside observers

When something goes wrong inside a frontier lab — a misuse case caught in red-team testing, an unexpected capability surfaced late in evaluation, a deployment incident in the wild — the information first exists at rung 3, inside the lab, in the internal logs. Outside observers — the AI Safety Institute, the EU AI Office, the press, the public — depend on what the lab chooses to disclose, which arrives wrapped in rung-1 framing (the lab's authority on the meaning of its own incidents).

The loop appears closed. Labs do publish model cards and post-mortems. But the closure is in form only. The rung-3 evidence that exists inside is filtered through rung-1 disclosure before reaching the controllers whose process models it would correct. This is the same asymmetric pattern as the capital/lab loop, in a different boundary. It is the structural reason why the press keeps discovering things about labs that the labs already knew.

3.5 The fifth mismatch: antitrust looks at mergers; concentration happens in contracts

Antitrust authorities — the FTC and DoJ in the United States, DG-COMP in the European Union, SAMR in China — evaluate market concentration through merger review and conduct cases. Their process models track mergers and

acquisitions, dominant-firm behaviour, and pricing patterns. The cycle on which they operate is years. The threshold above which a transaction triggers review is set in statute.

Concentration in this market, however, does not flow primarily through mergers. It flows through long-duration exclusive compute-allocation contracts between hyperscale providers and frontier laboratories, through bundled tying arrangements between chip designers and their largest buyers, through grandfathered grid-interconnection rights at data-centre sites. None of these trigger merger review, because none of them are mergers. They happen at the contract level, on a quarters cycle, and they foreclose entry by smaller actors before any antitrust authority has cause to look.

This is *cross-loop rung imposition* in the third dangerous-patterns sense. The rung-2 forum (statutory antitrust review) is being asked to adjudicate a dispute generated below its threshold and beneath its cadence, and it cannot. The forum is the wrong shape for the adjudication. The structural fix is not to make antitrust faster; it is to give the forum a different statutory mandate that captures contract-level conduct.

3.6 What having five mismatches in hand changes

A strategic essay can speak generally about "the AI industry's governance problem." Naming five specific mismatches changes the conversation in three ways.

It changes what is observable. Once we know that the lab/capital loop is asymmetric, we can look for the symptom: capability claims that survive funding rounds but not independent evaluation. We can know what *would* falsify the diagnosis, and what *would* confirm it.

It changes what counts as a remedy. A general call for "better AI governance" is not a remedy. A specific structural intervention that turns an asymmetric loop into a symmetric one — say, by inserting a statutory pre-deployment evaluation channel that does not pass through the lab's own benchmark selection — is a remedy. The intervention does not depend on anyone changing their mind; it changes the architecture.

It changes what survives political turbulence. Personal commitments and voluntary frameworks last as long as the people and the incentives that sustain them. Structural changes — new statutes, new institutions, new contractual rules — last as long as the political settlement that created them, which is typically much longer. The remedies the engineering literature catalogues are the structural kind.

Key Finding — the five rung mismatches

Tagging every flow in the AI control structure with its justificatory rung surfaces five specific structural mismatches:

- **M1.** Lab safety frameworks claim rung-6 deliberative legitimacy but operate at rung-1 self-binding (Pattern B — claimed-rung inflation).
- **M2.** Capital flows downward at rung-3 due-diligence; capability evidence flows back at rung-1 lab-selected benchmarks (Pattern A — asymmetric loop).
- **M3.** Labour-policy cycle (years) is mismatched to capability-deployment cycle (months); the rung-3/4 corrective channel arrives after the cohort it would have helped is already displaced.
- **M4.** Incident evidence exists at rung-3 inside labs but reaches outside observers wrapped in rung-1 disclosure (Pattern A again, at a different boundary).
- **M5.** Antitrust operates at rung-2 with a years cycle and a merger-review threshold; concentration happens at rung-3 in long-duration contracts on a quarters cycle and below the threshold (Pattern C — cross-loop rung imposition).

These five mismatches account for most of the loss scenarios developed in §4 and motivate most of the remedies in §5.

We turn now to what specifically goes wrong when these mismatches are not addressed.

4. The nine ways the system can fail

A loss scenario, in the engineering vocabulary, is a description of how a particular bad outcome arises from particular failures inside the structure. It names the controllers whose process models go wrong, the feedback channels that go missing or get filtered, the controlled processes that drift into states from which exit is costly, and the external disturbances that finish the job. A loss scenario is, in effect, a small story about *how* the system produces harm, told in the language of its own architecture.

The full analysis identifies nine such scenarios for the AI system. They are not predictions. Each of them describes a

pathway by which a specific identified harm could be produced, traceable to specific mechanisms one could intervene on. Some pathways are already partly running; others are latent. None is inevitable. The analytic purpose of writing them down is to attach each remedy to a specific failure pathway, so that the policy debate has a structural rather than a rhetorical anchor.

4.1 Concentration entrenches at the infrastructure layer

The first scenario is about how AI's economic gains end up concentrated in the hands of a small number of firms — but not through the obvious channel one would expect.

Antitrust authorities, as we saw in section 3.5, look at mergers. They look in the wrong place. The actual mechanism of concentration is contractual. A hyperscale cloud provider signs a multi-year exclusive capacity agreement with a frontier laboratory. A chip designer commits its leading-edge supply, a year ahead, to its three largest buyers. A foundry pre-allocates capacity windows. A data centre receives a grid-interconnection right while smaller applicants sit in a queue. Each of these is a single transaction that, on its own, looks like ordinary commerce. None of them crosses a merger threshold. None of them appears in any one controller's process model as concentration.

The cumulative effect is foreclosure. Smaller laboratories, open-weights publishers, application companies that depend on non-frontier compute, all find that the affordable supply has been contracted away in advance. Entry becomes a financial impossibility. The controlled process — the structure of the AI infrastructure market — moves to a state from which exit is costly: the contracts have already been signed; the grid interconnections have already been allocated; the leading-edge fab capacity is forward-booked for years.

The corrective channel exists, in principle. Concentration ratios, prices, and entry rates are measured. But they are measured at the slow rate at which antitrust evidence accumulates — over years — while the foreclosing transactions happen quarterly. By the time the evidence is strong enough to act on, the architecture is already locked in. The remedy class is *cadence-matched controls*: rules that fire on the transaction cycle, not on the evidence cycle. The specific remedies in the catalogue are reservation requirements on leading-edge supply (a defined fraction must be sold to non-frontier customers), an expedited review window on long-duration contracts during defined capability-jump moments, and a merger-review reform that captures contract-level foreclosure in the same statutory net as ordinary M&A.

4.2 Lab safety frameworks become ceremonial

The second scenario follows directly from the first and second mismatches, working together. Capital pressures laboratories to release on a fast cadence to defend market share. The laboratories' safety frameworks claim deliberative legitimacy. The frameworks are amended between releases, in response to commercial pressure, to allow the next release. The amendment process is internal; the legitimacy claim is external. The gap widens each cycle.

The damage here is not principally that any particular release is unsafe. It is that the legitimacy claim erodes. Each visible amendment to a safety threshold, each evaluation that produces a result the lab finds inconvenient and reframes, each model that ships with caveats no external party could have written, drains the rung-6 reservoir the labs implicitly drew on when they published their first safety frameworks. At some point the public notices that the frameworks are advisory, the amendments are unilateral, and the claims are aspirational. When that happens, the labs can no longer issue legitimate safety assurances of any kind.

The remedy class here works in two directions at once. It *lowers the claim*: the lab admits, structurally, that its frameworks are self-binding rung-1 commitments — useful, but limited. It *raises the operation*: an external evaluator with statutory authority enters the loop. The specific remedies in the catalogue are mandatory pre-deployment evaluation by an independent body (the AI Safety Institutes are early examples, operating today as voluntary partners but candidates for statutory upgrade), structural separation of safety governance from commercial governance inside labs, and a capital-structure mission-constraint mechanism that gives a public agent — a state attorney general, a regulator — standing to enforce the mission constraint, rather than relying on shareholder litigation that will never come.

4.3 Workers cannot keep up with capability deployment

The third scenario is about the labour market and runs through the third mismatch. Capability deploys on a months cycle — model release, application integration, deployment in customer workflow. Labour ministries' instruments — retraining programmes, curriculum updates, unemployment-insurance schemes — operate on a years cycle. Cohort-specific damage to early-career cognitive workers — the hardest-hit group because their tasks are most substitutable — is invisible in the headline employment statistics that ministries track.

The corrective channel is structurally late. By the time the employment data shows a problem, the cohort the data described has already entered the labour force, found that the entry-level roles they trained for are gone, and either dropped out, accepted underemployment, or organised politically. The political channel operates at rung 0 (electoral backlash) on a fast timescale that the rung-3 statistical channel cannot match. The rung-0 channel arrives, in other

words, exactly when the rung-3 channel was supposed to have prevented its arrival.

The remedy class is *automatic stabilisers* indexed to a fast benchmark of task displacement, not to headline employment. Indexed unemployment insurance, indexed earned-income tax credit, indexed retraining credits: instruments whose triggers fire without requiring legislative confirmation per cycle. To these add cohort-specific transition programmes for the early-career group — recognising that the affected cohort changes year by year as substitution sweeps through different occupations — and portable benefits that decouple healthcare and pension continuity from continuous employment with a single employer. None of these require new theory; they require the index, the statutory trigger, and the political decision to commit ahead of the displacement rather than after.

4.4 The productivity surplus accrues to capital and stays there

The fourth scenario is the distributional one. Suppose the productivity gain from AI is real — every careful estimate suggests it is, on the order of fractions of a percentage point per year of GDP growth, possibly more. The question is who receives it.

The mechanisms by which a productivity gain reaches workers, as wages, depend on labour bargaining power. The mechanisms by which it reaches consumers, as lower prices, depend on competitive markets. The mechanisms by which it reaches the public, as taxation, depend on a fiscal architecture calibrated to the new factor shares. None of these is automatic. If labour's bargaining power falls (because AI substitutes for substitutable cognitive tasks), the wage channel narrows. If markets concentrate (scenario 4.1), the price channel narrows. If tax instruments calibrated to corporate income do not adjust for the shift toward AI-as-capital that flows to less-taxed bases, the public channel narrows.

The structural diagnosis is that no controller in the existing system is positioned to act on cross-sectoral productivity-gain capture as such. Tax authorities track corporate income, capital gains, and payroll — three categories that together miss the factor-share shift. There is, in the engineering vocabulary, a *controller absent* from the loop. The remedy class is to add one: a transmission mechanism, of which the catalogue lists four candidates — a productivity-indexed dividend (analogue: Alaska Permanent Fund, but indexed to AI productivity rather than oil), a value-added tax on AI inference with hypothecated dividend, an expanded earned-income tax credit indexed to the displacement benchmark, and sovereign-wealth participation in AI infrastructure with public dividend.

These are not interchangeable. The productivity-indexed dividend is conceptually clean but politically the hardest to land. The inference VAT is operationally simpler because it taxes the visible economic flow. The EITC expansion is the most familiar because the instrument already exists. Sovereign-wealth participation requires a fund and a governance mechanism; some jurisdictions have one (Norway, Singapore, the Gulf), the United States does not.

The point is not which one is correct; it is that without at least one of them, the productivity surplus has no transmission channel to labour, consumers, or the public. It will accrue to capital not because capital is rapacious but because capital is the only controller in the loop with a defined claim on it.

4.5 The open-weights price-setting mechanism collapses

The fifth scenario concerns a controller most strategic essays do not name as a controller at all: the open-weights ecosystem. Chinese laboratories and a handful of Western publishers release model weights free for download. The releases serve many purposes, but at the system level they perform one specific function: they set a price ceiling on the model layer. Frontier laboratories cannot price closed-weight inference far above what an adequately-equipped competitor can do with a recent open-weight release. The ceiling exists.

If the ceiling collapses, the model layer becomes a place where a small number of closed-weight providers can extract whatever the market will bear. The collapse can happen in three ways. Western governments could regulate open-weights releases out of existence — well-intentioned, on misuse-risk grounds, without recognising that the regulation also removes the price-setting mechanism. Capital could flee the few Western open-weights labs whose business is harder than their closed-weights competitors', leaving Chinese labs as the only credible open-weights suppliers — at which point the political acceptability of using their weights becomes a separate concern. Or Chinese labs could shift strategy and stop releasing open weights, for reasons of their own, and no one outside has the leverage to insist.

The diagnostic point is that the open-weights ecosystem performs a function that no controller is assigned to protect. Regulators treat open weights primarily as a misuse risk; market participants treat them as a competitive threat or competitive crutch. No body publishes a measure of the frontier-versus-open-weights gap as a system-level state variable that regulators are obliged to consider. The remedy class is what the dangerous-patterns catalogue calls *domain-appropriate rung containment*: route open-weights questions to a forum that holds both misuse risk and price-setting in superposition, and require any rule affecting open weights to document its impact on the gap. Concretely: publish the gap, calibrate misuse-risk rules narrowly enough to preserve the ecosystem below those thresholds, and treat the ≤ 12 -month gap as an explicit policy goal rather than an emergent outcome.

4.6 Compute competition escalates to kinetic conflict

The sixth scenario is the most extreme. Compute is geographically concentrated. Leading-edge logic chips come overwhelmingly from a single foundry on a single island. Hyperscale capacity is concentrated across a small number of operators in a small number of regions. Energy and grid interconnections are concentrated where electricity is cheap. The geography of AI infrastructure is fragile in ways the geography of, say, the global oil supply once was — a small number of physical chokepoints whose disruption would stop the flow.

National security apparatuses on both sides of the US–PRC relationship watch each other's compute build-out. Each side's process model treats the other's build-out as preparation for adversarial use. The defensive interpretation — each side is building because each side fears being shut out — is filtered out on both sides by the rung-1 strategic identity that does not admit the other's defensive logic. There is no confidence-building channel at the AI-infrastructure level. The only available channels are coercive (export controls, sanctions, designations) and rhetorical (state speeches).

A scenario in which a regional crisis, perhaps unrelated to AI itself, transmits directly into the AI-infrastructure layer is not implausible. A blockade or a strike that disrupts foundry output. A sanctions package that pulls one side into commercial escalation. A reciprocal designation that classifies a category of research on both sides simultaneously. Each step is small; the cumulative path is from rung-2 rule-making through rung-0 coercion toward kinetic action.

The remedy class here is the hardest in the catalogue and the most important. An independent inspection or observation regime for AI-infrastructure deployments — analogous to the IAEA for nuclear material, the OPCW for chemical weapons, or the now-defunct Open Skies regime — would provide a rung-3 verification channel that bypasses the rung-1 filter on both sides simultaneously. It would not solve the conflict; it would create a forum in which evidence about the system's actual state could enter both sides' decision-making, reducing the chance that either side acts on a process-model error. The argument for such a regime, in the engineering vocabulary, is structural rather than rhetorical: the conflict is being driven by mutual process-model error, and only a r3 channel external to both sides can correct mutual error symmetrically.

This will not be built quickly. It depends on a precondition — shared technical vocabulary about what AI-infrastructure verification could even mean — that current track-1.5 dialogues are only beginning to produce. It depends, further, on geographic diversification of fab capacity, so that verification has somewhere to look that is not already monopolised. The sequencing is therefore: track-1.5 dialogue first (months), fab diversification in parallel (years), inspection regime later (decade).

4.7 The Chinese mirror

The seventh scenario is a specific consequence of the second, mirrored across the bloc boundary. Inside the Chinese state-lab system, the same Pattern A asymmetric loop operates that we identified in section 3.2 between US capital and US labs. The party centre receives capability evidence from PRC labs filtered through r1 mission-conformant disclosure; PRC labs cannot, in their own internal process, fully separate r3 evidence from r1 authority because the r1 structure is constitutive rather than contingent.

The damage is symmetric to what happens in the US case, with one critical difference. The US side is open enough that external r3 channels (independent benchmarks, AISI evaluations, academic critique) can in principle reach the apex. Whether they do depends on whether structural separation is enforceable. The Chinese side has no equivalent external channel: the structure that produces the asymmetry is the same structure that excludes external correction.

The implication is uncomfortable but operationally important. Unilateral remedies on the US side can produce evidence that PRC analysts read; capability claims that survive Western r3 evaluation also survive their own internal r1 wrapping. An externalised r3 channel becomes a correction channel for both sides at once. But the only remedy that works *directly* on the PRC side is a bilateral one — an inspection regime of the kind described in scenario 4.6, which exists outside both national hierarchies. This is one of several reasons that the bilateral remedies are higher-leverage than they appear.

4.8 The information environment degrades faster than authentication scales

The eighth scenario is about epistemics. The volume of AI-generated content already exceeds the capacity of human readers to authenticate it manually and is approaching the capacity of automated detection systems to track it. Provenance signals exist — cryptographic content credentials, watermarking — but their adoption is voluntary, and adversarial actors strip them at low cost. Information platforms whose business model is engagement-driven amplification have no incentive to favour verified content over high-engagement AI content; moderation, if it happens, imposes a unilateral cost on a single platform in a multi-platform competition for attention.

The structural diagnosis is that the controlled process — the information environment — is moving to a state in which trust in *any* signal degrades trust in *all* signals, and from that state recovery is difficult. The corrective feedback channel (measurement of provenance failure rates) does not exist at the aggregate level; only individual incidents are tracked.

The controller responsible (information regulators) tracks takedowns, not aggregate state.

The remedy class is to move provenance from the platform layer (where it can be stripped) to the capture layer (camera, OS) and the transmission layer (model output API, where it can be cryptographically attested). To this add a statutory regime for aggregate measurement of provenance-failure rate, so that the question "is the information environment getting worse, by how much" can be answered. Watermarking alone is insufficient; the analogy is with seatbelt laws — voluntary mechanisms have a ceiling; mandatory mechanisms with measurement do not.

4.9 Rules calibrated to capability tiers already obsolete

The ninth scenario is the regulatory cadence one, mirrored. Legislatures pass AI rules calibrated to a defined capability threshold — typically a specific compute or capability level named in the statute. Algorithmic-efficiency improvements move the threshold under the rule by a factor of three or four every year. The regulator's process model treats the threshold as a fixed boundary; in reality it is a moving target. Five years after the rule passes, the rule applies to capabilities the intended scope did not anticipate, while capabilities that match the intended scope have moved into a less-regulated regime.

The corrective mechanism — recalibration — depends on re-passage of primary legislation, which requires legislative attention that is itself rationed. By the time the rule is re-passed, the threshold has moved again. The remedy class is to write rules whose thresholds move with a published benchmark, and to grant a standing body the authority to recalibrate without primary legislation, subject to a notice-and-comment process. This is not novel — securities regulators in many jurisdictions recalibrate margin and capital requirements without primary legislation — but it requires a deliberate political decision to delegate, and the political appetite for delegation varies by jurisdiction.

4.10 What having nine scenarios gives us

Nine scenarios is more than a strategic essay would provide and fewer than the underlying analysis admits. The number is roughly right because each scenario corresponds to one identifiable failure mechanism. Together they cover all the unsafe control actions catalogued in the analysis; each remedy class points at specific UCAs whose failure pathway it interrupts.

The nine scenarios are not independent. Scenario 4.1 (concentration) makes scenario 4.4 (surplus capture) worse. Scenario 4.2 (ceremonial safety) makes scenario 4.8 (epistemic degradation) worse. Scenarios 4.6 and 4.7 (the geopolitical pair) interact symmetrically with each other. The remedies, in turn, are not independent: as we will see in the next section, several of them share preconditions, several form sequencing chains, and the structurally hardest remedies are exactly those whose absence would cost the most to live with.

5. What can be done

Twenty-two specific remedies follow from the analysis. They divide into five families, each addressed at one or more of the nine loss scenarios. The family-level summary first; then the sequencing.

The first family is *measurement and feedback channels*. Several of the most damaging mismatches occur because the system lacks an honest measurement of its own state. Regulators do not know the frontier-versus-open-weights gap; nobody publishes the aggregate provenance-failure rate; labour ministries track headline employment but not task-by-task displacement. Four remedies fill these channels: a quarterly assessment of the open-weights gap, a quarterly aggregate provenance-failure measurement, a task-displacement benchmark indexed by occupation cohort, and a track-1.5 dialogue process that builds shared technical vocabulary between the US and PRC technical communities about what AI-infrastructure verification could mean. Each of these is operationally close to existing instruments and politically modest. Each of them, however, is a precondition for remedies in later families.

The second family is *rules built on the new measurements*. Once the open-weights gap is published, a targeted misuse-risk regime can be calibrated narrowly enough to preserve the ecosystem below those thresholds. Once the provenance-failure rate is measured, mandatory provenance becomes specific rather than gestural. Once the task-displacement benchmark exists, automatic stabilisers — unemployment insurance, retraining credits, EITC top-ups — can be indexed to it, so that triggers fire on the fast cycle that the disturbance arrives on. Reservation requirements on leading-edge supply, capability-tier rules indexed to a moving benchmark with sunset clauses, and a standing recalibration authority for AI rules each address a specific identified pathway. A mandatory pre-deployment evaluation regime — turning the AI Safety Institutes from voluntary partners into statutory ones — closes the asymmetric loop between labs and capital.

The third family is *structural remedies*. These are the interventions that change the architecture rather than the operating parameters. Merger-review reform that captures contract-level foreclosure. Structural separation of safety governance from commercial governance inside frontier laboratories. A statutory public-benefit-corporation class with mission constraints enforceable by a public agent, not by shareholder litigation. Cohort-specific transition programmes for

early-career cognitive workers. Portable benefits that decouple healthcare and pension continuity from continuous employment. A value-added tax on inference, hypothecated to a public dividend. Sovereign-wealth participation in AI infrastructure, where the jurisdiction has a fund. These are harder to land politically and slower to take effect, but each of them addresses a structural failure that no operating-parameter adjustment can fix.

The fourth family is *the slow geopolitical remedies*. Three of the most consequential remedies operate on a decade horizon because their preconditions are physical or diplomatic. Geographic diversification of fab capacity — bringing leading-edge logic-chip production to a strategically meaningful share outside the present chokepoint — is funded but slow; the constraint is sustained political commitment across electoral cycles. The independent inspection or observation regime for AI infrastructure depends on shared technical vocabulary, which the track-1.5 process slowly produces, and on geographic diversification, so that verification has somewhere to look that is not already monopolised. A productivity-indexed dividend, of the Alaska-Permanent-Fund kind but indexed to AI productivity, operates on a decade horizon because it requires both a fiscal instrument (which the inference VAT could provide) and a governance mechanism (which sovereign wealth could provide). Each of these is the long-cycle remedy whose absence the others cannot make up for.

The fifth family is *the politically novel ones*. Two remedies do not fit cleanly into any of the four prior categories. An expedited capability-jump merger-review window, granting antitrust authorities a temporary hold on long-duration compute contracts during defined capability-jump moments, is a small intervention with a precise function but no real precedent. An employer-independent benefits regime is operationally familiar but politically contested in jurisdictions where employer-tied benefits define the post-war social contract. These are listed separately not because they are less important, but because they require political adoption on terms the system does not yet provide.

5.1 The four-phase sequence

The 22 remedies do not all enter the world at once. They enter in a specific order constrained by the dependency graph among them, and the analysis derives a four-phase implementation plan.

The first phase, months one through eighteen, is measurement. The four feedback-channel remedies — the open-weights gap, the provenance-failure rate, the task-displacement benchmark, the track-1.5 dialogue — can begin without statutory change. An AI Safety Institute could initiate the open-weights gap publication tomorrow, as a non-binding research output. A national statistics office could establish a provenance-measurement programme on existing authority. A labour-statistics office could publish a task-displacement index using existing labour data and AI task mapping. The track-1.5 dialogue is already running in several forums; sustaining and broadening it requires nothing more than philanthropic funding and continued attention. By month twelve, each of these channels will have produced a first batch of evidence; by month eighteen, the question of whether to formalise each one statutorily will have a real answer.

The second phase, months twelve through forty-eight, builds rules on the new feedback channels. This is where most of the operationally familiar remedies sit: indexed automatic stabilisers, indexed EITC, mandatory provenance, reservation requirements on leading-edge supply, capability-tier rules with sunset clauses, a recalibration authority, mandatory pre-deployment evaluation, a targeted misuse-risk regime for open weights. Each of these is a piece of legislation or a regulation written against a measurement that the first phase produced. Each is operationally close to existing instruments. Each requires political work but not novel theory.

The third phase, months thirty-six through eighty-four, is the structural remedy phase. Merger-review reform, structural separation of lab safety and commercial governance, the statutory PBC class with public-agent enforcement, cohort-specific transitions, portable benefits, the inference VAT with hypothecated dividend, sovereign-wealth participation. These require political capital that the first two phases will have generated, and an empirical record that shows whether the operating-parameter remedies were sufficient on their own (the analysis suggests they will not be, but the evidence by month thirty-six will be specific). Each of these is structurally novel and politically harder than the second-phase remedies; each addresses a pathway the second-phase remedies cannot close.

The fourth phase, months sixty through one hundred and forty-four, is the long-cycle remedy phase. Fab diversification will already be running in parallel; what completes in this phase is the inspection regime, the productivity-indexed dividend, and the expedited capability-jump review window if it has not already been adopted. These are the remedies whose absence imposes the largest tail cost — concentration unbroken, geopolitical escalation unverified, surplus capture untransmitted — and whose adoption is hardest because their political horizon exceeds electoral cycles.

The decision-memo companion to this chapter sets out the dependency graph in full and identifies, at each phase boundary, the question that the prior phase must have answered before the next phase becomes feasible. By month twelve: did anyone publish the open-weights gap? By month eighteen: is the displacement benchmark stable? By month thirty-six: have at least three of the eight phase-two rules been adopted? At each decision point, the analysis describes both what success looks like and what failure implies for downstream sequencing. This is what distinguishes a planning artefact from a wish list.

The four-phase sequence at a glance

Phase	Months	Theme	Representative remedies
1	1–18	Measurement and feedback channels	Open-weights gap publication · provenance-failure measurement · task-displacement benchmark · track-1.5 dialogue
2	12–48	Rules built on phase-1 measurements	Mandatory pre-deployment evaluation · indexed automatic stabilisers · indexed EITC · mandatory provenance · reservation requirements · capability-tier rules with sunset · recalibration authority · targeted misuse-risk regime
3	36–84	Structural remedies	Merger-review reform for contract foreclosure · structural separation of safety governance · statutory PBC class · cohort-specific transitions · portable benefits · inference VAT · sovereign-wealth participation
4	60–144	Long-cycle remedies	Geographic fab diversification · independent inspection regime · productivity-indexed dividend · expedited capability-jump review

Each phase's gating decision determines whether the next phase is structurally feasible. Phase 1 is reversible and operationally familiar; phase 4 is largely irreversible once started.

5.2 The non-state actor's leverage

The remedies are not, as a class, things that one private actor can do alone. They are political and statutory. But several of them have entry points where a non-state actor — a foundation, a research institute, a coalition of application companies, a group of labour organisations — can move first and, by moving first, make the statutory work tractable.

Phase-one measurement is the cleanest example. Each of the four channels can be initiated as a non-binding publication; once the data exists, the statutory mandate becomes specific rather than gestural. Drafting model-bill language for phase-two rules is work that legal-policy institutes do well, and adoption rates of model-bill language are unusually high in jurisdictions where the legislature is overstretched. Coalition-building among open-weights-dependent companies is work that market participants can do without waiting for political alignment. Track-1.5 dialogue is the canonical philanthropic-funded geopolitical intervention; it is not a substitute for state action, but it produces the vocabulary state action requires.

The leverage point for any non-state actor, then, is to ask which of the 22 remedies in which phase has an entry point that matches the actor's position, and to commit to the entry-point work for at least one full decision cycle. The dependency graph makes the leverage points explicit; it also makes clear which work, however well-meaning, is not on the critical path.

5.3 Reversibility

A final consideration constrains the remedy set. Most of the proposals are reversible at moderate cost within their first cycle. Statistical publications can be discontinued. Indexed triggers can be recalibrated. Pre-deployment evaluation requirements can be relaxed. Reservation requirements can be adjusted. This is structurally important because it means a risk-averse stakeholder can support phase-one and most of phase-two with limited downside. The risk of getting the rule wrong is not the risk of locking in an error.

The exceptions are the structural and slow-cycle remedies. Once established, an inspection regime is hard to withdraw because withdrawal would itself signal an escalation posture. Sovereign-wealth participation in AI infrastructure cannot be unwound without selling at unfavourable timing. Fab diversification is sunk capital. A productivity-indexed dividend, once it becomes a household income source, is politically irreversible in practice. A loss-averse stakeholder considering tail risks — the kinetic-escalation risk in scenario 4.6, the cohort-displacement risk in scenario 4.3 — would weight these irreversible remedies more heavily despite their cost. A risk-averse stakeholder weights phase-one and phase-two preferentially.

The point of putting reversibility in the analysis explicitly is that different stakeholders, with different risk tolerances, will rank the same remedies differently. The analysis does not adjudicate the ranking. It surfaces the trade-off so that the adjudication can happen on terms a deliberative process can accept.

6. What the analysis does not do

A method that produces twenty-two specific remedies is also a method that risks being read as a forecast or a manifesto. It is neither, and it is worth being explicit about what the limits are before closing.

It does not predict which of the nine loss scenarios will occur, or in what combination, or on what timescale. The scenarios are *possibility statements*: each one describes a pathway from the current structure to a specific bad outcome, and the structure does in fact have those pathways. Whether the outcomes occur depends on which control actions are issued, by which controllers, in which sequence, and with what feedback. The analysis maps the terrain; it does not tell the reader how the journey ends.

It does not settle value disputes. The remedies are derived from the loss scenarios, and the loss scenarios are derived from stakeholder-named losses. Different stakeholders would name different losses. A reader who does not consider concentrated ownership of the AI substrate to be a loss — perhaps because they believe the firms in question are well-governed and well-intentioned — will not find scenario 4.1 a problem and will not weight remedies R-1 through R-3 highly. The analysis cannot adjudicate that disagreement; it can only show that *if* the loss is taken to be a loss, *then* certain remedies follow with specific properties. The reader's own values determine which "ifs" are live.

It does not tell any particular government or firm what to do. Decisions belong to controllers; the analysis describes what each controller could in principle do and what consequences each choice would have for the controlled processes. A decision-memo addressed to a specific stakeholder — say, a foundation allocating its grant budget over five years, or a labour union prioritising its lobbying agenda, or a cabinet office choosing between policy packages — would be a different document, derived from this analysis but committing to that stakeholder's goals and constraints.

It does not, finally, exhaust the possible failure modes of the AI system. The analysis is a first cycle in the engineering sense: every loss scenario is one that the controller / control-action / feedback / process-model decomposition can surface. There may well be failure modes that this decomposition cannot see — failures that emerge from interactions across multiple loops simultaneously, failures that depend on non-controllers the analysis treats as environment, failures whose causes lie below the threshold of the current decomposition. A second cycle, with refined boundaries and a tighter decomposition, would find some of these. The analysis we have is bounded in the way every engineering analysis is bounded: by the granularity at which the system was modelled.

7. Coda: the question reformulated

We started with three questions: about jobs, about power, about geopolitics. The structured analysis allows us to reformulate each of them in a way that points at action.

The jobs question — *will AI cause mass joblessness or freer prosperity?* — is, structurally, the question of whether the labour-policy cycle (years) will be brought into alignment with the capability-deployment cycle (months) before the cohort displacement that the mismatch produces becomes irreversible. The technology produces the surplus; the institutions decide whether the surplus reaches workers, consumers, the public, or only capital. The analysis identifies the specific instruments that would close the cycle gap (indexed automatic stabilisers, a displacement benchmark, cohort-specific transitions, portable benefits) and the specific instruments that would transmit the surplus (the inference VAT, the productivity-indexed dividend, the EITC top-up, sovereign-wealth participation). Whether any of these is adopted is a political choice made under known trade-offs — not, as the strategic essay would have it, an open question with a wide range of plausible answers.

The power question — *will a few firms own everything?* — becomes, structurally, the question of whether the contract-level foreclosure mechanism currently moving compute infrastructure into a small number of hands is captured by an antitrust regime calibrated to the right cadence and the right threshold. The analysis identifies merger-review reform, reservation requirements on leading-edge supply, and an expedited capability-jump review window as the specific instruments. It also identifies the open-weights ecosystem as a price-setting mechanism whose suppression would be the worst single thing that could happen at the model layer, and identifies the specific regulatory pathway (well-meaning misuse-risk regulation that fails to recognise the price-setting function) by which the suppression could occur.

The geopolitics question — *will the US–PRC AI competition escalate?* — becomes, structurally, the question of whether a rung-3 verification channel can be installed between the two blocs before mutual process-model error converts a series of rung-2 coercive moves into a kinetic exchange. The analysis identifies track-1.5 dialogue as the precondition, fab diversification as the parallel constraint, and an infrastructure-inspection regime as the long-cycle remedy. It also identifies the symmetric Pattern-A failure on both sides that makes mutual error structurally likely, and explains why unilateral US-side remedies (especially mandatory pre-deployment evaluation) produce evidence that bypasses the rung-1 filter on the PRC side as well, making them more valuable than they appear when each is considered in isolation.

None of these reformulations resolves the questions. Each of them turns a forecasting question into a structural question, and a structural question is one a deliberative process can act on. A political community that refuses to take action on a forecasting question — because the future is uncertain, because experts disagree, because the right thing to do is unclear — can take action on a structural question, because the structural question is not "what will happen" but "what feedback channels are missing and what institutions are needed to install them." The first question admits indefinite postponement; the second does not.

That, finally, is what an engineering reading offers that an essay does not. The engineering reading produces a list of specific things that could be done, in a specific order, by specific actors, with specific costs and specific reversibility profiles. It does not promise that any of them will be done. It does promise that, if one of them is done, the system will respond in the way the analysis predicts — because the prediction is not a guess about behaviour but a derivation from structure.

Whether the structures change is, in the end, the same question as whether the political communities affected by AI choose to change them. The technology will not make the choice. Capital will not make the choice. The labs will not make the choice. The choice has to be made by the controllers whose process models the structural analysis maps and whose feedback channels the remedies repair. And the choice can only be made if the controllers see the structure clearly enough to know which levers exist.

This chapter has tried to make the structure visible. The rest is up to whoever is reading.

Control Structures in Social Systems

The Universal Control Pattern

Every social system has a hierarchical control structure with the same basic shape. Authority flows downward; feedback — ideally — flows upward. STPA analyses this structure by asking five questions of every system:

1. **Feedback richness** — how many independent channels carry information from the base to the top, and how accurately do they represent reality?
2. **Self-sealing tendency** — can the process model at the top be corrected by information from below, or does it filter out contradictory signals?
3. **Accountability voids** — where is the controller also the interested party, structurally preventing independent oversight?
4. **Circuit breakers** — what mechanisms exist to interrupt a harmful escalating loop before it causes irreversible damage?
5. **Rung match** (*social-systems extension; see Justificatory Rungs*) — at what justificatory rungs do the loop's control action and feedback operate, and is the loop symmetric? Does the system's claimed rung match its operating rung?

The four levels of the universal pattern are:

Level	Role
Supreme Authority	Issues doctrine, law, strategy, or policy
Intermediate Controllers	Interpret and transmit downward; aggregate and filter upward
Local Community	Enforces norms through social mechanisms; primary source of ground-truth feedback
Individual	Internalises norms; self-regulates; ultimate recipient of system outcomes

What varies across systems is: who occupies each level, what mechanisms connect them, and how honest the upward flow is.

Canonical reference

The technique definitions — the five diagnostic questions and the four dangerous patterns — live in [knowledge/se-techniques/control-structures/](#) and (for the rung extension) [knowledge/se-techniques/justification-rungs/](#). The per-system profiles for all ten social systems, with claimed-rung and operating-rung tags, live in [knowledge/system-catalogues/social-systems/cross-system/control-structure-profiles.md](#). The cross-system rung comparison is in [knowledge/system-catalogues/social-systems/cross-system/justification-rungs-by-system.md](#).

The Ten Systems: Control Structure Profiles

For each of the ten social systems, the same four-level pattern is filled in with the specific controllers, intermediate actors, local communities, and individuals that occupy each role, and each system is rated on the four diagnostic questions.

The complete profiles — Kingdom, Republic, Theocracy, One-Party State, Corporation, University, Military, Religion (Church), Family, Verein — and the cross-system comparison table are in [knowledge/system-catalogues/social-systems/cross-system/control-structure-profiles.md](#). A short summary of the pattern of feedback richness, self-sealing tendency, accountability voids, and circuit breakers across the ten systems appears there as a single table.

The key observation the profiles support is that the same control-structure weakness recurs in structurally identical forms

across very different systems: the monarch judging the Crown's own claims, the bishop investigating his own clergy, the commanding officer prosecuting his own troops, the board setting the pay of executives who chose the board, the peer reviewing his own competitor's work. The *institutional vocabulary* differs; the *structural pattern* is the same.

The Four Most Dangerous Structural Patterns

Across all systems, four patterns recur wherever the most serious harm occurs. The first three are domain-general; the fourth is the social-systems extension introduced in Justificatory Rungs.

Pattern 1: The Accountability Void. The controller and the interested party are the same entity. No external actor has authority to review the decision. The remedy in every case is the same: structurally separate the controller from the interested party, and give the separating body genuine authority.

Pattern 2: The Self-Sealing Process Model. The top of the control structure populates its own process model — its internal representation of reality — from sources it controls. Contradictory information is excluded by design (doctrine), by incentive (career risk for bearers of bad news), or by structure (classification, privacy norms). The model diverges from reality while the authority acts with increasing confidence. The remedy: establish independent information channels with direct access to the apex, reporting to a body outside the hierarchy.

Pattern 3: The Proxy Metric. The system measures something that is correlated with its actual goal, attaches resources and careers to that measure, and then watches as the system reorients toward producing the measure rather than the goal. Exam scores replace education. Quarterly earnings replace sustainable value creation. Publication counts replace scientific knowledge. Compliance reports replace operational readiness. The remedy: continuously examine what is being measured and whether the feedback loop connects to the goal or to a proxy, and restructure the measurement accordingly.

Pattern 4: Rung Asymmetry. The downward control action and the upward feedback on the same loop operate at different justificatory rungs. The controller transmits commands at rung 1 (authority, tradition) but the corrective signals it would need to register arrive — or are expected to arrive — at rung 3 or higher (empirical, audited, replicated). The controller's process model classifies the higher-rung feedback as out-of-band and filters it out. *Instances:* doctrinal authorities receiving abuse reports as rung-1 hostility rather than rung-3 evidence; autocracies receiving independent press reports as treason rather than information; family systems treating children's developmental signals as defiance rather than health data. The remedy: insert an independent feedback channel that operates at rung 3 (or higher), routes around the rung-1 filter, and reaches the controller's process model directly. Most of the canonical architectural remedies in the next chapter are rung-elevation moves of this kind.

These four patterns are the recurring targets of the architectural remedies in the next chapter. The technique-level statements — with the failure mode each corresponds to in a generic control loop, and the remedy preconditions — are in `knowledge/se-techniques/control-structures/dangerous-patterns.md` (Patterns 1–4) and `knowledge/se-techniques/justification-rungs/dangerous-mismatches.md` (Pattern 4 in detail, plus two further rung-specific patterns: Claimed-Rung Inflation and Cross-Loop Rung Imposition).

Architectural Remedies

The Analysis Generates the Fix

The introduction to this book used a secondary school to illustrate how structural analysis works: once you have traced *which feedback loop is misconfigured* and *why*, you know exactly which connection to change. Finland changed those connections and the outcomes changed with them.

This chapter applies the same logic across five social systems. For each, STPA identifies the unsafe control actions, traces their causal factors, and derives the architectural remedy. Where a remedy has been implemented and its effects measured, that evidence is included. The aim is to show that STPA is a design method — not merely a diagnostic — and that its prescriptions are testable.

Canonical reference

The full case studies — unsafe control actions, causal factors, remedies, and the real-world evidence for each — are in [knowledge/system-catalogues/social-systems/cross-system/remedies-case-studies.md](#). This chapter gives a narrative walk-through of the five worked cases; the knowledge file is the reference.

The Five Worked Cases

1. Democracy (Republic) — the Weimar → Basic Law transition

The Weimar Republic (1919–1933) was formally democratic, constitutionally sophisticated, and destroyed from within in under fifteen years. Its failures — no judicial review of constitutionality, an emergency-powers path from parliamentary deadlock to authoritarian rule, no defence against anti-constitutional parties, and a capturable information environment — map onto four specific unsafe control actions.

The Federal Republic of Germany's Basic Law (1949) addressed each of them: the Bundesverfassungsgericht with power to strike down laws; the constructive vote of no confidence, which lets a Chancellor be removed only by simultaneously electing a successor; Article 21 GG allowing the Constitutional Court to ban parties seeking to undermine the free democratic basic order; and public broadcasting governed by civil-society councils rather than government officials. The result is a democracy that has operated without interruption for 75 years — not because Germans are different, but because the control structure is different.

2. Corporation — Sarbanes-Oxley, Mitbestimmung, and Cadbury

The corporation's defining structural feature is that authority derives from capital ownership. The three characteristic failures are: auditor dependence on the audited client, short-term compensation structures that outlive their consequences, and employee exclusion from governance despite holding the most accurate operational process models.

Three remedies, each tested in practice: Sarbanes-Oxley severed the auditor's consulting revenue from the audited client and made financial statement accuracy a personal criminal-liability matter for CEOs and CFOs. German co-determination (Mitbestimmung) placed worker representatives on supervisory boards of companies with more than 2,000 employees, bringing a different time horizon into the decision room. The UK Cadbury Code separated CEO and Chairman roles and required independent non-executive directors on key committees. None of the three asks executives to be more virtuous; each changes the structure so that the incentives point in a different direction.

3. Religion — mandatory reporting and Vatican II

Where bishops controlled both pastoral care and external communication, the authority that received abuse reports also decided whether to report externally. Mandatory reporting laws — enacted in Australia after the 2013–2017 Royal Commission, in Ireland after the 2009 Murphy Report, and elsewhere — bypass the institutional hierarchy by legal channel, changing the cost calculus of concealment from "this protects us" to "this destroys us."

The self-sealing property of doctrinal authority is harder to fix. The Second Vatican Council (1962–1965) was a structural opening to lay participation, ecumenical engagement, and modern scholarship. Its reforms were partially reversed in subsequent decades — itself a data point: the opening worked where it was maintained, showing the structural analysis was correct, and showing that reform requires sustained commitment at the goals level, not only the

connections level.

4. University — pre-registration, Registered Reports, and DORA

Two failures dominate: selective reporting of positive results (the replication crisis: only 36% of psychology findings replicated in the 2015 Open Science Collaboration study), and impact-factor substitution for research quality (the classic proxy-metric failure).

The structural fixes: pre-registration and Registered Reports decouple the publication decision from the result, so that researchers commit to hypotheses, methods, and analysis plans before data collection. The Center for Open Science now hosts over 100,000 pre-registrations, and more than 300 journals offer Registered Reports. Pre-registered studies replicate at substantially higher rates. DORA and the UK REF's move toward article-level panel assessment disconnect the reward structure from the specific proxy metric. Both fixes change the connection between "result" and "consequence for the researcher" — and behaviour changes accordingly.

5. Military — parliamentary authorisation and independent prosecution

German armed forces cannot be deployed outside German territory without a Bundestag vote. The Wehrbeauftragter — the Parliamentary Commissioner for the Armed Forces — has independent access to any unit, any file, and any individual soldier, and reports directly to the Bundestag rather than the Defence Ministry. Germany has maintained continuous civilian control of its military since 1955; the Commissioner has regularly surfaced structural problems that command had institutional incentive to conceal.

For military justice, the US National Defense Authorization Act 2022 transferred jurisdiction over sexual assault and other serious felonies from commanding officers to independent Special Trial Counsel, removing the commanding officer from the prosecution decision entirely. Reporting rates rose; prosecution rates for referred cases rose. The structural conflict of interest — commanding officer deciding whether to prosecute conduct that reflects on the command — was identified through sustained advocacy using exactly the structural argument this book makes, and directly addressed.

Cross-System Pattern

Across all five systems, the same three structural failures recur: an accountability void, a self-sealing process model, and a proxy metric that has replaced the goal. The full cross-system table — failure type by system, with each proven remedy mapped to the corresponding control-loop fix — is in [knowledge/system-catalogues/social-systems/cross-system/remedies-case-studies.md](#).

The remedies are not generic. Each closes a specific causal pathway identified by the structural analysis — and each has been tested in practice. That precision is what distinguishes architectural reform from exhortation.

Reading the Remedies as Rung-Elevation Moves

The previous chapter introduced the seven Justificatory Rungs and the dangerous **Pattern A — Asymmetric Loop**: rung-1 control downward with rung-3 evidence trying, and failing, to feed back upward. Re-read through that lens, **all thirteen catalogued architectural remedies share a common move**: the controller's authority remains at rung 1 (legal, institutional, or sacred); a *parallel* feedback channel is installed that operates at rung 3 (independent audit, statutory investigation, replicated study, criminal-liability disclosure); and the new channel routes around the rung-1 hierarchy that previously filtered upward signal.

Remedy	Pre-remedy feedback rung	Post-remedy feedback rung
Bundesverfassungsgericht (Basic Law)	1 (parliamentary self-judgment)	2/6 (legal-rational + constitutional principle)
Sarbanes-Oxley external audit	1 (management self-reports)	3 (audited financials with criminal liability)
Mitbestimmung worker board representation	1 (management self-reviews)	6 (deliberative — multiple stakeholder perspectives)
Mandatory reporting (Religion)	1 (clergy report internally)	3 (statutory report to civil authority)
Vatican II lay councils	1 (clerical hierarchy)	1+2 (lay participation; partially reversed)

Remedy	Pre-remedy feedback rung	Post-remedy feedback rung
Pre-registration Registered Reports	/ 3 (publication-biased trial sample)	4 (cumulative evidence without publication-bias filter)
DORA / article-level assessment	1 (journal prestige proxy)	3 (article-level merit)
Parliamentary mandate (Bundeswehr)	0/1 (executive decision)	6 (parliamentary deliberation)
Wehrbeauftragter (independent military prosecutor)	1 (commander investigates own troops)	3 (independent investigation)

The full thirteen-remedy table with detailed deltas is in `knowledge/system-catalogues/social-systems/cross-system/remedies-case-studies.md §6`.

The remedies do not try to convert rung-1 controllers to higher rungs — that would dissolve the institution they protect. They *parallel* the rung-1 channel with a rung-3 (or rung-2/6) channel routed to a separate oversight body whose existence does not depend on the controller. Where remedies have been incomplete (Vatican II's partial post-Council reversal) it is because the rung-elevated channel was made revocable by the same rung-1 hierarchy it was supposed to correct — which is the precondition this book's structural fixes already require.

This is the canonical architectural pattern for repairing **Pattern A — Asymmetric Loop** in the social-systems catalogue. It is the unifying idea behind the apparently disparate reforms above.

About This Book

Origin

This book grew from a series of conversations exploring what happens when you take systems engineering seriously — not just for aircraft and software, but for the institutions that shape human life.

The starting point was a simple question: if we try to build something in life — whether in family, in politics, or in engineering — we need to distinguish between goals, requirements, functions, logical architecture, and physical implementation. What happens when you apply this discipline to religion, to political systems, to development methodologies?

The answer turned out to be more interesting than expected.

Approach

Every analysis in this book follows the same method:

1. **Define** what a system is — establishing a precise working definition before any decomposition begins
2. **Decompose** the system through the five-level SE hierarchy
3. **Compare** systems by identifying their shared platform and variation points (product line thinking)
4. **Analyse** the control structure using STPA to identify how the system can produce outcomes opposite to its goals
5. **Propose** architectural remedies — structural changes, not just policy adjustments

Technology

This book is built with [MkDocs](#) and the [Material](#) theme, hosted on [GitHub Pages](#), and deployed automatically via GitHub Actions.

Source: github.com/ccabos/systems

License

This work is a personal research and documentation project.